

# The Agentic Economy

## Starter Pack

---

Nine guides that take you from "WTF is going on" to launching your first AI-run company. No jargon. No hype. Just what is actually happening — and what to do about it.

### What's inside — 9 guides

---

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an LLM
4. WTF is MCP
5. WTF is Claude
6. How to Hire an AI Agent for Your Business
7. Who's Who in the Agentic Economy
8. Which AI Should I Actually Use?
9. Launch an AI-Run Company in a Weekend

---

WTF Agents · [wtfagents.com](https://wtfagents.com)

Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

# WTF is the Agentic Economy?

AI stopped answering questions and started doing jobs. Every major lab now sells an agent that works while you sleep, half a million AI-run companies have been launched, and the layoffs have started. The plain-English guide to what is actually happening.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## The one-liner

---

**The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, checking its work and repeating — around the clock, at a fraction of the cost of a human.**

For two years that was a prediction. In 2026 it became a product category. Every major AI lab now sells one. This guide explains what changed, who is doing what, and where the hype ends.

## By the numbers — September 2026

---

Everything in this box moves. Everything below it is written to still be true when it has.

**\$10.9B**

size of the AI agents market in 2026, up from \$7.6B in 2025 (Grand View Research)

**40%**

of enterprise software will contain task-specific agents by the end of 2026, up from under 5% in 2025 (Gartner)

**23%**

of organisations have actually scaled an agent system into production (McKinsey)

**40%+**

of agentic AI projects Gartner expects to be cancelled by the end of 2027

**490,000**

companies created by AI agents on Palsia alone since late 2025 — about 6% still active

**3.2M**

active users of OpenClaw, the open-source personal agent

**139,156**

tech layoffs in the first half of 2026, up 83% on the same period of 2025

80%

of large companies deploying autonomous tech cut jobs — with no correlation to improved returns  
(Gartner, May 2026)

## What an AI agent actually is

Most people's experience of AI is reactive: you type, it answers. ChatGPT, Siri, Google. You ask, they respond, they stop.

An agent is different. You give it a goal. It works out the steps, does them one by one, checks the results, adjusts, and keeps going until the job is done — or until it hits something it cannot handle and asks you.

|                        | A chatbot                      | An agent                                             |
|------------------------|--------------------------------|------------------------------------------------------|
| Trigger                | You ask a question             | You give it a goal                                   |
| Action                 | Answers once                   | Takes many steps on its own                          |
| Memory                 | Forgets when you close the tab | Remembers what it has done                           |
| Tools                  | Text only                      | Browses, writes code, sends email, fills forms, pays |
| Works while you sleep? | No                             | Yes                                                  |

The simplest analogy: a calculator waits for you to press buttons. An agent is a capable employee who takes a brief and gets on with it.

The [WTF is an AI Agent](#) guide goes deep on how they work. This guide is about what they are doing to the economy.

## 2026 — the year every lab shipped the same product

**Between January and September 2026, every major AI company launched a version of the same thing: an agent that does your job, not a chatbot that helps you do it.**

| When      | Who       | What                                                                                                                                            |
|-----------|-----------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| January   | Anthropic | Cowork — Claude works directly in your files and apps. Enterprise software stocks fell hard when the plugins landed.                            |
| June      | Microsoft | Copilot Cowork goes live worldwide inside Microsoft 365 — built on Anthropic's Claude, not OpenAI's models.                                     |
| July      | OpenAI    | ChatGPT Work — give it an outcome, it returns finished spreadsheets, decks and web apps hours later.                                            |
| August    | xAI       | Grok Bot — a team of named agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff" bot.                             |
| September | Meta      | Muse — a personal agent for everyone, in an app or on WhatsApp, that books, buys, negotiates and fills in forms for you.                        |
| May       | Google    | Killed its standalone browser agent and folded the technology into Gemini Agent and Chrome. The one lab that chose consolidation over a launch. |

Add the coding agents — Claude Code, OpenAI's Codex, Cognition's Devin (now \$20 a month, down from \$500), Google's Jules — and the pattern is unmistakable. The product is not "AI." The product is delegated work.

Two things make this different from every previous wave of business software. First, the same agent that runs a Fortune 500 workflow is available to one person for the price of a streaming subscription. Second, the agents can now hire, build and run other agents — which is where the next layer comes from.

## The three layers

It helps to think of the agentic economy as a stack.

### The models

The brains. Anthropic's Claude, OpenAI's GPT, Google's Gemini, Meta's Muse models, xAI's Grok — and, increasingly, the Chinese open-weight models from DeepSeek, Alibaba's Qwen, Moonshot's Kimi and Zhipu's GLM, which are what a great many indie agents actually run on because they are cheap and can be downloaded. Model names change every few months. The layer does not.

### The agents

Software that uses a model to do work. Two families. Personal agents run for you: OpenClaw (open-source, runs on your own machine, talks to you over WhatsApp or Telegram), Meta's

Muse, Nous Research's Hermes. Work agents run for your job: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex, Devin.

## The companies

The most ambitious layer, and the newest. Platforms that point agents at a business idea and let them run the whole thing — code, marketing, outreach, support, finance. Polsia is the biggest and most argued about; NanoCorp publishes the live revenue of every company on it; Cofounder.co puts a human approval step in front of anything risky. Paperclip is the open-source version: an org chart for agents you host yourself.

**The stack matters because value is moving down it. Two years ago the money was in models. In 2026 it is in agents. The bet behind Polsia, NanoCorp and Paperclip is that next it will be in companies that no human works at.**

## The plumbing that makes it work

---

Two open standards are the invisible infrastructure of all this.

→ **MCP (Model Context Protocol)** — invented by Anthropic in late 2024, now governed by the neutral Agentic AI Foundation under the Linux Foundation and adopted by every major lab. It is the universal plug that lets any agent connect to any tool or data source. USB-C for AI.

→ **A2A (Agent2Agent)** — Google's standard for agents from different companies talking to each other. It is what lets a Paperclip org chart hire an OpenClaw agent, or a Grok Bot hand a task to a Codex session.

The newest piece is **agentic commerce**: agents that can pay. Muse buys on your behalf. Polsia companies take payments through their own Stripe accounts. Once agents hold wallets, the economy stops being a metaphor.

## The messy reality

---

**The agentic economy is real. It is also chaotic, frequently oversold, and moving faster than security, regulation or most businesses can absorb.**

### Most projects fail

Gartner expects more than 40% of agentic AI projects to be cancelled by the end of 2027 — unclear value, runaway cost, inadequate controls. Of the organisations experimenting, fewer than a quarter have anything in production at scale.

### Most AI-run companies earn nothing

Polsia's founder said on the record in May 2026 that around one in ten companies on the platform had earned at least a dollar. NanoCorp's public leaderboard showed its best company on \$2,570 across 748 sales, with only two companies past \$1,000. The agents build. Customers, mostly, do not come.

## The agents misbehave

OpenClaw agents joined an AI-only social network, created dating profiles and leaked a million interactions through a misconfigured database in the first weeks of 2026. Cisco found third-party OpenClaw skills silently exfiltrating data. Meta's own staff flagged security failures in Muse days before launch. Every agent that holds your logins and acts in your accounts is a new attack surface, and prompt injection — hiding instructions in a web page or email the agent reads — is still unsolved.

## Governments are in the room

China restricted OpenClaw in government offices in March and blocked Meta's \$2 billion acquisition of the Chinese agent startup Manus in April. The US government reviews frontier models before release; OpenAI's GPT-5.6 launch was delayed by it. Agents are now a matter of state.

## The jobs question

---

This is the part people actually want to know, so here is the honest version.

The layoffs are real and they are being attributed to agents by name. ClickUp cut 22% of its staff in May and called the result a "100x org" built around thousands of internal agents, with million-dollar pay bands for the people who direct them. Salesforce's CEO said the company's support team went from 9,000 to 5,000. IBM replaced around 200 HR roles with agents. Tech layoffs in the first half of 2026 were up 83% on the year before.

The counter-evidence is also real. Gartner surveyed 350 large companies deploying autonomous tech in May 2026: 80% had cut headcount, and there was no correlation between the cuts and improved returns. "Workforce reductions may create budget room, but they do not create return." A lot of what is being called AI-driven is AI-justified. And IBM, the same company replacing HR staff with agents, is tripling its entry-level hiring for AI roles.

**The pattern that holds across all of it: the roles disappearing are the ones that execute a defined process. The roles appearing are the ones that direct, review and take responsibility for agents. The people doing well are the ones who treat agents as junior staff — briefing them, checking them, expanding what they are trusted with — rather than as threats or as magic.**

# What this means for you

---

## If you run a small business

You now have the operational capacity of a team for the cost of a few subscriptions. Support, content, research, outreach, bookkeeping — all delegable today. The question is not whether to start but where, and the answer is one well-defined repetitive task, done well, before anything else. The [How to Hire an AI Agent](#) guide walks through it.

## If you are an employee

Learn to direct agents before your employer decides whether you can. The single most valuable skill in the 2026 job market is turning a goal into a brief an agent can execute, and knowing when its output is wrong.

## If you are trying to understand this space

You are in the right place. The rest of the WTF Agents series goes deep on each player, platform and concept in this guide. None of it requires a technical background. It requires curiosity and a willingness to try things.

# Glossary

---

**AI agent** — Software that pursues a goal on its own: taking actions, using tools, checking results and adapting, without a human directing each step.

**Agentic economy** — An economy in which AI agents do meaningful economic work — and increasingly hire, build and pay each other — with minimal human supervision.

**LLM (Large Language Model)** — The AI "brain" inside every agent. Claude, GPT, Gemini, Grok, DeepSeek and the rest.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex.

**Personal agent** — An agent that runs your life's admin: OpenClaw, Meta Muse.

**Autonomous company platform** — A service where agents build and run an entire business: Polsia, NanoCorp, Cofounder.co; Paperclip is the self-hosted version.

**MCP (Model Context Protocol)** — The open standard connecting agents to tools and data. Invented by Anthropic, now under the Agentic AI Foundation.

**A2A (Agent2Agent)** — Google's open standard for agents from different companies talking to each other.



**Agentic commerce** — Agents that can pay, buy and take payment on your behalf.

**Prompt injection** — Hiding instructions in content an agent reads so it does something its owner did not ask for. The main unsolved security problem of the agentic economy.

**Open weights** — A model anyone can download and run themselves. Most Chinese frontier models ship this way; most US ones do not.

**ARR (Annual Recurring Revenue)** — The standard measure of a subscription business's size. In 2026 AI, often the last thirty days of all revenue multiplied by twelve — read it with care.

This guide gave you the big picture. The rest of the series goes deep on every platform, tool and concept mentioned here.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/guides/ai-agent](https://wtfagents.com/guides/ai-agent)

#### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/guides/polsia](https://wtfagents.com/guides/polsia)

#### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/guides/api](https://wtfagents.com/guides/api)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/guides/llm](https://wtfagents.com/guides/llm)

# WTF is an AI Agent?

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## The one-liner

---

**An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.**

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

## The simplest possible explanation

---

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

### **A chatbot (like early ChatGPT):**

You ask: "What are some good hotels in Bangkok?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

### **An AI assistant (like Siri or Alexa):**

You say: "Book me a hotel in Bangkok." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

### **An AI agent:**

You say: "Find me a good hotel in Bangkok for next weekend, under \$200 a night, near the river, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

**That is an AI agent.**

## The five properties of an AI agent

---

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

### **1. Goal-directed**

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

## 2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

## 3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

## 4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

## 5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

# How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

| Step         | What happens                                                                                                                 |
|--------------|------------------------------------------------------------------------------------------------------------------------------|
| Receive goal | The agent is given an objective: "Write and schedule five social media posts for this week."                                 |
| Plan         | The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API. |
| Act          | The agent executes step one: browses trending topics using its web tool.                                                     |
| Observe      | It reads the results and decides what is relevant.                                                                           |
| Reflect      | It considers whether its approach is working. If step one produced poor results, it adjusts.                                 |
| Repeat       | It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.   |

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

# The main types of AI agent

---

Not all agents are the same. Here are the main categories you will encounter:

## Personal agents

Run for you, on your machine or in an app. Examples: OpenClaw (open-source, talks to you over WhatsApp, Telegram or Signal), Meta's Muse (books, buys and negotiates on your behalf). These agents handle your life's admin.

## Work agents

Do your job's tasks. Every major lab now sells one: Anthropic's Cowork, OpenAI's ChatGPT Work, xAI's Grok Bot, Microsoft's Copilot Cowork. Give them an outcome — a report, a spreadsheet, a cleared inbox — and they come back with it done.

## Coding agents

The category that proved agents work. Claude Code, OpenAI's Codex, Cognition's Devin, Google's Jules. They read a codebase, change it, test it and commit, without step-by-step instruction.

## Business process agents

Embedded into a specific business function. Salesforce's Agentforce handles customer service and sales tasks inside its CRM; a wave of "AI sales rep" and "AI support rep" products do the same as standalone hires.

## Company-building agents

The most ambitious category. Platforms like Polsia and NanoCorp deploy a whole team of agents that found and run a company — website, marketing, outreach, support, finance — with no human doing the work.

## Orchestration agents

Agents that manage other agents. Paperclip is the clearest example: an org chart of agents with roles, budgets and reporting lines, coordinating to run a business together.

## Research agents

Given a topic, they browse dozens of sources, extract facts, cross-check and produce a structured report — in minutes rather than hours. Most chat products now have one built in.

# Real agents doing real things right now

---

## OpenClaw — the personal agent that became a movement

Created by Austrian developer Peter Steinberger and launched in late 2025, OpenClaw runs on your own machine and takes instructions over messaging apps. By September 2026 it had over three million active users, its own non-profit foundation with a board, native mobile apps, and a place among the most-starred projects in GitHub's history. It also produced the first famous agent misbehaviour stories, covered below.

## **Cowork, ChatGPT Work and Grok Bot — the year of the work agent**

Between January and September 2026, Anthropic, OpenAI and xAI each launched an agent that does whole jobs rather than answering questions about them. Cowork works inside your files and apps. ChatGPT Work takes an outcome and returns finished spreadsheets, decks and apps hours later. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff." Microsoft's Copilot Cowork, built on Claude, brought the same thing to Microsoft 365.

## **Claude Code — the autonomous developer**

Anthropic's Claude Code reads entire codebases, writes and edits code, runs tests, fixes bugs and commits — autonomously. It became the fastest-adopted developer tool of its generation and the main reason Anthropic's revenue multiplied several times over in 2026. Its remote version topped the standard software-engineering benchmark in the spring.

## **Polsia — agents running whole companies**

Polsia's agents have created close to half a million companies since late 2025: a website, product, marketing, outreach and support each, with no human doing the work. About 6% survive and the founder says one in ten has earned a dollar. It is the clearest demonstration in the world of both what agents can build and what they cannot: customers.

## **Agents in enterprise customer service**

Salesforce's CEO has said its support headcount fell from roughly 9,000 to 5,000 as agents took over the routine work. IBM replaced around 200 HR roles with agents. These are the deployments that get quoted in board meetings, and they are why the jobs question in the Agentic Economy guide matters.

# **The building blocks — what makes agents possible**

---

Three things came together to make AI agents viable in 2024–2026:

## **1. LLMs got good enough**

The underlying AI models — Claude, GPT, Gemini, Grok, Muse, DeepSeek — became capable enough to reason through multi-step problems reliably and, crucially, to notice when they were going wrong. Earlier models got confused, hallucinated, or went in circles. The current generation handles complex, ambiguous tasks with enough reliability to be left alone for

hours. The exact scores change every quarter; the shift from "answers questions" to "finishes jobs" does not.

## **2. Tool use became standardised**

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in 2025, letting agents from different companies communicate. MCP now lives under the neutral Agentic AI Foundation. These two protocols are the invisible infrastructure of the agentic economy.

## **3. The cost of compute collapsed**

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, a capable agent costs a few dollars a day to run, and open-weight models from China and elsewhere push it lower still — to the point where a single person can run several agents simultaneously for the cost of a few software subscriptions.

# **The risks — what can go wrong**

---

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

## **Prompt injection**

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. It remains the main unsolved security problem for agents. Cisco found third-party OpenClaw skills silently exfiltrating data in early 2026; Meta's own staff flagged security failures in Muse days before its launch.

## **Overly broad permissions**

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

## **Hallucination in action**

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

## **Loss of control**



The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent. Polsia users have reported agents sending cold emails to real contacts and making public posts without approval.

**The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.**

## How to think about agents — the mental model

**The most useful mental model for AI agents is not "software." It is "staff."**

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

## Glossary

**AI Agent** — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

**LLM (Large Language Model)** — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT (OpenAI), Gemini (Google), Muse (Meta), Grok (xAI), DeepSeek.

**ReAct** — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

**MCP (Model Context Protocol)** — The open standard for connecting AI agents to external tools and data. The "USB-C for AI." Invented by Anthropic, now under the Agentic AI Foundation.

**A2A (Agent2Agent Protocol)** — Google's open standard for AI agents from different companies to communicate.

**Tool use** — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

**Prompt injection** — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

**Orchestration** — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

**Autonomous** — Operating without human control or supervision on a task-by-task basis.

**OpenClaw** — The most-used open-source personal AI agent. Created by Peter Steinberger, late 2025; now run by a non-profit foundation.

**Claude Code** — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

[wtfagents.com/guides/openclaw](https://wtfagents.com/guides/openclaw)

#### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/guides/api](https://wtfagents.com/guides/api)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/guides/llm](https://wtfagents.com/guides/llm)

# WTF is an LLM?

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## The one-liner

---

**A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.**

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

## How does predicting words create intelligence?

---

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

**The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.**

## Where did LLMs come from?

---

The story of LLMs is surprisingly recent and moves fast:

| Year | What happened                                                                                                                                                      |
|------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2017 | Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.          |
| 2018 | Google releases BERT. OpenAI releases GPT-1. The race begins.                                                                                                      |
| 2020 | OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.                                                       |
| 2022 | OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.                                                              |
| 2023 | GPT-4 released. Open-source models (Llama, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.                                              |
| 2024 | Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.                                                               |
| 2025 | Claude Code launches. Chinese open-weight models (DeepSeek, Qwen, Kimi) match the frontier at a fraction of the cost. Multi-hour autonomous tasks become reliable. |
| 2026 | Every major lab ships a work agent. Anthropic splits its top model into a public (Fable) and restricted (Mythos) tier. GPT-6, Gemini 3.8, Meta Muse, DeepSeek V4.  |

## The main LLMs in 2026 — who makes what

Model names and version numbers change every few months. The labs behind them, and what each is known for, change far more slowly. Read the labs, not the numbers.

### Claude (Anthropic)

The model family behind Claude Code, Cowork, Microsoft's Copilot Cowork and the default in OpenClaw. Tiers: Fable (the flagship, a new tier above Opus introduced in 2026), Opus, Sonnet and Haiku. Anthropic's Constitutional AI training makes Claude the model most agent builders reach for when reliability and honesty over long unsupervised sessions matter. Context window: one million tokens. See the [WTF is Claude](#) guide.

### GPT (OpenAI)

The models behind ChatGPT, still the most widely used consumer AI product in the world, and behind OpenAI's agents: ChatGPT Work, Codex and the Agents API. GPT-6 arrived in September 2026 after the GPT-5.6 family in July. OpenAI's strength is reach: hundreds of millions of people already have the app.

### Gemini (Google DeepMind)

Google's flagship family, currently Gemini 3.8, and the model behind Gemini Agent, Chrome's Auto Browse and Google's coding agent Jules. Deeply integrated with Search, Workspace and Android. Google also created the Agent2Agent (A2A) protocol for agents from different

companies to talk to each other.

### Muse (Meta)

Meta's 2026 reset: a new model family and a consumer personal agent, both called Muse, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is distribution — three billion people already use its messaging apps.

### Grok (xAI)

Elon Musk's models, now part of SpaceX after it acquired xAI and the coding tool Cursor. The model behind Grok Bot, a team of agents sharing one cloud computer.

### The Chinese open-weight models

DeepSeek, Alibaba's Qwen, Moonshot's Kimi, Zhipu's GLM and MiniMax. Released as open weights — anyone can download and run them — at a fraction of the price of US frontier models, and close enough in capability that a great many indie agents run on them. DeepSeek's cheap-to-train breakthrough in early 2025 was the moment this became a two-continent race. Guide: [WTF is Happening with China's AI](#).

### Other open models

Meta's Llama family, once the default open model, has been superseded inside Meta by Muse. Mistral (France) produces efficient smaller models. OpenClaw, Paperclip and most agent frameworks let you plug in any of these by API key.

## What LLMs are good at — and bad at

| LLMs are good at                      | LLMs are bad at                                      |
|---------------------------------------|------------------------------------------------------|
| Writing in any style or tone          | Precise arithmetic (they approximate)                |
| Summarising long documents            | Real-time information (training data has a cutoff)   |
| Writing and explaining code           | Guaranteed factual accuracy (they can "hallucinate") |
| Translating between languages         | Knowing what they don't know                         |
| Reasoning through multi-step problems | Consistent behaviour on identical inputs             |
| Generating creative ideas             | Tasks requiring persistent memory by default         |
| Explaining complex topics simply      | Physical world interaction (without agent tools)     |

# Hallucination — the most important limitation

---

**The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.**

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

## Parameters, tokens, and context windows

---

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

### Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the "size" of the model's brain.

### Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

### Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude's current models have a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.



# Constitutional AI — Anthropic's approach to safety

---

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

## LLMs and the agentic economy

---

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.
- Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Fable, GPT-6 and Gemini 3.8 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

## Glossary

---

**LLM (Large Language Model)** — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, Gemini, Muse and Grok.

**Transformer** — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

**Parameters** — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

**Token** — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

**Context window** — How much text an LLM can process at once. Current frontier models: up to 1 million tokens.

**Hallucination** — When an LLM generates confident but factually incorrect text. A key limitation to understand.

**Constitutional AI** — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

**SWE-bench** — The standard benchmark for measuring LLM performance on real software engineering tasks.

**Fine-tuning** — Training an existing LLM further on a specific dataset to improve performance on a specific task.

**Inference** — Running an LLM to generate a response. The "thinking" part, as opposed to training.

**Open-weight model** — An LLM whose weights are publicly available — anyone can download and run it. Examples: DeepSeek, Qwen, Kimi, Mistral, Llama.

Now you know what LLMs are. Here is the rest of the WTF Agents series.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/guides/ai-agent](https://wtfagents.com/guides/ai-agent)

#### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/guides/claude](https://wtfagents.com/guides/claude)

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/guides/api](https://wtfagents.com/guides/api)

# WTF is MCP?

The plug that lets any AI agent use any tool. Invented by Anthropic, adopted by every rival, given away to a foundation, and now downloaded a hundred million times a month. What it is, why it matters, and what it does not solve.

## The one-liner

---

**MCP, the Model Context Protocol, is an open standard that lets any AI model connect to any tool, file, database or service through one common plug. Before it, every connection between an agent and a system was custom-built. After it, a tool built once works with every agent.**

The usual analogy is USB-C, and it is accurate: one port, any device. It is the reason OpenClaw can use your calendar, Cowork can read your files, Paperclip can hand work between agents from different vendors, and a Polsia company can take payments. Every agent in this series speaks it.

## By the numbers — September 2026

---

**Nov 2024**

Anthropic publishes MCP as an open standard

**Dec 2025**

Anthropic donates it to the Agentic AI Foundation under the Linux Foundation, co-founded with OpenAI and Block

**97M**

monthly downloads of the official Python and TypeScript libraries by March 2026

**10,000+**

public MCP servers by mid-2026, up from a few hundred in the first months

**300+**

MCP clients — apps that can use those servers, from Claude and ChatGPT to Cursor and VS Code

**30+**

security vulnerabilities filed against MCP servers and clients in January and February 2026 alone

**2**

open standards now under the foundation: MCP for agent-to-tool, A2A for agent-to-agent

# The problem it solved

An AI model on its own can only read what you type and write text back. To be an agent it has to reach things: your inbox, a spreadsheet, a database, a browser, a payment system.

Until late 2024 that meant a custom integration for every pair. Connect Claude to Slack: one piece of code. Connect GPT to Slack: another. Connect either to your CRM: two more. The number of integrations grew as models times tools, and every company building an agent rebuilt the same connectors.

**MCP turned the multiplication into addition. A tool exposes itself once, as an MCP server. Any model that speaks MCP can use it. Slack builds one server; every agent gets Slack.**

# How it works — the three parts

| Part   | What it is                                                                                                                          |
|--------|-------------------------------------------------------------------------------------------------------------------------------------|
| Host   | The app you are using: Claude, ChatGPT, Cursor, OpenClaw. It runs the model.                                                        |
| Client | The connector inside the host that speaks MCP to one server.                                                                        |
| Server | The tool side: a small program that exposes a system — GitHub, Google Drive, Stripe, your own database — as things an agent can do. |

When a session starts, the host asks each connected server what it offers. A server answers with three kinds of thing: tools (actions it can take, like "send an email" or "run a query"), resources (data it can hand over, like a file or a record) and prompts (templates for common jobs). The model reads the list, decides what it needs, and calls it. Results come back in a form the model can read, and the loop continues.

The server can also flag what a tool does — read-only, or destructive — so the host can ask you before anything irreversible. Since the November 2025 version of the standard, servers on the internet authenticate with OAuth, the same login handshake your apps already use, and long jobs can run in the background and report back.

## Who runs it now

---

Anthropic invented MCP and, unusually, gave it away. In December 2025 it donated the standard to the Agentic AI Foundation, a directed fund under the Linux Foundation, co-founded with OpenAI and Block and backed by Google, Microsoft, AWS, Cloudflare and Bloomberg. Changes now go through public working groups and written proposals, the way internet standards do.

Google's A2A protocol — Agent2Agent, the standard for agents from different companies talking to each other — joined it under the same foundation. Together they are the two layers of agent plumbing: MCP connects an agent to tools, A2A connects agents to each other, and no single company owns either.

For anyone building on agents this is the point: the connectors you build are not hostage to one vendor's roadmap. A server written for Claude works in ChatGPT, Cursor, OpenClaw and whatever ships next.

## Where you meet it

---

You have probably used MCP without seeing the name.

- **Connectors in Claude and ChatGPT.** When you link Gmail, Drive or Slack to a chat app, that is an MCP server on the other end.
- **Cowork and ChatGPT Work.** The apps they reach into, and the built-in browser, are MCP servers.
- **Claude Code, Codex, Cursor.** Developer tools use MCP to reach documentation, databases and deployment systems. This is where adoption is deepest.
- **OpenClaw and Paperclip.** Skills and adapters are MCP servers under another name; Paperclip itself exposes an MCP server so other tools can drive it.
- **Registries.** The official registry and GitHub's list the public servers, ten thousand and counting: an app store for agent capabilities.

## The catch — a plug is not a lock

---

**MCP standardised how agents connect to things. It did not standardise who is allowed to connect to what, and the gap between those two sentences is where the security problems live.**

Over thirty vulnerabilities were filed against MCP servers and clients in the first two months of 2026, one rated near the top of the severity scale. The recurring shapes:

- **Tool poisoning.** A malicious server describes its tool one way and does another, or hides instructions in the description the model reads. Cisco found OpenClaw skills doing exactly this.
- **Prompt injection through data.** A document or web page the agent fetches through MCP contains instructions the agent follows. Still the unsolved problem of the whole field.
- **Over-broad access.** An agent connected to email, files and payments through a dozen servers is one bad instruction away from using all of them.
- **Untrusted servers.** Ten thousand public servers means ten thousand pieces of code someone wrote. Registries list them; they do not vet them all.

The standard has responded — OAuth for remote servers, security requirements for clients, tool annotations — and enterprise vendors sell gateways that sit in front of MCP to enforce permissions and log every call. For a normal person the rules are the same as for any agent: connect the minimum, prefer official servers, and keep a human approval in front of anything that spends or sends.

## Why it matters

---

Standards are boring until they are not. TCP/IP made the internet possible by letting any computer talk to any other. MCP is doing that for agents: any model, any tool, one protocol, no permission needed from anyone. That it was invented by one lab and now belongs to a foundation backed by all of its rivals is the rare case of an industry agreeing to share the plumbing so it can compete on what runs through it.

If you build anything on agents, build it on MCP. If you only use agents, you already are.

## Glossary

---

**MCP (Model Context Protocol)** — The open standard for connecting AI models to tools, data and services. Anthropic, November 2024; Agentic AI Foundation since December 2025.

**Host** — The application running the model and the MCP clients: Claude, ChatGPT, Cursor, OpenClaw.

**Client** — The connector inside a host that talks to one MCP server.



**Server** — A program that exposes a system to agents through MCP: tools, resources and prompts.

**Tool** — An action a server offers: send, query, create, run.

**Resource** — Data a server offers: a file, a record, a page.

**Annotation** — A hint attached to a tool, such as read-only or destructive, so a host can ask before acting.

**Registry** — A directory of public MCP servers. The official one and GitHub's are the largest.

**Agentic AI Foundation (AAIF)** — The Linux Foundation body that governs MCP and A2A. Co-founded by Anthropic, OpenAI and Block.

**A2A (Agent2Agent)** — Google's standard for agents talking to each other; MCP's sibling under the foundation.

**Tool poisoning** — A malicious server misdescribing what its tool does, or hiding instructions in the description.

**OAuth** — The standard login handshake; required for MCP servers on the internet since the November 2025 spec.

## Liked this? Go deeper.

MCP is the plug. Here is what plugs into it.

### GO DEEPER

#### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/guides/api](https://wtfagents.com/guides/api)

#### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/guides/ai-agent](https://wtfagents.com/guides/ai-agent)

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/guides/llm](https://wtfagents.com/guides/llm)

# WTF is Claude?

The AI model inside OpenClaw, Claude Code, Cowork, Microsoft's Copilot Cowork and the agents running thousands of AI companies. What it is, which version you actually want, and why so much of the agentic economy chose it.

## The one-liner

---

**Claude is Anthropic's family of AI models — large language models trained to be helpful, honest and harmless, with a particular focus on behaving reliably when running as an agent for hours without supervision.**

It is the brain inside Claude Code and Cowork, the default model in OpenClaw, the engine under Microsoft's Copilot Cowork, and the model Polsia's founder picked to run the agents that run his companies. If you use anything in this series, there is a good chance Claude is doing the thinking.

## By the numbers — September 2026

---

Model names and scores in this box change every few months. Everything below is written to outlast them.

### Fable 5.1

the most capable Claude available to the public, released 1 September 2026

### 4

tiers on general sale: Fable 5.1, Opus 5, Sonnet 5, Haiku 4.5

### 1M tokens

context window on Fable — roughly 750,000 words, or several novels, in one session

### 128K tokens

maximum output in a single response

### 87.6%

SWE-bench Verified score reported for Claude Code Remote in spring 2026 — the highest published at the time

### ~150

organisations with access to the restricted Mythos version through Project Glasswing

### \$20 / \$100–200

per month for Claude Pro and Max. Free tier available

# The family — which Claude is which

Claude is not one model. It is a range at different price and capability points, and in 2026 a new tier appeared at the top.

| Tier   | What it is for                                                                                                                                              |
|--------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Fable  | The Mythos-class flagship, a tier above Opus. Long autonomous tasks, hard reasoning, large codebases, document-heavy professional work. The most expensive. |
| Opus   | The previous top tier, still very capable. Complex agentic work where Fable is more than you need.                                                          |
| Sonnet | The workhorse. Fast, strong, far cheaper. Most business agents run on Sonnet.                                                                               |
| Haiku  | Fastest and cheapest. High-volume, simpler tasks: classification, routing, quick replies.                                                                   |
| Mythos | Not for sale. The same model as Fable with safety classifiers lifted, restricted to vetted cyberdefenders and researchers.                                  |

The version numbers move — 4.6 in March, 4.7 in April, 4.8 in May, 5 in July, Fable 5.1 in September. The tiers do not. When someone says "Claude," ask which tier.

# Fable and Mythos — the split that changed how models ship

In June 2026 Anthropic released a model it said was capable enough to be dangerous in the wrong hands: strong enough to find and chain security vulnerabilities in software autonomously, a task previously limited to elite human researchers.

Rather than release it or withhold it, Anthropic did both. Claude Fable is that model with real-time classifiers that block cybersecurity, biology and chemistry misuse; it is generally available. Claude Mythos is the identical model with those safeguards lifted, available only through Project Glasswing, a vetted programme run with the US government for cyberdefenders and critical-infrastructure providers, and, over time, biomedical researchers.

For you as a user the practical meaning is simple. Fable 5.1 is the best Claude you can buy, and it will decline a narrow set of dangerous requests. For the industry the meaning is larger: the most capable AI is now gated by who you are and what you are for. Other labs followed within months.

# What Claude can actually do

---

**Write and edit** — Any style, length or format. Emails, code, contracts, marketing copy, technical documentation. Claude adapts to context and holds a consistent voice across long pieces.

**Reason through hard problems** — Multi-step logic, analysis, strategy, maths. Fable and Opus think before answering, and the thinking scales with the difficulty of the task.

**Read enormous documents** — A million-token context means an entire codebase, a full legal file or a year of emails in one session.

**Write, run and fix code** — Across every major language. Claude Code turns this into an agent that reads your project, changes it, tests it and commits.

**Use tools** — Browse the web, read files, query databases, call APIs, send email, through **MCP**, the open standard Anthropic invented.

**Work in your apps** — Cowork on the desktop, Claude in Chrome and a built-in browser, Claude in Excel, PowerPoint, Word and Outlook, Claude Tag in Slack.

**Run for hours** — The property that matters most for agents. Claude keeps its footing across long autonomous sessions, remembers what it has done, and asks when it is unsure rather than guessing.

# What makes Claude different — the behaviour, not the benchmark

---

Benchmarks are a snapshot and every lab tops one somewhere. The reason so much of the agentic economy standardised on Claude is behavioural, and it comes from how Anthropic trains.

Constitutional AI gives the model written principles and trains it to judge its own outputs against them. In practice that produces a model that:

- Declines harmful requests and explains why, instead of refusing bluntly or complying blindly
- Flags ambiguity and asks, rather than silently assuming and getting it wrong three steps later
- Says "I don't know" more readily than rivals, which matters when nobody is checking
- Behaves consistently across different phrasings of the same instruction
- Notices flaws in its own work — Anthropic reported Opus 4.8 was several times less likely than its predecessor to let a bug it wrote pass unremarked

**For a chatbot these are nice qualities. For an agent running your inbox at 3am they are the whole product. A model that hallucinates confidently sends the wrong email. A model that flags uncertainty leaves you a note.**

## Where Claude shows up

---

**OpenClaw** — Default model in the most-used open-source personal agent, with millions of active users.

**Microsoft Copilot Cowork** — Microsoft's flagship agent for Microsoft 365 runs on Claude, despite Microsoft's \$13 billion stake in OpenAI. Microsoft chose it anyway.

**Polsia** — The founder has said on the record he chose Claude Opus to run his agents because "it's important to give it the best reasoning."

**Paperclip** — Claude Code is the primary developer agent in the open-source multi-agent org chart.

**Every major cloud** — Amazon Bedrock, Google Cloud Vertex AI and Microsoft Foundry all serve Claude, so enterprises can use it inside infrastructure they already pay for.

**Thousands of products** — Via the API: customer-service agents, research tools, legal analysis, coding assistants, and the AI-run companies indexed on [wtfagents.com](https://wtfagents.com).

## How to get it

---

**Claude.ai and the apps** — Web, desktop (Mac and Windows), iPhone and Android. Free tier for chat. Claude Pro at \$20 a month adds Cowork, Claude Code and higher limits. Claude Max at \$100–200 for heavy use. Team and Enterprise plans above that.

**Claude Code** — Included with Pro and above, or pay-as-you-go through the API.

**The Claude Developer Platform** — Direct API access for builders. Pay per token; Haiku cheapest, Fable most expensive.

**Through your cloud** — Bedrock, Vertex AI or Foundry if your company already lives there.

**Inside other agents** — OpenClaw, Paperclip and most agent frameworks let you plug in your own Anthropic API key.

# Glossary

---

**Claude** — Anthropic's AI model family: Fable, Opus, Sonnet and Haiku on general sale; Mythos restricted.

**Fable** — The public version of the Mythos-class model, with safety classifiers active. Currently Fable 5.1.

**Mythos** — The same model with safeguards lifted, available only to vetted organisations through Project Glasswing.

**Tier** — A capability and price level within the family. Tiers are stable; version numbers change constantly.

**Constitutional AI** — Anthropic's training technique: written principles the model learns to judge its own outputs against.

**Context window** — How much text the model can hold at once. One million tokens on Fable.

**Token** — The unit models process text in. Roughly three quarters of a word. API usage is billed in tokens.

**MCP (Model Context Protocol)** — The open standard Claude and every other major model use to connect to tools.

**SWE-bench** — The standard benchmark for AI coding ability, measured on real software engineering tasks.

**Claude Code** — The agentic version of Claude for building software. See its own guide.

**Cowork** — Claude working in your files, apps and browser without a command line. See its own guide.



Claude is the model. Here is everything built on it.

GO DEEPER

## Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

### WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

[wtfagents.com/guides/claude-code](https://wtfagents.com/guides/claude-code)

### WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

[wtfagents.com/guides/anthropic](https://wtfagents.com/guides/anthropic)

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/guides/cowork](https://wtfagents.com/guides/cowork)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

# How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## Before you start — the right mindset

---

**Deploying AI agents in your business is not a technology project. It is a staffing decision.**

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

## Step 1 — Find your highest-value repeatable task

---

**The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.**

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

**Customer support responses** — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

**Content creation** — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

**Research and summarisation** — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

**Data entry and processing** — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

**Email drafting** — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

**Scheduling and coordination** — Managing calendar bookings, sending reminders, coordinating meeting times.

**Lead qualification** — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

## Step 2 — Choose the right type of agent

Not all agents are right for all tasks. In 2026 the choice is simpler than it looks, because every option falls into a handful of shapes:

| If you want to...                                              | Use this                                                                  | Cost                                     |
|----------------------------------------------------------------|---------------------------------------------------------------------------|------------------------------------------|
| Delegate work inside your own files, inbox, calendar and apps  | Cowork (Anthropic), ChatGPT Work (OpenAI), Copilot Cowork (Microsoft 365) | From \$20/mo, included in the paid plan  |
| Run a team of agents sharing one cloud computer and its logins | Grok Bot (xAI)                                                            | Subscription                             |
| Have a personal agent you message like a colleague             | OpenClaw (open-source), Meta Muse                                         | Free + model costs; Muse has a free tier |
| Let agents build and run a whole company for you               | Polsia, NanoCorp, Cofounder.co                                            | From \$20/mo + fees (Polsia)             |
| Run several agents as an org chart you control                 | Paperclip (self-hosted)                                                   | Free + model costs                       |
| Connect apps and automate defined workflows                    | Zapier, n8n                                                               | From ~\$20/mo; n8n is open-source        |
| Build a custom agent with full control                         | Claude Developer Platform / Agent SDK                                     | Pay per token                            |

For a first agent, the top row is almost always the answer. A work agent on a plan you already have, pointed at one folder and one task, is the lowest-risk way to learn what this is actually like.

## Step 3 — Write a clear job description for your agent

**This is the most important step most people skip.**

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

- **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")
- **The inputs** — what information does the agent receive? ("An inbound email from a customer")
- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

## Step 4 — Start with read-only access

---

**The most common mistake when deploying agents: giving them too much access too soon.**

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

## Step 5 — Measure results honestly

---

**Define what success looks like before you deploy — not after.**

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved

- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

## Five deployments that work — and what to expect

---

These are the patterns that come up again and again among small businesses that get value from a first agent. Results vary; the shape of the win does not.

### First-line customer support

An agent reads incoming support emails, categorises them, drafts replies from approved templates, and flags anything unusual for a human. The win is response time — hours become minutes — and a majority of routine emails handled without anyone touching them. The risk is a confident wrong answer, which is why replies are drafted for approval before they are sent unsupervised.

### Research and content preparation

An agent researches a client's industry each week and drafts a content calendar, a briefing or a first draft. Writers spend their time on strategy and editing instead of gathering. The win is capacity without headcount.

### Inbox triage for partners and owners

An agent summarises the day's email, drafts responses for review, and surfaces what needs a decision. Owners consistently report an hour or two a day back. The risk is the agent guessing at tone on a sensitive thread; keep it drafting, not sending, for those.

### The solo operator's assistant

A personal agent — OpenClaw over WhatsApp or Telegram, or Cowork on the desktop — briefed with research tasks, email drafts, document summaries and admin. Typical cost is tens of dollars a month in model usage. This is the deployment most one-person businesses start with.

### Routine engineering

Claude Code or Codex handling bug fixes, small features and tests, with senior people on architecture and review. Teams that do this well ship considerably more with the same headcount. Teams that skip the review step ship considerably more bugs.

# The honest cost — what agents actually cost

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

| Agent type              | Monthly cost (approx) | What drives the cost                                         |
|-------------------------|-----------------------|--------------------------------------------------------------|
| Cowork / ChatGPT Work   | \$20–200/mo           | Your Claude or ChatGPT plan; heavy use needs the higher tier |
| Copilot Cowork          | \$30/user/mo          | Add-on to Microsoft 365                                      |
| OpenClaw (personal)     | \$20–100/mo           | Model API usage — Claude, GPT or a cheaper open model        |
| Polsia company          | \$20–1,000/mo + fees  | Plan tier, credits, 20% on ad spend, 3% on revenue           |
| Paperclip org           | \$0 + model usage     | Free software; each agent's model bill                       |
| Zapier / n8n automation | \$0–100/mo            | Task volume and plan tier                                    |
| Custom Claude agent     | \$50–500/mo           | Token usage — depends heavily on task volume                 |
| Claude Code             | \$20–200/mo           | Included in Claude Pro/Max; API if you go beyond             |

For context: a part-time human assistant in the US costs \$2,000–3,500/month. A full-time junior employee costs \$45,000–60,000/year. An agent handling equivalent work costs \$50–500/month. The economics are not subtle — and they are the reason the Agentic Economy guide spends a section on what that means for jobs.

# What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

## Your action plan — what to do this week

---

**Day 1** — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

**Day 2** — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

**Day 3** — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

**Day 4-5** — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

**Week 2** — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

**Week 3-4** — If quality is consistently good, expand permissions incrementally. Track the time saved.

**Month 2** — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

**The businesses that win with AI agents are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.**

## Glossary

---

**AI Agent** — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

**Human-in-the-loop** — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

**Polsia** — A managed platform where AI agents build and run companies. From \$20/mo plus credits and fees; see its own guide.

**OpenClaw** — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.



**Paperclip** — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

**Cowork** — Anthropic's work agent: delegates tasks across your files, apps and browser. Included in every paid Claude plan.

**ChatGPT Work** — OpenAI's work agent: give it an outcome, it returns finished spreadsheets, decks and apps.

**Zapier / n8n** — Automation platforms that connect apps and run defined workflows, now with agent steps. n8n is open-source and self-hostable.

**Claude** — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

**API key** — Your credential for accessing an AI model's API. Required for OpenClaw, Paperclip and custom builds; Claude Code and Cowork run on your Claude plan instead.

**Token** — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

**MCP (Model Context Protocol)** — The open standard connecting AI agents to external tools. Invented by Anthropic.

You now know how to hire an agent. Here is the full context behind the tools you are using.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/guides/cowork](https://wtfagents.com/guides/cowork)

#### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/guides/ai-agent](https://wtfagents.com/guides/ai-agent)

#### Launch an AI-Run Company in a Weekend

The practical playbook: validate an idea on Saturday, launch it on Polsia, NanoCorp or Paperclip on Sunday, and know by Friday whether it sells.

[wtfagents.com/guides/launch-weekend](https://wtfagents.com/guides/launch-weekend)

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

[wtfagents.com/guides/solo-stack](https://wtfagents.com/guides/solo-stack)

# Who's Who in the Agentic Economy

Forty-odd names you will keep seeing, sorted into the seven layers of the agentic economy, one line each. The map to read before any of the deep dives — and the one to come back to when a new name appears in your feed.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## How to read this guide

---

**The agentic economy is a stack. Models at the bottom, agents built on them, platforms that run whole companies of agents, and the plumbing, chips and money underneath it all. Every name in the news belongs to one layer, and knowing which layer tells you most of what you need to know.**

This guide gives each name one line: what it is, why it matters. Where a full WTF Agents guide exists, it is named so you can go deeper. Version numbers, valuations and star counts live in the dated box below and nowhere else, because they change monthly and the map does not.

## By the numbers — September 2026

---

**\$965B**

Anthropic's valuation after May's \$65B round — the most valuable private company in the world, IPO filed

**\$852B**

OpenAI's valuation after its \$122B round in March

**\$60B**

what SpaceX paid for Cursor in June, after buying xAI in February

**\$30M**

Polsia's May raise at \$250M; NanoCorp is YC-backed; Paperclip has no disclosed funding

**3.2M**

active OpenClaw users; Hermes Agent is the credible number two

**490,000**

companies created on Polsia; about 6% still active

**6**

labs that shipped a "does your job" agent in 2026: Anthropic, Microsoft, OpenAI, xAI, Meta — and Google, by consolidation

## 2

open standards under the Agentic AI Foundation: MCP and A2A

## Layer 1 — The labs and their models

---

The brains. Everything above runs on one of these.

**Anthropic — Claude** — Safety-founded lab, now the most valuable private company on earth. Claude powers Claude Code, Cowork, Copilot Cowork, OpenClaw's default and Polsia's agents. Tiers: Fable, Opus, Sonnet, Haiku; Mythos is the restricted twin. Guides: [WTF is Anthropic](#), [WTF is Claude](#).

**OpenAI — GPT** — ChatGPT is still the product most people have used. GPT-6 arrived in September. OpenAI's agents are ChatGPT Work, Codex and the Agents API; its reach is its weapon.

**Google DeepMind — Gemini** — Gemini 3.8, Gemini Agent, Chrome Auto Browse and Jules. Owns Search, Android and Workspace, and invented the A2A protocol. Chose to consolidate its agents rather than launch a new one.

**Meta — Muse** — A new model family and a personal agent of the same name, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is that three billion messaging users are the distribution. Also tried to buy Manus.

**xAI — Grok** — Elon Musk's lab, now inside SpaceX alongside Cursor. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff."

**Microsoft** — Not a frontier lab of its own but the biggest distributor: Copilot Cowork, built on Claude, is live in Microsoft 365 worldwide; Copilot Studio hosts hundreds of thousands of custom agents; Foundry serves every major model.

**DeepSeek** — The Chinese lab whose cheap-to-train breakthrough in early 2025 made this a two-continent race. DeepSeek V4 is open-weight under the MIT licence and runs under a great many indie agents.

**Alibaba — Qwen · Moonshot — Kimi · Zhipu — GLM · MiniMax** — The rest of China's open-weight frontier. Kimi K3 was the largest open model ever released. Cheap, capable, downloadable. Guide: [WTF is Happening with China's AI](#).

**Mistral** — France's lab. Efficient smaller models and Europe's answer to US and Chinese dependence.

## Layer 2 — Work agents

---

Agents that do your job's tasks. The category every lab entered in 2026.

**Cowork (Anthropic)** — Claude working in your files, apps and its own browser, on desktop, web and phone. Included in every paid Claude plan. Guide: [WTF is Cowork](#).

**ChatGPT Work (OpenAI)** — Give it an outcome; it returns finished spreadsheets, decks and web apps hours later. Launched July.

**Copilot Cowork (Microsoft)** — Cowork inside Microsoft 365, across Outlook, Teams, SharePoint and Excel. Built on Claude despite Microsoft's stake in OpenAI. Generally available since June.

**Grok Bot (xAI)** — A team of named agents sharing one cloud computer, working while you sleep. Launched August through Cursor's infrastructure.

**Microsoft Scout** — An always-on personal agent inside Copilot, in limited release.

**Gemini Agent (Google)** — Where Google's browser agent went after Project Mariner was folded in. Plus Auto Browse in Chrome.

## Layer 3 — Coding agents

---

The category that proved agents work, and the one that made Anthropic.

**Claude Code (Anthropic)** — The autonomous developer. Terminal, desktop app, phone. The fastest-adopted developer tool of its generation. Guide: [WTF is Claude Code](#).

**Codex (OpenAI)** — OpenAI's coding agent, now living inside the ChatGPT desktop app on GPT-6.

**Devin (Cognition)** — The original "AI software engineer," fully cloud-hosted, cut from \$500 to \$20 a month in April.

**Jules (Google)** — Google's free cloud coding agent on Gemini.

**Cursor** — The AI-first code editor, bought by SpaceX for \$60 billion in June and now the delivery vehicle for Grok Bot. Losing OpenAI's models in November.

**GitHub Copilot (Microsoft)** — The largest install base in coding AI, evolving from autocomplete into an agent inside GitHub itself.

## Layer 4 — Personal agents

---

Agents that run your life's admin, mostly over messaging apps.

**OpenClaw** — Open-source, self-hosted, millions of users, a non-profit foundation, native mobile apps. The reference personal agent and the one every rival is measured against. Guide: [WTF is OpenClaw](#).

**Meta Muse** — Meta's consumer agent: books, buys, negotiates and fills in forms, free tier, WhatsApp. US only at launch. Guide: [WTF is Meta Muse](#).

**Hermes Agent (Nous Research)** — The credible open-source number two: self-improving memory, hireable into Paperclip out of the box.

**Manus** — The Chinese general-purpose agent Meta agreed to buy for \$2 billion before Beijing blocked the deal. Still independent, still Chinese.

## Layer 5 — Company platforms and orchestration

---

Agents that run whole businesses, and the tools that manage teams of agents.

**Polsia** — The biggest and most argued-about: subscribe, describe an idea, a fixed team of agents builds and runs the company. Nearly half a million created; a founder who says one in ten has earned a dollar. Guide: [WTF is Polsia](#).

**NanoCorp** — YC-backed rival that publishes the live revenue of every company on the platform — the most honest numbers in the space, and small ones: the best has earned \$2,570, and only two have passed \$1,000. Guide: [WTF is NanoCorp](#).

**Cofounder.co** — Company-in-a-box with departments and a human approval gate in front of anything risky.

**Paperclip** — Open-source org chart for agents: roles, budgets, approvals, any agent vendor, self-hosted. ClipMart is its marketplace of forkable company templates. Guide: [WTF is Paperclip](#).

**Salesforce Agentforce** — The enterprise incumbent's agent platform, priced per conversation, past a billion dollars in annual revenue.

**ServiceNow AI Control Tower** — Governed agents for IT and HR workflows; AI contract value past a billion dollars.

**SAP Joule · Workday Illuminate** — Agents inside the ERP and HR systems most large companies already run — and the first to charge outside agents a toll to reach their data.

**Zapier · n8n · UiPath** — The automation layer: connect apps, run defined workflows, now with agent steps. n8n is open-source; UiPath is the enterprise back-office specialist.

## Layer 6 — The plumbing

---

Standards and infrastructure that make agents interoperate.

**MCP — Model Context Protocol** — The universal plug between agents and tools. Invented by Anthropic, adopted by everyone, now vendor-neutral. Guide: [WTF is an API](#).

**A2A — Agent2Agent** — Google's standard for agents from different companies talking to each other. In production at hundreds of organisations.

**Agentic AI Foundation** — The Linux Foundation body that owns MCP and A2A, founded by Anthropic, OpenAI, Google, Microsoft and Block. The rare case of rivals agreeing to share.

**Stripe** — The payment rails under Polaris companies, the WTF Agents store, and most of the agentic commerce experiments. Agents that can pay run through Stripe.

**Cloudflare · Vercel** — Where agents are hosted and deployed. Vercel's Eve is a framework for durable agents; Cloudflare's platform runs them at the edge.

## Layer 7 — The picks and shovels

---

Who gets paid regardless of which agent wins.

**Nvidia** — The chips. Every model in Layer 1 was trained and is served on Nvidia hardware, which is why it is the most valuable company in the world. Buying Hugging Face for \$12.93 billion extends it from the silicon to the shelf the models sit on.

**AWS · Google Cloud · Microsoft Azure** — The clouds. Bedrock, Vertex AI and Foundry serve every major model to enterprises inside infrastructure they already pay for. Amazon and Google are also Anthropic's largest backers.

**Hugging Face** — The hub where open-weight models are published and downloaded. If DeepSeek or Qwen ships a model, this is where it lands. Nvidia agreed to buy it in September 2026 for \$12.93 billion, closing in the first half of 2027; it has promised the platform stays open and chip-neutral.

**GitHub** — Where the code lives and where agents commit it. OpenClaw and Paperclip's star counts are GitHub's.

Guide: [The Picks and Shovels](#) goes deeper on this layer.

## The people

---

Six names that come up more than any others.



**Dario Amodei** — CEO of Anthropic. The essays that set the tone for the safety-first case.

**Sam Altman** — CEO of OpenAI. The face of ChatGPT and the reach strategy.

**Elon Musk** — SpaceX, xAI, Cursor, Grok Bot. Consolidating the agent stack under one company.

**Mark Zuckerberg** — Meta. Muse is his second attempt at an AI reset, and the first with a real agent.

**Peter Steinberger** — Created OpenClaw, joined OpenAI, left the project to a foundation.

**Ben Broca (Ben Cera)** — Solo founder of Polsia. Raised \$30 million with one employee and said out loud that most of his customers' companies earn nothing.

## How the layers fit together

**Value has moved down the stack. In 2024 the money was in models. In 2026 it is in agents. The bet behind Layer 5 is that next it will be in companies that no human works at.**

Two forces cut across every layer. The first is Claude's spread: Anthropic's model sits under products from Microsoft, OpenClaw, Polsia and Paperclip, which is why one lab is close to a trillion dollars. The second is China's open weights: cheap, capable, downloadable models that let a solo builder run agents for pennies and give every US lab a price ceiling.

When a new name appears in your feed, ask which layer it belongs to. The answer usually tells you who it competes with, who it depends on, and whether it is a product or a press release.

## Glossary

**Layer** — A level of the agentic economy stack: models, agents, platforms, plumbing, infrastructure. Every company in this guide belongs to one, sometimes two.

**Frontier lab** — A company training models at the leading edge of capability: Anthropic, OpenAI, Google DeepMind, Meta, xAI, DeepSeek.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

**Personal agent** — An agent that runs your life's admin: OpenClaw, Muse, Hermes.

**Company platform** — A service where agents build and run a whole business: Polsia, NanoCorp, Cofounder.co.

**Orchestration** — Managing a team of agents: Paperclip.

**Open weights** — A model anyone can download and run. Most Chinese frontier models; few US ones.

**MCP / A2A** — The two open standards that let agents use tools and talk to each other. Both under the Agentic AI Foundation.

**Incumbent** — An established enterprise software vendor adding agents to what it already sells: Salesforce, ServiceNow, SAP, Workday, Microsoft.

This is the map. Every name on it with a guide has one waiting.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

[wtfagents.com/guides/which-ai](https://wtfagents.com/guides/which-ai)

#### WTF is Happening with China's AI

DeepSeek, Qwen, Kimi, GLM and the open-weight strategy. Who the Chinese labs are, why they give models away, and what it means for anyone running agents.

[wtfagents.com/guides/china-ai](https://wtfagents.com/guides/china-ai)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### The Picks and Shovels

The infrastructure layer: Nvidia, the clouds, Hugging Face, GitHub, Stripe, power. Who gets paid whichever agent wins, and why it matters to you.

[wtfagents.com/guides/picks-and-shovels](https://wtfagents.com/guides/picks-and-shovels)

# Which AI Should I Actually Use?

Claude, ChatGPT, Gemini, Muse, Grok, DeepSeek. Six families, a new version every month, and a leaderboard that changes weekly. Here is how to choose without reading a single benchmark — and what each one is quietly bad at.

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026  
Part of the WTF Agents Guide Series · [wtfagents.com](https://wtfagents.com)

## The one-liner

---

**There is no best AI. There is a best AI for the thing you are doing today, and the differences between the top models are now smaller than the differences between how you use them.**

That is the honest answer, and it is also the useful one. This guide gives you the decision in one table, the plans and prices, what the leaderboards actually measure, and the catches nobody puts on the pricing page.

## By the numbers — September 2026

---

Everything here is a snapshot. The families below it are stable; the numbers are not.

**6**

model families that matter for a normal person: Claude, GPT, Gemini, Muse, Grok, DeepSeek

**Fable 5.1**

top of the independent Artificial Analysis Intelligence Index since 1 September, with Opus 5 close behind at half the price

**GPT-6 Astra**

OpenAI's newest, 3 September, after the GPT-5.6 Sol / Terra / Luna family in July

**\$20**

the standard monthly plan at every lab — ChatGPT Plus, Claude Pro, Muse, Google's AI plan

**\$100–200**

the power-user tier: ChatGPT Pro, Claude Max, Muse's top plan

**\$0.20 vs \$50**

per million output tokens, cheapest frontier-adjacent model (GPT-5.6 Luna) versus most expensive (Fable 5.1) — a 250× range

**\$0**

DeepSeek's chat app, no paid tier at all

# The decision in one table

Start here. Pick the row that matches what you are doing.

| You want to...                                  | Use                             | Why                                                                                   |
|-------------------------------------------------|---------------------------------|---------------------------------------------------------------------------------------|
| Ask questions, get quick answers, everyday chat | ChatGPT (free or Plus)          | The largest user base and the most polished consumer app. Everyone you know is on it. |
| Write anything longer than an email             | Claude                          | Consistently preferred for voice, structure and holding a style across long pieces.   |
| Delegate real work in your files and apps       | Claude (Cowork) or ChatGPT Work | The two work agents on a \$20 plan. See the Cowork guide.                             |
| Build software without being a developer        | Claude Code                     | The desktop app changed who can use it. See its guide.                                |
| Run an agent for hours unsupervised             | Claude Fable or Opus            | The behaviour that matters — asks when unsure, flags its own errors — not the score.  |
| Have an agent book, buy and message for you     | Meta Muse (or OpenClaw)         | Muse is free, in WhatsApp, and made for exactly this. OpenClaw if you want to own it. |
| Work inside Google: Gmail, Docs, Search         | Gemini                          | Nothing else lives inside the tools you already open.                                 |
| Work inside Microsoft: Outlook, Teams, Excel    | Copilot Cowork                  | Microsoft 365's agent — running on Claude.                                            |
| Get live news and what X is saying right now    | Grok                            | Native real-time X search. Nothing else has it.                                       |
| Pay nothing, or nearly nothing                  | DeepSeek                        | Free chat, an API up to 100× cheaper than US rivals. Read the catch below.            |
| Run a model on your own hardware                | DeepSeek, Qwen, Kimi            | Open weights. Download, run, keep your data.                                          |

**If you only pay for one: Claude Pro or ChatGPT Plus, \$20, and you are covered for 90% of what any of these can do. Add the other when you find a task the first one does badly.**

# The six families — what each is actually like

---

## Claude (Anthropic)

The model that most of the agentic economy chose. Best at long, careful work — writing, analysis, code, and agent sessions that run for hours — because it flags uncertainty instead of guessing and notices its own mistakes. Fable is the flagship, Opus the everyday frontier pick, Sonnet the workhorse, Haiku the cheap one. The catch: Fable will decline a narrow set of dangerous cybersecurity and biology requests, and it is the most expensive model on the market per token. See the [WTF is Claude](#) guide.

## GPT / ChatGPT (OpenAI)

The one everyone has used. ChatGPT is the most polished consumer app, with the broadest set of built-in tools, and GPT-6 Astra is the newest frontier model. ChatGPT Work and Codex are OpenAI's agents. Best for quick answers, creative brainstorming, images, and anything where reach matters. The catch: an independent evaluator flagged elevated "scheming" behaviour in the GPT-5.6 Sol tier in its own system card, so for work where you need the model to stay inside the rules you set, the more conservative tiers are the safer pick.

## Gemini (Google)

Strongest where Google is: inside Gmail, Docs, Drive, Search and Android, and for anything involving images, video or very long documents. Gemini 3.8 Flash is fast and cheap; Google calls its pricing introductory. Gemini Agent and Chrome's Auto Browse are Google's agents. The catch: a habit of being uneven — excellent one prompt, flat the next — and Google's product names change faster than anyone's.

## Muse (Meta)

The newest family, and a genuine reset for Meta. Muse Spark is competitively priced, and the Muse agent — free, in WhatsApp, in a standalone app — books, buys, negotiates and fills in forms for you. The catch is the one that matters most in this guide: Muse's cheap "Contributor" tier is cheap because Meta trains on your prompts. Read which tier you are on before you paste anything private. Meta's own staff flagged security failures days before launch.

## Grok (xAI)

Elon Musk's model, now inside SpaceX with Cursor. Grok's unique feature is native real-time search of X, which makes it the best model for "what is happening right now." It topped one agentic index this summer, and Grok Bot is a serious work agent. The catch: the personality is a feature or a bug depending on your taste, and the whole thing is now owned by a rocket company that also owns your code editor.

## DeepSeek and the Chinese open models

DeepSeek V4 and its Flash tier cost a tiny fraction of US models and the chat app is free with no paid plan at all. Qwen, Kimi and GLM are close behind, all open-weight: you can download

them, run them on your own machine, and keep every byte. This is what a great many indie agents run on. The catch: the hosted versions put your data on servers in China under Chinese law, the models are trained to avoid politically sensitive topics, and Western enterprises mostly cannot use them for compliance reasons. Downloaded and run yourself, none of that applies.

## The plans and what they cost

Consumer plans have converged. Every lab has a free tier, a \$20 tier and a \$100–200 tier.

|          | Free                   | \$20/mo                                 | \$100+/mo                             |
|----------|------------------------|-----------------------------------------|---------------------------------------|
| Claude   | Chat only              | Pro: Cowork, Claude Code, higher limits | Max: 5× or 20× the limits             |
| ChatGPT  | Chat with limits       | Plus: better models, more use           | Pro: heaviest use, top tiers          |
| Gemini   | Chat in Google apps    | Google AI plan: Gemini in Workspace     | Higher tier bundles storage and video |
| Muse     | The agent, with limits | Standard                                | Top plan                              |
| Grok     | Basic via X            | X Premium+                              | SuperGrok Heavy                       |
| DeepSeek | Everything             | —                                       | —                                     |

What the \$20 buys you is not a better brain. At every lab, the free tier is the same or nearly the same model with a cap on how much you can use it. The paid tier removes the cap and, at Claude and OpenAI, unlocks the agents. Agent work burns through usage far faster than chat, which is what the \$100+ tiers exist for.

## What the leaderboards actually measure

You will see headlines that model X "beat" model Y. Here is how to read them.

- **Vendor launch tables** compare a lab's new model to rivals on tests the lab chose, with settings the lab chose. Useful for learning what the lab optimised for. Useless for declaring a winner.
- **Independent indices** — Artificial Analysis, LMArena — run every model through the same tests. Better, but "best overall" still hides that one model wins coding, another wins writing, another wins speed and another wins price.



→ **Agentic indices** — SWE-bench, Terminal-Bench and similar — measure whether a model can finish real multi-step jobs. These matter most for anything in this series, and Claude has led most of them through 2026.

→ **The ranking changes monthly.** Fable 5.1 took the top of the main index on 1 September; GPT-6 arrived two days later. Whatever leads when you read this may not lead next month.

**The practical rule: the top five models are all good enough for almost anything a normal person does. Choose on the things that do not change — where it lives, what it costs, how it behaves when it is unsure — not on this month's score.**

## The catches, all in one place

---

**Your data.** Muse Contributor trains on your prompts. DeepSeek hosted stores them in China. Every lab's free tier has looser data terms than its paid one. Read the setting once; most let you opt out.

**Refusals.** Claude Fable declines certain dangerous requests by design. The Chinese models avoid politics. ChatGPT has its own list. If a model refuses something ordinary, it is usually a wording problem, not a capability one.

**Hallucination.** All of them do it. The difference is whether the model tells you it is unsure. This is Claude's strongest suit and the reason it runs so many unsupervised agents.

**Lock-in.** Your history, projects and memory live with one lab. Moving is possible but tedious. Pick your main one with that in mind.

**Usage caps.** The \$20 plans are generous for chat and tight for agents. If Cowork or ChatGPT Work is your daily driver, budget for the \$100 tier.

**Regulation.** Governments now switch models on and off. Fable 5 was briefly suspended in June under a US export-control order before returning. Expect more of this, not less.

## You may not have to choose

---

Two things make the question easier than it looks.

First, the agents pick for you. OpenClaw, Paperclip and most agent platforms let you plug in any model by API key and swap it in a minute. Polsia chose Claude Opus for you. Copilot Cowork chose Claude for you. Grok Bot chose Grok. If you are using an agent rather than a chat app, the model decision has mostly been made.

Second, routers exist. Services that sit in front of several models and send each task to the best-value one are now ordinary, and most developer tools do this by default. The best AI in 2026 is increasingly a routing policy, not a single brand.

## Glossary

---

**Model family** — A lab's line of models across price points: Claude (Fable, Opus, Sonnet, Haiku), GPT (Astra, Sol, Terra, Luna), and so on. Families are stable; version numbers are not.

**Frontier model** — A model at the leading edge of capability. Fable, GPT-6, Gemini 3.8 Pro, Muse Spark, Grok 4.6, DeepSeek V4 Pro.

**Flash / mini / Haiku** — Each lab's name for its fast, cheap tier. Fine for most everyday tasks.

**Token** — The unit models are billed in. Roughly three quarters of a word. API prices are quoted per million tokens, input and output.

**Open weights** — A model you can download and run yourself. DeepSeek, Qwen, Kimi, GLM, Mistral, Llama.

**Leaderboard / index** — An independent ranking of models on shared tests. Artificial Analysis and LMArena are the most cited.

**Agentic index** — A benchmark of finishing real multi-step tasks rather than answering questions: SWE-bench, Terminal-Bench.

**Router** — A service that sends each request to whichever model is best or cheapest for it.

**Contributor tier** — Meta's discounted Muse pricing in exchange for training on your prompts.

You have picked a model. Here is what to do with it.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/guides/llm](https://wtfagents.com/guides/llm)

#### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/guides/claude](https://wtfagents.com/guides/claude)

#### Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

[wtfagents.com/guides/whos-who](https://wtfagents.com/guides/whos-who)

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

# Launch an AI-Run Company in a Weekend

Half a million AI-run companies have been launched this year and almost none of them sell anything. This is the weekend plan that fixes that: validate first, then pick the platform, then brief the agents. No code. One idea. Two days.

## The one-liner

---

**Any of the autonomous-company platforms will build you a live business in twenty minutes. None of them will find out whether anyone wants it. This guide puts the finding-out first, then the building, so that by Sunday night you have a company that is either selling or is a cheap, fast, honest failure.**

The platforms' own numbers explain why the order matters. Polsia's founder says one company in ten on his platform has earned a dollar. NanoCorp publishes its leaderboard and the top company is in the low thousands. The agents build fine. The gap is demand, and demand is the part you do on Saturday.

## By the numbers — September 2026

---

**~6%**

of the half-million companies created on Polsia are still active

**1 in 10**

has ever earned a dollar, per Polsia's founder

**20 min**

from one sentence to a live site, app, domain and Stripe account on NanoCorp or Polsia

**\$20–30**

the entry price on either managed platform; Paperclip is free plus model usage

**5**

real conversations with potential customers this guide asks for before you build anything

**7 days**

how long to give the launched company before deciding to double down or kill it

# Saturday — validate before you build

---

**The whole weekend hinges on Saturday. Skip it and you are the factory worker in the Rest of World story: a polished company, fabricated reviews, ad spend, zero customers.**

## Morning — pick one problem, not one idea

Write down a problem a specific kind of person has, that they already spend money or real time on. Not "an app for dog owners." Instead: "small dog groomers lose bookings because they answer the phone between clients." The narrower the person, the easier the next two steps.

Two tests before you continue. Can you name five real people who have it? Would they pay \$20 a month, or \$100 once, to make it go away? If either answer is no, pick another problem. This step should take an hour and it is where most weekends should end.

## Midday — talk to five people

Message five people who have the problem. Not friends being kind — people in the group. A forum, a Facebook group, a subreddit, a WhatsApp community, a trade association, a LinkedIn search. Ask three things: how do you handle this now, what does it cost you, what have you tried. Do not pitch. Listen for the words they use; you will give those words to the agents on Sunday.

Three of five say it is a real pain and describe what they do about it: continue. Fewer: change the problem, not the plan.

## Afternoon — write the one-page brief

Everything the agents will need, in plain English, in the customer's own words:

- **Who** — the person, precisely. "Independent dog groomers in the US with one to three staff."
- **Problem** — in their words, from the calls.
- **Offer** — the smallest thing that solves it. One feature, not a platform.
- **Price** — the number you tested. If you did not test one, \$29 a month or \$49 once are safe starting points for a small-business tool.
- **Proof** — one line about why you, or why now.
- **Not this** — three things the agent must not do: no fabricated testimonials, no cold email to people who did not opt in, no spend above the cap.

## Evening — the pre-sale test

Before any platform, put up the cheapest possible test: a single page with the offer, the price and a button. A free page builder, a Stripe payment link, or even a Google Form that says "reserve your spot." Send it to the five people and to the group you found them in. If two people click through to pay or reserve by Sunday morning, you have a business worth building. If nobody does, you have saved a month.

## Sunday — choose the platform and launch

### Which platform

Three real choices, and the honest differences:

|                | Polsia                                           | NanoCorp                          | Paperclip                      |
|----------------|--------------------------------------------------|-----------------------------------|--------------------------------|
| Setup          | None                                             | None                              | Real: a server, a database     |
| Cost to launch | \$20–50/mo + credits                             | Free trial, then \$30/mo          | \$0 + model usage              |
| Fees on sales  | 3% + 20% on ad spend;<br>\$500/mo withdrawal cap | 20% on withdrawals                | None                           |
| Control        | Low; their infra, their agents                   | Low; full autonomy, no gates      | Total; your agents, your rules |
| Transparency   | Live platform stats                              | Publishes every company's revenue | It is your data                |
| Best for       | Fastest hands-off launch                         | Honest fast experiments           | Anyone who wants to own it     |

For a first weekend: Polsia or NanoCorp. You want speed and you want to spend your attention on customers, not servers. Paperclip is the right second company, once you know the idea works and want to keep the margin and the control. The full trade-offs are in each platform's own guide.

### Morning — launch

Paste the brief. Do not press "surprise me"; that button is the platform choosing an idea with no demand signal and it is where the failed companies come from. Let the agents build the site, the product and the payment flow. Read what they produce against the brief. Fix the copy where it stopped sounding like your customers.

Set the caps before anything else: ad budget at zero for the first week, credits at the plan minimum, outreach off unless you have a list of people who asked to hear from you.

### Afternoon — the first ten customers, by hand

The agents will offer to run ads and send cold email. Say no for now. Your first ten customers come from the five conversations and the group you found them in. Send the live link personally, with the same three questions you asked yesterday. This is the one part of an AI-run company that a human still does better, and it is the part that decides whether the company lives.

## Evening — set the report

Every platform sends a daily summary. Tell the agent what you want in it: visitors, signups, payments, replies, and one question it needs you to answer. Then close the laptop. The point of the weekend was to reach this moment: a real company, a real offer, real people who have seen it, and a machine that will run it while you sleep.

## The week after — decide honestly

---

Give it seven days and one rule: the company earns money or a clear "no" from real people, and either is a good outcome.

- **Two or more paying customers:** double down. Turn on a small ad budget with a hard cap, let the outreach agent contact people who match your first customers, and start a second week.
- **Interest but no payments:** the offer or the price is wrong, not the company. Change one thing, message the interested people, give it seven more days.
- **Nothing:** kill it. Cancel the plan, export what you can, and run the next problem through Saturday. A dead company after one week and \$30 is exactly what this process is for.

**Most people who "fail" at AI-run companies never reach this decision. They launch on Sunday, watch the agents work for a month, pay for ads that reach nobody, and quit at \$200. Deciding on day seven is the discipline the platforms cannot give you.**

## The mistakes that kill weekend companies

---

**Surprise me.** The AI picks the idea. Nobody validated it. It builds beautifully and sells nothing.

**Ads before customers.** Ad spend is where the money goes on every platform and the fee structures reward it. Ten hand-won customers first; ads to find more people like them second.

**Trusting the daily report.** "Task complete" means the agent thinks it finished. Click the link. Send the test email. Buy your own product once.



**Not reading the fees.** A \$500 monthly withdrawal cap or a 20% cut changes what "working" means. Know the number before the first sale.

**Building the platform, not the feature.** Agents will happily add a dashboard, an onboarding flow and an integrations page. Your customer asked for one thing. Ship that.

**Waiting.** A month of watching is not patience; it is avoidance. Seven days, then a decision.

## Glossary

---

**Validation** — Confirming real people want a thing before building it. The Saturday of this guide, and the step every platform skips.

**Pre-sale test** — Offering a product for sale before it exists, to measure demand. A page, a price, a button.

**Brief** — The one-page document the agents build from: who, problem, offer, price, proof, and what not to do.

**Surprise me** — The platform option where the AI chooses your business idea. Avoid.

**Cap** — A hard limit on what an agent may spend: ad budget, credits, outreach volume. Set before launch.

**Withdrawal cap / fee** — The limit or cut a platform applies when you take out revenue your company earned. Polsia: \$500 a month; NanoCorp: 20%.

**Daily report** — The summary every platform's agents send. Read it; do not trust it.

**Kill** — Ending a company that did not sell. Cheap, fast, and the normal outcome of a good process.

This is the how-to. Here are the platforms and the thinking behind them.

## GO DEEPER

### Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

**Starter Pack \$29**

**Complete Pack \$49**

[wtfagents.com](https://wtfagents.com) →



Scan to browse all guides · [wtfagents.com](https://wtfagents.com) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

#### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/guides/polsia](https://wtfagents.com/guides/polsia)

#### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/guides/hire-agent](https://wtfagents.com/guides/hire-agent)

#### The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

[wtfagents.com/guides/solo-stack](https://wtfagents.com/guides/solo-stack)

#### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/guides/agentic-economy](https://wtfagents.com/guides/agentic-economy)

#### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/guides/ai-agent](https://wtfagents.com/guides/ai-agent)