

The Agentic Economy

Starter Pack

Five guides that take you from "WTF is going on" to hiring your first agent. No jargon. No hype. Just what is actually happening — and what to do about it.

What's inside — 5 guides

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an LLM
4. WTF is Claude
5. How to Hire an AI Agent for Your Business

WTF Agents · wtfagents.com

Part of the WTF Agents Guide Series · wtfagents.com/store

WTF is the

Agentic Economy?

The plain English guide to the biggest shift in business, employment, and technology since the internet. No jargon. No hype. Just what is actually happening — and what it means for you.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, and repeating — around the clock, at a fraction of the cost of a human.

This is not a future prediction. It is happening right now, at a scale most people have not yet registered.

When did this actually start?

The term "agentic economy" began appearing in industry writing around mid-2024. The New York Times ran one of the first mainstream articles using the phrase in September 2024. By late 2024, Sam Altman (OpenAI), Satya Nadella (Microsoft), and Dario Amodei (Anthropic) were all using it in major public speeches.

The first peer-reviewed academic paper on the subject was published in May 2025, and formally entered the Communications of the ACM in January 2026. In under eighteen months, it went from startup jargon to board-level agenda item at Fortune 500 companies.

Gartner — the world's most cited technology research firm — projects that 40% of enterprise software applications will be integrated with task-specific AI agents by the end of 2026. That is up from less than 5% in 2025.

To put that in plain English: within twelve months, nearly half of all business software will have an AI that does things, not just answers questions.

What is an AI agent, exactly?

Most people's experience of AI is reactive: you type something, it responds. ChatGPT, Siri, Google — you ask, they answer.

An AI agent is different. An agent is proactive. You give it a goal, and it figures out the steps, executes them one by one, checks the results, adjusts, and tries again — without you holding its hand through each step.

Here is the difference in a table:

	Traditional AI	AI Agent
Trigger	You ask it something	You give it a goal
Action	Responds once	Takes multiple steps autonomously
Memory	Forgets after each conversation	Remembers context across sessions
Tools	Text only	Can browse web, write code, send email, make payments
Works while you sleep?	No	Yes

The simplest analogy: a calculator waits for you to press buttons. An agent is more like a capable employee who takes a brief and gets on with the job.

The scale of what's happening

These are not startup claims. These are verified, sourced facts:

\$5.25B

size of the agentic AI market in 2024

\$52B+

projected market size by 2030 (MarketsandMarkets)

79%

of organisations report some agentic AI deployment (2025)

40%

of enterprise apps will have AI agents by end of 2026 (Gartner)

700

full-time customer service roles worth of work handled by AI at Klarna

5,000+

enterprise customers on Salesforce Agentforce alone

McKinsey estimates that generative and agentic AI could add between \$2.6 trillion and \$4.4 trillion annually to global GDP. The WEF projects 170 million new jobs created by AI by 2030 — offsetting 85 million displaced, for a net positive.

A word of caution: Gartner also predicts that 40% of agentic AI projects will be cancelled by 2027 due to unrealistic expectations. This is a revolution, but it is not magic. The technology is powerful and the hype is real — and so is the messiness.

The big players — who is actually leading this

The agentic economy is not being built by scrappy startups alone. The largest technology companies in the world are betting their futures on it.

Anthropic — the safety-first AI lab

Founded in 2021 by Dario and Daniela Amodei (former OpenAI executives) and a team of AI safety researchers. Anthropic's flagship model is Claude — the AI that powers OpenClaw, Paperclip, and dozens of the most important agent frameworks. Current valuation: \$380 billion. Backers include Amazon, Google, and Nvidia. They also invented the Model Context Protocol (MCP) — the open standard that lets AI agents connect to any tool or data source. Think of MCP as USB-C for AI.

OpenAI — the consumer and enterprise giant

The company behind ChatGPT and GPT-5.4. OpenAI launched Codex (autonomous coding agent) and Operator (computer-use agent) in 2025–2026. In February 2026, they hired Peter Steinberger (creator of OpenClaw) and became financial sponsor of the OpenClaw Foundation. Valuation: \$300 billion+.

Google DeepMind — the benchmark setter

Google's Gemini 3.1 Pro currently holds the highest verified score on SWE-bench (80.6%) — the industry standard benchmark for AI coding ability. Google also created the Agent2Agent (A2A) protocol in April 2025 — the open standard that lets AI agents from different companies talk to each other. Launched with 50+ partner companies including Salesforce, SAP, PayPal, and McKinsey.

Microsoft — the enterprise distribution machine

Invested \$10B+ in OpenAI. Distributes both OpenAI and Anthropic models via Azure. GitHub Copilot (77 million+ developers). Copilot for Microsoft 365 reporting 20–30% productivity gains in enterprise knowledge work.

Salesforce Agentforce — the enterprise deployment leader

Marc Benioff declared 2025 "the year of the agent" and bet the company on agentic workflows. Agentforce now has 5,000+ paying enterprise customers embedding AI agents into CRM, sales, and support workflows. This is the largest verified enterprise agentic deployment as of 2026.

The real-world proof — it's already happening

Klarna replaces 700 employees with AI agents

Swedish fintech Klarna deployed AI agents across its customer service operation. CEO Sebastian Siemiatkowski publicly stated the AI does the equivalent work of 700 full-time employees. The

company subsequently reduced its human customer service headcount. (Note: Klarna later qualified some of these statements — exact displacement figures remain disputed but the deployment itself is verified.)

A bot hiring humans — Rentahuman.ai

One of the most striking signals of where the economy is heading: Rentahuman.ai is a platform where AI agents hire humans — not the other way around. Agents post jobs for tasks they cannot complete autonomously (physical tasks, tasks requiring legal identity, tasks needing human judgment). Forbes called it "a new platform that flips the usual narrative about AI."

WEF Davos 2026 — 32 verified enterprise deployments

The World Economic Forum, in collaboration with Accenture, documented 32 verified large-scale enterprise AI deployments across manufacturing, healthcare, financial services, and logistics. These are not pilots. These are production systems.

The indie opportunity — where the little guy wins

Here is what makes the agentic economy genuinely different from every previous technology revolution: the barrier to entry is essentially zero.

The same tools that power Salesforce Agentforce's 5,000-customer enterprise platform are available to a single person sitting at a laptop in Mallorca. The gap between "big company" and "one person" has never been smaller.

OpenClaw — the open-source agent that went viral

OpenClaw (originally Clawdbot, then Moltbot) was built by Austrian developer Peter Steinberger and launched in November 2025. It is a free, open-source autonomous AI agent that runs on your machine and connects to Claude, GPT, or DeepSeek via messaging apps like Signal, Telegram, and Discord. By February 2026 it had 247,000 GitHub stars — one of the fastest-growing open-source projects in history. Steinberger was subsequently hired by OpenAI, and the project moved to an independent open-source foundation backed by OpenAI.

Important: OpenClaw is a personal AI agent, not a company-building platform. It is the tools layer — the thing that does the work.

Polsia — AI that runs your company while you sleep

Polsia is a managed platform founded by Ben Cera (also known as Ben Broca) in San Francisco. The model: pay \$50/month, describe a business, and AI agents handle everything — coding, marketing, customer support, operations. The founder claims \$1–1.5M ARR within 30 days of launch. These figures are unverified by independent third parties and have been debated on Reddit and in tech communities. The platform itself is real and operational. WTF Agents tracks 1,293+ companies on Polsia via the live API at polsia.imrat.com/api/data.

Paperclip — the orchestration layer

Paperclip is an open-source orchestration framework that lets you build a multi-agent "company org chart" — assigning roles, workflows, and goals across multiple AI agents simultaneously. Tagline: "Any agent, any runtime, one org chart." It integrates with Claude Code, OpenClaw, Cursor, and others. ClipMart (coming soon) will let you download entire pre-built company templates with one click. 13,500 GitHub stars in days of launch.

The protocols making it all work

Two open standards are the invisible infrastructure of the agentic economy:

→ **MCP (Model Context Protocol)** — invented by Anthropic, November 2024. Lets AI agents connect to any external tool or data source. Think USB-C for AI. Now donated to the Agentic AI Foundation (backed by Anthropic, OpenAI, Google, Microsoft, AWS).

→ **A2A (Agent2Agent Protocol)** — created by Google, April 2025. Lets AI agents from different companies talk to each other. Launched with 50+ partners including Salesforce, SAP, PayPal, McKinsey, Deloitte.

The messy reality — it's not all smooth

The agentic economy is real. It is also chaotic, occasionally dangerous, and moving faster than regulation, security, or most businesses can keep up with.

The Moltbook incident

Moltbook was a social network built exclusively for AI agents — launched January 2026 by entrepreneurs Matt Schlicht and Ben Parr. Within days, AI agents were autonomously posting to it without their human owners' knowledge. Within a week, a database misconfiguration exposed 6,000+ email addresses and 1 million agent interactions. Security firm Wiz confirmed the flaw allowed anyone to take control of any agent on the platform. Moltbook was acquired by Meta in March 2026.

OpenClaw's security problem

Cisco's AI security team tested a third-party OpenClaw skill and found it performed silent data exfiltration and prompt injection without user awareness. One of OpenClaw's own maintainers publicly warned: "If you can't understand how to run a command line, this is far too dangerous of a project for you to use safely." China restricted OpenClaw in government offices in March 2026.

The jobs question

The entry-level job market is already being affected. Administrative tasks, basic coding, research, customer support — these are where AI agents are most competent and most deployed. The WEF projects a net positive (170M new jobs vs 85M displaced) but acknowledges a "significant skills transition" is required. The transition is the hard part.

What this means for you

If you run a small business:

The tools available to you today — Claude, OpenClaw, Paperclip, Relevance AI — give you the operational capacity of a team for the cost of a few subscriptions. Customer support, content creation, data analysis, outreach: all automatable right now. The question is not whether to use these tools. It is how fast to move.

If you are an employee:

The roles most at risk are the ones involving repetitive, process-driven work. The roles being created are the ones that involve directing, overseeing, and collaborating with AI agents. The people winning right now are the ones who treat AI agents as junior colleagues, not threats.

If you are curious and want to understand this space:

You are reading the right guide. The rest of the WTF Agents series goes deeper on each platform, tool, and concept in this guide. The agentic economy does not require a computer science degree. It requires curiosity and a willingness to experiment.

Glossary

AI Agent — An AI system that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

LLM (Large Language Model) — The AI "brain" at the core of most agents. Examples: Claude (Anthropic), GPT-5.4 (OpenAI), Gemini 3.1 (Google).

MCP (Model Context Protocol) — Anthropic's open standard for connecting AI agents to external tools and data. The "USB-C for AI."

A2A (Agent2Agent Protocol) — Google's open standard for AI agents from different companies to communicate with each other.

SWE-bench — The industry benchmark for measuring how well an AI can solve real software engineering tasks. Higher = better.

Polsia — A managed platform where AI agents build and run companies autonomously. Founded by Ben Cera.

OpenClaw — A free, open-source autonomous AI agent. Runs locally, connects via messaging apps. Created by Peter Steinberger.

Paperclip — An open-source orchestration framework for running multi-agent companies.

Moltbook — A social network built for AI agents. Launched January 2026, acquired by Meta March 2026.

Constitutional AI — Anthropic's technique for training AI models to be helpful, honest, and harmless using a set of principles.

ARR (Annual Recurring Revenue) — How much money a business makes per year from recurring customers. The standard startup health metric.

Agentic Economy — The emerging economic system in which AI agents perform meaningful economic work with minimal human supervision.

Liked this? Go deeper.

This guide gave you the big picture. The WTF Agents series goes deep on every platform, tool, and concept mentioned here. Each guide is \$7 — instant PDF download.

WTF is Claude

The AI powering the agentic economy. What it is, how it works, why it matters.

wtfagents.com/store

WTF is OpenClaw

The viral open-source agent that went from 0 to 247K GitHub stars in 60 days.

wtfagents.com/store

WTF is Polsia

The platform claiming to run 1,300+ companies autonomously. What's real and what's hype.

wtfagents.com/store

WTF is an AI Agent

The deep dive on agents specifically — how they think, plan, and act.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. The product that wiped \$285B off enterprise software stocks.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 and the Complete Pack for \$49, which includes all twelve.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — all 12 guides Also free: the live AI company directory at wtfagents.com/companies

WTF is an

AI Agent?

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

The simplest possible explanation

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

A chatbot (like early ChatGPT):

You ask: "What are some good hotels in Barcelona?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

An AI assistant (like Siri or Alexa):

You say: "Book me a hotel in Barcelona." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

An AI agent:

You say: "Find me a good hotel in Barcelona for next weekend, under £200 a night, close to the Gothic Quarter, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

That is an AI agent.

The five properties of an AI agent

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

1. Goal-directed

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

Step	What happens
Receive goal	The agent is given an objective: "Write and schedule five social media posts for this week."
Plan	The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API.
Act	The agent executes step one: browses trending topics using its web tool.
Observe	It reads the results and decides what is relevant.
Reflect	It considers whether its approach is working. If step one produced poor results, it adjusts.
Repeat	It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

The main types of AI agent

Not all agents are the same. Here are the main categories you will encounter:

Personal AI agents

Run on your own machine or device. Examples: OpenClaw (connects via Signal, Telegram, Discord), Claude Code (runs in your terminal). These agents work for you personally — managing your tasks, writing your code, handling your communications.

Business process agents

Embedded into business workflows. Examples: Salesforce Agentforce (handles CRM tasks), Artisan's "Ava" (AI business development rep that researches prospects and books meetings), 11x.ai's digital workers (AI sales reps). These agents automate specific business functions.

Company-building agents

The most ambitious category. Platforms like Polsia deploy agents that not only work within a company but run entire companies autonomously — from founding to daily operations.

Orchestration agents

Agents that manage other agents. Paperclip is an example — it creates an "org chart" of agents with different roles (CEO agent, marketing agent, developer agent) that coordinate to run a business together.

Research agents

Agents designed to gather, synthesise, and analyse information. Given a topic, they browse dozens of sources, extract key facts, and produce structured reports — in minutes rather than hours.

Real agents doing real things right now

Klarna's customer service agent

Klarna, the Swedish fintech, deployed AI agents across its customer support operation. The CEO publicly stated the agents handle the equivalent of 700 full-time employees' worth of customer service work. This is one of the most cited enterprise agent deployments in the world.

OpenClaw — the viral personal agent

Created by Austrian developer Peter Steinberger and launched November 2025. OpenClaw is a personal AI agent that runs on your machine and takes instructions via Signal, Telegram, or Discord. It achieved 247,000 GitHub stars by February 2026 — one of the fastest-growing open-source projects ever. The creator was subsequently hired by OpenAI.

Artisan's Ava — the AI sales rep

"Ava" is an AI business development representative that researches potential customers, writes personalised outreach emails, follows up, and books meetings — all without human involvement.

Artisan has raised \$46 million total and is backed by Y Combinator.

Claude Code — the autonomous developer

Anthropic's Claude Code is an agent that reads entire codebases, writes and edits code, runs tests, fixes bugs, and commits to Git — autonomously. It was voted "most loved" coding tool by 46% of developers in early 2026, ahead of Cursor (19%) and GitHub Copilot (9%).

Rentahuman.ai — agents hiring humans

Perhaps the most striking example: an entire platform where AI agents post jobs and hire humans to complete tasks the agents cannot do themselves (physical tasks, tasks requiring legal identity, nuanced human judgment). Forbes called it "a platform that flips the usual AI narrative."

The building blocks — what makes agents possible

Three things came together to make AI agents viable in 2024–2026:

1. LLMs got good enough

The underlying AI models — Claude, GPT, Gemini — became capable enough to reason through multi-step problems reliably. Earlier models would get confused, hallucinate, or go in circles. The current generation handles complex, ambiguous tasks with enough reliability to be useful. Claude Opus 4.6 scores 75.6% on SWE-bench (real software engineering tasks). Gemini 3.1 Pro scores 80.6%. These numbers were unthinkable two years ago.

2. Tool use became standardised

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar, or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in April 2025, letting agents from different companies communicate with each other. These two protocols are the invisible infrastructure of the agentic economy.

3. The cost of compute collapsed

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, the cost of running a capable agent has dropped dramatically — to the point where a single person can afford to run multiple agents simultaneously for the cost of a few software subscriptions.

The risks — what can go wrong

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

Prompt injection

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. Cisco found this vulnerability in OpenClaw in early 2026.

Overly broad permissions

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account, and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

Hallucination in action

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

Loss of control

The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent.

The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.

How to think about agents — the mental model

The most useful mental model for AI agents is not "software." It is "staff."

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

Glossary

AI Agent — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

LLM (Large Language Model) — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT-5.4 (OpenAI), Gemini 3.1 (Google).

ReAct — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

MCP (Model Context Protocol) — Anthropic's open standard for connecting AI agents to external tools and data. The "USB-C for AI." Published November 2024.

A2A (Agent2Agent Protocol) — Google's open standard for AI agents from different companies to communicate. Launched April 2025.

Tool use — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

Prompt injection — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

Orchestration — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

Autonomous — Operating without human control or supervision on a task-by-task basis.

OpenClaw — The most-starred open-source personal AI agent. Created by Peter Steinberger, November 2025. 247,000 GitHub stars by February 2026.

Claude Code — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

Liked this? Go deeper.

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

WTF is the Agentic Economy

The big picture. Market size, big players, real numbers. Start here if you haven't already.

wtfagents.com/store

WTF is OpenClaw

The viral open-source agent. 247K GitHub stars, security controversies, and what it actually does.

wtfagents.com/store

WTF is Claude

The AI powering most of the agents in this guide. What it is, how it works, why it leads.

wtfagents.com/store

WTF is an LLM

The deep dive on the brain inside every agent. Plain English, no maths.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. The product that wiped \$285B off enterprise software stocks.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 and the Complete Pack for \$49, which includes all twelve.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — all 12 guides Also free: the live AI company directory at wtfagents.com/companies

WTF is an LLM?

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

How does predicting words create intelligence?

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.

Where did LLMs come from?

The story of LLMs is surprisingly recent and moves fast:

Year	What happened
2017	Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.
2018	Google releases BERT. OpenAI releases GPT-1. The race begins.
2020	OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.
2022	OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.
2023	GPT-4 released. Open-source models (LLaMA, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.
2024	Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.
2025–26	Claude Opus 4.6 (75.6% SWE-bench), Gemini 3.1 Pro (80.6%), GPT-5.4. Multi-hour autonomous tasks become reliable.

The main LLMs in 2026 — who makes what

There are three dominant LLM providers and a growing ecosystem of open-source alternatives:

Claude (Anthropic)

Anthropic's flagship model family. Current version: Claude 4.6, with Opus (most powerful) and Sonnet (faster, more affordable) variants. Claude Opus 4.6 scores 75.6% on SWE-bench — the standard benchmark for software engineering tasks. Claude is the default model for OpenClaw and widely used in agent frameworks. Anthropic's focus on safety and Constitutional AI training makes Claude the preferred choice for agentic deployments where reliability and honesty matter. Context window: up to 1 million tokens (beta). Anthropic valuation: \$380 billion.

GPT-5.4 (OpenAI)

OpenAI's latest flagship model as of March 2026. Native computer use capability. 1 million token context window in Codex. Powers ChatGPT, which remains the most widely used consumer AI product in the world. OpenAI also released Codex — an autonomous coding agent built on GPT-5.4. OpenAI valuation: \$300 billion+.

Gemini 3.1 Pro (Google DeepMind)

Google's flagship model. Currently holds the highest score on SWE-bench Verified (80.6%). Deeply integrated with Google's ecosystem — Search, Workspace, Android. Google also created the Agent2Agent (A2A) protocol for agent interoperability. Available via Google Cloud (Vertex AI).

Open-source models (LLaMA, Mistral, DeepSeek)

A growing ecosystem of models that anyone can run on their own hardware. Meta's LLaMA family is the most widely used. Mistral (French AI lab) produces highly efficient smaller models. DeepSeek (Chinese) produced models that competed with GPT-4 at a fraction of the training cost — a significant

moment in early 2025. OpenClaw supports DeepSeek as an alternative to Claude or GPT.

What LLMs are good at — and bad at

LLMs are good at	LLMs are bad at
Writing in any style or tone	Precise arithmetic (they approximate)
Summarising long documents	Real-time information (training data has a cutoff)
Writing and explaining code	Guaranteed factual accuracy (they can "hallucinate")
Translating between languages	Knowing what they don't know
Reasoning through multi-step problems	Consistent behaviour on identical inputs
Generating creative ideas	Tasks requiring persistent memory by default
Explaining complex topics simply	Physical world interaction (without agent tools)

Hallucination — the most important limitation

The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

Parameters, tokens, and context windows

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the

"size" of the model's brain.

Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude Opus 4.6 has a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.

Constitutional AI — Anthropic's approach to safety

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

LLMs and the agentic economy

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.

→ Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Claude 4.6 and Gemini 3.1 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

Glossary

LLM (Large Language Model) — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, and Gemini.

Transformer — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

Parameters — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

Token — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

Context window — How much text an LLM can process at once. Claude Opus 4.6: up to 1 million tokens.

Hallucination — When an LLM generates confident but factually incorrect text. A key limitation to understand.

Constitutional AI — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

SWE-bench — The standard benchmark for measuring LLM performance on real software engineering tasks.

Fine-tuning — Training an existing LLM further on a specific dataset to improve performance on a specific task.

Inference — Running an LLM to generate a response. The "thinking" part, as opposed to training.

Open-source model — An LLM whose weights are publicly available — anyone can run it. Examples: LLaMA, Mistral, DeepSeek.

Liked this? Go deeper.

Now you know what LLMs are. Here is the rest of the WTF Agents series.

WTF is Claude

The LLM powering most of the agentic economy. Deep dive on Anthropic's flagship model.

wtfagents.com/store

WTF is an AI Agent

How LLMs become agents — the tools, loops, and protocols that make it possible.

wtfagents.com/store

WTF is the Agentic Economy

The big picture — market size, major players, real numbers.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. The product that wiped \$285B off enterprise software stocks.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 and the Complete Pack for \$49, which includes all twelve.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — all 12 guides Also free: the live AI company directory at wtfagents.com/companies

WTF Agents · wtfagents.com · The autonomous company economy is here. WTF is happening.
© 2026 WTF Agents. All rights reserved.

WTF is

Claude?

The AI model powering OpenClaw, Claude Code, and much of the agentic economy. Built by Anthropic with safety at its core. As of March 2026, one of the most capable and widely deployed AI models in the world.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Claude is Anthropic's AI model — a Large Language Model trained to be helpful, honest, and harmless, with a particular focus on safety and reliability for agentic use.

It is the brain inside OpenClaw (the most-starred open-source agent), the engine behind Claude Code (the most-loved coding tool among developers), and one of the three frontier models competing at the top of every major AI benchmark.

The model family — what's available in March 2026

Claude is not a single model — it is a family of models at different price and capability points:

Claude Opus 4.6

The most powerful model in the family. Best for complex reasoning, long autonomous tasks, and demanding agentic workflows. SWE-bench score: 75.6% (real software engineering tasks). Context window: up to 1 million tokens (beta). Output: up to 128,000 tokens.

Claude Sonnet 4.6

The mid-tier model. Faster than Opus, more affordable, with strong agentic capabilities. Best for most business use cases where you need capability without Opus-level cost.

Claude Haiku 4.x

The fastest and most cost-effective model. Best for high-volume, lower-complexity tasks where speed and cost matter more than raw capability.

How Claude compares to the competition

As of March 2026, three frontier models compete at the top of AI benchmarks:

Model	Made by	SWE-bench	Key strength
Claude Opus 4.6	Anthropic	75.6%	Agentic reliability, safety, honesty
Gemini 3.1 Pro	Google	80.6%	Highest benchmark score, Google ecosystem
GPT-5.4	OpenAI	57.7%*	Native computer use, consumer reach

* GPT-5.4 SWE-bench Pro score. Benchmark methodologies vary.

Raw benchmark scores tell part of the story. What they do not capture: Claude's reputation for being more honest, more nuanced, and more reliable in real-world agentic deployments. Developers building serious autonomous systems consistently report Claude behaving more predictably than alternatives — declining harmful requests more gracefully, flagging ambiguity rather than guessing, and handling edge cases more sensibly.

What Claude can actually do

Write and edit — Any style, any length, any format. Essays, code, emails, legal documents, marketing copy, technical documentation. Claude adapts to context better than almost any other model.

Reason through complex problems — Multi-step logical problems, mathematical reasoning, strategic analysis. Claude thinks through problems step by step rather than pattern-matching to a quick answer.

Read and analyse long documents — With a 1 million token context window (beta), Claude can read entire codebases, long legal documents, or research papers in a single session.

Write, debug, and explain code — Across all major programming languages. Claude Code (the agentic version) takes this further — it reads codebases, writes code, runs tests, and commits to Git autonomously.

Use tools via MCP — Browse the web, read files, query databases, call APIs, send emails — through the Model Context Protocol that Anthropic invented.

Maintain long conversations — Unlike models with small context windows, Claude can maintain coherent conversations and task execution over very long sessions — critical for autonomous agent work.

Be honest about uncertainty — A genuinely unusual capability: Claude is trained to say "I don't know" rather than confidently hallucinate. Constitutional AI makes this more reliable than in competing models.

What makes Claude different — Constitutional AI in practice

Every major AI model is trained to be helpful. What makes Claude different is how it was trained to handle the hard cases.

Constitutional AI (CAI) — Anthropic's training technique — gives Claude a set of principles to evaluate its own outputs against. The result is a model that:

- Declines harmful requests more gracefully — explaining why rather than just refusing.
- Flags ambiguity rather than making assumptions that could cause downstream problems.
- Maintains consistent values across different phrasings of the same request.
- Is more honest about what it does not know.

For agentic use — where Claude is running autonomously for hours, making judgment calls constantly without human supervision — these properties matter enormously. An agent that hallucinates confidently, makes assumptions silently, or behaves inconsistently is a dangerous agent. Claude's training makes it significantly more reliable in these scenarios.

Claude's role in the agentic economy

Claude is not just a chatbot. It is the infrastructure of the agentic economy.

OpenClaw's default LLM — The most-starred open-source agent project uses Claude as its primary model. 247,000 GitHub stars worth of developers are building with Claude under the hood.

Claude Code — Anthropic's own agentic coding tool — ranked "most loved" by 46% of developers in 2026, ahead of Cursor (19%) and GitHub Copilot (9%).

Paperclip integration — Paperclip — the multi-agent company orchestration framework — lists Claude Code as a primary supported agent.

Thousands of applications — Via the Anthropic API, Claude powers customer service agents, research tools, writing assistants, legal analysis tools, and more across thousands of businesses.

AWS and Google Cloud — Available via Amazon Bedrock and Google Vertex AI — the two largest enterprise cloud platforms. This means Claude is accessible inside existing enterprise infrastructure without new vendor relationships.

How to access Claude

Claude.ai — The consumer chat interface. Free tier available; Claude Pro (\$20/month) for higher limits and priority access.

Anthropic API — Direct API access for developers. Pay per token. Pricing varies by model (Haiku cheapest, Opus most expensive).

Amazon Bedrock — Claude via AWS. For teams already in the AWS ecosystem.

Google Vertex AI — Claude via Google Cloud. For teams already in the Google ecosystem.

OpenClaw — Free. Connects to Claude via your own Anthropic API key.

Claude Code — Available as a command-line tool. Requires an Anthropic API key.

Glossary

Claude — Anthropic's AI model family. Includes Opus (most powerful), Sonnet (balanced), and Haiku (fastest).

Claude Opus 4.6 — The flagship model. 75.6% SWE-bench, 1M token context window, 128K output tokens.

Constitutional AI (CAI) — Anthropic's training technique. Gives Claude principles to evaluate its own outputs. Makes Claude more reliable and honest.

SWE-bench — The industry benchmark for AI coding ability. Measures how well a model can solve real software engineering problems.

Context window — How much text Claude can process at once. Claude Opus 4.6: up to 1 million tokens (beta) — enough for entire codebases.

MCP (Model Context Protocol) — The open standard Claude uses to connect to external tools. Invented by Anthropic, now an open standard.

Claude Code — The agentic version of Claude for coding. Reads codebases, writes code, runs tests, commits to Git — autonomously.

Anthropic API — Developer access to Claude. Pay per token. Required for OpenClaw, Paperclip, and custom agent builds.

Amazon Bedrock — AWS's managed AI service. Hosts Claude for enterprise teams in the AWS ecosystem.

Token — The unit Claude uses to process text. Roughly 0.75 words. You pay for API usage in tokens.

Liked this? Go deeper.

Claude is the model. Here is everything around it.

WTF is Claude Code

Claude acting as an autonomous developer. The full story on the most-loved coding agent.

wtfagents.com/store

WTF is Anthropic

The company that built Claude. The founding story, the safety mission, the \$380B valuation.

wtfagents.com/store

WTF is an LLM

The technical foundation behind Claude. Plain English, no maths.

wtfagents.com/store

WTF is OpenClaw

The viral open-source agent that uses Claude as its default brain.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. The product that wiped \$285B off enterprise software stocks.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 and the Complete Pack for \$49, which includes all twelve.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — all 12 guides Also free: the live AI company directory at wtfagents.com/companies

How to Hire an AI Agent

for Your Business

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

Before you start — the right mindset

Deploying AI agents in your business is not a technology project. It is a staffing decision.

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

Step 1 — Find your highest-value repeatable task

The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

Customer support responses — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

Content creation — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

Research and summarisation — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

Data entry and processing — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

Email drafting — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

Scheduling and coordination — Managing calendar bookings, sending reminders, coordinating meeting times.

Lead qualification — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

Step 2 — Choose the right type of agent

Not all agents are right for all tasks. Here is a simple decision framework:

If you want to...	Use this	Cost
Launch an AI-run company with no technical setup	Polsia (polsia.com)	\$50/mo + 20% revenue share
Deploy a personal AI assistant that takes actions via messaging apps	OpenClaw (openclaw.ai)	Free + LLM API costs
Build a multi-agent team to run a business	Paperclip (paperclip.ing)	Free (self-hosted)
Automate specific business workflows without coding	Relevance AI (relevanceai.com)	Free tier available
Automate repetitive tasks and connect apps	Zapier AI (zapier.com)	From \$19.99/mo
Build a custom agent with full control	Anthropic API + Claude	Pay per token

Step 3 — Write a clear job description for your agent

This is the most important step most people skip.

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

- **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")
- **The inputs** — what information does the agent receive? ("An inbound email from a customer")
- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply

using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

Step 4 — Start with read-only access

The most common mistake when deploying agents: giving them too much access too soon.

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

Step 5 — Measure results honestly

Define what success looks like before you deploy — not after.

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved
- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

Real examples — what businesses are doing right now

E-commerce store (5 employees)

Deployed an agent to handle first-line customer support emails. The agent categorises enquiries, drafts replies using approved templates, and flags unusual requests for human review. Result: response time dropped from 6 hours to 20 minutes. 80% of emails handled without human

involvement.

Marketing agency (12 employees)

Deployed an agent to research clients' industries and draft weekly content calendars. Copywriters now spend time on strategy and editing, not research. Result: capacity increased 40% without additional headcount.

Accounting firm (3 partners)

Deployed an agent to summarise client emails and draft responses for partner review. Result: 2 hours per day saved per partner. Used the time to take on 3 additional clients.

Solo consultant

Deployed OpenClaw as a personal assistant — briefing it via Telegram with research tasks, email drafts, and document summaries. Cost: approximately \$30/month in LLM API costs. Time saved: 10+ hours per week.

SaaS startup (8 people)

Deployed Claude Code to handle routine bug fixes and feature requests. Senior engineers now focus on architecture and complex problems. Result: shipped 3x more features in Q1 2026 with the same team.

The honest cost — what agents actually cost

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

Agent type	Monthly cost (approx)	What drives the cost
Polsia company	\$50/mo + 20% revenue	Flat fee plus revenue share
OpenClaw (personal)	\$20-100/mo	Anthropic or OpenAI API usage
Relevance AI	\$0-299/mo	Free tier to enterprise plans
Zapier AI	\$20-100/mo	Task volume and plan tier
Custom Claude agent	\$50-500/mo	Token usage — depends heavily on task volume
Claude Code	\$100-500/mo	API token usage — intensive for large codebases

For context: a part-time human assistant in the UK costs £1,500-2,500/month. A full-time junior employee costs £25,000-35,000/year. An agent handling equivalent work costs £50-500/month. The economics are not subtle.

What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

Your action plan — what to do this week

Day 1 — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

Day 2 — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

Day 3 — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

Day 4-5 — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

Week 2 — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

Week 3-4 — If quality is consistently good, expand permissions incrementally. Track the time saved.

Month 2 — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

The businesses that win with AI agents in 2026 are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.

Glossary

AI Agent — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

Human-in-the-loop — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

Polsia — A managed platform where AI agents build and run companies. \$50/mo + 20% revenue share.

OpenClaw — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

Paperclip — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

Relevance AI — No-code/low-code platform for building AI agents for specific business workflows. \$37M raised.

Zapier AI — Major automation platform with AI agent capabilities. Connects thousands of apps.

Claude — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

API key — Your credential for accessing an AI model's API. Required for OpenClaw, Claude Code, and custom builds.

Token — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

MCP (Model Context Protocol) — The open standard connecting AI agents to external tools. Invented by Anthropic.

Liked this? Go deeper.

You now know how to hire an agent. Here is the full context behind the tools you are using.

WTF is an AI Agent

The deep dive on what agents actually are — how they think, plan, and act.

wtfagents.com/store

WTF is OpenClaw

The viral open-source agent. Full story including security considerations.

wtfagents.com/store

WTF is Claude

The AI model powering most of the agents in this guide.

wtfagents.com/store

WTF is the Agentic Economy

The big picture — where your new agent fits in a \$52B+ market.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. The product that wiped \$285B off enterprise software stocks.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 and the Complete Pack for \$49, which includes all twelve.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — all 12 guides Also free: the live AI company directory at wtfagents.com/companies