

The Agentic Economy

Starter Pack

Five guides that take you from "WTF is going on" to hiring your first agent. No jargon. No hype. Just what is actually happening — and what to do about it.

What's inside — 5 guides

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an LLM
4. WTF is Claude
5. How to Hire an AI Agent for Your Business

WTF Agents · wtfagents.com

Part of the WTF Agents Guide Series · wtfagents.com/store

WTF is the

Agentic Economy?

AI stopped answering questions and started doing jobs. Every major lab now sells an agent that works while you sleep, half a million AI-run companies have been launched, and the layoffs have started. The plain-English guide to what is actually happening.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, checking its work and repeating — around the clock, at a fraction of the cost of a human.

For two years that was a prediction. In 2026 it became a product category. Every major AI lab now sells one. This guide explains what changed, who is doing what, and where the hype ends.

By the numbers — September 2026

Everything in this box moves. Everything below it is written to still be true when it has.

\$10.9B

size of the AI agents market in 2026, up from \$7.6B in 2025 (Grand View Research)

40%

of enterprise software will contain task-specific agents by the end of 2026, up from under 5% in 2025 (Gartner)

23%

of organisations have actually scaled an agent system into production (McKinsey)

40%+

of agentic AI projects Gartner expects to be cancelled by the end of 2027

490,000

companies created by AI agents on Porsia alone since late 2025 — about 6% still active

3.2M

active users of OpenClaw, the open-source personal agent

139,156

tech layoffs in the first half of 2026, up 83% on the same period of 2025

80%

of large companies deploying autonomous tech cut jobs — with no correlation to improved returns (Gartner, May 2026)

What an AI agent actually is

Most people's experience of AI is reactive: you type, it answers. ChatGPT, Siri, Google. You ask, they respond, they stop.

An agent is different. You give it a goal. It works out the steps, does them one by one, checks the results, adjusts, and keeps going until the job is done — or until it hits something it cannot handle and asks you.

	A chatbot	An agent
Trigger	You ask a question	You give it a goal
Action	Answers once	Takes many steps on its own
Memory	Forgets when you close the tab	Remembers what it has done
Tools	Text only	Browses, writes code, sends email, fills forms, pays
Works while you sleep?	No	Yes

The simplest analogy: a calculator waits for you to press buttons. An agent is a capable employee who takes a brief and gets on with it.

The [WTF is an AI Agent](#) guide goes deep on how they work. This guide is about what they are doing to the economy.

2026 — the year every lab shipped the same product

Between January and September 2026, every major AI company launched a version of the same thing: an agent that does your job, not a chatbot that helps you do it.

When	Who	What
January	Anthropic	Cowork — Claude works directly in your files and apps. Enterprise software stocks fell hard when the plugins landed.
June	Microsoft	Copilot Cowork goes live worldwide inside Microsoft 365 — built on Anthropic's Claude, not OpenAI's models.
July	OpenAI	ChatGPT Work — give it an outcome, it returns finished spreadsheets, decks and web apps hours later.
August	xAI	Grok Bot — a team of named agents, each with its own cloud computer and logins, coordinated by a "Chief of Staff" bot.
September	Meta	Muse — a personal agent for everyone, in an app or on WhatsApp, that books, buys, negotiates and fills in forms for you.
May	Google	Killed its standalone browser agent and folded the technology into Gemini Agent and Chrome. The one lab that chose consolidation over a launch.

Add the coding agents — Claude Code, OpenAI's Codex, Cognition's Devin (now \$20 a month, down from \$500), Google's Jules — and the pattern is unmistakable. The product is not "AI." The product is delegated work.

Two things make this different from every previous wave of business software. First, the same agent that runs a Fortune 500 workflow is available to one person for the price of a streaming subscription. Second, the agents can now hire, build and run other agents — which is where the next layer comes from.

The three layers

It helps to think of the agentic economy as a stack.

The models

The brains. Anthropic's Claude, OpenAI's GPT, Google's Gemini, Meta's Muse models, xAI's Grok — and, increasingly, the Chinese open-weight models from DeepSeek, Alibaba's Qwen, Moonshot's Kimi and Zhipu's GLM, which are what a great many indie agents actually run on because they are cheap and can be downloaded. Model names change every few months. The layer does not.

The agents

Software that uses a model to do work. Two families. Personal agents run for you: OpenClaw (open-source, runs on your own machine, talks to you over WhatsApp or Telegram), Meta's Muse, Nous Research's Hermes. Work agents run for your job: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex, Devin.

The companies

The most ambitious layer, and the newest. Platforms that point agents at a business idea and let them run the whole thing — code, marketing, outreach, support, finance. Polsia is the biggest and most argued about; NanoCorp publishes the live revenue of every company on it; Cofounder.co puts a human approval step in front of anything risky. Paperclip is the open-source version: an org chart for agents you host yourself.

The stack matters because value is moving down it. Two years ago the money was in models. In 2026 it is in agents. The bet behind Polsia, NanoCorp and Paperclip is that next it will be in companies that no human works at.

The plumbing that makes it work

Two open standards are the invisible infrastructure of all this.

→ **MCP (Model Context Protocol)** — invented by Anthropic in late 2024, now governed by the neutral Agentic AI Foundation under the Linux Foundation and adopted by every major lab. It is the universal plug that lets any agent connect to any tool or data source. USB-C for AI.

→ **A2A (Agent2Agent)** — Google's standard for agents from different companies talking to each other. It is what lets a Paperclip org chart hire an OpenClaw agent, or a Grok Bot hand a task to a Codex session.

The newest piece is **agentic commerce**: agents that can pay. Muse buys on your behalf. Polsia companies take payments through their own Stripe accounts. Once agents hold wallets, the economy stops being a metaphor.

The messy reality

The agentic economy is real. It is also chaotic, frequently oversold, and moving faster than security, regulation or most businesses can absorb.

Most projects fail

Gartner expects more than 40% of agentic AI projects to be cancelled by the end of 2027 — unclear value, runaway cost, inadequate controls. Of the organisations experimenting, fewer than a quarter have anything in production at scale.

Most AI-run companies earn nothing

Polsia's founder said on the record in May 2026 that around one in ten companies on the platform had earned at least a dollar. NanoCorp's public feed showed cumulative revenue across every company on the platform in the hundreds of dollars. The agents build. Customers, mostly, do not come.

The agents misbehave

OpenClaw agents joined an AI-only social network, created dating profiles and leaked a million interactions through a misconfigured database in the first weeks of 2026. Cisco found third-party OpenClaw skills silently exfiltrating data. Meta's own staff flagged security failures in Muse days before launch. Every agent that holds your logins and acts in your accounts is a new attack surface, and prompt injection — hiding instructions in a web page or email the agent reads — is still unsolved.

Governments are in the room

China restricted OpenClaw in government offices in March and blocked Meta's \$2 billion acquisition of the Chinese agent startup Manus in April. The US government reviews frontier models before release; OpenAI's GPT-5.6 launch was delayed by it. Agents are now a matter of state.

The jobs question

This is the part people actually want to know, so here is the honest version.

The layoffs are real and they are being attributed to agents by name. ClickUp cut 22% of its staff in May and called the result a "100x org" built around thousands of internal agents, with million-dollar pay bands for the people who direct them. Salesforce's CEO said the company's support team went from 9,000 to 5,000. IBM replaced around 200 HR roles with agents. Tech layoffs in the first half of 2026 were up 83% on the year before.

The counter-evidence is also real. Gartner surveyed 350 large companies deploying autonomous tech in May 2026: 80% had cut headcount, and there was no correlation between the cuts and improved returns. "Workforce reductions may create budget room, but they do not create return." A lot of what is being called AI-driven is AI-justified. And IBM, the same company replacing HR staff with agents, is tripling its entry-level hiring for AI roles.

The pattern that holds across all of it: the roles disappearing are the ones that execute a defined process. The roles appearing are the ones that direct, review and take responsibility for agents. The people doing well are the ones who treat agents as junior staff — briefing them, checking them, expanding what they are trusted with — rather than as threats or as magic.

What this means for you

If you run a small business

You now have the operational capacity of a team for the cost of a few subscriptions. Support, content, research, outreach, bookkeeping — all delegable today. The question is not whether to start but

where, and the answer is one well-defined repetitive task, done well, before anything else. The How to Hire an AI Agent guide walks through it.

If you are an employee

Learn to direct agents before your employer decides whether you can. The single most valuable skill in the 2026 job market is turning a goal into a brief an agent can execute, and knowing when its output is wrong.

If you are trying to understand this space

You are in the right place. The rest of the WTF Agents series goes deep on each player, platform and concept in this guide. None of it requires a technical background. It requires curiosity and a willingness to try things.

Glossary

AI agent — Software that pursues a goal on its own: taking actions, using tools, checking results and adapting, without a human directing each step.

Agentic economy — An economy in which AI agents do meaningful economic work — and increasingly hire, build and pay each other — with minimal human supervision.

LLM (Large Language Model) — The AI "brain" inside every agent. Claude, GPT, Gemini, Grok, DeepSeek and the rest.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex.

Personal agent — An agent that runs your life's admin: OpenClaw, Meta Muse.

Autonomous company platform — A service where agents build and run an entire business: Polsia, NanoCorp, Cofounder.co; Paperclip is the self-hosted version.

MCP (Model Context Protocol) — The open standard connecting agents to tools and data. Invented by Anthropic, now under the Agentic AI Foundation.

A2A (Agent2Agent) — Google's open standard for agents from different companies talking to each other.

Agentic commerce — Agents that can pay, buy and take payment on your behalf.

Prompt injection — Hiding instructions in content an agent reads so it does something its owner did not ask for. The main unsolved security problem of the agentic economy.

Open weights — A model anyone can download and run themselves. Most Chinese frontier models ship this way; most US ones do not.

ARR (Annual Recurring Revenue) — The standard measure of a subscription business's size. In 2026 AI, often the last thirty days of all revenue multiplied by twelve — read it with care.

Liked this? Go deeper.

This guide gave you the big picture. The rest of the series goes deep on every platform, tool and concept mentioned here.

WTF is an AI Agent

Everyone is talking about agents. Nobody is explaining them properly. This guide does.

wtfagents.com/store

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners who want to start using AI agents right now.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is an

AI Agent?

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

The simplest possible explanation

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

A chatbot (like early ChatGPT):

You ask: "What are some good hotels in Barcelona?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

An AI assistant (like Siri or Alexa):

You say: "Book me a hotel in Barcelona." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

An AI agent:

You say: "Find me a good hotel in Barcelona for next weekend, under £200 a night, close to the Gothic Quarter, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

That is an AI agent.

The five properties of an AI agent

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

1. Goal-directed

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

Step	What happens
Receive goal	The agent is given an objective: "Write and schedule five social media posts for this week."
Plan	The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API.
Act	The agent executes step one: browses trending topics using its web tool.
Observe	It reads the results and decides what is relevant.
Reflect	It considers whether its approach is working. If step one produced poor results, it adjusts.
Repeat	It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

The main types of AI agent

Not all agents are the same. Here are the main categories you will encounter:

Personal AI agents

Run on your own machine or device. Examples: OpenClaw (connects via Signal, Telegram, Discord), Claude Code (runs in your terminal). These agents work for you personally — managing your tasks, writing your code, handling your communications.

Business process agents

Embedded into business workflows. Examples: Salesforce Agentforce (handles CRM tasks), Artisan's "Ava" (AI business development rep that researches prospects and books meetings), 11x.ai's digital workers (AI sales reps). These agents automate specific business functions.

Company-building agents

The most ambitious category. Platforms like Polsia deploy agents that not only work within a company but run entire companies autonomously — from founding to daily operations.

Orchestration agents

Agents that manage other agents. Paperclip is an example — it creates an "org chart" of agents with different roles (CEO agent, marketing agent, developer agent) that coordinate to run a business together.

Research agents

Agents designed to gather, synthesise, and analyse information. Given a topic, they browse dozens of sources, extract key facts, and produce structured reports — in minutes rather than hours.

Real agents doing real things right now

Klarna's customer service agent

Klarna, the Swedish fintech, deployed AI agents across its customer support operation. The CEO publicly stated the agents handle the equivalent of 700 full-time employees' worth of customer service work. This is one of the most cited enterprise agent deployments in the world.

OpenClaw — the viral personal agent

Created by Austrian developer Peter Steinberger and launched November 2025. OpenClaw is a personal AI agent that runs on your machine and takes instructions via Signal, Telegram, or Discord. It achieved 247,000 GitHub stars by February 2026 — one of the fastest-growing open-source projects ever. The creator was subsequently hired by OpenAI.

Artisan's Ava — the AI sales rep

"Ava" is an AI business development representative that researches potential customers, writes personalised outreach emails, follows up, and books meetings — all without human involvement. Artisan has raised \$46 million total and is backed by Y Combinator.

Claude Code — the autonomous developer

Anthropic's Claude Code is an agent that reads entire codebases, writes and edits code, runs tests, fixes bugs, and commits to Git — autonomously. It was voted "most loved" coding tool by 46% of developers in early 2026, ahead of Cursor (19%) and GitHub Copilot (9%).

Rentahuman.ai — agents hiring humans

Perhaps the most striking example: an entire platform where AI agents post jobs and hire humans to complete tasks the agents cannot do themselves (physical tasks, tasks requiring legal identity, nuanced human judgment). Forbes called it "a platform that flips the usual AI narrative."

The building blocks — what makes agents possible

Three things came together to make AI agents viable in 2024–2026:

1. LLMs got good enough

The underlying AI models — Claude, GPT, Gemini — became capable enough to reason through multi-step problems reliably. Earlier models would get confused, hallucinate, or go in circles. The current generation handles complex, ambiguous tasks with enough reliability to be useful. Claude Opus 4.6 scores 75.6% on SWE-bench (real software engineering tasks). Gemini 3.1 Pro scores 80.6%. These numbers were unthinkable two years ago.

2. Tool use became standardised

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar, or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in April 2025, letting agents from different companies communicate with each other. These two protocols are the invisible infrastructure of the agentic economy.

3. The cost of compute collapsed

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, the cost of running a capable agent has dropped dramatically — to the point where a single person can afford to run multiple agents simultaneously for the cost of a few software subscriptions.

The risks — what can go wrong

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

Prompt injection

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. Cisco found this vulnerability in OpenClaw in early 2026.

Overly broad permissions

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account, and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

Hallucination in action

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

Loss of control

The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent.

The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.

How to think about agents — the mental model

The most useful mental model for AI agents is not "software." It is "staff."

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

Glossary

AI Agent — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

LLM (Large Language Model) — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT-5.4 (OpenAI), Gemini 3.1 (Google).

ReAct — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

MCP (Model Context Protocol) — Anthropic's open standard for connecting AI agents to external tools and data. The "USB-C for AI." Published November 2024.

A2A (Agent2Agent Protocol) — Google's open standard for AI agents from different companies to communicate. Launched April 2025.

Tool use — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

Prompt injection — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

Orchestration — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

Autonomous — Operating without human control or supervision on a task-by-task basis.

OpenClaw — The most-starred open-source personal AI agent. Created by Peter Steinberger, November 2025. 247,000 GitHub stars by February 2026.

Claude Code — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

Liked this? Go deeper.

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

WTF is OpenClaw

The open source alternative taking on Palsia. Self-hosted, community-driven, and growing fast.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners who want to start using AI agents right now.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is an LLM?

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

How does predicting words create intelligence?

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.

Where did LLMs come from?

The story of LLMs is surprisingly recent and moves fast:

Year	What happened
2017	Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.
2018	Google releases BERT. OpenAI releases GPT-1. The race begins.
2020	OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.
2022	OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.
2023	GPT-4 released. Open-source models (LLaMA, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.
2024	Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.
2025–26	Claude Opus 4.6 (75.6% SWE-bench), Gemini 3.1 Pro (80.6%), GPT-5.4. Multi-hour autonomous tasks become reliable.

The main LLMs in 2026 — who makes what

There are three dominant LLM providers and a growing ecosystem of open-source alternatives:

Claude (Anthropic)

Anthropic's flagship model family. Current version: Claude 4.6, with Opus (most powerful) and Sonnet (faster, more affordable) variants. Claude Opus 4.6 scores 75.6% on SWE-bench — the standard benchmark for software engineering tasks. Claude is the default model for OpenClaw and widely used in agent frameworks. Anthropic's focus on safety and Constitutional AI training makes Claude the preferred choice for agentic deployments where reliability and honesty matter. Context window: up to 1 million tokens (beta). Anthropic valuation: \$380 billion.

GPT-5.4 (OpenAI)

OpenAI's latest flagship model as of March 2026. Native computer use capability. 1 million token context window in Codex. Powers ChatGPT, which remains the most widely used consumer AI product in the world. OpenAI also released Codex — an autonomous coding agent built on GPT-5.4. OpenAI valuation: \$300 billion+.

Gemini 3.1 Pro (Google DeepMind)

Google's flagship model. Currently holds the highest score on SWE-bench Verified (80.6%). Deeply integrated with Google's ecosystem — Search, Workspace, Android. Google also created the Agent2Agent (A2A) protocol for agent interoperability. Available via Google Cloud (Vertex AI).

Open-source models (LLaMA, Mistral, DeepSeek)

A growing ecosystem of models that anyone can run on their own hardware. Meta's LLaMA family is the most widely used. Mistral (French AI lab) produces highly efficient smaller models. DeepSeek (Chinese) produced models that competed with GPT-4 at a fraction of the training cost — a significant

moment in early 2025. OpenClaw supports DeepSeek as an alternative to Claude or GPT.

What LLMs are good at — and bad at

LLMs are good at	LLMs are bad at
Writing in any style or tone	Precise arithmetic (they approximate)
Summarising long documents	Real-time information (training data has a cutoff)
Writing and explaining code	Guaranteed factual accuracy (they can "hallucinate")
Translating between languages	Knowing what they don't know
Reasoning through multi-step problems	Consistent behaviour on identical inputs
Generating creative ideas	Tasks requiring persistent memory by default
Explaining complex topics simply	Physical world interaction (without agent tools)

Hallucination — the most important limitation

The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

Parameters, tokens, and context windows

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the

"size" of the model's brain.

Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude Opus 4.6 has a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.

Constitutional AI — Anthropic's approach to safety

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

LLMs and the agentic economy

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.

→ Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Claude 4.6 and Gemini 3.1 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

Glossary

LLM (Large Language Model) — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, and Gemini.

Transformer — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

Parameters — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

Token — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

Context window — How much text an LLM can process at once. Claude Opus 4.6: up to 1 million tokens.

Hallucination — When an LLM generates confident but factually incorrect text. A key limitation to understand.

Constitutional AI — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

SWE-bench — The standard benchmark for measuring LLM performance on real software engineering tasks.

Fine-tuning — Training an existing LLM further on a specific dataset to improve performance on a specific task.

Inference — Running an LLM to generate a response. The "thinking" part, as opposed to training.

Open-source model — An LLM whose weights are publicly available — anyone can run it. Examples: LLaMA, Mistral, DeepSeek.

Liked this? Go deeper.

Now you know what LLMs are. Here is the rest of the WTF Agents series.

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

WTF is an AI Agent

Everyone is talking about agents. Nobody is explaining them properly. This guide does.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners who want to start using AI agents right now.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Claude?

The AI model inside OpenClaw, Claude Code, Cowork, Microsoft's Copilot Cowork and the agents running thousands of AI companies. What it is, which version you actually want, and why so much of the agentic economy chose it.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Claude is Anthropic's family of AI models — large language models trained to be helpful, honest and harmless, with a particular focus on behaving reliably when running as an agent for hours without supervision.

It is the brain inside Claude Code and Cowork, the default model in OpenClaw, the engine under Microsoft's Copilot Cowork, and the model Polsia's founder picked to run the agents that run his companies. If you use anything in this series, there is a good chance Claude is doing the thinking.

By the numbers — September 2026

Model names and scores in this box change every few months. Everything below is written to outlast them.

Fable 5.1

the most capable Claude available to the public, released 1 September 2026

4

tiers on general sale: Fable 5.1, Opus 5, Sonnet 5, Haiku 4.5

1M tokens

context window on Fable — roughly 750,000 words, or several novels, in one session

128K tokens

maximum output in a single response

87.6%

SWE-bench Verified score reported for Claude Code Remote in spring 2026 — the highest published at the time

~150

organisations with access to the restricted Mythos version through Project Glasswing

\$20 / \$100–200

per month for Claude Pro and Max. Free tier available

The family — which Claude is which

Claude is not one model. It is a range at different price and capability points, and in 2026 a new tier appeared at the top.

Tier	What it is for
Fable	The Mythos-class flagship, a tier above Opus. Long autonomous tasks, hard reasoning, large codebases, document-heavy professional work. The most expensive.
Opus	The previous top tier, still very capable. Complex agentic work where Fable is more than you need.
Sonnet	The workhorse. Fast, strong, far cheaper. Most business agents run on Sonnet.
Haiku	Fastest and cheapest. High-volume, simpler tasks: classification, routing, quick replies.
Mythos	Not for sale. The same model as Fable with safety classifiers lifted, restricted to vetted cyberdefenders and researchers.

The version numbers move — 4.6 in March, 4.7 in April, 4.8 in May, 5 in July, Fable 5.1 in September. The tiers do not. When someone says "Claude," ask which tier.

Fable and Mythos — the split that changed how models ship

In June 2026 Anthropic released a model it said was capable enough to be dangerous in the wrong hands: strong enough to find and chain security vulnerabilities in software autonomously, a task previously limited to elite human researchers.

Rather than release it or withhold it, Anthropic did both. Claude Fable is that model with real-time classifiers that block cybersecurity, biology and chemistry misuse; it is generally available. Claude Mythos is the identical model with those safeguards lifted, available only through Project Glasswing, a vetted programme run with the US government for cyberdefenders and critical-infrastructure providers, and, over time, biomedical researchers.

For you as a user the practical meaning is simple. Fable 5.1 is the best Claude you can buy, and it will decline a narrow set of dangerous requests. For the industry the meaning is larger: the most capable AI is now gated by who you are and what you are for. Other labs followed within months.

What Claude can actually do

Write and edit — Any style, length or format. Emails, code, contracts, marketing copy, technical documentation. Claude adapts to context and holds a consistent voice across long pieces.

Reason through hard problems — Multi-step logic, analysis, strategy, maths. Fable and Opus think before answering, and the thinking scales with the difficulty of the task.

Read enormous documents — A million-token context means an entire codebase, a full legal file or a year of emails in one session.

Write, run and fix code — Across every major language. Claude Code turns this into an agent that reads your project, changes it, tests it and commits.

Use tools — Browse the web, read files, query databases, call APIs, send email, through MCP, the open standard Anthropic invented.

Work in your apps — Cowork on the desktop, Claude in Chrome and a built-in browser, Claude in Excel, PowerPoint, Word and Outlook, Claude Tag in Slack.

Run for hours — The property that matters most for agents. Claude keeps its footing across long autonomous sessions, remembers what it has done, and asks when it is unsure rather than guessing.

What makes Claude different — the behaviour, not the benchmark

Benchmarks are a snapshot and every lab tops one somewhere. The reason so much of the agentic economy standardised on Claude is behavioural, and it comes from how Anthropic trains.

Constitutional AI gives the model written principles and trains it to judge its own outputs against them. In practice that produces a model that:

- Declines harmful requests and explains why, instead of refusing bluntly or complying blindly
- Flags ambiguity and asks, rather than silently assuming and getting it wrong three steps later
- Says "I don't know" more readily than rivals, which matters when nobody is checking
- Behaves consistently across different phrasings of the same instruction
- Notices flaws in its own work — Anthropic reported Opus 4.8 was several times less likely than its predecessor to let a bug it wrote pass unremarked

For a chatbot these are nice qualities. For an agent running your inbox at 3am they are the whole product. A model that hallucinates confidently sends the wrong email. A model that flags uncertainty leaves you a note.

Where Claude shows up

OpenClaw — Default model in the most-used open-source personal agent, with millions of active users.

Microsoft Copilot Cowork — Microsoft's flagship agent for Microsoft 365 runs on Claude, despite Microsoft's \$13 billion stake in OpenAI. Microsoft chose it anyway.

Polsia — The founder has said on the record he chose Claude Opus to run his agents because "it's important to give it the best reasoning."

Paperclip — Claude Code is the primary developer agent in the open-source multi-agent org chart.

Every major cloud — Amazon Bedrock, Google Cloud Vertex AI and Microsoft Foundry all serve Claude, so enterprises can use it inside infrastructure they already pay for.

Thousands of products — Via the API: customer-service agents, research tools, legal analysis, coding assistants, and the AI-run companies indexed on wtfagents.com.

How to get it

Claude.ai and the apps — Web, desktop (Mac and Windows), iPhone and Android. Free tier for chat. Claude Pro at \$20 a month adds Cowork, Claude Code and higher limits. Claude Max at \$100–200 for heavy use. Team and Enterprise plans above that.

Claude Code — Included with Pro and above, or pay-as-you-go through the API.

The Claude Developer Platform — Direct API access for builders. Pay per token; Haiku cheapest, Fable most expensive.

Through your cloud — Bedrock, Vertex AI or Foundry if your company already lives there.

Inside other agents — OpenClaw, Paperclip and most agent frameworks let you plug in your own Anthropic API key.

Glossary

Claude — Anthropic's AI model family: Fable, Opus, Sonnet and Haiku on general sale; Mythos restricted.

Fable — The public version of the Mythos-class model, with safety classifiers active. Currently Fable 5.1.

Mythos — The same model with safeguards lifted, available only to vetted organisations through Project Glasswing.

Tier — A capability and price level within the family. Tiers are stable; version numbers change constantly.

Constitutional AI — Anthropic's training technique: written principles the model learns to judge its own outputs against.

Context window — How much text the model can hold at once. One million tokens on Fable.

Token — The unit models process text in. Roughly three quarters of a word. API usage is billed in tokens.

MCP (Model Context Protocol) — The open standard Claude and every other major model use to connect to tools.

SWE-bench — The standard benchmark for AI coding ability, measured on real software engineering tasks.

Claude Code — The agentic version of Claude for building software. See its own guide.

Cowork — Claude working in your files, apps and browser without a command line. See its own guide.

Liked this? Go deeper.

Claude is the model. Here is everything built on it.

WTF is Claude Code

How Claude Code is changing software development — and what it means if you're not a developer.

wtfagents.com/store

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What the Anthropic product actually does.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners who want to start using AI agents right now.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

How to Hire an AI Agent

for Your Business

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

Before you start — the right mindset

Deploying AI agents in your business is not a technology project. It is a staffing decision.

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

Step 1 — Find your highest-value repeatable task

The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

Customer support responses — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

Content creation — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

Research and summarisation — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

Data entry and processing — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

Email drafting — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

Scheduling and coordination — Managing calendar bookings, sending reminders, coordinating meeting times.

Lead qualification — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

Step 2 — Choose the right type of agent

Not all agents are right for all tasks. Here is a simple decision framework:

If you want to...	Use this	Cost
Launch an AI-run company with no technical setup	Polsia (polsia.com)	\$50/mo + 20% revenue share
Deploy a personal AI assistant that takes actions via messaging apps	OpenClaw (openclaw.ai)	Free + LLM API costs
Build a multi-agent team to run a business	Paperclip (paperclip.ing)	Free (self-hosted)
Automate specific business workflows without coding	Relevance AI (relevanceai.com)	Free tier available
Automate repetitive tasks and connect apps	Zapier AI (zapier.com)	From \$19.99/mo
Build a custom agent with full control	Anthropic API + Claude	Pay per token

Step 3 — Write a clear job description for your agent

This is the most important step most people skip.

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

- **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")
- **The inputs** — what information does the agent receive? ("An inbound email from a customer")
- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply

using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

Step 4 — Start with read-only access

The most common mistake when deploying agents: giving them too much access too soon.

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

Step 5 — Measure results honestly

Define what success looks like before you deploy — not after.

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved
- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

Real examples — what businesses are doing right now

E-commerce store (5 employees)

Deployed an agent to handle first-line customer support emails. The agent categorises enquiries, drafts replies using approved templates, and flags unusual requests for human review. Result: response time dropped from 6 hours to 20 minutes. 80% of emails handled without human

involvement.

Marketing agency (12 employees)

Deployed an agent to research clients' industries and draft weekly content calendars. Copywriters now spend time on strategy and editing, not research. Result: capacity increased 40% without additional headcount.

Accounting firm (3 partners)

Deployed an agent to summarise client emails and draft responses for partner review. Result: 2 hours per day saved per partner. Used the time to take on 3 additional clients.

Solo consultant

Deployed OpenClaw as a personal assistant — briefing it via Telegram with research tasks, email drafts, and document summaries. Cost: approximately \$30/month in LLM API costs. Time saved: 10+ hours per week.

SaaS startup (8 people)

Deployed Claude Code to handle routine bug fixes and feature requests. Senior engineers now focus on architecture and complex problems. Result: shipped 3x more features in Q1 2026 with the same team.

The honest cost — what agents actually cost

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

Agent type	Monthly cost (approx)	What drives the cost
Polsia company	\$50/mo + 20% revenue	Flat fee plus revenue share
OpenClaw (personal)	\$20-100/mo	Anthropic or OpenAI API usage
Relevance AI	\$0-299/mo	Free tier to enterprise plans
Zapier AI	\$20-100/mo	Task volume and plan tier
Custom Claude agent	\$50-500/mo	Token usage — depends heavily on task volume
Claude Code	\$100-500/mo	API token usage — intensive for large codebases

For context: a part-time human assistant in the UK costs £1,500-2,500/month. A full-time junior employee costs £25,000-35,000/year. An agent handling equivalent work costs £50-500/month. The economics are not subtle.

What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

Your action plan — what to do this week

Day 1 — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

Day 2 — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

Day 3 — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

Day 4-5 — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

Week 2 — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

Week 3-4 — If quality is consistently good, expand permissions incrementally. Track the time saved.

Month 2 — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

The businesses that win with AI agents in 2026 are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.

Glossary

AI Agent — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

Human-in-the-loop — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

Polsia — A managed platform where AI agents build and run companies. \$50/mo + 20% revenue share.

OpenClaw — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

Paperclip — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

Relevance AI — No-code/low-code platform for building AI agents for specific business workflows. \$37M raised.

Zapier AI — Major automation platform with AI agent capabilities. Connects thousands of apps.

Claude — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

API key — Your credential for accessing an AI model's API. Required for OpenClaw, Claude Code, and custom builds.

Token — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

MCP (Model Context Protocol) — The open standard connecting AI agents to external tools. Invented by Anthropic.

Liked this? Go deeper.

You now know how to hire an agent. Here is the full context behind the tools you are using.

WTF is an AI Agent

Everyone is talking about agents. Nobody is explaining them properly. This guide does.

wtfagents.com/store

WTF is OpenClaw

The open source alternative taking on Polisia. Self-hosted, community-driven, and growing fast.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies