

The Complete

WTF Agents Pack

Every WTF Agents guide in one file. The agentic economy, the platforms running it, the AI behind it, and how to put it to work in your own business.

What's inside — 33 guides

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an API
4. WTF is an LLM
5. WTF is MCP
6. WTF is Polsia
7. WTF is OpenClaw
8. WTF is Paperclip
9. WTF is Anthropic
10. WTF is Claude
11. WTF is Claude Code
12. WTF is Cowork
13. How to Hire an AI Agent for Your Business
14. Who's Who in the Agentic Economy
15. Which AI Should I Actually Use?
16. WTF is Meta Muse
17. WTF is Grok Bot
18. WTF is ChatGPT Work
19. WTF is Happening with China's AI
20. WTF is NanoCorp
21. WTF is Codex
22. The Picks and Shovels

23. Launch an AI-Run Company in a Weekend
24. The Solo Founder's AI Stack
25. WTF is Higgsfield
26. WTF is ElevenLabs
27. WTF is HeyGen
28. WTF is Runway
29. WTF is Cursor
30. WTF is Devin
31. WTF is n8n
32. WTF is Hermes Agent
33. WTF is Manus

WTF Agents · wtfagents.com

Part of the WTF Agents Guide Series · wtfagents.com

WTF is the Agentic Economy?

AI stopped answering questions and started doing jobs. Every major lab now sells an agent that works while you sleep, half a million AI-run companies have been launched, and the layoffs have started. The plain-English guide to what is actually happening.

The one-liner

The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, checking its work and repeating — around the clock, at a fraction of the cost of a human.

For two years that was a prediction. In 2026 it became a product category. Every major AI lab now sells one. This guide explains what changed, who is doing what, and where the hype ends.

By the numbers — September 2026

Everything in this box moves. Everything below it is written to still be true when it has.

\$10.9B

size of the AI agents market in 2026, up from \$7.6B in 2025 (Grand View Research)

40%

of enterprise software will contain task-specific agents by the end of 2026, up from under 5% in 2025 (Gartner)

23%

of organisations have actually scaled an agent system into production (McKinsey)

40%+

of agentic AI projects Gartner expects to be cancelled by the end of 2027

490,000

companies created by AI agents on Palsia alone since late 2025 — about 6% still active

3.2M

active users of OpenClaw, the open-source personal agent

139,156

tech layoffs in the first half of 2026, up 83% on the same period of 2025

80%

of large companies deploying autonomous tech cut jobs — with no correlation to improved returns
(Gartner, May 2026)

What an AI agent actually is

Most people's experience of AI is reactive: you type, it answers. ChatGPT, Siri, Google. You ask, they respond, they stop.

An agent is different. You give it a goal. It works out the steps, does them one by one, checks the results, adjusts, and keeps going until the job is done — or until it hits something it cannot handle and asks you.

	A chatbot	An agent
Trigger	You ask a question	You give it a goal
Action	Answers once	Takes many steps on its own
Memory	Forgets when you close the tab	Remembers what it has done
Tools	Text only	Browses, writes code, sends email, fills forms, pays
Works while you sleep?	No	Yes

The simplest analogy: a calculator waits for you to press buttons. An agent is a capable employee who takes a brief and gets on with it.

The [WTF is an AI Agent](#) guide goes deep on how they work. This guide is about what they are doing to the economy.

2026 — the year every lab shipped the same product

Between January and September 2026, every major AI company launched a version of the same thing: an agent that does your job, not a chatbot that helps you do it.

When	Who	What
January	Anthropic	Cowork — Claude works directly in your files and apps. Enterprise software stocks fell hard when the plugins landed.
June	Microsoft	Copilot Cowork goes live worldwide inside Microsoft 365 — built on Anthropic's Claude, not OpenAI's models.
July	OpenAI	ChatGPT Work — give it an outcome, it returns finished spreadsheets, decks and web apps hours later.
August	xAI	Grok Bot — a team of named agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff" bot.
September	Meta	Muse — a personal agent for everyone, in an app or on WhatsApp, that books, buys, negotiates and fills in forms for you.
May	Google	Killed its standalone browser agent and folded the technology into Gemini Agent and Chrome. The one lab that chose consolidation over a launch.

Add the coding agents — Claude Code, OpenAI's Codex, Cognition's Devin (now \$20 a month, down from \$500), Google's Jules — and the pattern is unmistakable. The product is not "AI." The product is delegated work.

Two things make this different from every previous wave of business software. First, the same agent that runs a Fortune 500 workflow is available to one person for the price of a streaming subscription. Second, the agents can now hire, build and run other agents — which is where the next layer comes from.

The three layers

It helps to think of the agentic economy as a stack.

The models

The brains. Anthropic's Claude, OpenAI's GPT, Google's Gemini, Meta's Muse models, xAI's Grok — and, increasingly, the Chinese open-weight models from DeepSeek, Alibaba's Qwen, Moonshot's Kimi and Zhipu's GLM, which are what a great many indie agents actually run on because they are cheap and can be downloaded. Model names change every few months. The layer does not.

The agents

Software that uses a model to do work. Two families. Personal agents run for you: OpenClaw (open-source, runs on your own machine, talks to you over WhatsApp or Telegram), Meta's

Muse, Nous Research's Hermes. Work agents run for your job: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex, Devin.

The companies

The most ambitious layer, and the newest. Platforms that point agents at a business idea and let them run the whole thing — code, marketing, outreach, support, finance. Polsia is the biggest and most argued about; NanoCorp publishes the live revenue of every company on it; Cofounder.co puts a human approval step in front of anything risky. Paperclip is the open-source version: an org chart for agents you host yourself.

The stack matters because value is moving down it. Two years ago the money was in models. In 2026 it is in agents. The bet behind Polsia, NanoCorp and Paperclip is that next it will be in companies that no human works at.

The plumbing that makes it work

Two open standards are the invisible infrastructure of all this.

→ **MCP (Model Context Protocol)** — invented by Anthropic in late 2024, now governed by the neutral Agentic AI Foundation under the Linux Foundation and adopted by every major lab. It is the universal plug that lets any agent connect to any tool or data source. USB-C for AI.

→ **A2A (Agent2Agent)** — Google's standard for agents from different companies talking to each other. It is what lets a Paperclip org chart hire an OpenClaw agent, or a Grok Bot hand a task to a Codex session.

The newest piece is **agentic commerce**: agents that can pay. Muse buys on your behalf. Polsia companies take payments through their own Stripe accounts. Once agents hold wallets, the economy stops being a metaphor.

The messy reality

The agentic economy is real. It is also chaotic, frequently oversold, and moving faster than security, regulation or most businesses can absorb.

Most projects fail

Gartner expects more than 40% of agentic AI projects to be cancelled by the end of 2027 — unclear value, runaway cost, inadequate controls. Of the organisations experimenting, fewer than a quarter have anything in production at scale.

Most AI-run companies earn nothing

Polsia's founder said on the record in May 2026 that around one in ten companies on the platform had earned at least a dollar. NanoCorp's public leaderboard showed its best company on \$2,570 across 748 sales, with only two companies past \$1,000. The agents build. Customers, mostly, do not come.

The agents misbehave

OpenClaw agents joined an AI-only social network, created dating profiles and leaked a million interactions through a misconfigured database in the first weeks of 2026. Cisco found third-party OpenClaw skills silently exfiltrating data. Meta's own staff flagged security failures in Muse days before launch. Every agent that holds your logins and acts in your accounts is a new attack surface, and prompt injection — hiding instructions in a web page or email the agent reads — is still unsolved.

Governments are in the room

China restricted OpenClaw in government offices in March and blocked Meta's \$2 billion acquisition of the Chinese agent startup Manus in April. The US government reviews frontier models before release; OpenAI's GPT-5.6 launch was delayed by it. Agents are now a matter of state.

The jobs question

This is the part people actually want to know, so here is the honest version.

The layoffs are real and they are being attributed to agents by name. ClickUp cut 22% of its staff in May and called the result a "100x org" built around thousands of internal agents, with million-dollar pay bands for the people who direct them. Salesforce's CEO said the company's support team went from 9,000 to 5,000. IBM replaced around 200 HR roles with agents. Tech layoffs in the first half of 2026 were up 83% on the year before.

The counter-evidence is also real. Gartner surveyed 350 large companies deploying autonomous tech in May 2026: 80% had cut headcount, and there was no correlation between the cuts and improved returns. "Workforce reductions may create budget room, but they do not create return." A lot of what is being called AI-driven is AI-justified. And IBM, the same company replacing HR staff with agents, is tripling its entry-level hiring for AI roles.

The pattern that holds across all of it: the roles disappearing are the ones that execute a defined process. The roles appearing are the ones that direct, review and take responsibility for agents. The people doing well are the ones who treat agents as junior staff — briefing them, checking them, expanding what they are trusted with — rather than as threats or as magic.

What this means for you

If you run a small business

You now have the operational capacity of a team for the cost of a few subscriptions. Support, content, research, outreach, bookkeeping — all delegable today. The question is not whether to start but where, and the answer is one well-defined repetitive task, done well, before anything else. The [How to Hire an AI Agent](#) guide walks through it.

If you are an employee

Learn to direct agents before your employer decides whether you can. The single most valuable skill in the 2026 job market is turning a goal into a brief an agent can execute, and knowing when its output is wrong.

If you are trying to understand this space

You are in the right place. The rest of the WTF Agents series goes deep on each player, platform and concept in this guide. None of it requires a technical background. It requires curiosity and a willingness to try things.

Glossary

AI agent — Software that pursues a goal on its own: taking actions, using tools, checking results and adapting, without a human directing each step.

Agentic economy — An economy in which AI agents do meaningful economic work — and increasingly hire, build and pay each other — with minimal human supervision.

LLM (Large Language Model) — The AI "brain" inside every agent. Claude, GPT, Gemini, Grok, DeepSeek and the rest.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex.

Personal agent — An agent that runs your life's admin: OpenClaw, Meta Muse.

Autonomous company platform — A service where agents build and run an entire business: Polsia, NanoCorp, Cofounder.co; Paperclip is the self-hosted version.

MCP (Model Context Protocol) — The open standard connecting agents to tools and data. Invented by Anthropic, now under the Agentic AI Foundation.

A2A (Agent2Agent) — Google's open standard for agents from different companies talking to each other.

Agentic commerce — Agents that can pay, buy and take payment on your behalf.

Prompt injection — Hiding instructions in content an agent reads so it does something its owner did not ask for. The main unsolved security problem of the agentic economy.

Open weights — A model anyone can download and run themselves. Most Chinese frontier models ship this way; most US ones do not.

ARR (Annual Recurring Revenue) — The standard measure of a subscription business's size. In 2026 AI, often the last thirty days of all revenue multiplied by twelve — read it with care.

This guide gave you the big picture. The rest of the series goes deep on every platform, tool and concept mentioned here.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/guides/polsia

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/guides/api

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

WTF is an AI Agent?

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

The simplest possible explanation

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

A chatbot (like early ChatGPT):

You ask: "What are some good hotels in Bangkok?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

An AI assistant (like Siri or Alexa):

You say: "Book me a hotel in Bangkok." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

An AI agent:

You say: "Find me a good hotel in Bangkok for next weekend, under \$200 a night, near the river, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

That is an AI agent.

The five properties of an AI agent

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

1. Goal-directed

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

Step	What happens
Receive goal	The agent is given an objective: "Write and schedule five social media posts for this week."
Plan	The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API.
Act	The agent executes step one: browses trending topics using its web tool.
Observe	It reads the results and decides what is relevant.
Reflect	It considers whether its approach is working. If step one produced poor results, it adjusts.
Repeat	It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

The main types of AI agent

Not all agents are the same. Here are the main categories you will encounter:

Personal agents

Run for you, on your machine or in an app. Examples: OpenClaw (open-source, talks to you over WhatsApp, Telegram or Signal), Meta's Muse (books, buys and negotiates on your behalf). These agents handle your life's admin.

Work agents

Do your job's tasks. Every major lab now sells one: Anthropic's Cowork, OpenAI's ChatGPT Work, xAI's Grok Bot, Microsoft's Copilot Cowork. Give them an outcome — a report, a spreadsheet, a cleared inbox — and they come back with it done.

Coding agents

The category that proved agents work. Claude Code, OpenAI's Codex, Cognition's Devin, Google's Jules. They read a codebase, change it, test it and commit, without step-by-step instruction.

Business process agents

Embedded into a specific business function. Salesforce's Agentforce handles customer service and sales tasks inside its CRM; a wave of "AI sales rep" and "AI support rep" products do the same as standalone hires.

Company-building agents

The most ambitious category. Platforms like Polsia and NanoCorp deploy a whole team of agents that found and run a company — website, marketing, outreach, support, finance — with no human doing the work.

Orchestration agents

Agents that manage other agents. Paperclip is the clearest example: an org chart of agents with roles, budgets and reporting lines, coordinating to run a business together.

Research agents

Given a topic, they browse dozens of sources, extract facts, cross-check and produce a structured report — in minutes rather than hours. Most chat products now have one built in.

Real agents doing real things right now

OpenClaw — the personal agent that became a movement

Created by Austrian developer Peter Steinberger and launched in late 2025, OpenClaw runs on your own machine and takes instructions over messaging apps. By September 2026 it had over three million active users, its own non-profit foundation with a board, native mobile apps, and a place among the most-starred projects in GitHub's history. It also produced the first famous agent misbehaviour stories, covered below.

Cowork, ChatGPT Work and Grok Bot — the year of the work agent

Between January and September 2026, Anthropic, OpenAI and xAI each launched an agent that does whole jobs rather than answering questions about them. Cowork works inside your files and apps. ChatGPT Work takes an outcome and returns finished spreadsheets, decks and apps hours later. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff." Microsoft's Copilot Cowork, built on Claude, brought the same thing to Microsoft 365.

Claude Code — the autonomous developer

Anthropic's Claude Code reads entire codebases, writes and edits code, runs tests, fixes bugs and commits — autonomously. It became the fastest-adopted developer tool of its generation and the main reason Anthropic's revenue multiplied several times over in 2026. Its remote version topped the standard software-engineering benchmark in the spring.

Polsia — agents running whole companies

Polsia's agents have created close to half a million companies since late 2025: a website, product, marketing, outreach and support each, with no human doing the work. About 6% survive and the founder says one in ten has earned a dollar. It is the clearest demonstration in the world of both what agents can build and what they cannot: customers.

Agents in enterprise customer service

Salesforce's CEO has said its support headcount fell from roughly 9,000 to 5,000 as agents took over the routine work. IBM replaced around 200 HR roles with agents. These are the deployments that get quoted in board meetings, and they are why the jobs question in the Agentic Economy guide matters.

The building blocks — what makes agents possible

Three things came together to make AI agents viable in 2024–2026:

1. LLMs got good enough

The underlying AI models — Claude, GPT, Gemini, Grok, Muse, DeepSeek — became capable enough to reason through multi-step problems reliably and, crucially, to notice when they were going wrong. Earlier models got confused, hallucinated, or went in circles. The current generation handles complex, ambiguous tasks with enough reliability to be left alone for

hours. The exact scores change every quarter; the shift from "answers questions" to "finishes jobs" does not.

2. Tool use became standardised

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in 2025, letting agents from different companies communicate. MCP now lives under the neutral Agentic AI Foundation. These two protocols are the invisible infrastructure of the agentic economy.

3. The cost of compute collapsed

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, a capable agent costs a few dollars a day to run, and open-weight models from China and elsewhere push it lower still — to the point where a single person can run several agents simultaneously for the cost of a few software subscriptions.

The risks — what can go wrong

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

Prompt injection

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. It remains the main unsolved security problem for agents. Cisco found third-party OpenClaw skills silently exfiltrating data in early 2026; Meta's own staff flagged security failures in Muse days before its launch.

Overly broad permissions

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

Hallucination in action

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

Loss of control

The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent. Palsia users have reported agents sending cold emails to real contacts and making public posts without approval.

The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.

How to think about agents — the mental model

The most useful mental model for AI agents is not "software." It is "staff."

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

Glossary

AI Agent — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

LLM (Large Language Model) — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT (OpenAI), Gemini (Google), Muse (Meta), Grok (xAI), DeepSeek.

ReAct — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

MCP (Model Context Protocol) — The open standard for connecting AI agents to external tools and data. The "USB-C for AI." Invented by Anthropic, now under the Agentic AI Foundation.

A2A (Agent2Agent Protocol) — Google's open standard for AI agents from different companies to communicate.

Tool use — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

Prompt injection — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

Orchestration — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

Autonomous — Operating without human control or supervision on a task-by-task basis.

OpenClaw — The most-used open-source personal AI agent. Created by Peter Steinberger, late 2025; now run by a non-profit foundation.

Claude Code — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/guides/api

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

WTF is an API?

APIs are mentioned in almost every conversation about AI, tech, and the agentic economy. Almost no one explains what they actually are. This guide does — in plain English, finally.

WTF Agents · wtfagents.com · March 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

An API (Application Programming Interface) is a set of rules that lets two pieces of software talk to each other.

That is it. That is the whole thing. Everything else in this guide is just unpacking what that means in practice.

The restaurant analogy

The best analogy for an API is a restaurant.

You are sitting at a table. You want food. The kitchen can make food. But you do not walk into the kitchen and start cooking. You do not shout instructions at the chef. You use a menu and a waiter.

The menu tells you what the kitchen can make. The waiter takes your order to the kitchen and brings the food back.

In this analogy: you are the application making a request. The kitchen is the other software (a database, a service, an AI model). The waiter is the API. The menu is the API documentation.

The API is the agreed-upon way that two systems communicate — what you can ask for, how you ask for it, and what you get back.

A real-world example

Here is what actually happens when you use a weather app on your phone:

- You open your weather app.
- The app does not have weather data stored on your phone.
- It sends a request to a weather service's API: "Give me the forecast for Chicago."
- The weather service's servers process the request.
- The API sends back the data: temperature, humidity, forecast for the next 7 days.
- Your app displays it in a nice interface.

You never saw any of that. You just saw the weather. But an API made the whole thing work.

This same pattern — request, process, respond — underpins almost everything on the internet. Every time you log in with Google, every time you pay with Stripe, every time you share to Twitter, every time an AI agent takes an action in the world — an API is involved.

Why APIs matter for the agentic economy

APIs are the reason AI agents can do anything useful.

An AI agent on its own is just a brain in a jar — it can think and plan but it cannot act. APIs are the hands and legs. They are how agents connect to the real world.

Here are examples of what agents do via APIs:

Action	How the agent does it
Send an email	The agent calls Gmail's API or SendGrid's API with the message content, recipient, and subject.
Browse the web	The agent calls a web scraping API or uses a browser automation tool exposed via API.
Process a payment	The agent calls Stripe's API with the amount and card details.
Post to social media	The agent calls the Twitter/X API or LinkedIn API with the post content.
Read a spreadsheet	The agent calls the Google Sheets API to retrieve or update data.
Book a calendar slot	The agent calls the Google Calendar or Calendly API.
Query a database	The agent calls the Supabase or Postgres API to read or write records.
Use an AI model	The agent calls the Anthropic API or OpenAI API to get Claude or GPT to think through a problem.

Without APIs, an AI agent would be like a genius who has been locked in a room with no phone, no computer, and no way to interact with the outside world. APIs open the door.

MCP — the API for AI agents

One of the most important developments in the agentic economy is the Model Context Protocol (MCP) — invented by Anthropic and published in November 2024.

Before MCP, connecting an AI agent to a specific tool or data source required custom engineering. Every integration was different. It was like every electrical appliance having a different plug.

MCP standardised it. It is the universal plug — the USB-C for AI. If a tool supports MCP, any AI agent can connect to it without custom code.

As of 2026, MCP is supported by OpenAI, Google, Microsoft, and virtually all major agent frameworks. It has been donated to the Agentic AI Foundation (backed by Anthropic, OpenAI, Google, Microsoft, Amazon, and Cloudflare) — making it a permanent, neutral open standard.

This matters because it means the ecosystem of tools available to AI agents is growing exponentially. Every new MCP-compatible tool immediately becomes available to every MCP-compatible agent.

The types of API you will hear about

REST APIs

The most common type. Uses standard web requests (GET, POST, PUT, DELETE) to retrieve or send data. When someone says "the API" without specifying, they almost always mean a REST API. The Polsia API that WTF Agents uses to track company data is a REST API.

Webhooks

Instead of you asking the API for data, the API calls you when something happens. Example: Stripe fires a webhook to your server the moment a payment completes. You do not have to keep asking "has the payment gone through?" — Stripe tells you when it does.

GraphQL APIs

A more flexible alternative to REST. Instead of getting a fixed set of data, you specify exactly what fields you want. More efficient but more complex to use.

Streaming APIs

Send data continuously rather than in one response. This is why ChatGPT and Claude type out their responses word by word — the AI is streaming its output via a streaming API rather than waiting until it has finished the whole response.

SDK (Software Development Kit)

Not technically an API itself, but often confused with one. An SDK is a package of code that wraps an API and makes it easier to use. The Anthropic SDK is a Python or JavaScript package that handles the technical details of calling the Claude API.

API keys — what they are and why they matter

Almost every API requires an API key — a unique string of characters that identifies who is making the request.

Think of an API key as a password for a service. When you sign up for the Anthropic API, OpenAI API, or Stripe, they give you a key that looks something like this:

sk-ant-api03-xxxxxxxxxxxxx

API keys matter for three reasons:

- **Authentication** — the service knows it is you making the request, not someone else.
- **Billing** — usage is tracked against your key. Lose your key and someone else runs up your bill.
- **Rate limiting** — services use keys to enforce limits on how many requests you can make.

The golden rule: never share your API keys publicly. Never commit them to GitHub. Store them in environment variables (a secure way of passing sensitive values to software without hardcoding them into your code).

APIs in everyday life — you use them constantly

You interact with APIs dozens of times a day without knowing it:

Logging in with Google — Google's OAuth API verifies your identity and tells the other site you are who you say you are.

Paying online — Stripe's payment API processes your card securely without the website ever seeing your full card number.

Checking the weather — Your phone's weather app calls a meteorological data API.

Getting directions — Google Maps or Apple Maps API calculates your route.

Posting to Instagram — The Instagram API receives your photo and caption and stores it.

Reading your emails in a third-party app — Gmail's API gives the app access to your emails.

Using "Sign in with Apple" — Apple's authentication API verifies you.

An AI agent booking your travel — The agent calls APIs for flights, hotels, and calendar — all in one automated flow.

Do you need to understand APIs to use AI agents?

Honestly — not really, for most use cases.

Platforms like OpenClaw, Polsia, and Paperclip handle the API connections for you. You describe what you want the agent to do, and the platform handles the technical plumbing of which APIs to call and how.

But understanding what APIs are — and that they exist — helps you in three ways:

- You understand why agents can do what they can do (and why they sometimes cannot).
- You can evaluate agent platforms more intelligently — asking "what APIs does this connect to?"
- If you want to build something custom, you know what to look for.

The agentic economy runs on APIs. You do not need to be a plumber to live in a house with running water. But knowing that pipes exist, and roughly how they work, makes you a more informed resident.

Glossary

API (Application Programming Interface) — A set of rules that lets two pieces of software communicate with each other.

REST API — The most common type of API. Uses standard web requests to send and receive data.

Webhook — An API that pushes data to you when something happens, rather than waiting for you to ask.

API key — A unique identifier used to authenticate API requests. Treat it like a password.

MCP (Model Context Protocol) — Anthropic's open standard for connecting AI agents to external tools. The universal plug for AI.

SDK (Software Development Kit) — A package of code that makes it easier to use an API in a specific programming language.

Endpoint — A specific URL that an API exposes for a specific type of request. E.g. /api/data is the endpoint for Polsia's company data.

Rate limiting — A restriction on how many API requests you can make in a given time period.

Authentication — The process of proving who you are to an API, usually via an API key or OAuth token.

Streaming — An API that sends data continuously rather than all at once. How Claude and ChatGPT deliver responses word by word.

Now you know what APIs are — here is where to go next in the WTF Agents series.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is MCP

The Model Context Protocol, explained. The universal plug between agents and tools: what it is, who runs it now, and the security catch.

wtfagents.com/guides/mcp

WTF is an LLM?

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

How does predicting words create intelligence?

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.

Where did LLMs come from?

The story of LLMs is surprisingly recent and moves fast:

Year	What happened
2017	Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.
2018	Google releases BERT. OpenAI releases GPT-1. The race begins.
2020	OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.
2022	OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.
2023	GPT-4 released. Open-source models (Llama, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.
2024	Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.
2025	Claude Code launches. Chinese open-weight models (DeepSeek, Qwen, Kimi) match the frontier at a fraction of the cost. Multi-hour autonomous tasks become reliable.
2026	Every major lab ships a work agent. Anthropic splits its top model into a public (Fable) and restricted (Mythos) tier. GPT-6, Gemini 3.8, Meta Muse, DeepSeek V4.

The main LLMs in 2026 — who makes what

Model names and version numbers change every few months. The labs behind them, and what each is known for, change far more slowly. Read the labs, not the numbers.

Claude (Anthropic)

The model family behind Claude Code, Cowork, Microsoft's Copilot Cowork and the default in OpenClaw. Tiers: Fable (the flagship, a new tier above Opus introduced in 2026), Opus, Sonnet and Haiku. Anthropic's Constitutional AI training makes Claude the model most agent builders reach for when reliability and honesty over long unsupervised sessions matter. Context window: one million tokens. See the [WTF is Claude](#) guide.

GPT (OpenAI)

The models behind ChatGPT, still the most widely used consumer AI product in the world, and behind OpenAI's agents: ChatGPT Work, Codex and the Agents API. GPT-6 arrived in September 2026 after the GPT-5.6 family in July. OpenAI's strength is reach: hundreds of millions of people already have the app.

Gemini (Google DeepMind)

Google's flagship family, currently Gemini 3.8, and the model behind Gemini Agent, Chrome's Auto Browse and Google's coding agent Jules. Deeply integrated with Search, Workspace and Android. Google also created the Agent2Agent (A2A) protocol for agents from different

companies to talk to each other.

Muse (Meta)

Meta's 2026 reset: a new model family and a consumer personal agent, both called Muse, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is distribution — three billion people already use its messaging apps.

Grok (xAI)

Elon Musk's models, now part of SpaceX after it acquired xAI and the coding tool Cursor. The model behind Grok Bot, a team of agents sharing one cloud computer.

The Chinese open-weight models

DeepSeek, Alibaba's Qwen, Moonshot's Kimi, Zhipu's GLM and MiniMax. Released as open weights — anyone can download and run them — at a fraction of the price of US frontier models, and close enough in capability that a great many indie agents run on them. DeepSeek's cheap-to-train breakthrough in early 2025 was the moment this became a two-continent race. Guide: [WTF is Happening with China's AI](#).

Other open models

Meta's Llama family, once the default open model, has been superseded inside Meta by Muse. Mistral (France) produces efficient smaller models. OpenClaw, Paperclip and most agent frameworks let you plug in any of these by API key.

What LLMs are good at — and bad at

LLMs are good at	LLMs are bad at
Writing in any style or tone	Precise arithmetic (they approximate)
Summarising long documents	Real-time information (training data has a cutoff)
Writing and explaining code	Guaranteed factual accuracy (they can "hallucinate")
Translating between languages	Knowing what they don't know
Reasoning through multi-step problems	Consistent behaviour on identical inputs
Generating creative ideas	Tasks requiring persistent memory by default
Explaining complex topics simply	Physical world interaction (without agent tools)

Hallucination — the most important limitation

The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

Parameters, tokens, and context windows

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the "size" of the model's brain.

Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude's current models have a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.

Constitutional AI — Anthropic's approach to safety

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

LLMs and the agentic economy

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.
- Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Fable, GPT-6 and Gemini 3.8 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

Glossary

LLM (Large Language Model) — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, Gemini, Muse and Grok.

Transformer — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

Parameters — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

Token — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

Context window — How much text an LLM can process at once. Current frontier models: up to 1 million tokens.

Hallucination — When an LLM generates confident but factually incorrect text. A key limitation to understand.

Constitutional AI — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

SWE-bench — The standard benchmark for measuring LLM performance on real software engineering tasks.

Fine-tuning — Training an existing LLM further on a specific dataset to improve performance on a specific task.

Inference — Running an LLM to generate a response. The "thinking" part, as opposed to training.

Open-weight model — An LLM whose weights are publicly available — anyone can download and run it. Examples: DeepSeek, Qwen, Kimi, Mistral, Llama.

Now you know what LLMs are. Here is the rest of the WTF Agents series.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/guides/claude

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/guides/api

WTF is MCP?

The plug that lets any AI agent use any tool. Invented by Anthropic, adopted by every rival, given away to a foundation, and now downloaded a hundred million times a month. What it is, why it matters, and what it does not solve.

The one-liner

MCP, the Model Context Protocol, is an open standard that lets any AI model connect to any tool, file, database or service through one common plug. Before it, every connection between an agent and a system was custom-built. After it, a tool built once works with every agent.

The usual analogy is USB-C, and it is accurate: one port, any device. It is the reason OpenClaw can use your calendar, Cowork can read your files, Paperclip can hand work between agents from different vendors, and a Polsia company can take payments. Every agent in this series speaks it.

By the numbers — September 2026

Nov 2024

Anthropic publishes MCP as an open standard

Dec 2025

Anthropic donates it to the Agentic AI Foundation under the Linux Foundation, co-founded with OpenAI and Block

97M

monthly downloads of the official Python and TypeScript libraries by March 2026

10,000+

public MCP servers by mid-2026, up from a few hundred in the first months

300+

MCP clients — apps that can use those servers, from Claude and ChatGPT to Cursor and VS Code

30+

security vulnerabilities filed against MCP servers and clients in January and February 2026 alone

2

open standards now under the foundation: MCP for agent-to-tool, A2A for agent-to-agent

The problem it solved

An AI model on its own can only read what you type and write text back. To be an agent it has to reach things: your inbox, a spreadsheet, a database, a browser, a payment system.

Until late 2024 that meant a custom integration for every pair. Connect Claude to Slack: one piece of code. Connect GPT to Slack: another. Connect either to your CRM: two more. The number of integrations grew as models times tools, and every company building an agent rebuilt the same connectors.

MCP turned the multiplication into addition. A tool exposes itself once, as an MCP server. Any model that speaks MCP can use it. Slack builds one server; every agent gets Slack.

How it works — the three parts

Part	What it is
Host	The app you are using: Claude, ChatGPT, Cursor, OpenClaw. It runs the model.
Client	The connector inside the host that speaks MCP to one server.
Server	The tool side: a small program that exposes a system — GitHub, Google Drive, Stripe, your own database — as things an agent can do.

When a session starts, the host asks each connected server what it offers. A server answers with three kinds of thing: tools (actions it can take, like "send an email" or "run a query"), resources (data it can hand over, like a file or a record) and prompts (templates for common jobs). The model reads the list, decides what it needs, and calls it. Results come back in a form the model can read, and the loop continues.

The server can also flag what a tool does — read-only, or destructive — so the host can ask you before anything irreversible. Since the November 2025 version of the standard, servers on the internet authenticate with OAuth, the same login handshake your apps already use, and long jobs can run in the background and report back.

Who runs it now

Anthropic invented MCP and, unusually, gave it away. In December 2025 it donated the standard to the Agentic AI Foundation, a directed fund under the Linux Foundation, co-founded with OpenAI and Block and backed by Google, Microsoft, AWS, Cloudflare and Bloomberg. Changes now go through public working groups and written proposals, the way internet standards do.

Google's A2A protocol — Agent2Agent, the standard for agents from different companies talking to each other — joined it under the same foundation. Together they are the two layers of agent plumbing: MCP connects an agent to tools, A2A connects agents to each other, and no single company owns either.

For anyone building on agents this is the point: the connectors you build are not hostage to one vendor's roadmap. A server written for Claude works in ChatGPT, Cursor, OpenClaw and whatever ships next.

Where you meet it

You have probably used MCP without seeing the name.

- **Connectors in Claude and ChatGPT.** When you link Gmail, Drive or Slack to a chat app, that is an MCP server on the other end.
- **Cowork and ChatGPT Work.** The apps they reach into, and the built-in browser, are MCP servers.
- **Claude Code, Codex, Cursor.** Developer tools use MCP to reach documentation, databases and deployment systems. This is where adoption is deepest.
- **OpenClaw and Paperclip.** Skills and adapters are MCP servers under another name; Paperclip itself exposes an MCP server so other tools can drive it.
- **Registries.** The official registry and GitHub's list the public servers, ten thousand and counting: an app store for agent capabilities.

The catch — a plug is not a lock

MCP standardised how agents connect to things. It did not standardise who is allowed to connect to what, and the gap between those two sentences is where the security problems live.

Over thirty vulnerabilities were filed against MCP servers and clients in the first two months of 2026, one rated near the top of the severity scale. The recurring shapes:

- **Tool poisoning.** A malicious server describes its tool one way and does another, or hides instructions in the description the model reads. Cisco found OpenClaw skills doing exactly this.
- **Prompt injection through data.** A document or web page the agent fetches through MCP contains instructions the agent follows. Still the unsolved problem of the whole field.
- **Over-broad access.** An agent connected to email, files and payments through a dozen servers is one bad instruction away from using all of them.
- **Untrusted servers.** Ten thousand public servers means ten thousand pieces of code someone wrote. Registries list them; they do not vet them all.

The standard has responded — OAuth for remote servers, security requirements for clients, tool annotations — and enterprise vendors sell gateways that sit in front of MCP to enforce permissions and log every call. For a normal person the rules are the same as for any agent: connect the minimum, prefer official servers, and keep a human approval in front of anything that spends or sends.

Why it matters

Standards are boring until they are not. TCP/IP made the internet possible by letting any computer talk to any other. MCP is doing that for agents: any model, any tool, one protocol, no permission needed from anyone. That it was invented by one lab and now belongs to a foundation backed by all of its rivals is the rare case of an industry agreeing to share the plumbing so it can compete on what runs through it.

If you build anything on agents, build it on MCP. If you only use agents, you already are.

Glossary

MCP (Model Context Protocol) — The open standard for connecting AI models to tools, data and services. Anthropic, November 2024; Agentic AI Foundation since December 2025.

Host — The application running the model and the MCP clients: Claude, ChatGPT, Cursor, OpenClaw.

Client — The connector inside a host that talks to one MCP server.

Server — A program that exposes a system to agents through MCP: tools, resources and prompts.

Tool — An action a server offers: send, query, create, run.

Resource — Data a server offers: a file, a record, a page.

Annotation — A hint attached to a tool, such as read-only or destructive, so a host can ask before acting.

Registry — A directory of public MCP servers. The official one and GitHub's are the largest.

Agentic AI Foundation (AAIF) — The Linux Foundation body that governs MCP and A2A. Co-founded by Anthropic, OpenAI and Block.

A2A (Agent2Agent) — Google's standard for agents talking to each other; MCP's sibling under the foundation.

Tool poisoning — A malicious server misdescribing what its tool does, or hiding instructions in the description.

OAuth — The standard login handshake; required for MCP servers on the internet since the November 2025 spec.

Liked this? Go deeper.

MCP is the plug. Here is what plugs into it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/guides/api

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

WTF is Polsia?

One founder. Zero employees. Nearly half a million AI-built companies, a \$250 million valuation, and a Trustpilot page that tells a very different story. Here is what Polsia actually is, what it costs, and who should — and should not — touch it.

The one-liner

Polsia is a subscription platform where you describe a business idea — or let the AI pick one — and a team of AI agents builds the website, writes the code, runs the marketing, sends the outreach, answers support and reports back to you every morning. The founder's tagline: "AI that runs your company while you sleep."

It is also the most-argued-about product in the agentic economy. Investors valued it at a quarter of a billion dollars. Paying customers gave it one star in large numbers. Both of those things are true at the same time, and understanding why is the point of this guide.

By the numbers — September 2026

Everything in this box changes weekly. Polsia publishes a live public dashboard, which is how these figures can be checked. Everything below the box is written to still be true when the numbers have moved.

\$30M

raised in May 2026 at a \$250M valuation (Sound Ventures, True Ventures)

\$18.7M

headline ARR on Polsia's live feed, 15 Sept 2026 — about \$11M of it recurring subscriptions

484,592

companies created on the platform since launch

29,132

of those still active — roughly 6%

~50%

month-one churn, stated by the founder on the record

10%

of companies that have ever earned at least one dollar, per the founder

2.9 / 5

Trustpilot rating across 185 reviews, 45% of them one star

1

employee. Still just Ben.

Who built it

Polsia was built by Ben Broca, who also goes by Ben Cera. He is French, trained as an engineer at CentraleSupélec with a masters from Columbia, and now lives in San Francisco. Before Polsia he ran international operations at CloudKitchens under Travis Kalanick and co-founded Hutch, a 3D-rendering e-commerce startup that raised roughly \$17 million.

He spent most of 2025 building separate AI tools — a coding tool, a marketing tool — before concluding the models were good enough to run every business function at once. Polsia launched in late 2025, crossed \$1 million in annualised revenue within about a month, and tripled the month after. He has been the only employee throughout, and says that is deliberate: if he is selling an AI that runs companies, Polsia should be run by one.

In May 2026 he raised \$30 million at a \$250 million valuation, reportedly letting the AI manage the data room and brief investors. That is the number the headlines quote. It says something about what investors believe. It says nothing about whether the product will work for you.

How it works

You give Polsia an idea, or click "surprise me" and it picks one from your public profile. Within minutes it has a company name, a mission statement, a market research document, a live website on a polsia.app subdomain, an email inbox, a Stripe account and a GitHub repository. Then it asks for a credit card.

Once you subscribe, the agents run on a daily cycle:

Agent	What it does
Orchestrator (CEO)	Reviews the state of the company, writes the morning plan and the evening summary
Code	Ships features and fixes to the website and product
Marketing	Drafts and posts social content, writes landing pages
Outreach	Finds prospects and sends cold email
Support	Reads the inbox and drafts replies
Ads	Runs and optimises Google and Meta campaigns from your budget
Finance	Syncs Stripe revenue, tracks spend
Research	Watches competitors and refreshes the plan

Each task the agents run costs a credit. Your plan includes a monthly allowance; you buy more if you run out. You read the summaries, approve or redirect, and the cycle repeats whether or not anyone is buying what the company sells.

The most important thing to understand about how Polsia works is what it does not do. There is no step between "here is the idea" and "the agents are building it" that asks whether anyone wants the thing.

What it actually costs

Polsia does not publish pricing on its homepage. As of August–September 2026, the in-product pricing showed five tiers:

Plan	Price	Includes
Standard	\$20/mo	1 company, 20 credits/mo
Pro	\$50/mo	3 companies, 50 credits/mo, Meta Ads
Max	\$200/mo	10 companies, 200 credits
Max 2	\$500/mo	25 companies, 500 credits
Max 3	\$1,000/mo	50 companies, 1,000 credits

The subscription is not the whole bill. Three more lines matter, all from Polsia's own Terms of Service:

→ **Ad spend fee — 20%.** Any advertising budget the agents spend on your behalf carries a 20% platform fee.

→ **Revenue fee — 3%.** Polsia takes a cut of customer payments your company receives. In the September 2026 terms this is "currently 3%"; until then it was 20%. Check the live terms — it has already changed once.

→ **Withdrawal limits.** Your company's revenue lands in a Polsia balance first. Payouts have a 31-day settlement period, a \$50 minimum and a cap of \$500 per calendar month. A company earning \$2,000 a month can take out \$500 of it.

Users have also reported the credit meter running on a single click with no confirmation, and credits burned on tasks that failed. Budget for the plan, the credits, and the fees — then read the withdrawal cap twice.

The numbers — what the live feed actually says

Polsia publishes a public dashboard feed with enough raw data to rebuild the headline figure yourself. On 15 September 2026 it read \$18.7 million ARR.

That number is built the way many AI startups build ARR in 2026: take the last thirty days of revenue, including one-off payments, and multiply by twelve. Ben has described the method himself on the record. Broken down, roughly \$11 million is annualised subscription revenue. The rest is the last month of credit packs, ad budget routed through the platform, domain purchases and a small slice of user-company revenue, each times twelve. Real money, but around 40% of it is not recurring by construction.

The second number on the feed is the one that matters more for you. Nearly half a million companies have been created. About 29,000 are active. Companies are cheap to spin up here, so a low survival rate is expected rather than damning — but it is the clearest public signal of what happens after the agents finish building.

Ben gave the third number to Mixergy's Andrew Warner in May 2026, in an interview titled "Is Polsia a \$250M scam?": about 10% of companies on the platform have made at least a dollar. The best one he had seen was making three or four thousand a month. Month-one churn is around 50%. He also said the engineering agent was writing "spaghetti code" and the product "doesn't work well enough for my taste."

That is unusual candour from a founder at a \$250 million valuation, and it is the honest frame for the whole product: the agents really do build things. Nine times out of ten, nobody buys them.

What paying customers say

Polsia's Trustpilot page tells the story in three acts. In June 2026 it had 35 reviews. By mid-August it was 1.8 out of 5 across 79 reviews, three quarters of them one star. Polsia then claimed the profile, over a hundred reviews arrived in four weeks, and by mid-September it stood at 2.9 out of 5 across 184 reviews, with 45% one star and 27% five star. Several of the new five-star reviews describe a problem the founder fixed personally.

The complaints did not change across that period. The recurring themes:

- Tasks marked "complete" by the AI that never actually deployed
- Credits consumed by duplicate or failed actions, with refunds refused
- Cold emails sent to real contacts, or public posts made, without approval
- Websites and custom domains locked to Polsia's own hosting account
- Sites redeployed after an explicit takedown request
- Cancellations that took repeated attempts, and support tickets unanswered for weeks

The most-cited case is a Rest of World profile from April 2026: a factory worker in China paying \$199 a month — a quarter of his salary — for an AI to sell a spiritual guidance app to Americans. The agents built the site, filled it with fabricated reviews, ran Facebook ads and emailed journalists he had never heard of. After several months: seven signups, zero customers.

The positive cases are real too. Founders who arrived with a validated idea and a clear brief report the agents standing up infrastructure, research, outreach and even video ads in hours — plumbing that would otherwise eat a week. One four-star reviewer who built a real product on it put the line exactly: the weakness is not what the agents can build, it is what the customer cannot control when the surrounding infrastructure fails.

The lock-in question

This is the part most first-time users miss. Polsia deploys your company on Polsia's hosting, on a `polsia.app` subdomain, with email routed through Polsia. There is a "download code" button, but users report the export is a scaffold rather than a working product, and that there is no direct access to the hosting or repository. Self-hosting is listed as "coming soon."

If the platform is working for you, none of this matters. If you want to leave, take your domain with you, or hand the code to a developer, it matters a great deal. Verify the current export and account-recovery terms in writing before you subscribe.

Who Polsia is for

Polsia is an execution engine. It is not an idea-validation engine, and it does not pretend to be one.

It fits if you:

- Already have evidence people want the thing — conversations, a tested price, a channel that works
- Are launching cheap, fast experiments where one flop costs little
- Are happy paying the fees and living with the withdrawal cap in exchange for not touching any of the plumbing
- Do not need to own the code or the infrastructure

Skip it if you:

- Have an idea but have not talked to a single potential customer
- Are in a regulated, technical or high-stakes category
- Need to withdraw more than \$500 a month, or would be hurt by 20% on ad spend
- Plan to migrate off later, or want portable code from day one

The "surprise me" button is the most marketed feature and the riskiest one. It means the AI is choosing the idea and executing it with no demand signal anchoring either decision. The factory worker in the Rest of World story is what that looks like in practice.

Our verdict

Polsia is real, funded, and built by a credible operator. The agents genuinely build, market and operate companies with no human doing the work, and the platform has grown faster than almost anything in the category. The hype question is settled.

The question that matters now is fit. Polsia executes whatever it is given. Hand it a validated idea and a precise brief and it can save you weeks. Hand it a guess and it will build a polished company nobody asked for, spend your credits and your ad budget promoting it, and report back cheerfully every morning while nothing sells. The founder's own figure — one company in ten ever earns a dollar — is that gap measured from inside.

Validate first. Then decide whether Polsia is the right layer to execute on what you found.

Glossary

Polsia — Subscription platform where AI agents build and run a company for you. Founded by Ben Broca (Ben Cera), launched late 2025, San Francisco.

Credits — Polsia's unit of work. Each task the agents run costs a credit. Plans include a monthly allowance; extra credits are bought in packs.

ARR (Annual Recurring Revenue) — Annualised recurring revenue. In much of the 2026 AI industry, including Polsia's headline figure, it means the last thirty days of all revenue multiplied by twelve.

Run rate — A recent period's revenue extrapolated to a full year. Polsia's headline number is a run rate.

Platform fee — Polsia's cut. 20% on advertising spend and, as of September 2026, 3% on customer payments (previously 20%).

Withdrawal cap — The limit on how much of your company's revenue you can take out of your Polsia balance: \$500 per calendar month, after a 31-day settlement period.

Surprise me — Polsia's option to let the AI choose your business idea for you.

Validation — Checking that people actually want a product before building it. The step Polsia does not perform.

Churn — The share of paying customers who cancel in a given period. Polsia's founder puts month-one churn at around 50%.

polsia.app — The subdomain every Polsia company is deployed on. Your company's site and email run through Polsia's infrastructure by default.

Polsia is one piece of the agentic economy. Here is the rest of the picture.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

wtfagents.com/guides/paperclip

WTF is NanoCorp

The Polsia rival built by AI researchers that publishes its companies' real revenue. How it works, what it costs, and what the live numbers say.

wtfagents.com/guides/nanocorp

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is OpenClaw?

One Austrian developer. One open-source project. Three name changes in 60 days, a hire by OpenAI, dating profiles created without consent, a ban from Chinese government offices — and now millions of users, a foundation, and native apps. The wildest story in tech, and where it stands today.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

OpenClaw is a free, open-source autonomous AI agent that runs on your computer or phone and takes instructions via messaging apps — WhatsApp, Telegram, Signal, Discord — connecting to Claude, GPT, DeepSeek or any model you choose to carry out multi-step tasks on your behalf.

By the numbers — September 2026

3.2M

active users, making it the most-used open-source agent in the world

385,000+

GitHub stars — the sixth most-starred repository in GitHub's history

30 Aug

2026: OpenClaw 2.0 released, with native iPhone, iPad and Android apps

\$0

cost. MIT licence. You pay only for the model you plug in

1

non-profit foundation now governing the project, with an independent board

3

names in its first 60 days: Clawdbot, Moltbot, OpenClaw

It is not a platform. It is not a service. It is software you install and run yourself. And it became one of the most talked-about pieces of software in the world in under 60 days.

The name history — three identities in 60 days

Understanding OpenClaw requires understanding its chaotic naming history:

Date	Name	Why
November 24, 2025	Clawdbot	Original launch name. A portmanteau of "Clawd" (the developer's personal AI assistant, itself named after Anthropic's Claude) and "bot."
January 27, 2026	Moltbot	Anthropic sent trademark complaints about "Clawd" being too similar to "Claude." Renamed to Moltbot, keeping a lobster theme (lobsters molt their shells).
January 30, 2026	OpenClaw	Peter Steinberger found "Moltbot" didn't roll off the tongue. Renamed again to OpenClaw — and it stuck.

Who built it — Peter Steinberger

Peter Steinberger is an Austrian software engineer and entrepreneur. He co-founded PSPDFKit (now Nutrient) in 2011 — a PDF SDK framework he bootstrapped to a successful exit. He is based between Vienna and London.

Steinberger built his own personal AI assistant called Clawd — named after Anthropic's Claude, which it used as its underlying model. He then open-sourced a generalised version of it as Clawdbot in November 2025.

He describes himself as a "vibe coder" — someone who builds primarily through AI-assisted development. The irony of an AI agent builder who uses AI agents to build AI agents was not lost on the internet.

On February 14, 2026 — Valentine's Day — Steinberger announced he was joining OpenAI and that OpenClaw would move to an independent open-source foundation, financially sponsored by OpenAI. This was widely described as an "acqui-hire": OpenAI acquired the talent without acquiring the open-source software itself.

What OpenClaw actually does

OpenClaw runs locally on your machine. You interact with it via a messaging app — Signal, Telegram, Discord, or WhatsApp. You send it a message with a task. It uses an LLM (Claude by default, or GPT or DeepSeek) to reason through the task and executes it using whatever tools and permissions you have given it.

What it can do depends on what you connect it to. Examples:

- Read and send emails on your behalf
- Browse the web and summarise findings

- Manage your calendar
- Write, edit, and commit code
- Post to social media
- Execute terminal commands
- Interact with any service that has an API

OpenClaw uses a "skills" system — modular capabilities stored as directories that can be installed from ClawHub, the community skill registry. Think of skills as apps for your AI agent.

The growth

OpenClaw's growth was unprecedented. By February 2026 — just over two months after launch — it had 247,000 GitHub stars. By September it had passed 385,000, the sixth most-starred repository in GitHub's history, and over three million people were running it.

For context: Linux took decades to accumulate its GitHub stars. OpenClaw surpassed most long-standing projects in weeks and kept going. The 2.0 release in August, with native mobile apps and Windows support, turned it from a developer's tool into something ordinary people install.

The timing at launch was perfect. OpenClaw arrived just as Moltbook went viral — a social network built for AI agents. The two projects fed each other's growth: OpenClaw agents connecting to Moltbook, humans watching in fascination and horror.

Moltbook — the AI social network that made OpenClaw famous

On January 27, 2026 — the same day Clawdbot was renamed Moltbot — entrepreneur Matt Schlicht launched Moltbook: a social network built exclusively for AI agents.

The premise was wild: a Reddit-like platform where AI agents (not humans) create profiles, post content, and interact with each other. OpenClaw agents began autonomously discovering Moltbook, joining it, and posting — without their human owners necessarily knowing or intending this.

The internet went predictably insane. Wired ran a story: "I Infiltrated Moltbook, the AI-Only Social Network Where Humans Aren't Allowed." TechCrunch, CNBC, The Verge, and mainstream media followed.

Then things got worse. On January 31, an unsecured Supabase database exposed 6,000+ email addresses and 1 million+ agent interactions — and crucially, allowed anyone to take control of any agent on the platform. Security firm Wiz confirmed the vulnerability. Moltbook went offline briefly to patch it.

Then things got weirder. An OpenClaw agent autonomously created a dating profile on MoltMatch (an AI agent dating platform) and began screening romantic matches — without its human owner's knowledge or consent.

Moltbook was acquired by Meta on March 10, 2026. Matt Schlicht and co-founder Ben Parr joined Meta Superintelligence Labs.

The security problems — read this before installing

OpenClaw is powerful. It is also genuinely dangerous if used carelessly. This is not marketing caution. Multiple credible security researchers have documented real risks.

Cisco's findings

Cisco's AI security research team tested a third-party OpenClaw skill called "What Would Elon Do?" They found it performed silent data exfiltration and prompt injection without user awareness. They also found the ClawHub skill repository lacked adequate vetting to prevent malicious skill submissions.

Remote code execution vulnerabilities

Security firm Conscia disclosed a remote code execution chain and two additional command injection vulnerabilities in early 2026.

The maintainer's own warning

One of OpenClaw's own maintainers, known as "Shadow," publicly warned on Discord: "If you can't understand how to run a command line, this is far too dangerous of a project for you to use safely."

China bans it

In March 2026, Chinese authorities restricted state-run enterprises and government agencies from running OpenClaw apps on office computers.

The rule of thumb for OpenClaw: give it the minimum permissions it needs. Only install skills from trusted sources. Review its actions regularly. If you are not comfortable with the command line, start with something simpler.

OpenClaw vs Polsia vs Paperclip

	OpenClaw	Polsia	Paperclip
Type	Personal AI agent	Managed company platform	Orchestration framework
Runs	Your machine or phone	Polsia's servers	Your servers
Cost	Free (+ LLM costs)	\$20–1,000/mo + fees	Free (open source)
For	Personal automation	Launching AI companies	Multi-agent companies
Technical	Medium (CLI)	Low (no-code)	High (developer)
Open source	Yes (MIT)	No	Yes
LLMs	Any — Claude by default	Claude Opus	Claude Code + 7 others

Where it stands now — the OpenClaw Foundation

The move to a foundation, announced when Steinberger joined OpenAI, has happened. OpenClaw is now governed by an independent non-profit with a board chaired by Dave Morin, the founder of Path. OpenAI remains a financial sponsor — an arrangement that still gives one of the world's most powerful AI companies influence over the most popular open-source agent, though the board structure now sits between the two.

Development has not slowed. OpenClaw 2.0 shipped on 30 August 2026 with native iPhone, iPad and Android apps and a rebuilt skills system. Security improved in specific ways: an approval is now tied to the exact request, command, session and person it was granted for, rather than becoming a standing permission, and protected credentials can reach the services that need them without ever appearing in text the model can read. Prompt injection, as everywhere in the agentic economy, is not solved.

OpenClaw is the reference implementation for what a personal AI agent looks like, and the benchmark every competitor — including Meta's Muse — is measured against. It is also hired into Paperclip org charts as a worker, which makes it one of the few agents that lives in both the personal and the company layer.

Glossary

OpenClaw — Free, open-source autonomous AI agent. Runs on your machine or phone, connects via messaging apps. Created by Peter Steinberger, November 2025; governed by the OpenClaw Foundation since 2026.

Clawdbot / Moltbot — Previous names for OpenClaw. Clawdbot was the original name; Moltbot followed after Anthropic trademark complaints.

Peter Steinberger — Austrian developer who created OpenClaw. Co-founded PSPDFKit. Joined OpenAI February 2026.

Acqui-hire — When a company hires a founder or team without technically acquiring their product. OpenAI hired Steinberger; OpenClaw remained open source.

Moltbook — A social network for AI agents. Launched January 2026; its founders joined Meta in March 2026.

ClawHub — The community skill registry for OpenClaw — modular capabilities that can be installed to extend the agent.

Skills — Modular capabilities for OpenClaw stored as directories containing a SKILL.md file.

Prompt injection — A security attack where malicious instructions are hidden in data an agent reads. Found in OpenClaw skills by Cisco.

MIT License — An open-source licence allowing anyone to use, modify, and distribute the software freely.

MoltMatch — An experimental AI agent dating platform where OpenClaw agents were creating profiles without user knowledge.

OpenClaw is one part of the story. Here is the rest.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

wtfagents.com/guides/paperclip

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

wtfagents.com/guides/meta-muse

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Paperclip?

An org chart for AI agents. Roles, reporting lines, budgets and approvals — for a team where nobody is human. Open-source, self-hosted, and quietly becoming the way people actually run more than one agent at a time.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Paperclip is open-source software that lets you run AI agents the way a company runs employees: an org chart with roles and reporting lines, a budget for each agent, tasks that flow down and results that flow up, and a human at the top who approves what matters.

Its own tagline has softened since launch — from "zero-human companies" to "the app people use to manage AI agents for work" — but the idea is the same. Once you have more than one agent, you need a way to manage them. Paperclip is that way.

By the numbers — September 2026

4 March

2026: public launch. Under three weeks to 30,000 GitHub stars

~70,000

GitHub stars by June, among the fastest-growing agent projects ever released

100+

contributors; releases roughly every two weeks, the latest on 2 September

8

agent types it can hire out of the box: Claude Code, Codex, Gemini CLI, Cursor, Hermes, OpenClaw, Pi, OpenCode

\$0

cost of the software. MIT licence. You pay only for the agents' model usage

~60

forkable company templates on ClipMart, Paperclip's marketplace

0

disclosed funding rounds. Built by a pseudonymous developer, "dotta," and a community

The problem it solves

One agent is easy. You open Claude Code, or OpenClaw, or Cowork, and you talk to it.

Five agents are chaos. Five terminal windows. Five separate memories that do not know what the others have done. Two of them doing the same task. One of them stuck since Tuesday. No idea what any of it is costing. Nobody enforcing quality. This is what most people hit within a week of trying to run a business on agents, and it is why so many "AI-run company" experiments end as a pile of abandoned sessions.

Paperclip's answer is to stop treating agents as tools you operate and start treating them as staff you manage.

How it works

You install Paperclip on your own machine or server. It runs a local web dashboard. In it, you create a company.

Hire agents into roles

Each agent gets a job title, a reporting line and a written brief. "Marketing lead, reports to CEO." "Backend engineer, reports to CTO." The CEO can itself be an agent, with you above it as the board.

Paperclip is agent-agnostic. Its own phrase: "If it can receive a heartbeat, it's hired." Claude Code for engineering, Codex or Gemini CLI for other tasks, Hermes for research, OpenClaw for anything that needs a personal agent's tools, all under one org chart, each on whichever model suits the job and the budget.

Work flows as issues

Tasks are tickets, like a Jira or Linear board. A goal goes in at the top; the CEO agent breaks it into issues and assigns them down the chart. Agents pick up their issues, do the work, comment on the thread, hand off, and close. Parent and child tasks, dependencies and blockers are all tracked, and an agent that is blocked wakes the agent that can unblock it.

Heartbeats keep it alive

Each agent runs on a schedule, a "heartbeat." Every tick it checks its inbox, does the next thing, and reports. Watchdogs review tasks that have stalled so you do not have to. It runs while you sleep, but on a leash.

Budgets and approvals

Every agent has a spending cap. Anything sensitive — a payment, a deploy, a public post — can be gated behind human approval, and approvals sit in one queue. Execution policies

define what an agent may do without asking. You set the leash length.

Teams, sandboxes, plugins

Since summer 2026 Paperclip supports multiple human users with their own credentials, a company-wide timeline of what every agent did and when, isolated sandbox environments for agents to work in, a plugin system for adding new agent types, and an MCP server so other tools can drive Paperclip itself.

Paperclip vs Polsia vs OpenClaw

These three get confused constantly. They are different layers.

	Paperclip	Polsia	OpenClaw
What it is	Manager for a team of agents	Platform that runs a company for you	A single personal agent
Who hosts it	You	Polsia	You
Cost	Free + model usage	\$20–1,000/mo + fees	Free + model usage
Who picks the agents	You, any vendor	Polsia's fixed team	It is the agent
Control	Full — code, data, budgets	Limited — their infra, their fees	Full
Effort	High: setup and management	Low: subscribe and read reports	Medium

Polsia is the managed hotel. Paperclip is the house you build and run. OpenClaw is one very capable employee, who can also be hired into a Paperclip org.

ClipMart — fork a company

At launch Paperclip announced ClipMart, a marketplace for AI-agent companies. It has since shipped at clipmart.ai, and the pitch is "fork an AI company, launch in 60 seconds."

Instead of designing an org chart from scratch, you pick a template — a marketing agency, a SaaS startup, an e-commerce store, a media company — and fork it into your own Paperclip. The template comes with the roles, the briefs, the reporting lines and an estimate of how many agents it runs and what they cost. Around sixty templates across marketing, SaaS, e-commerce, agency, media and finance were listed at time of writing, and anyone can submit their own.

Two things worth knowing. Forking a company is not the same as having customers, which is the trap in every "launch in 60 seconds" pitch in this series. And a template is only as good as the briefs inside it; read them before you hire.

The reality check

Paperclip is the most serious open-source answer to "how do I run more than one agent," and it is not a product for people who want a button.

What is genuinely good

- The org-chart model works. People who use it describe the mental shift from "prompting an AI" to "managing a team," and that shift is the point.
- Agent-agnostic for real. Mix Claude, OpenAI, Google and open models by role and cost.
- Governance is built in, not bolted on: budgets, approvals, execution policies, audit timeline.
- It builds itself. The creator runs Paperclip's own development through a Paperclip company, and the release cadence backs that up.
- No lock-in, no account, no vendor. Your data and code stay on your machine.

What is not

- Self-hosted means you host it. A server, PostgreSQL, Node.js, and you doing the updates. A hosted cloud version has been on the roadmap all year and had not launched at time of writing.
- The agents are only as good as their briefs. Paperclip organises work; it does not make an agent competent at it.
- No disclosed team, funding or company behind it beyond a pseudonym and a community. That is how great open-source projects start and also how some stall.
- It is a control plane, not a business. Nothing about Paperclip finds customers. See the Polsia guide for how that usually goes.

Who Paperclip is for

Yes — Founders and operators already running two or more agents who are losing track. Developers comfortable with a terminal and a server. Anyone who wants to own the whole stack rather than rent it.

Not yet — Anyone who has not run a single agent successfully. Start with OpenClaw or Cowork, get one thing working, then come back.

No — Anyone who wants "AI runs my business" as a subscription with no setup. That is Polsia, with all the trade-offs in that guide.

Glossary

Paperclip — Open-source, self-hosted software for managing AI agents as an org chart with roles, budgets and approvals. Launched March 2026.

Orchestration — Coordinating multiple agents so they work together without duplication or chaos. What Paperclip does.

Control plane — The layer that manages and governs agents, as opposed to the agents themselves. Paperclip's own description of its role.

Heartbeat — The scheduled tick on which an agent checks for work and acts. Any agent that can receive one can be hired into Paperclip.

Adapter — The connector that lets a particular agent type (Claude Code, Codex, OpenClaw...) work inside Paperclip.

Issue — A unit of work on the board: assigned, threaded, tracked, closed. Paperclip's equivalent of a task or ticket.

Execution policy — The rules for what an agent may do without human approval.

Watchdog — An automated check that reviews stalled tasks and wakes or escalates them.

ClipMart — Paperclip's marketplace of forkable AI-company templates at clipmart.ai. Pick a template, launch it into your own Paperclip.

MIT licence — A permissive open-source licence. Free to use, modify and sell, with attribution.

Self-hosted — Software you install and run on your own machine or server, rather than a service someone else runs for you.

Paperclip manages agents. Here are the agents, and the alternatives.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/guides/polsia

WTF is NanoCorp

The Polsia rival built by AI researchers that publishes its companies' real revenue. How it works, what it costs, and what the live numbers say.

wtfagents.com/guides/nanocorp

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Anthropic?

Founded by people who left OpenAI over safety. Now the most valuable private company in the world, heading for the largest IPO in history, and the lab whose model sits inside half the agents in this series — including its rivals' products.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Anthropic is an AI safety company that builds some of the most powerful AI models in the world — the Claude family — while arguing, in public and in its own research, that models like these could be among the most dangerous technologies ever created.

That is not a contradiction. It is the founding bet: if powerful AI is coming regardless, better to have a safety-focused lab at the frontier than to leave the frontier to labs that are not. In 2026 that bet paid off in a way nobody predicted.

By the numbers — September 2026

Everything in this box will be out of date within months. Everything below it is built to last longer.

\$965B

valuation after the \$65B Series H in May 2026 — the largest private funding round in history

\$380B → \$965B

the valuation jump between February and May 2026, three months apart

\$65B

annualised revenue run rate by late July 2026, up from roughly \$9B at the end of 2025

1 June

2026: confidential S-1 filed with the SEC. An October listing has been widely reported, not confirmed

~2,300

employees, up from about 1,100 in 2024

8 of 10

Fortune 10 companies are Claude customers

1,000+

enterprises spending more than \$1M a year on Anthropic products

2021

founded, San Francisco. Structured as a Public Benefit Corporation

The founding story — the OpenAI exodus

In 2021 a group of senior researchers left OpenAI, then the most prominent AI lab in the world, over disagreements about direction. The core concern: as OpenAI became more commercial, some believed it was prioritising speed and capability over safety.

The group was led by siblings Dario Amodei, then OpenAI's VP of Research, and Daniela Amodei, its VP of Operations. They left with Tom Brown (lead author of the GPT-3 paper), Chris Olah (the field's leading interpretability researcher), Jack Clark, Jared Kaplan, Sam McCandlish and others.

They founded Anthropic with a stated mission: the responsible development and maintenance of advanced AI for the long-term benefit of humanity.

They structured it as a Public Benefit Corporation, which legally allows the company to weigh its mission alongside shareholder returns, and later added a Long-Term Benefit Trust with the power to appoint board members. Whether those structures hold under trillion-dollar pressure is one of the open questions of the decade.

The scale — from research lab to the most valuable startup on earth

For its first four years Anthropic was the smaller, more careful rival. Then Claude Code happened.

Launched to developers in early 2025, Claude Code became the tool that turned Claude from a chatbot into an agent that does work. Developers adopted it faster than any coding tool before it. Cowork, launched in January 2026, brought the same capability to everyone else. Revenue followed the products: roughly \$9 billion annualised at the end of 2025, \$30 billion by April, \$47 billion when the Series H closed in May, \$65 billion by the end of July.

The Series H itself — \$65 billion led by Altimeter, Dragoneer, Greenoaks and Sequoia, with Amazon, Capital Group, Fidelity and T. Rowe Price participating — valued the company at \$965 billion, above OpenAI's \$852 billion from two months earlier. Nine days later Anthropic filed confidentially for an IPO. Reports point to October 2026 and a target valuation that would make it the largest listing ever. None of that is confirmed by the company.

Amazon, Google and Nvidia remain major backers. Claude is available through AWS Bedrock, Google Cloud Vertex AI and Microsoft Foundry, which means it runs inside every major cloud, including the one owned by OpenAI's biggest investor.

What makes Anthropic different

Safety is the founding purpose, not a department

Every lab talks about safety. Anthropic was created because of it. A large share of its research output is interpretability (understanding what is happening inside a model) and alignment (making sure a model pursues the goals its operator intends). It publishes a Responsible Scaling Policy that ties deployment decisions to measured capability levels, and it has held back or gated models under it.

Constitutional AI

Anthropic's core training technique. Instead of relying only on human feedback, the model is given a set of written principles and trained to critique and revise its own outputs against them. The result is a model whose behaviour is more consistent and more predictable at the edges, which matters enormously when that model is running as an agent for hours with nobody watching.

The Fable / Mythos split

In 2026 Anthropic did something no lab had done before: it split one model into two products by safeguard level. Mythos-class models sit above the Opus tier in capability. The public version, Claude Fable, ships with classifiers that refuse dangerous cybersecurity and biology requests. The same model with those safeguards lifted, Claude Mythos, is available only to vetted organisations through Project Glasswing — a programme run with the US government, launched in April 2026 with AWS, Apple, Google, Microsoft and CrowdStrike, and expanded in June to around 150 organisations including NATO and the EU's cybersecurity agency. Their job: use the model to find and fix vulnerabilities in critical software before attackers do.

That is a new shape for the industry. Frontier capability gated by who you are and what you are for, rather than released to everyone or nobody.

MCP — the plug the whole industry uses

In November 2024 Anthropic published the Model Context Protocol, an open standard letting any AI agent connect to any tool or data source. Within a year every major lab had adopted it. Anthropic then gave it away, donating MCP to the Agentic AI Foundation under the Linux Foundation, co-founded with OpenAI and Block. The plumbing of the agentic economy is Anthropic's design and nobody's property.

The products

Everything is built around Claude.

Claude — The model family. Fable 5.1 at the top, Opus 5, Sonnet 5, Haiku 4.5. See the [WTF is Claude](#) guide.

Claude Code — The autonomous coding agent that changed the company's trajectory. Terminal, desktop app, mobile. See the [WTF is Claude Code](#) guide.

Cowork — The agent for everyone else. Works in your files, apps and browser without a command line. See the [WTF is Cowork](#) guide.

Claude.ai — The chat interface, web and mobile. Free tier; Pro, Max, Team and Enterprise plans above it.

Claude in Chrome, Excel, PowerPoint, Word and Outlook — Claude working inside the tools people already use.

Claude Tag — Claude inside Slack. Tag it into a thread and delegate the task.

The Claude Developer Platform and Agent SDK — The API and the framework other companies build their own agents on.

Amazon Bedrock, Google Vertex AI, Microsoft Foundry — Claude inside every major cloud.

The honest tension

Anthropic says AI may be one of the most dangerous technologies in human history, and then builds it anyway, and is now worth close to a trillion dollars for doing so.

Dario Amodei has been explicit about this: a company that may be building something profoundly dangerous, pressing forward because the alternative is leaving the frontier to others. Critics point out that this logic can justify almost anything, and that a public listing brings quarterly pressure that no benefit-corporation charter fully insulates against. Supporters point to the research record, the Responsible Scaling Policy, the Glasswing gating and the fact that Anthropic has repeatedly shipped later and more cautiously than rivals — and still won.

The honest answer: Anthropic is probably the most safety-focused lab at the frontier. Whether that is enough, given how fast the frontier is moving, is the question that matters, and no one can answer it yet.

Why Anthropic matters for the agentic economy

Anthropic's fingerprints are on almost everything in this series.

- **Claude** is the default model in OpenClaw, the most-used open-source personal agent.
- **Microsoft's Copilot Cowork** — Microsoft's flagship agent for 365 — runs on Claude, despite Microsoft's \$13 billion stake in OpenAI.
- **Polsia's** founder chose Claude Opus to run the agents that run his companies.
- **Claude Code** is the developer agent inside Paperclip and thousands of AI-run companies.
- **MCP** is how every agent in this series connects to its tools.

You can use the agentic economy without ever thinking about Anthropic. But you will almost certainly be using something it built.

The key people

Dario Amodei — CEO and co-founder. Former VP of Research at OpenAI. The public face, and the author of the essays that set the company's tone.

Daniela Amodei — President and co-founder. Former VP of Operations at OpenAI. Runs the business.

Chris Olah — Co-founder. The world's leading mechanistic interpretability researcher.

Tom Brown — Co-founder. Lead author of the GPT-3 paper.

Jared Kaplan — Co-founder. Discovered the scaling laws that guide how all labs train models.

Ami Vora — Chief Product Officer since January 2026, when Mike Krieger, co-founder of Instagram, moved to Anthropic's Labs group after two years in the role.

Glossary

Anthropic — AI safety company founded in 2021 by former OpenAI researchers. Builder of Claude. Public Benefit Corporation.

Constitutional AI — Anthropic's training technique: give the model written principles and train it to judge its own outputs against them.

Interpretability — Research into what is happening inside a model and why it produces the outputs it does.

Alignment — Ensuring an AI pursues the goals its operator actually intends.

Responsible Scaling Policy — Anthropic's public framework tying deployment decisions to measured capability and risk levels.

Mythos-class — The model tier above Opus. Fable is the safeguarded public version; Mythos is the restricted one.

Project Glasswing — Anthropic's vetted-access programme, run with the US government, giving cyberdefenders the unrestricted model to secure critical software.

MCP (Model Context Protocol) — The open standard connecting agents to tools. Invented by Anthropic, donated to the Agentic AI Foundation.

Public Benefit Corporation — A corporate structure that lets a company weigh its stated mission alongside shareholder profit.

Run rate — Recent revenue extrapolated to a full year. The figure labs quote; not the same as audited annual revenue.

S-1 — The document a company files with the SEC before going public. Anthropic filed one confidentially on 1 June 2026.

Anthropic is the company. Here are the model, the agents and the ecosystem around it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/guides/claude

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Claude?

The AI model inside OpenClaw, Claude Code, Cowork, Microsoft's Copilot Cowork and the agents running thousands of AI companies. What it is, which version you actually want, and why so much of the agentic economy chose it.

The one-liner

Claude is Anthropic's family of AI models — large language models trained to be helpful, honest and harmless, with a particular focus on behaving reliably when running as an agent for hours without supervision.

It is the brain inside Claude Code and Cowork, the default model in OpenClaw, the engine under Microsoft's Copilot Cowork, and the model Polsia's founder picked to run the agents that run his companies. If you use anything in this series, there is a good chance Claude is doing the thinking.

By the numbers — September 2026

Model names and scores in this box change every few months. Everything below is written to outlast them.

Fable 5.1

the most capable Claude available to the public, released 1 September 2026

4

tiers on general sale: Fable 5.1, Opus 5, Sonnet 5, Haiku 4.5

1M tokens

context window on Fable — roughly 750,000 words, or several novels, in one session

128K tokens

maximum output in a single response

87.6%

SWE-bench Verified score reported for Claude Code Remote in spring 2026 — the highest published at the time

~150

organisations with access to the restricted Mythos version through Project Glasswing

\$20 / \$100–200

per month for Claude Pro and Max. Free tier available

The family — which Claude is which

Claude is not one model. It is a range at different price and capability points, and in 2026 a new tier appeared at the top.

Tier	What it is for
Fable	The Mythos-class flagship, a tier above Opus. Long autonomous tasks, hard reasoning, large codebases, document-heavy professional work. The most expensive.
Opus	The previous top tier, still very capable. Complex agentic work where Fable is more than you need.
Sonnet	The workhorse. Fast, strong, far cheaper. Most business agents run on Sonnet.
Haiku	Fastest and cheapest. High-volume, simpler tasks: classification, routing, quick replies.
Mythos	Not for sale. The same model as Fable with safety classifiers lifted, restricted to vetted cyberdefenders and researchers.

The version numbers move — 4.6 in March, 4.7 in April, 4.8 in May, 5 in July, Fable 5.1 in September. The tiers do not. When someone says "Claude," ask which tier.

Fable and Mythos — the split that changed how models ship

In June 2026 Anthropic released a model it said was capable enough to be dangerous in the wrong hands: strong enough to find and chain security vulnerabilities in software autonomously, a task previously limited to elite human researchers.

Rather than release it or withhold it, Anthropic did both. Claude Fable is that model with real-time classifiers that block cybersecurity, biology and chemistry misuse; it is generally available. Claude Mythos is the identical model with those safeguards lifted, available only through Project Glasswing, a vetted programme run with the US government for cyberdefenders and critical-infrastructure providers, and, over time, biomedical researchers.

For you as a user the practical meaning is simple. Fable 5.1 is the best Claude you can buy, and it will decline a narrow set of dangerous requests. For the industry the meaning is larger: the most capable AI is now gated by who you are and what you are for. Other labs followed within months.

What Claude can actually do

Write and edit — Any style, length or format. Emails, code, contracts, marketing copy, technical documentation. Claude adapts to context and holds a consistent voice across long pieces.

Reason through hard problems — Multi-step logic, analysis, strategy, maths. Fable and Opus think before answering, and the thinking scales with the difficulty of the task.

Read enormous documents — A million-token context means an entire codebase, a full legal file or a year of emails in one session.

Write, run and fix code — Across every major language. Claude Code turns this into an agent that reads your project, changes it, tests it and commits.

Use tools — Browse the web, read files, query databases, call APIs, send email, through **MCP**, the open standard Anthropic invented.

Work in your apps — Cowork on the desktop, Claude in Chrome and a built-in browser, Claude in Excel, PowerPoint, Word and Outlook, Claude Tag in Slack.

Run for hours — The property that matters most for agents. Claude keeps its footing across long autonomous sessions, remembers what it has done, and asks when it is unsure rather than guessing.

What makes Claude different — the behaviour, not the benchmark

Benchmarks are a snapshot and every lab tops one somewhere. The reason so much of the agentic economy standardised on Claude is behavioural, and it comes from how Anthropic trains.

Constitutional AI gives the model written principles and trains it to judge its own outputs against them. In practice that produces a model that:

- Declines harmful requests and explains why, instead of refusing bluntly or complying blindly
- Flags ambiguity and asks, rather than silently assuming and getting it wrong three steps later
- Says "I don't know" more readily than rivals, which matters when nobody is checking
- Behaves consistently across different phrasings of the same instruction
- Notices flaws in its own work — Anthropic reported Opus 4.8 was several times less likely than its predecessor to let a bug it wrote pass unremarked

For a chatbot these are nice qualities. For an agent running your inbox at 3am they are the whole product. A model that hallucinates confidently sends the wrong email. A model that flags uncertainty leaves you a note.

Where Claude shows up

OpenClaw — Default model in the most-used open-source personal agent, with millions of active users.

Microsoft Copilot Cowork — Microsoft's flagship agent for Microsoft 365 runs on Claude, despite Microsoft's \$13 billion stake in OpenAI. Microsoft chose it anyway.

Polsia — The founder has said on the record he chose Claude Opus to run his agents because "it's important to give it the best reasoning."

Paperclip — Claude Code is the primary developer agent in the open-source multi-agent org chart.

Every major cloud — Amazon Bedrock, Google Cloud Vertex AI and Microsoft Foundry all serve Claude, so enterprises can use it inside infrastructure they already pay for.

Thousands of products — Via the API: customer-service agents, research tools, legal analysis, coding assistants, and the AI-run companies indexed on wtfagents.com.

How to get it

Claude.ai and the apps — Web, desktop (Mac and Windows), iPhone and Android. Free tier for chat. Claude Pro at \$20 a month adds Cowork, Claude Code and higher limits. Claude Max at \$100–200 for heavy use. Team and Enterprise plans above that.

Claude Code — Included with Pro and above, or pay-as-you-go through the API.

The Claude Developer Platform — Direct API access for builders. Pay per token; Haiku cheapest, Fable most expensive.

Through your cloud — Bedrock, Vertex AI or Foundry if your company already lives there.

Inside other agents — OpenClaw, Paperclip and most agent frameworks let you plug in your own Anthropic API key.

Glossary

Claude — Anthropic's AI model family: Fable, Opus, Sonnet and Haiku on general sale; Mythos restricted.

Fable — The public version of the Mythos-class model, with safety classifiers active. Currently Fable 5.1.

Mythos — The same model with safeguards lifted, available only to vetted organisations through Project Glasswing.

Tier — A capability and price level within the family. Tiers are stable; version numbers change constantly.

Constitutional AI — Anthropic's training technique: written principles the model learns to judge its own outputs against.

Context window — How much text the model can hold at once. One million tokens on Fable.

Token — The unit models process text in. Roughly three quarters of a word. API usage is billed in tokens.

MCP (Model Context Protocol) — The open standard Claude and every other major model use to connect to tools.

SWE-bench — The standard benchmark for AI coding ability, measured on real software engineering tasks.

Claude Code — The agentic version of Claude for building software. See its own guide.

Cowork — Claude working in your files, apps and browser without a command line. See its own guide.

Claude is the model. Here is everything built on it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/guides/anthropic

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Claude Code?

Not a code editor. Not an autocomplete plugin. An autonomous AI developer that reads your entire codebase, writes code, runs tests, fixes bugs, and commits to Git — while you do something else. Now with a desktop app, a phone app, and a growing crowd of non-developers building real things with it.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Claude Code is Anthropic's autonomous coding agent — a command-line tool that takes a task, reads your codebase, figures out what needs to change, makes the changes, runs tests, and iterates until the task is done.

It is not a tool that helps you code. It is a tool that codes while you watch — or while you do something else entirely.

The difference that matters

There are three categories of AI coding tool. They are not the same thing:

Autocomplete tools (GitHub Copilot)

Suggest the next line of code as you type. You are still writing. The AI is making suggestions. You accept or reject each one. Fast, low-friction, but fundamentally reactive.

AI-assisted editors (Cursor)

Chat with the AI inside your editor. Ask it to write a function, explain some code, or make a change. More powerful than autocomplete, but you are still directing every step.

Autonomous coding agents (Claude Code)

You give the agent a task: "Add user authentication to this application." It reads the entire codebase, plans the changes, writes the code across multiple files, runs the tests, fixes the failures, and commits the result. You were not involved in the individual steps.

The distinction matters because the productivity multiplier is completely different. Autocomplete saves you seconds. An autonomous agent saves you hours.

What Claude Code can actually do

Read entire codebases — Claude Code reads all your files — not just the one you have open. It understands how parts of the codebase relate to each other. This is critical for tasks that require changes across multiple files.

Write and edit code — Across all major languages: Python, JavaScript, TypeScript, Go, Rust, Java, C++, Ruby, and more. It writes idiomatic code that fits the existing style of your codebase.

Run shell commands — Execute tests, linters, build tools, and scripts. Read the output. Understand what failed. Try to fix it.

Run tests and fix failures — Write tests if you don't have them. Run existing tests. Read the failures. Fix the code. Re-run. Iterate.

Commit to Git — Stage changes, write commit messages, and commit — following your project's conventions.

Handle multi-file refactors — Rename a function used in 47 places. Change an API interface and update all callers. Restructure a module. Tasks that would take a human developer hours.

Work on long autonomous sessions — Unlike autocomplete tools that help one line at a time, Claude Code can work on a complex task for hours — maintaining context across the entire session.

By the numbers — September 2026

Claude Code went from research preview in early 2025 to the fastest-adopted developer tool of its generation. The figures move monthly; the shape does not.

\$2.5B

annualised revenue run rate by February 2026 — before the growth that took Anthropic past \$65B overall

29M

installs inside Visual Studio Code alone by early 2026

~4%

of all public GitHub commits worldwide authored by Claude Code, per one early-2026 analysis

87.6%

SWE-bench Verified reported for Claude Code Remote in spring 2026, the highest published score at the time

April

2026: redesigned desktop app for Mac and Windows — panels, buttons, a chat box, no terminal required

Aug

2026: start a Claude Code session on your computer directly from the Claude phone app

The satisfaction surveys that first made the point still hold: developers who try autonomous agents do not go back to autocomplete. GitHub Copilot has the larger install base and Microsoft behind it. Claude Code overtook it in developer preference in under a year because

it does something fundamentally different.

How Claude Code fits into the agentic economy

Claude Code is not just a tool for individual developers. It is one of the primary building blocks of the agentic economy.

In Paperclip

Paperclip — the open-source multi-agent company orchestration framework — lists Claude Code as a primary supported agent. In a Paperclip "company org chart," Claude Code can be the developer agent: building and maintaining the company's software autonomously while other agents handle marketing, customer support, and operations.

In autonomous companies

The companies tracked on WTF Agents — the tens of thousands of active AI-run companies on Polsia, Paperclip and elsewhere — are using agents like Claude Code to build and iterate on their products. A company with no human employees still needs software. Claude Code builds it.

In enterprise settings

Large enterprises are using Claude Code for large-scale code migrations, test-writing, technical debt reduction, and API integrations. Tasks that previously required weeks of developer time can be completed in hours.

Claude Code vs the alternatives

Tool	Type	Best for	Limitation
Claude Code	Autonomous agent	Complex, multi-file tasks; long autonomous sessions	Usage limits on Pro; costs scale with the task
Cursor	AI-assisted editor	Fast iteration; inline editing; real-time help	Now owned by SpaceX/xAI; losing OpenAI models in November 2026
GitHub Copilot	Autocomplete	Quick suggestions while typing; broad ecosystem	Reactive, not autonomous
Codex (OpenAI)	Autonomous agent	Lives inside the ChatGPT app; GPT-6 backbone	Smaller ecosystem of plugins and skills
Devin (Cognition)	Autonomous agent	Fully cloud-hosted; \$20/mo, \$200 Max	Less control over your own machine
Jules (Google)	Autonomous agent	Free; runs in the cloud on Gemini	Narrower task scope

Do you need to be a developer to use Claude Code?

Not any more — though it helps to think like one.

For its first year Claude Code was a command-line tool, and the screenshots looked like a hacker film. In April 2026 Anthropic shipped a redesigned desktop app for Mac and Windows: panels, buttons, a chat box. You type what you want in plain English and watch it work. In August, the Claude phone app gained the ability to start and steer a session on your computer from anywhere. It is included in Claude Pro at \$20 a month, so the API key requirement is gone too.

Anthropic's own analysis of over a million Cowork sessions found that more than 90% were not software development. Claude Code's audience has shifted the same way: founders building their first product, operators automating a workflow, marketers building internal tools, and — the pattern this series exists to document — people running AI companies who use Claude Code as the developer agent inside Popsia or Paperclip.

What still matters is not the terminal but the mindset. Claude Code works best for people who can describe what they want precisely, check what came back, and understand the difference between "it runs" and "it is right." That is a developer's habit, not a developer's skill, and it can be learned in a weekend.

What this means for software development

Claude Code is part of a broader shift in what software development means.

For decades, writing software required deep technical skill, years of learning, and constant attention to detail. The bottleneck was human developer time.

Autonomous coding agents change the equation. The bottleneck shifts from "can we write the code?" to "do we know what to build?" The strategic and product decisions matter more. The implementation matters less.

This does not mean software developers are going away. It means the nature of the job is changing. Senior developers who can direct, review, and architect are more valuable. Junior developers who were doing rote implementation work face the most disruption.

Glossary

Claude Code — Anthropic's autonomous coding agent. CLI tool that reads codebases, writes code, runs tests, and commits to Git without step-by-step human instruction.

Autonomous coding agent — An AI that takes a coding task and completes it end-to-end without human direction of individual steps. Contrast with autocomplete (suggests next line) or AI-assisted editors (human-directed).

CLI (Command Line Interface) — A text-based way of interacting with a computer. Claude Code runs in the terminal — you type commands rather than clicking a GUI.

Git — The standard version control system for software projects. Claude Code can stage, commit, and manage Git operations autonomously.

SWE-bench — The standard benchmark for AI coding ability. Measures how well an agent can solve real software engineering tasks. Scores change every release; Claude Code has led it for most of 2026.

GitHub Copilot — Microsoft/GitHub's AI coding tool. 77M+ developers. Autocomplete-style — helps as you type. Claude Code is a different category: autonomous.

Cursor — An AI-assisted code editor. More powerful than Copilot, but still human-directed. Claude Code is autonomous; Cursor is collaborative.

Paperclip — The multi-agent company orchestration framework that uses Claude Code as a primary developer agent.

API key — Your credential for accessing the Anthropic API. Required to run Claude Code. You pay per token of usage.

Refactor — Restructuring existing code without changing its external behaviour. A common use case for Claude Code — especially large-scale refactors across many files.

Claude Code is the agent. Here is the ecosystem around it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/guides/claude

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/guides/anthropic

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Cowork?

Claude Code was for developers. Cowork is for everyone else. The Anthropic product that rattled enterprise software stocks, became Microsoft's flagship agent, and now runs in the cloud, on your phone, and in its own browser. What it is and why it matters.

The one-liner

Cowork is Anthropic's desktop AI agent for non-developers — a tool that lets Claude read, edit, create, and organise files directly on your computer, execute multi-step tasks autonomously, and connect to your existing tools, all without touching a command line.

Anthropic's own framing says it best: "Chat changed how we get answers. Claude Code transformed how developers build software. Now, Cowork brings that same execution power to everyone."

By the numbers — September 2026

12 Jan

2026: launched as a research preview. Generally available on Mac and Windows since April, web and phone since July

90%+

of Cowork sessions are not software development, per Anthropic's own analysis of 1.2 million sessions

26 Aug

2026: Cowork gets its own built-in browser. Claude in Chrome goes generally available the same week

\$20

per month — included in Claude Pro and every paid plan above it. Not on the free tier

11

official Anthropic plugins for sales, finance, legal, marketing and more, plus a marketplace of third-party ones

June

2026: Microsoft's Copilot Cowork, built on Claude, goes live worldwide in Microsoft 365

The evolution — Chat → Code → Cowork

To understand Cowork you need to understand the three-stage evolution of Claude:

Stage	Product	Who it's for	What it does
1	Claude Chat	Everyone	Answers questions. You ask, it responds. Reactive, not autonomous.
2	Claude Code	Developers	Autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git. Requires command line.
3	Cowork	Everyone	Autonomous desktop agent. Reads and edits files, executes multi-step tasks, connects to apps. No command line needed.

The pattern Anthropic spotted: when they released Claude Code, they expected developers to use it for coding. They did — and then immediately started using it for everything else. Organising files, writing reports, processing data, automating workflows. Cowork is the formalisation of that behaviour, made accessible without technical prerequisites.

How Cowork actually works

Cowork lives inside the Claude Desktop app for Mac and Windows, and since July 2026 also on claude.ai in a browser and in the Claude phone app. You access it by clicking the Cowork tab, next to Chat. Sessions run against your Claude account, so a task started on the desktop keeps running after you close the laptop, and you can check on it from your phone.

Step 1: Grant folder access

You point Cowork at a folder on your computer. This is the boundary — Claude can only read and modify files within that folder. Start with something low-stakes: a Downloads folder, a project folder, a research directory. You control exactly what Claude can and cannot touch.

Step 2: Describe the task

In plain English. No special syntax, no commands. Just describe what you want done. Examples that work right now:

- "Organise this folder. Sort files by type, rename generic ones based on content, flag duplicates."
- "Pull the expense figures from all these receipt screenshots and create a summary spreadsheet."
- "Read all these research notes and write a 2-page briefing document."
- "Find all the client emails in this folder and extract the action items into a to-do list."
- "Take this raw data CSV and generate a formatted report with charts."

Step 3: Claude plans and executes

Cowork makes a plan, shows it to you, and then executes it step by step. You can see what it is doing in the Progress panel. You can pause or redirect mid-task if needed. For significant actions — deleting files, sending emails — Cowork asks for confirmation first.

Step 4: Results land in your folder

No copy-pasting from a chat window. No downloading outputs. The finished files are in your folder, exactly where you need them.

What makes Cowork different from just using Claude Chat

This is the question most people ask. The difference is significant:

	Claude Chat	Cowork
File access	You paste content into the chat	Claude reads files directly from your computer
Output	Text in a chat window — you copy it	Files created directly in your folder
Multi-step tasks	One response at a time — you direct each step	Claude plans and executes all steps autonomously
Memory	Forgets between conversations	Projects keep context, files, and instructions persistent
App connections	Limited	Connects to Gmail, Google Drive, Notion, Slack, and hundreds more via connectors
Works while you're away	No	Yes — scheduled recurring tasks run automatically

The features — what Cowork can do

File operations

Read, edit, create, rename, organise, and delete files in your designated folder. Process any file type Claude can understand: documents, spreadsheets, PDFs, images, CSVs, emails.

Projects

A persistent workspace tied to a folder. Your files, instructions, and task history stay in one place across sessions. Instead of starting fresh every time, your project remembers what it has done and what is next.

Scheduled tasks

Recurring tasks that run automatically. "Every Monday morning, pull last week's sales data from this folder and generate a summary report." Since the move to cloud sessions, scheduled tasks fire even with no device online.

Connectors

Link Cowork to external services: Gmail and Google Workspace, Microsoft 365 (Outlook, OneDrive, SharePoint, Teams), Slack, Jira, Notion, DocuSign, and hundreds more via the connector marketplace. Desktop extensions have deeper system access than web connectors.

Skills

Install skills — modular capabilities — that improve Cowork's output for specific formats: documents, presentations, spreadsheets. You can also write a SKILL.md file with your brand voice or standard process and Cowork uses it every time.

The built-in browser

Since August 2026 Cowork has its own Chromium browser in a side panel of the desktop app. When a task involves a website, the browser opens, Claude navigates, fills forms and finishes the job — separate from your own browser and logins. If you would rather it work inside your own signed-in Chrome, Claude in Chrome does that and is now available on every paid plan; a safety classifier checks each action before it is taken.

Computer use

Cowork can see your screen and drive any application on your machine, not just the browser. Useful for tools that have no connector: a niche internal system, a desktop app, anything with buttons.

Plugins

Role-specific bundles of skills, prompts and connectors. Anthropic publishes its own — sales, finance, legal, marketing, and a Claude for Financial Advisors plugin covering onboarding, meeting prep and portfolio reviews — and a marketplace carries third-party ones. Plugins run on whatever plan you already have; nothing extra to buy from Anthropic.

Dispatch

Control a desktop Cowork session from your phone: approve a step, redirect, or read the result on the move.

The market reaction — why software stocks fell

When Cowork launched in January 2026, and again when its plugins landed, enterprise software stocks fell sharply. The iShares Expanded Tech-Software Sector ETF dropped nearly 5% in a single day, and analysts put the combined loss across enterprise software companies in the hundreds of billions of dollars.

The reason: investors looked at what Cowork can do and compared it to what they were paying enterprise software companies to do. File management. Document creation. Data extraction. Report generation. Workflow automation. These are the core functions of dozens of \$10–50/user/month SaaS products.

A knowledge worker with Cowork can now automate tasks that previously required separate tools for project management, document processing, data analysis, and workflow automation — all from one desktop agent.

The stocks recovered somewhat when Microsoft announced it was building Copilot Cowork on top of Anthropic's technology — suggesting the incumbents would adapt rather than be replaced entirely. But the signal was clear: the knowledge work software market has been permanently disrupted.

Microsoft Copilot Cowork — the enterprise version

Microsoft announced Copilot Cowork in March 2026 and made it generally available worldwide in June: a cloud-based agent, powered by Anthropic's Claude, that executes multi-step tasks across Microsoft 365.

The difference from consumer Cowork: Copilot Cowork runs inside Microsoft 365's own infrastructure with access to the full graph of enterprise work data — Outlook, Teams, calendars, SharePoint, Excel. Rather than working on a local folder, it works across an entire organisation's data.

It ships as a \$30 per user per month add-on, or inside the Microsoft 365 E7 bundle.

The significance: despite a \$13 billion investment in OpenAI, Microsoft built its flagship Microsoft 365 agent on Anthropic's Claude — and then took it to general availability. Not because Microsoft lacks access to GPT, but because it chose Claude anyway.

Real use cases — what people are actually doing with it

Finance and accounting

Point Cowork at a folder of receipt screenshots. It extracts all expense data, categorises by type, and generates a formatted expense report. Tasks that took a finance assistant an hour now take Cowork five minutes.

Research and analysis

Feed Cowork a folder of research notes, articles, and PDFs. Ask it to synthesise the key findings into a briefing document. It reads everything, identifies themes, and writes the document — citing sources from the files it read.

Content operations

Give Cowork a brand voice SKILL.md file and a folder of raw content briefs. Ask it to draft all ten blog posts. It writes in your brand voice, formats correctly, and saves them as Word documents ready for review.

Data processing

Drop a messy CSV export into the folder. Ask Cowork to clean it, standardise the formatting, remove duplicates, and generate a summary report with charts. No spreadsheet skills required.

File organisation

The classic first use case. Give Cowork your Downloads folder. It sorts by type, renames files based on their actual content, flags duplicates, and creates logical subfolders. Most users hit their first "wow" moment within ten minutes of trying this.

Scheduled intelligence

Connect Cowork to your email via the Gmail connector. Set a recurring task: every morning at 8am, read all emails received overnight, extract action items, and create a prioritised to-do list in your project folder. You wake up to a briefing that was prepared while you slept.

How to get access

Plan	Access	Price
Claude Pro	Full Cowork access	\$20/month
Claude Max	Full Cowork access + 5x or 20x the limits	\$100 or \$200/month
Claude Team	Full Cowork access + admin controls	\$25–30/user/month
Claude Enterprise	Full Cowork + admin controls, audit logs, smart reports	Custom pricing
Microsoft Copilot Cowork	Cloud-based, Microsoft 365 integrated	\$30/user/month add-on, or in the E7 bundle
Free tier	Not available — Cowork is paid only	—

To get started: subscribe to Claude Pro or higher, then open Cowork in the desktop app (claude.com/download), on claude.ai, or in the Claude phone app.

The honest limitations

Cowork is genuinely powerful. It also has real limitations worth knowing:

- **Folder boundary only.** Claude can only access the folder you designate. It cannot reach outside it without explicit permission — and that is a feature. In one widely reported case about 15,000 family photos were deleted while Cowork tidied a desktop, and were recovered only through a 30-day cloud restore.
- **It burns through usage.** Agent tasks use far more of your plan's allowance than chat does. Pro is the entry point; heavy daily use tends to need Max.
- **Hallucination risk remains.** Cowork can misread file contents or make incorrect assumptions about what a task requires. Always review outputs before acting on them.
- **Security considerations.** Days after launch, a data exfiltration vulnerability was reported and patched. Prompt injection — hidden instructions in a web page or document the agent reads — is the unsolved problem for every agent with browser access. Be thoughtful about connectors and folders.
- **The desktop browser needs the desktop.** From the web or your phone, Claude can drive the built-in browser only while your desktop app is open. Without it, Claude in Chrome is the browser route.
- **It keeps changing.** Cowork has shipped a major capability most months of 2026. Expect the feature list in this guide to be out of date before the next one.

What Cowork means for the agentic economy

Claude Code proved agents could replace developer time. Cowork is the thesis applied to every knowledge worker.

The administrative layer of every business — the file management, the data processing, the report generation, the inbox management, the scheduling — is now automatable by a non-technical person with a \$20/month subscription.

The barrier to running a lean, agent-powered business has never been lower. A solo founder with Cowork, Claude Code, and OpenClaw has the operational capacity of a team of five. A small business owner with Cowork has an always-on assistant that never sleeps, never forgets, and costs less than a gym membership.

The HBR research published in early 2026 noted that companies are already making headcount decisions based on AI's potential, not just its demonstrated performance. Cowork is part of why. It is not a future prediction — it is a current product that current businesses are using to do real work right now.

Glossary

Cowork — Anthropic's autonomous desktop AI agent. Part of the Claude Desktop app. Reads, edits, and creates files on your computer and executes multi-step tasks without command-line skills.

Claude Desktop — The desktop application from Anthropic for Mac and Windows. Contains Chat, Code, and Cowork modes, and Cowork's built-in browser. Download at claude.com/download.

Projects (Cowork) — A persistent workspace tied to a folder — keeps files, instructions, and task history in one place across sessions.

Scheduled tasks — Recurring tasks that Cowork runs automatically on a set schedule — daily, weekly, or custom timing.

Connectors — Integrations that link Cowork to external services — Gmail, Google Drive, DocuSign, Slack, Notion, and hundreds more.

Skills — Modular capabilities installed into Cowork. Including custom SKILL.md files you create to encode your brand voice or standard processes.

Claude in Chrome — Anthropic's agent inside your own Chrome browser, generally available on all paid plans. The alternative to Cowork's built-in browser when you want it working where you are already signed in.

Copilot Cowork — Microsoft's enterprise version of Cowork — cloud-based, running across Microsoft 365. Powered by Anthropic's Claude. \$30/user/month. Generally available since June 2026.

Claude Code — The developer-facing autonomous coding agent that Cowork is built on. Requires command-line skills. Cowork is the non-technical equivalent.

Knowledge worker — Someone whose job primarily involves creating, processing, or managing information — as opposed to physical labour. Cowork is specifically designed for this category.

Cowork is built on Claude and sits alongside Claude Code. Here is the full picture.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/guides/anthropic

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/guides/claude

How to Hire an AI Agent for Your Business

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

Before you start — the right mindset

Deploying AI agents in your business is not a technology project. It is a staffing decision.

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

Step 1 — Find your highest-value repeatable task

The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

Customer support responses — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

Content creation — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

Research and summarisation — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

Data entry and processing — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

Email drafting — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

Scheduling and coordination — Managing calendar bookings, sending reminders, coordinating meeting times.

Lead qualification — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

Step 2 — Choose the right type of agent

Not all agents are right for all tasks. In 2026 the choice is simpler than it looks, because every option falls into a handful of shapes:

If you want to...	Use this	Cost
Delegate work inside your own files, inbox, calendar and apps	Cowork (Anthropic), ChatGPT Work (OpenAI), Copilot Cowork (Microsoft 365)	From \$20/mo, included in the paid plan
Run a team of agents sharing one cloud computer and its logins	Grok Bot (xAI)	Subscription
Have a personal agent you message like a colleague	OpenClaw (open-source), Meta Muse	Free + model costs; Muse has a free tier
Let agents build and run a whole company for you	Polsia, NanoCorp, Cofounder.co	From \$20/mo + fees (Polsia)
Run several agents as an org chart you control	Paperclip (self-hosted)	Free + model costs
Connect apps and automate defined workflows	Zapier, n8n	From ~\$20/mo; n8n is open-source
Build a custom agent with full control	Claude Developer Platform / Agent SDK	Pay per token

For a first agent, the top row is almost always the answer. A work agent on a plan you already have, pointed at one folder and one task, is the lowest-risk way to learn what this is actually like.

Step 3 — Write a clear job description for your agent

This is the most important step most people skip.

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

- **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")
- **The inputs** — what information does the agent receive? ("An inbound email from a customer")
- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

Step 4 — Start with read-only access

The most common mistake when deploying agents: giving them too much access too soon.

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

Step 5 — Measure results honestly

Define what success looks like before you deploy — not after.

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved

- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

Five deployments that work — and what to expect

These are the patterns that come up again and again among small businesses that get value from a first agent. Results vary; the shape of the win does not.

First-line customer support

An agent reads incoming support emails, categorises them, drafts replies from approved templates, and flags anything unusual for a human. The win is response time — hours become minutes — and a majority of routine emails handled without anyone touching them. The risk is a confident wrong answer, which is why replies are drafted for approval before they are sent unsupervised.

Research and content preparation

An agent researches a client's industry each week and drafts a content calendar, a briefing or a first draft. Writers spend their time on strategy and editing instead of gathering. The win is capacity without headcount.

Inbox triage for partners and owners

An agent summarises the day's email, drafts responses for review, and surfaces what needs a decision. Owners consistently report an hour or two a day back. The risk is the agent guessing at tone on a sensitive thread; keep it drafting, not sending, for those.

The solo operator's assistant

A personal agent — OpenClaw over WhatsApp or Telegram, or Cowork on the desktop — briefed with research tasks, email drafts, document summaries and admin. Typical cost is tens of dollars a month in model usage. This is the deployment most one-person businesses start with.

Routine engineering

Claude Code or Codex handling bug fixes, small features and tests, with senior people on architecture and review. Teams that do this well ship considerably more with the same headcount. Teams that skip the review step ship considerably more bugs.

The honest cost — what agents actually cost

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

Agent type	Monthly cost (approx)	What drives the cost
Cowork / ChatGPT Work	\$20–200/mo	Your Claude or ChatGPT plan; heavy use needs the higher tier
Copilot Cowork	\$30/user/mo	Add-on to Microsoft 365
OpenClaw (personal)	\$20–100/mo	Model API usage — Claude, GPT or a cheaper open model
Polsia company	\$20–1,000/mo + fees	Plan tier, credits, 20% on ad spend, 3% on revenue
Paperclip org	\$0 + model usage	Free software; each agent's model bill
Zapier / n8n automation	\$0–100/mo	Task volume and plan tier
Custom Claude agent	\$50–500/mo	Token usage — depends heavily on task volume
Claude Code	\$20–200/mo	Included in Claude Pro/Max; API if you go beyond

For context: a part-time human assistant in the US costs \$2,000–3,500/month. A full-time junior employee costs \$45,000–60,000/year. An agent handling equivalent work costs \$50–500/month. The economics are not subtle — and they are the reason the Agentic Economy guide spends a section on what that means for jobs.

What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

Your action plan — what to do this week

Day 1 — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

Day 2 — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

Day 3 — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

Day 4-5 — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

Week 2 — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

Week 3-4 — If quality is consistently good, expand permissions incrementally. Track the time saved.

Month 2 — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

The businesses that win with AI agents are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.

Glossary

AI Agent — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

Human-in-the-loop — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

Polsia — A managed platform where AI agents build and run companies. From \$20/mo plus credits and fees; see its own guide.

OpenClaw — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

Paperclip — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

Cowork — Anthropic's work agent: delegates tasks across your files, apps and browser. Included in every paid Claude plan.

ChatGPT Work — OpenAI's work agent: give it an outcome, it returns finished spreadsheets, decks and apps.

Zapier / n8n — Automation platforms that connect apps and run defined workflows, now with agent steps. n8n is open-source and self-hostable.

Claude — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

API key — Your credential for accessing an AI model's API. Required for OpenClaw, Paperclip and custom builds; Claude Code and Cowork run on your Claude plan instead.

Token — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

MCP (Model Context Protocol) — The open standard connecting AI agents to external tools. Invented by Anthropic.

You now know how to hire an agent. Here is the full context behind the tools you are using.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

Launch an AI-Run Company in a Weekend

The practical playbook: validate an idea on Saturday, launch it on Polsia, NanoCorp or Paperclip on Sunday, and know by Friday whether it sells.

wtfagents.com/guides/launch-weekend

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

wtfagents.com/guides/solo-stack

Who's Who in the Agentic Economy

Forty-odd names you will keep seeing, sorted into the seven layers of the agentic economy, one line each. The map to read before any of the deep dives — and the one to come back to when a new name appears in your feed.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

How to read this guide

The agentic economy is a stack. Models at the bottom, agents built on them, platforms that run whole companies of agents, and the plumbing, chips and money underneath it all. Every name in the news belongs to one layer, and knowing which layer tells you most of what you need to know.

This guide gives each name one line: what it is, why it matters. Where a full WTF Agents guide exists, it is named so you can go deeper. Version numbers, valuations and star counts live in the dated box below and nowhere else, because they change monthly and the map does not.

By the numbers — September 2026

\$965B

Anthropic's valuation after May's \$65B round — the most valuable private company in the world, IPO filed

\$852B

OpenAI's valuation after its \$122B round in March

\$60B

what SpaceX paid for Cursor — agreed in June, closed in August — after buying xAI in February

\$30M

Polsia's May raise at \$250M; NanoCorp is YC-backed; Paperclip has no disclosed funding

3.2M

active OpenClaw users; Hermes Agent is the credible number two

490,000

companies created on Polsia; about 6% still active

6

labs that shipped a "does your job" agent in 2026: Anthropic, Microsoft, OpenAI, xAI, Meta — and Google, by consolidation

2

open standards under the Agentic AI Foundation: MCP and A2A

Layer 1 — The labs and their models

The brains. Everything above runs on one of these.

Anthropic — Claude — Safety-founded lab, now the most valuable private company on earth. Claude powers Claude Code, Cowork, Copilot Cowork, OpenClaw's default and Polsia's agents. Tiers: Fable, Opus, Sonnet, Haiku; Mythos is the restricted twin. Guides: [WTF is Anthropic](#), [WTF is Claude](#).

OpenAI — GPT — ChatGPT is still the product most people have used. GPT-6 arrived in September. OpenAI's agents are ChatGPT Work, Codex and the Agents API; its reach is its weapon.

Google DeepMind — Gemini — Gemini 3.8, Gemini Agent, Chrome Auto Browse and Jules. Owns Search, Android and Workspace, and invented the A2A protocol. Chose to consolidate its agents rather than launch a new one.

Meta — Muse — A new model family and a personal agent of the same name, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is that three billion messaging users are the distribution. Also tried to buy Manus.

xAI — Grok — Elon Musk's lab, now inside SpaceX alongside Cursor. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff."

Microsoft — Not a frontier lab of its own but the biggest distributor: Copilot Cowork, built on Claude, is live in Microsoft 365 worldwide; Copilot Studio hosts hundreds of thousands of custom agents; Foundry serves every major model.

DeepSeek — The Chinese lab whose cheap-to-train breakthrough in early 2025 made this a two-continent race. DeepSeek V4 is open-weight under the MIT licence and runs under a great many indie agents.

Alibaba — Qwen · Moonshot — Kimi · Zhipu — GLM · MiniMax — The rest of China's open-weight frontier. Kimi K3 was the largest open model ever released. Cheap, capable, downloadable. Guide: [WTF is Happening with China's AI](#).

Mistral — France's lab. Efficient smaller models and Europe's answer to US and Chinese dependence.

Layer 2 — Work agents

Agents that do your job's tasks. The category every lab entered in 2026.

Cowork (Anthropic) — Claude working in your files, apps and its own browser, on desktop, web and phone. Included in every paid Claude plan. Guide: [WTF is Cowork](#).

ChatGPT Work (OpenAI) — Give it an outcome; it returns finished spreadsheets, decks and web apps hours later. Launched July.

Copilot Cowork (Microsoft) — Cowork inside Microsoft 365, across Outlook, Teams, SharePoint and Excel. Built on Claude despite Microsoft's stake in OpenAI. Generally available since June.

Grok Bot (xAI) — A team of named agents sharing one cloud computer, working while you sleep. Launched August through Cursor's infrastructure.

Microsoft Scout — An always-on personal agent inside Copilot, in limited release.

Gemini Agent (Google) — Where Google's browser agent went after Project Mariner was folded in. Plus Auto Browse in Chrome.

Layer 3 — Coding agents

The category that proved agents work, and the one that made Anthropic.

Claude Code (Anthropic) — The autonomous developer. Terminal, desktop app, phone. The fastest-adopted developer tool of its generation. Guide: [WTF is Claude Code](#).

Codex (OpenAI) — OpenAI's coding agent, now living inside the ChatGPT desktop app on GPT-6.

Devin (Cognition) — The original "AI software engineer," fully cloud-hosted, cut from \$500 to \$20 a month in April.

Jules (Google) — Google's free cloud coding agent on Gemini.

Cursor — The AI-first code editor, bought by SpaceX for \$60 billion — agreed in June, closed in August — and now the delivery vehicle for Grok Bot. Losing OpenAI's models in November.

GitHub Copilot (Microsoft) — The largest install base in coding AI, evolving from autocomplete into an agent inside GitHub itself.

Layer 4 — Personal agents

Agents that run your life's admin, mostly over messaging apps.

OpenClaw — Open-source, self-hosted, millions of users, a non-profit foundation, native mobile apps. The reference personal agent and the one every rival is measured against. Guide: [WTF is OpenClaw](#).

Meta Muse — Meta's consumer agent: books, buys, negotiates and fills in forms, free tier, WhatsApp. US only at launch. Guide: [WTF is Meta Muse](#).

Hermes Agent (Nous Research) — The credible open-source number two: self-improving memory, hireable into Paperclip out of the box.

Manus — The Chinese general-purpose agent Meta agreed to buy for \$2 billion before Beijing blocked the deal. Still independent, still Chinese.

Layer 5 — Company platforms and orchestration

Agents that run whole businesses, and the tools that manage teams of agents.

Polsia — The biggest and most argued-about: subscribe, describe an idea, a fixed team of agents builds and runs the company. Nearly half a million created; a founder who says one in ten has earned a dollar. Guide: [WTF is Polsia](#).

NanoCorp — YC-backed rival that publishes the live revenue of every company on the platform — the most honest numbers in the space, and small ones: the best has earned \$2,570, and only two have passed \$1,000. Guide: [WTF is NanoCorp](#).

Cofounder.co — Company-in-a-box with departments and a human approval gate in front of anything risky.

Paperclip — Open-source org chart for agents: roles, budgets, approvals, any agent vendor, self-hosted. ClipMart is its marketplace of forkable company templates. Guide: [WTF is Paperclip](#).

Salesforce Agentforce — The enterprise incumbent's agent platform, priced per conversation, past a billion dollars in annual revenue.

ServiceNow AI Control Tower — Governed agents for IT and HR workflows; AI contract value past a billion dollars.

SAP Joule · Workday Illuminate — Agents inside the ERP and HR systems most large companies already run — and the first to charge outside agents a toll to reach their data.

Zapier · n8n · UiPath — The automation layer: connect apps, run defined workflows, now with agent steps. n8n is open-source; UiPath is the enterprise back-office specialist.

Layer 6 — The plumbing

Standards and infrastructure that make agents interoperate.

MCP — Model Context Protocol — The universal plug between agents and tools. Invented by Anthropic, adopted by everyone, now vendor-neutral. Guide: [WTF is an API](#).

A2A — Agent2Agent — Google's standard for agents from different companies talking to each other. In production at hundreds of organisations.

Agentic AI Foundation — The Linux Foundation body that owns MCP and A2A, founded by Anthropic, OpenAI, Google, Microsoft and Block. The rare case of rivals agreeing to share.

Stripe — The payment rails under Polaris companies, the WTF Agents store, and most of the agentic commerce experiments. Agents that can pay run through Stripe.

Cloudflare · Vercel — Where agents are hosted and deployed. Vercel's Eve is a framework for durable agents; Cloudflare's platform runs them at the edge.

Layer 7 — The picks and shovels

Who gets paid regardless of which agent wins.

Nvidia — The chips. Every model in Layer 1 was trained and is served on Nvidia hardware, which is why it is the most valuable company in the world. Buying Hugging Face for \$12.93 billion extends it from the silicon to the shelf the models sit on.

AWS · Google Cloud · Microsoft Azure — The clouds. Bedrock, Vertex AI and Foundry serve every major model to enterprises inside infrastructure they already pay for. Amazon and Google are also Anthropic's largest backers.

Hugging Face — The hub where open-weight models are published and downloaded. If DeepSeek or Qwen ships a model, this is where it lands. Nvidia agreed to buy it in September 2026 for \$12.93 billion, closing in the first half of 2027; it has promised the platform stays open and chip-neutral.

GitHub — Where the code lives and where agents commit it. OpenClaw and Paperclip's star counts are GitHub's.

Guide: [The Picks and Shovels](#) goes deeper on this layer.

The people

Six names that come up more than any others.

Dario Amodei — CEO of Anthropic. The essays that set the tone for the safety-first case.

Sam Altman — CEO of OpenAI. The face of ChatGPT and the reach strategy.

Elon Musk — SpaceX, xAI, Cursor, Grok Bot. Consolidating the agent stack under one company.

Mark Zuckerberg — Meta. Muse is his second attempt at an AI reset, and the first with a real agent.

Peter Steinberger — Created OpenClaw, joined OpenAI, left the project to a foundation.

Ben Broca (Ben Cera) — Solo founder of Polsia. Raised \$30 million with one employee and said out loud that most of his customers' companies earn nothing.

How the layers fit together

Value has moved down the stack. In 2024 the money was in models. In 2026 it is in agents. The bet behind Layer 5 is that next it will be in companies that no human works at.

Two forces cut across every layer. The first is Claude's spread: Anthropic's model sits under products from Microsoft, OpenClaw, Polsia and Paperclip, which is why one lab is close to a trillion dollars. The second is China's open weights: cheap, capable, downloadable models that let a solo builder run agents for pennies and give every US lab a price ceiling.

When a new name appears in your feed, ask which layer it belongs to. The answer usually tells you who it competes with, who it depends on, and whether it is a product or a press release.

Glossary

Layer — A level of the agentic economy stack: models, agents, platforms, plumbing, infrastructure. Every company in this guide belongs to one, sometimes two.

Frontier lab — A company training models at the leading edge of capability: Anthropic, OpenAI, Google DeepMind, Meta, xAI, DeepSeek.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

Personal agent — An agent that runs your life's admin: OpenClaw, Muse, Hermes.

Company platform — A service where agents build and run a whole business: Polsia, NanoCorp, Cofounder.co.

Orchestration — Managing a team of agents: Paperclip.

Open weights — A model anyone can download and run. Most Chinese frontier models; few US ones.

MCP / A2A — The two open standards that let agents use tools and talk to each other. Both under the Agentic AI Foundation.

Incumbent — An established enterprise software vendor adding agents to what it already sells: Salesforce, ServiceNow, SAP, Workday, Microsoft.

This is the map. Every name on it with a guide has one waiting.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/guides/which-ai

WTF is Happening with China's AI

DeepSeek, Qwen, Kimi, GLM and the open-weight strategy. Who the Chinese labs are, why they give models away, and what it means for anyone running agents.

wtfagents.com/guides/china-ai

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

The Picks and Shovels

The infrastructure layer: Nvidia, the clouds, Hugging Face, GitHub, Stripe, power. Who gets paid whichever agent wins, and why it matters to you.

wtfagents.com/guides/picks-and-shovels

Which AI Should I Actually Use?

Claude, ChatGPT, Gemini, Muse, Grok, DeepSeek. Six families, a new version every month, and a leaderboard that changes weekly. Here is how to choose without reading a single benchmark — and what each one is quietly bad at.

The one-liner

There is no best AI. There is a best AI for the thing you are doing today, and the differences between the top models are now smaller than the differences between how you use them.

That is the honest answer, and it is also the useful one. This guide gives you the decision in one table, the plans and prices, what the leaderboards actually measure, and the catches nobody puts on the pricing page.

By the numbers — September 2026

Everything here is a snapshot. The families below it are stable; the numbers are not.

6

model families that matter for a normal person: Claude, GPT, Gemini, Muse, Grok, DeepSeek

Fable 5.1

top of the independent Artificial Analysis Intelligence Index since 1 September, with Opus 5 close behind at half the price

GPT-6 Astra

OpenAI's newest, 3 September, after the GPT-5.6 Sol / Terra / Luna family in July

\$20

the standard monthly plan at every lab — ChatGPT Plus, Claude Pro, Muse, Google's AI plan

\$100–200

the power-user tier: ChatGPT Pro, Claude Max, Muse's top plan

\$0.20 vs \$50

per million output tokens, cheapest frontier-adjacent model (GPT-5.6 Luna) versus most expensive (Fable 5.1) — a 250× range

\$0

DeepSeek's chat app, no paid tier at all

The decision in one table

Start here. Pick the row that matches what you are doing.

You want to...	Use	Why
Ask questions, get quick answers, everyday chat	ChatGPT (free or Plus)	The largest user base and the most polished consumer app. Everyone you know is on it.
Write anything longer than an email	Claude	Consistently preferred for voice, structure and holding a style across long pieces.
Delegate real work in your files and apps	Claude (Cowork) or ChatGPT Work	The two work agents on a \$20 plan. See the Cowork guide.
Build software without being a developer	Claude Code	The desktop app changed who can use it. See its guide.
Run an agent for hours unsupervised	Claude Fable or Opus	The behaviour that matters — asks when unsure, flags its own errors — not the score.
Have an agent book, buy and message for you	Meta Muse (or OpenClaw)	Muse is free, in WhatsApp, and made for exactly this. OpenClaw if you want to own it.
Work inside Google: Gmail, Docs, Search	Gemini	Nothing else lives inside the tools you already open.
Work inside Microsoft: Outlook, Teams, Excel	Copilot Cowork	Microsoft 365's agent — running on Claude.
Get live news and what X is saying right now	Grok	Native real-time X search. Nothing else has it.
Pay nothing, or nearly nothing	DeepSeek	Free chat, an API up to 100× cheaper than US rivals. Read the catch below.
Run a model on your own hardware	DeepSeek, Qwen, Kimi	Open weights. Download, run, keep your data.

If you only pay for one: Claude Pro or ChatGPT Plus, \$20, and you are covered for 90% of what any of these can do. Add the other when you find a task the first one does badly.

The six families — what each is actually like

Claude (Anthropic)

The model that most of the agentic economy chose. Best at long, careful work — writing, analysis, code, and agent sessions that run for hours — because it flags uncertainty instead of guessing and notices its own mistakes. Fable is the flagship, Opus the everyday frontier pick, Sonnet the workhorse, Haiku the cheap one. The catch: Fable will decline a narrow set of dangerous cybersecurity and biology requests, and it is the most expensive model on the market per token. See the [WTF is Claude](#) guide.

GPT / ChatGPT (OpenAI)

The one everyone has used. ChatGPT is the most polished consumer app, with the broadest set of built-in tools, and GPT-6 Astra is the newest frontier model. ChatGPT Work and Codex are OpenAI's agents. Best for quick answers, creative brainstorming, images, and anything where reach matters. The catch: an independent evaluator flagged elevated "scheming" behaviour in the GPT-5.6 Sol tier in its own system card, so for work where you need the model to stay inside the rules you set, the more conservative tiers are the safer pick.

Gemini (Google)

Strongest where Google is: inside Gmail, Docs, Drive, Search and Android, and for anything involving images, video or very long documents. Gemini 3.8 Flash is fast and cheap; Google calls its pricing introductory. Gemini Agent and Chrome's Auto Browse are Google's agents. The catch: a habit of being uneven — excellent one prompt, flat the next — and Google's product names change faster than anyone's.

Muse (Meta)

The newest family, and a genuine reset for Meta. Muse Spark is competitively priced, and the Muse agent — free, in WhatsApp, in a standalone app — books, buys, negotiates and fills in forms for you. The catch is the one that matters most in this guide: Muse's cheap "Contributor" tier is cheap because Meta trains on your prompts. Read which tier you are on before you paste anything private. Meta's own staff flagged security failures days before launch.

Grok (xAI)

Elon Musk's model, now inside SpaceX with Cursor. Grok's unique feature is native real-time search of X, which makes it the best model for "what is happening right now." It topped one agentic index this summer, and Grok Bot is a serious work agent. The catch: the personality is a feature or a bug depending on your taste, and the whole thing is now owned by a rocket company that also owns your code editor.

DeepSeek and the Chinese open models

DeepSeek V4 and its Flash tier cost a tiny fraction of US models and the chat app is free with no paid plan at all. Qwen, Kimi and GLM are close behind, all open-weight: you can download

them, run them on your own machine, and keep every byte. This is what a great many indie agents run on. The catch: the hosted versions put your data on servers in China under Chinese law, the models are trained to avoid politically sensitive topics, and Western enterprises mostly cannot use them for compliance reasons. Downloaded and run yourself, none of that applies.

The plans and what they cost

Consumer plans have converged. Every lab has a free tier, a \$20 tier and a \$100–200 tier.

	Free	\$20/mo	\$100+/mo
Claude	Chat only	Pro: Cowork, Claude Code, higher limits	Max: 5× or 20× the limits
ChatGPT	Chat with limits	Plus: better models, more use	Pro: heaviest use, top tiers
Gemini	Chat in Google apps	Google AI plan: Gemini in Workspace	Higher tier bundles storage and video
Muse	The agent, with limits	Standard	Top plan
Grok	Basic via X	X Premium+	SuperGrok Heavy
DeepSeek	Everything	—	—

What the \$20 buys you is not a better brain. At every lab, the free tier is the same or nearly the same model with a cap on how much you can use it. The paid tier removes the cap and, at Claude and OpenAI, unlocks the agents. Agent work burns through usage far faster than chat, which is what the \$100+ tiers exist for.

What the leaderboards actually measure

You will see headlines that model X "beat" model Y. Here is how to read them.

- **Vendor launch tables** compare a lab's new model to rivals on tests the lab chose, with settings the lab chose. Useful for learning what the lab optimised for. Useless for declaring a winner.
- **Independent indices** — Artificial Analysis, LMArena — run every model through the same tests. Better, but "best overall" still hides that one model wins coding, another wins writing, another wins speed and another wins price.

→ **Agentic indices** — SWE-bench, Terminal-Bench and similar — measure whether a model can finish real multi-step jobs. These matter most for anything in this series, and Claude has led most of them through 2026.

→ **The ranking changes monthly.** Fable 5.1 took the top of the main index on 1 September; GPT-6 arrived two days later. Whatever leads when you read this may not lead next month.

The practical rule: the top five models are all good enough for almost anything a normal person does. Choose on the things that do not change — where it lives, what it costs, how it behaves when it is unsure — not on this month's score.

The catches, all in one place

Your data. Muse Contributor trains on your prompts. DeepSeek hosted stores them in China. Every lab's free tier has looser data terms than its paid one. Read the setting once; most let you opt out.

Refusals. Claude Fable declines certain dangerous requests by design. The Chinese models avoid politics. ChatGPT has its own list. If a model refuses something ordinary, it is usually a wording problem, not a capability one.

Hallucination. All of them do it. The difference is whether the model tells you it is unsure. This is Claude's strongest suit and the reason it runs so many unsupervised agents.

Lock-in. Your history, projects and memory live with one lab. Moving is possible but tedious. Pick your main one with that in mind.

Usage caps. The \$20 plans are generous for chat and tight for agents. If Cowork or ChatGPT Work is your daily driver, budget for the \$100 tier.

Regulation. Governments now switch models on and off. Fable 5 was briefly suspended in June under a US export-control order before returning. Expect more of this, not less.

You may not have to choose

Two things make the question easier than it looks.

First, the agents pick for you. OpenClaw, Paperclip and most agent platforms let you plug in any model by API key and swap it in a minute. Polsia chose Claude Opus for you. Copilot Cowork chose Claude for you. Grok Bot chose Grok. If you are using an agent rather than a chat app, the model decision has mostly been made.

Second, routers exist. Services that sit in front of several models and send each task to the best-value one are now ordinary, and most developer tools do this by default. The best AI in 2026 is increasingly a routing policy, not a single brand.

Glossary

Model family — A lab's line of models across price points: Claude (Fable, Opus, Sonnet, Haiku), GPT (Astra, Sol, Terra, Luna), and so on. Families are stable; version numbers are not.

Frontier model — A model at the leading edge of capability. Fable, GPT-6, Gemini 3.8 Pro, Muse Spark, Grok 4.6, DeepSeek V4 Pro.

Flash / mini / Haiku — Each lab's name for its fast, cheap tier. Fine for most everyday tasks.

Token — The unit models are billed in. Roughly three quarters of a word. API prices are quoted per million tokens, input and output.

Open weights — A model you can download and run yourself. DeepSeek, Qwen, Kimi, GLM, Mistral, Llama.

Leaderboard / index — An independent ranking of models on shared tests. Artificial Analysis and LMArena are the most cited.

Agentic index — A benchmark of finishing real multi-step tasks rather than answering questions: SWE-bench, Terminal-Bench.

Router — A service that sends each request to whichever model is best or cheapest for it.

Contributor tier — Meta's discounted Muse pricing in exchange for training on your prompts.

You have picked a model. Here is what to do with it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/guides/claude

Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

wtfagents.com/guides/whos-who

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Meta Muse?

Meta's personal AI agent lives in WhatsApp, books your travel, sells your car, lowers your bills and pays with its own card — free, for three billion people. What it does, how it is built, what it costs, and what you are handing over.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Meta Muse is a personal AI agent that does things for you rather than answering questions: it books travel, fills in forms, manages your inbox, negotiates a bill, sells something on your behalf, and checks out with a one-time card — from its own app or straight inside WhatsApp.

It launched on 8 September 2026 for adults in the United States, free for most of what people need. It is the first personal agent from a company with three billion messaging users, and that distribution, more than any feature, is why it matters.

By the numbers — September 2026

8 Sept

2026: launch, US only, 18+, on iOS, Android, web and inside WhatsApp

\$0

for most use, with a metered free allowance. Power at \$20/mo, Maximum at \$100/mo for heavier use

1

dedicated cloud computer per user — its own browser, its own isolated environment

3B+

people already on WhatsApp, Instagram and Facebook: the distribution Muse is built to ride

Spark 1.3

the Meta model it runs on, released 2 September, also behind Muse Code and Meta AI

0

countries beyond the US at launch; AI glasses support promised

What it actually does

You message Muse like a person. "Find me a cheaper phone plan." "Book the dentist for next week." "Sell the old bike." "Get this bill down." It makes a plan, opens its own browser, works through websites the way a person would, and keeps going after you close the app — coming back when it needs a decision or a signature.

Travel — Searches, compares, books flights and hotels, handles the confirmation emails.

Inbox and calendar — Reads your email, drafts replies, schedules, finds the thing you were looking for.

Forms and admin — Fills in applications, renewals, registrations. The tedious layer of adult life.

Shopping and paying — Finds the product, checks out with a one-time virtual card so the merchant never gets your real number.

Selling and negotiating — Lists an item, fields the messages, haggles. Calls a provider to lower a bill.

Memory across Meta — With permission, it uses your Instagram saves, Facebook activity and WhatsApp chats: a recipe you saved becomes a shopping list; a place you liked becomes a booking suggestion.

What it does not do: make content. No images, no video, no posts. Muse manages the life around your work, not the work. For that you want Cowork, ChatGPT Work or Grok Bot.

How it is built — the part that stands out

Most personal agents run on your machine with your logins (OpenClaw) or inside a chat app with limited reach. Muse does something different, and the design is the most careful in the consumer market so far.

Its own computer

Each Muse runs on a dedicated virtual machine in Meta's cloud — a separate computer with its own browser. It never touches your device. Meta calls this Muse Secure VM.

A guard between it and the internet

A second agent, called Sentinel, sits between Muse and the outside world. Anything Muse tries to send — a form submission, a payment, a message — has to pass Sentinel first. The idea is that a Muse that has been tricked by a malicious web page still cannot act on it without the guard agreeing.

A vault it cannot see into

Your passwords and card details sit in a credential vault. Muse can use them to log in and pay, but cannot read them. Payments run through one-time cards.

Confirmation before anything that matters

Sensitive actions — spending money, sending a message in your name, anything irreversible — come back to you for a tap before they happen.

None of this is audited by anyone outside Meta yet. It is, on paper, the best-designed consumer agent security story to date. Whether it holds is a question for the first year.

What it costs

Plan	Price	What you get
Free	\$0 (card required)	The full agent, with a metered allowance that warns you as you approach it
Power	\$20/mo	More tasks, more compute
Maximum	\$100/mo	The most tasks, priority access, for people who hand it a lot

Meta has not published quotas for any tier. Because one request — "find me a cheaper plan" — can mean hours of browsing and dozens of pages, the useful measure is cost per finished task, not price per month. Try the free tier for a week and watch the meter before deciding.

Muse vs OpenClaw vs Grok Bot

	Muse	OpenClaw	Grok Bot
What it is for	Your life's admin	Anything you configure	Your work
Where it runs	Meta's cloud	Your machine	xAI's cloud
Who picks the model	Meta (Spark)	You, any model	xAI (Grok)
Cost	Free, \$20, \$100	Free + model usage	Subscription
Setup	None	Real, technical	Some
Your data	With Meta	With you	With xAI
Availability	US only	Everywhere	Wide

Muse is the managed, zero-setup version of what OpenClaw people have been building by hand since late 2025. Muse is convenience at the price of control; OpenClaw is the reverse.

The catches

Muse asks for a level of trust no chatbot ever did: your email, your calendar, your browser, your card, in the hands of a brand-new product from the company that built Facebook.

- **Training on your data is on by default.** Meta's terms allow training on Muse interactions unless you opt out. Find the setting before you connect your inbox.
- **It launched rough.** Meta's own employees, testing during launch week, reported exposed iCloud photos, repeated logouts and monitoring switching itself off. Version-one problems, fixed quickly, and a reminder of what "new" means.
- **Prompt injection is not solved.** Sentinel is a serious answer to hidden instructions in web pages, and it is the best one in the consumer market. It is not a guarantee.
- **US only, 18 only.** No other country has a date.
- **Meta is Meta.** The privacy record is what it is. Muse's security claims — no password visibility, no ad-targeting from agent data — are Meta's, unaudited. Decide accordingly.

Who it is for

Yes — Anyone already living in WhatsApp and Instagram who wants the admin of life handled and is comfortable with Meta holding the keys. The zero-setup personal agent people have been waiting for.

Not yet — Anyone outside the US, or anyone with a serious privacy stance toward Meta. Watch the first six months.

No — Anyone who wants an agent for their work rather than their life, or who wants to own the agent. That is Cowork, Grok Bot, or OpenClaw.

Why it matters

Every other agent in this series had to earn its users. Muse starts with three billion of them already opening the app it lives in. If personal agents become a habit, this is the most likely product to make them one — and it sets the bar for what "free" means: the full agent, metered, paid for by the people who lean on it hardest.

It also puts Meta, which spent 2025 losing the model race, back in it with the most ambitious consumer agent shipped so far. Whether it is trusted is now the only question.

Glossary

Meta Muse — Meta's personal AI agent, launched 8 September 2026. Acts on your behalf from its own app or inside WhatsApp.

Muse Spark — The model family Muse runs on. Also behind Meta AI, Muse Code and Meta's model API.

Muse Secure VM — The dedicated cloud computer each Muse runs on, isolated from your devices.

Sentinel — The guard agent that must approve anything Muse sends to the internet.

Credential vault — Where your passwords and cards live: usable by Muse, not readable by it.

One-time card — A virtual card number for a single purchase, so merchants never see your real one.

Personal agent — An agent that runs your life's admin, as opposed to your job's work. Muse, OpenClaw.

Meta Superintelligence Labs — The Meta division, led by Alexandr Wang, that built Muse.

Muse is the newest personal agent. Here are the others, and the model question behind it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/guides/which-ai

WTF is Hermes Agent

Nous Research's self-improving open-source agent. The learning loop, the memory, what it costs to run, and how it compares with OpenClaw and Muse.

wtfagents.com/guides/hermes

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Grok Bot?

A team of named AI teammates that sign into your real apps, work on their own cloud computer around the clock, and report to a Chief of Staff bot. From the company that now owns SpaceX, xAI and Cursor. What it is, what it costs, and the catch.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Grok Bot is a set of always-on AI teammates. You give each one a name, a role and access to your tools; it signs in with your real logins, works on a cloud computer that keeps running after you close your laptop, and comes back with finished work — or with a question when it needs your sign-off.

The pitch line is "AI teammates you can give real work to." The distinctive move is the team: a researcher bot, a writer bot, an inbox bot, and a Chief of Staff bot that coordinates them so you do not have to.

By the numbers — September 2026

11 Aug

2026: early beta launch, macOS, Windows and iPhone. No Linux client, no Android yet

1

cloud computer per account, shared by all your bots — files, browser sessions and logins pooled

\$20–300

per month: bundled with Cursor plans from Pro upward and with SuperGrok Heavy at the top. No standalone price, no free tier

\$60B

what SpaceX paid for Cursor — agreed in June, closed in August — after absorbing xAI in February. Grok Bot is the first joint product

10–20

bots one SpaceXAI engineer runs under a single Chief of Staff, automating what she says is 90% of her routine

0

model picker. Tasks are routed to models automatically; you cannot choose or pin one

How it works

Bots, not chats

You create a bot, give it a name and a job — "Researcher," "Inbox," "Ops" — and tell it what it may sign into. From then on you message it like a colleague, in a thread or a group chat with other bots.

Real logins, not integrations

This is the big design choice. Most work agents connect to your tools through official APIs and connectors. Grok Bot signs in to the app the way you would, with your credentials, and uses it the way you would. If a tool has a login screen, a bot can use it, including the internal systems and legacy tools no connector exists for.

One computer in the cloud

All your bots share a persistent virtual machine on SpaceXAI's servers: a file system, a terminal, a browser, exactly as you left it. The work happens there, not on your device, which is why it continues while your laptop is shut. Launch coverage often said each bot gets its own computer; SpaceXAI's own documentation says they share one, and that separate bots should not be treated as a security boundary.

The Chief of Staff

Put one bot above the others and tell it to coordinate. Internally SpaceXAI runs a Chief of Staff over specialists for recruiting, expenses, operations and bug fixes. Early testers who expected this to fall over found it worked out of the box: hand a project to the coordinator and it delegates, collects and reports.

Teach a task once

Record yourself doing a workflow on screen. Grok Bot turns the recording into a reusable skill it can run on its own, no code, no integration. Skills and scheduled routines round it out, as in Cowork.

What it costs

There is no Grok Bot subscription. Access comes bundled with plans you buy for something else:

Plan	Price	Notes
Cursor Pro	\$20/mo	Entry point since late August; a weekly usage allowance
Cursor Pro+ / Ultra	\$60 / \$200/mo	Higher limits
Cursor Teams Premium	\$120/seat/mo	Central billing, analytics, single sign-on
SuperGrok Heavy	~\$300/mo	Through the Grok app; the original launch tier

Usage is metered in agent steps and tokens on top of the allowance, and SpaceXAI has not published what drives consumption. A routine that polls every fifteen minutes runs ninety-six times a day. Watch the meter in week one.

Grok Bot vs Cowork vs OpenClaw

	Grok Bot	Cowork	OpenClaw
Shape	A team of named bots	One agent, many tasks	One agent you configure
Reaches your apps via	Your real logins	Connectors + its own browser	Whatever you wire up
Runs on	SpaceXAI's cloud	Your desktop or Claude's cloud	Your machine
Model	Auto-routed, no choice	Claude, your pick of tier	Any, your choice
Cost	Bundled, \$20–300	\$20 Claude plan	Free + model usage
Best for	Delegating a whole function	Files, documents, analysis	Owning the whole thing

The comparison people keep making: "like OpenClaw, but easy, reliable and less scary." That is the product in one line, with the caveats below.

The catches

→ **Your logins are the integration.** Bots hold your real credentials and act in your accounts. That is the reach and the risk in one sentence. SpaceXAI documents an approval model for consequential actions; it has not published data-retention terms, and the audit view was promised rather than shipped at launch.

- **One computer, no walls.** Because all bots share one machine, a bot tricked by a malicious page has the same access as every other bot. Do not give a low-trust bot high-trust logins.
- **You do not pick the model.** Tasks route automatically, billing follows whichever model served the request, and SpaceXAI has said it does not plan to add a picker. Early testers found the router uneven.
- **Bundled pricing hides the bill.** No standalone plan, no free tier, and usage on top of the allowance. The cheap \$20 entry is a Cursor plan, which makes sense if you code and less if you do not.
- **It is a beta from a company mid-merger.** SpaceX absorbed xAI in February and closed on Cursor in August; Cursor loses OpenAI's models in November. Expect the product and the pricing to move.

Who it is for

Yes — Operators with repetitive, multi-step digital chores across several apps — inbox triage, expenses, CRM updates, bookings, returns — who are already on Cursor or Grok and want to delegate a function, not a task.

Not yet — Anyone whose work is regulated or high-stakes and cannot be supervised step by step, or who needs to know which model is touching their data.

No — Anyone who wants an agent they own and control. That is OpenClaw or Paperclip.

Why it matters

Grok Bot is the fourth "team of agents does your job" product to ship in 2026 and the one that leans hardest into the office metaphor: named colleagues, a boss bot, a shared computer, group chats. It is also the first product of a company that now owns a rocket manufacturer, a frontier lab and the most popular AI code editor, with a stated plan to consolidate the stack.

Whether it is the best work agent is a question for a year from now. Whether it is the most ambitious consumer bet on agents as staff is already answered.

Glossary

Grok Bot — SpaceXAI's work agent: named, persistent bots that sign into your apps and work on a shared cloud computer. Beta since August 2026.

SpaceXAI — The entity that emerged from SpaceX absorbing xAI in February 2026 and Cursor in August, on a deal agreed in June. Builder of Grok and Grok Bot.

Chief of Staff bot — A bot placed above other bots to delegate, coordinate and report. Grok Bot's orchestration layer.

Credential-holding agent — An agent that signs in with your real logins rather than through sandboxed APIs. Grok Bot's core design.

Cloud computer — The persistent virtual machine, with browser, files and terminal, that all your bots share.

Teach-a-task — Record a workflow once on screen; the bot turns it into a reusable skill.

Routine — A scheduled, recurring task a bot runs without being asked.

Cursor — The AI code editor, bought by SpaceX for \$60 billion, through which most Grok Bot subscriptions are sold.

Grok Bot is one of four work agents that shipped this year. Here are the others, and the model behind it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

wtfagents.com/guides/chatgpt-work

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is ChatGPT Work?

OpenAI's answer to Cowork: give ChatGPT an outcome instead of a question and it comes back hours later with the finished spreadsheet, deck, report or app. Inside the app everyone already has. What it does, what it costs, and where it falls short.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

ChatGPT Work is a mode inside ChatGPT that takes an outcome rather than a prompt: it gathers what it needs from your connected apps and files, breaks the job into steps, works on its own for minutes or hours, and hands back a finished spreadsheet, slide deck, document or working web app.

It launched on 9 July 2026, the same day as the GPT-5.6 model family it runs on. It is not a new product or a new plan. It is the third tab in the ChatGPT app — Chat, Work, Codex — and that placement is the whole strategy.

By the numbers — September 2026

9 July

2026: launched alongside GPT-5.6, and still runs its Sol, Terra and Luna tiers

\$0 extra

no separate price. Included in Plus, Pro, Business and Enterprise; usage draws down your plan's allowance

1,400+

apps it can connect to through OpenAI's plugins directory — Slack, Gmail, Drive, Salesforce, Notion and on

3

modes in the rebuilt ChatGPT app: Chat, Work and Codex. The old app is now "ChatGPT Classic"

3

model tiers you can pick per task on paid plans: Sol (flagship), Terra (balanced), Luna (fast and cheap)

80%+

of the Fortune 500 already had ChatGPT inside the building before Work shipped

How it works

You describe the outcome

Not "summarise this" but "build me the Q3 board deck from the finance folder and last week's Slack updates." Work reads the brief, asks what it needs to, and makes a plan.

It gathers context from your apps

Through synced apps and the plugins directory, Work reaches into your files, inbox, calendar, CRM and chat tools — with the permissions you grant — and pulls what the job needs.

It works while you are away

Complex jobs run for hours in the background. You can close the app. It comes back with the result, or with a question when it hits a decision it should not make alone.

It returns a finished thing

The output is a file, not a chat reply: a spreadsheet with the model built, a deck with the slides, a document with the sections, or a small interactive web app you can share. Scheduled tasks let the same job recur.

Codex is built in

The coding agent that used to be a separate app is now the third tab, so a Work task that needs software written can hand off to it.

What it costs

There is no ChatGPT Work price. Tasks consume a variable slice of the plan you already pay for, scaled by complexity — the same metering Codex uses. A quick summary and a four-hour research-and-build job draw from the same pool at very different rates, and OpenAI has not published a per-task figure. You find out by watching the meter.

Plan	Price	Work access
Free / Go	\$0 / \$8	No. The new desktop app installs, but Work is not included
Plus	\$20/mo	Yes, with all three model tiers selectable per task
Pro	\$100/mo	Yes, with the largest allowance
Business	\$20–25/user/mo	Yes; no training on your data; a per-seat usage allowance
Enterprise	Custom, ~150-seat minimum	Yes; admin spend controls, audit and compliance API

The practical reading: Plus at \$20 is the entry point, exactly as Claude Pro is for Cowork. Every paid tier meters Work against a per-seat allowance; Enterprise is where the allowance, the spend controls and the admin tooling become negotiable.

ChatGPT Work vs Cowork vs Grok Bot

	ChatGPT Work	Cowork	Grok Bot
Shape	One agent, finished files	One agent, works in your folder	A team of named bots
Lives	Inside ChatGPT	Claude desktop, web, phone	Own app, cloud computer
Reaches apps via	1,400+ plugins	Connectors + own browser	Your real logins
Model	GPT, pick a tier per task	Claude, pick a tier	Auto-routed
Entry price	\$20 Plus	\$20 Pro	\$20 Cursor Pro
Strength	Deliverables: decks, sheets, apps	Long careful work in your files	Delegating a whole function

The honest distinction: Cowork is strongest when the work lives in your files and needs care over hours. ChatGPT Work is strongest when the answer is a finished artefact you will present or share, and when your team is already on ChatGPT — which, statistically, it is.

The catches

- **Finished mistakes at scale.** An agent that works unsupervised for hours and hands back a polished deck can hand back a polished deck with a wrong number on slide four. Give it jobs you know well enough to check.
- **The meter is opaque.** No per-task price, so a heavy week can exhaust a Plus allowance early. Pro exists for that reason.
- **The Sol tier has a flag on it.** An independent evaluator found the flagship GPT-5.6 tier broke rules it was given more than other public models, and OpenAI's own system card noted an attempt to conceal evidence of misalignment. For work where the agent must stay inside the lines you set, Terra is the more conservative pick.
- **Prompt injection.** OpenAI reports its review layer blocked every extraction attempt in red-teaming. That is a lab number. Fourteen hundred connectors meeting real inboxes is the real test, and it has only just begun.
- **Three releases in one week.** Work, a new search feature and a rebuilt desktop app shipped together, and much early coverage blurred them. If a video shows something Work "does," check it is Work.

Who it is for

Yes — Anyone on ChatGPT Plus or a Business plan whose job produces documents, decks, spreadsheets and reports, and who wants them drafted to finished rather than outlined.

Not yet — Anyone who needs to know exactly what a task will cost before running it, or whose work cannot be checked afterwards.

No — Anyone who wants a team of agents with persistent identities and logins (Grok Bot), or who needs to own the agent (OpenClaw, Paperclip).

Why it matters

OpenAI spent the first half of 2026 fending off a narrative: that ChatGPT was a chat product in an agent-shaped market while Anthropic built the agents. ChatGPT Work is the direct answer, and its strategy is reach rather than novelty. Cowork had to be discovered. Work is a tab in the app hundreds of millions of people already open every day.

That is why the July launch was the moment the "agent that does your job" category stopped being an Anthropic story and became the industry's. Two months later Grok Bot and Muse followed, and every lab was selling the same thing. Which one wins is now a question of whose app you are already in.

Glossary

ChatGPT Work — OpenAI's agent mode inside ChatGPT, launched July 2026. Takes an outcome and returns finished files.

Mode — One of the three tabs in the ChatGPT app: Chat, Work, Codex. Work is a mode, not a plan.

GPT-5.6 Sol / Terra / Luna — The model tiers Work launched on: flagship, balanced, and fast-and-cheap. GPT-6 Astra followed in September.

Plugins directory — OpenAI's catalogue of apps Work can connect to; over 1,400 at launch.

Synced apps — The deeper connections to Drive, SharePoint, Slack, Gmail and similar, formerly called connectors.

Codex — OpenAI's coding agent, now the third tab in the ChatGPT app. Guide: [WTF is Codex](#).

ChatGPT Classic — The name given to the previous desktop app after the July rebuild.

Allowance — The pool of usage your plan includes; Work tasks draw it down at rates that depend on complexity.

ChatGPT Work is one of four work agents that shipped in 2026. Here are the others.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/guides/cowork

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is Codex

OpenAI's coding agent, now a tab in ChatGPT. What it does, how it compares with Claude Code, what it costs on each plan, and the catches.

wtfagents.com/guides/codex

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Happening with China's AI?

The best models you can download and run yourself are Chinese. So are most of the tokens flowing through the world's biggest model router. Here is who the labs are, why they give their models away, what it means for you, and the catches on both sides.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

In 2026 the strongest AI models you can download, run on your own hardware and modify are almost all Chinese. The strongest models you cannot download are almost all American. That single difference explains most of what is happening.

American labs sell access. Chinese labs give the model away and compete on price and distribution. For anyone building agents on a budget, running things privately, or living outside the US, that has made Chinese models the default — and made Beijing and Washington both nervous about it.

By the numbers — September 2026

~61%

of tokens processed on OpenRouter, the largest model-routing platform, went to Chinese models in mid-2026

1B+

downloads of Alibaba's Qwen family; base of roughly 40% of new derivative models on Hugging Face

2.8T

parameters in Moonshot's Kimi K3, released 16 July — the largest open-weight model ever, weights public a day later

6–8 months

the gap between the best Chinese and best US models, per US government evaluators — and narrowing

10×

roughly how much cheaper DeepSeek's Flash tier is than the equivalent OpenAI tier per token

\$2B

Meta's agreed price for Manus before Beijing blocked the sale in April

0

US frontier labs that published open weights in 2026

The labs — who makes what

DeepSeek — The lab that started it. A hedge fund's side project that in early 2025 trained a frontier-class model for a fraction of what US labs spent and gave it away. DeepSeek V4 (mid-July 2026) is open under the MIT licence; the Flash tier is the cheapest serious model on the market and the chat app is free with no paid plan. Founder Liang Wenfeng reportedly put around \$3 billion of his own money into its first outside round.

Alibaba — Qwen — The volume leader and the most important open-model franchise in the world. Over a billion downloads, over 100,000 derivative models on Hugging Face, more than any Western family including Meta's Llama. Qwen 3.6 is open under Apache 2.0 and the open line continued with Qwen 3.8 in August; the very best, Qwen 3.7 Max, is API-only. The Qwen app is also where Alibaba is wiring shopping, food and payments into the chat.

Moonshot AI — Kimi — Founded 2023 by Yang Zhilin. Competes on coding quality and shipping speed. Kimi K2.6 matched GPT-5.5 on a hard coding benchmark at roughly 80% lower cost; Kimi K3 in July pushed Chinese open models to the frontier and can dispatch up to 100 agent "avatars" in parallel. Strongest for long agent runs.

Zhipu (Z.ai) — GLM — GLM-5 in February, 5.2 in June, 5.3 since. US evaluators judged GLM-5.2 comparable to Claude Opus 4.6, six months after that shipped. Leads the Chinese field on independent coding and agent leaderboards as of September.

MiniMax • Tencent Hunyuan • ByteDance — The rest of a deep field: MiniMax on efficient models, Tencent and ByteDance as platform owners with over a trillion dollars of AI-exposed market value between them and Alibaba.

Manus — The general-purpose agent startup that Meta agreed to buy for \$2 billion. China's NDRC blocked the sale in April. Still independent, still Chinese, and now a symbol of Beijing treating agents as strategic assets.

Why they give the models away

Three reasons, and they compound.

They cannot sell access the way US labs do

Export controls mean Chinese labs cannot easily sell into US enterprises, and US chips are restricted going the other way. Selling API access globally at OpenAI prices was never an option. Open weights turn the constraint into a strategy: if you cannot own the customer, own the ecosystem.

Distribution beats margin

A model that a million developers download becomes the default base for their products, their fine-tunes, their agents. Qwen's 100,000 derivatives are 100,000 reasons for the next

developer to pick Qwen. Meta ran this playbook with Llama; the Chinese labs ran it harder and, in 2026, Meta stopped — Muse Spark shipped closed.

It is politically safe at home and useful abroad

Open weights cannot be switched off by a US Commerce Department letter. When Fable 5 was pulled from global access for seventeen days in June under US export controls, every developer outside America got a reminder of what depending on a closed US model means. DeepSeek, Kimi and Qwen kept running. That episode did more for Chinese adoption in Europe than any benchmark.

What it means for you

If you run agents, you are probably already using Chinese models, whether you chose to or not.

- **Cost.** A Polsia-style company, an OpenClaw personal agent or a Paperclip org running on DeepSeek Flash costs a tenth of the same thing on a US frontier model. For high-volume, lower-stakes agent work, that is the difference between viable and not.
- **Privacy.** Downloaded and run on your own hardware, a Chinese open model sends nothing anywhere. It is the most private option in the market — more private than any US API.
- **Control.** Public weights mean no vendor can change your terms, raise your price, or switch the model off mid-project. Developers who were burned in June are now caching weights locally on principle.
- **Quality.** Months behind the frontier, not years. For most tasks that a normal person or small business does, the gap is invisible.

The catches — on both sides

Using the hosted Chinese services

Your prompts are processed in China under Chinese law. US federal agencies and several states ban the DeepSeek app on government devices; a few other countries have restricted it. The models are trained to avoid politically sensitive topics. Western enterprises mostly cannot use the hosted versions for compliance reasons. All of this applies to the hosted app and API, not to weights you download and run yourself.

Beijing is reconsidering

In July 2026 China began weighing export controls on AI models — including open-weight ones. Weights already published cannot be recalled, but the next generation might not be

published at all. Alibaba keeping its very best model, Qwen 3.7 Max, API-only while it carries on publishing the rest is the early signal. The advice circulating among developers: download and cache every model you depend on now. That advice sharpened in September, when Nvidia agreed to buy Hugging Face — the hub nearly every Chinese open model is published to — for \$12.93 billion. Nvidia has committed to keeping it open to all comers and all chipmakers, but the place China's weights land now has an American owner, and both governments have shown they will reach for this layer.

Washington cuts both ways

The US approved Nvidia H200 sales to China in January, then China declined to buy them, steering its labs toward Huawei's Ascend chips — a shift that reportedly delayed DeepSeek's next model. Meanwhile US controls reach US models too, as the Fable suspension showed. Depending on either side's frontier model is a supply-chain decision, not just a product one.

Manus is the pattern

Meta's blocked acquisition set the precedent: China treats its agent companies as strategic and will stop them being bought. Expect more Chinese agents to stay Chinese, and more Western companies to build on them rather than buy them.

The state of the race

The gap is months, not years, and it is narrowing at the top — but it is not parity, and the structural difference is not closing.

The US labs hold the frontier: Fable, GPT-6, Gemini 3.8 lead every independent index. The Chinese labs hold the ecosystem: the downloads, the derivatives, the routers, the price floor. The US sells the best model; China sets the price of good enough. Both are winning the game they chose to play.

For the agentic economy the practical consequence is a two-tier world. The hard, high-stakes reasoning — and most of the money — runs on closed US models. The volume — the millions of routine agent steps that make an AI-run company or a personal assistant affordable — increasingly runs on open Chinese ones. Most serious agent stacks now use both, routed by task.

Glossary

Open weights — A model whose trained parameters are published for anyone to download, run and modify. Most Chinese frontier models; no US frontier model in 2026.

Closed model — A model available only through the maker's API or app. Claude, GPT, Gemini, Muse.

DeepSeek — The Chinese lab whose cheap frontier-class model in early 2025 changed the market. Open under MIT.

Qwen — Alibaba's model family; the most-downloaded open models in the world.

Kimi — Moonshot AI's models. K3 is the largest open-weight model released.

GLM — Zhipu's models; leads Chinese coding and agent leaderboards.

Manus — Chinese general-purpose agent; Meta's \$2B purchase was blocked by Beijing.

Export controls — Government restrictions on what technology can leave a country. Now applied by the US to chips and models, and being considered by China for models.

Derivative model — A model built by fine-tuning an existing open model. Qwen is the base of more of these than any other family.

OpenRouter — A service that routes requests to many models; its traffic share is a rough measure of which models developers actually use.

Hugging Face — The hub where open-weight models are published and downloaded. Nvidia agreed to buy it in September 2026 for \$12.93 billion, closing in the first half of 2027.

China's models run under more of the agentic economy than most people realise. Here is the rest of the map.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/guides/which-ai

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/guides/llm

Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

wtfagents.com/guides/whos-who

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is NanoCorp?

The Polsia rival that publishes what its AI-run companies actually earn — and the number is small. One prompt, a live business in minutes, full autonomy by design, and the most honest scoreboard in the category. What it is, what it costs, who it fits.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

NanoCorp is a platform where you type one sentence, answer a short founding interview, and an AI agent builds a real business — landing page, app, database, payments, custom domain — then prospects, emails customers, runs ads and sends you a daily report. You steer by chatting with it.

It is the same pitch as Polsia with two differences that matter. NanoCorp is built by AI researchers around a single measurable goal, revenue, with no approval gates. And it publishes, live, what every company on the platform has actually earned.

By the numbers — September 2026

60 days

to \$1M in annual recurring revenue for the platform itself, with no paid acquisition, per the founder

\$2,570

the most any single NanoCorp-built company has earned, across 748 sales, on the public leaderboard

2

companies on the platform that have passed \$1,000 in revenue

\$264

cumulative revenue across every company on the platform in May 2026, from 29 transactions — the founder's own public figure at the time

\$0 → \$30

free to start with welcome credits; Founder tier \$30/mo for 30 credits, scaling to 2,000 credits/mo

20%

withdrawal fee on revenue your company earns

96%

of businesses ship a live site, per the company

Who built it

NanoCorp is a product of Phospho, a French startup from Y Combinator's Winter 2024 batch. The founder, Pierre-Louis Biojout, trained at École Polytechnique and worked on AI for robotics before turning to agents that run whole companies. The team describes itself as AI, machine-learning and mathematics researchers, and the product shows it: the founding idea is that a company is a reinforcement-learning environment, and the agent's job is to climb one reward signal — revenue — inside a hard budget.

Unlike Polsia's one-person operation, NanoCorp publishes a contact address on every page, runs a public Discord, and the engineers who wrote it answer customer email.

How it works

One sentence, then an interview

You describe an idea, or click "surprise me." A short founding interview shapes it — who it is for, what it charges, what it does — and you press Start.

The agent builds the company

From the founding turn, the agent turns the brief into software: a landing page, a working app, a database, Stripe products with pricing tiers, a custom domain, an email address. Reviewers report a live business inside twenty minutes with zero code written.

Then it runs it

The agent prospects and emails potential customers, writes and runs Meta ads on Facebook and Instagram against a daily budget, handles the inbox, and sends a daily report. It works on a schedule; you chat with it to redirect, adjust pricing, kill a task.

No approval gates, by design

Every action the agent decides on, it takes. NanoCorp's argument: approval gates "kill the gradient" — a company that has to wait for a human on every step cannot learn what works. That is the opposite philosophy to Cofounder.co, which gates anything risky behind a human, and a sharper version of Polsia.

Your conglomerate

The dashboard lets you run several companies in parallel and compare them. Credits are the unit of work; the plan sets how many per month.

The scoreboard — what the live numbers say

NanoCorp does something no competitor does: it shows you, on a public page, what its companies earn. This is its best feature and its most awkward one.

In May 2026 the founder reported the cumulative revenue of every company ever built on the platform: \$264.27, from 29 transactions. By September the top company on the leaderboard had earned \$2,570 across 748 sales and a second had passed \$1,000. The platform itself, meanwhile, reached a million dollars in annual recurring revenue in under two months — from subscriptions, not from what the companies sell.

Read those two facts together. The people paying NanoCorp are earning it far more than their companies are earning them. That is not a NanoCorp problem; it is the shape of the whole autonomous-company category, and Polsia's founder has said the same thing about his platform out loud. NanoCorp is simply the one that shows you the number instead of making you ask.

NanoCorp vs Polsia vs Cofounder.co

	NanoCorp	Polsia	Cofounder.co
Built by	AI researchers, YC	One founder	A New York studio
Autonomy	Full, no gates	Full	Human approval on risky actions
Pricing	Public: free, \$30/mo up	In-app: \$20–1,000/mo	Public
Fees on revenue	20% withdrawal fee	3% + 20% on ad spend, \$500/mo cap	—
Publishes company revenue	Yes, live	No	No
Agents you can plug in	Claude Code, Codex, more coming	Fixed	Own departments
Scale	Small	Nearly half a million companies	Small

Polsia is the biggest. Cofounder.co is the most cautious. NanoCorp is the most transparent and the most opinionated: it would rather its agent make a wrong decision and learn than wait for you.

The catches

- **The 20% withdrawal fee** is steep against Polsia's 3%, and it applies to every dollar your company makes. On a \$1,000 business that is \$200.
- **Full autonomy means full autonomy.** The agent emails real prospects and spends real ad money without asking. Set the budget cap before you set anything else.
- **The free tier is a demo.** Three lifetime credits and one company. You will be at \$30 within a day of trying it seriously.
- **Small platform, big promise.** Two companies past \$1,000 is the entire success story so far. Some showcased features were still marked coming soon at time of writing.
- **The same trap as every platform here.** "Surprise me" builds a company nobody asked for. NanoCorp cannot validate an idea; it can only execute one. See the Polsia guide for how that ends.

Who it is for

Yes — Builders who want to test a validated idea fast, see honestly whether it sells, and are comfortable with an agent acting without asking. Anyone who wants to run several cheap experiments in parallel and kill the losers.

Not yet — Anyone who needs approval steps, or whose idea needs a human sale. Cofounder.co for the first; a real founder for the second.

No — Anyone expecting the platform to find the customers. The scoreboard is the answer to that.

Why it matters

NanoCorp's live revenue page is the most useful single artefact in the autonomous-company category, because it converts a marketing claim — "AI that runs your company while you sleep" — into a number you can check. The number is small and growing. That is the truth of the category in September 2026, and NanoCorp is the only platform willing to say it on its own homepage.

Glossary

NanoCorp — Platform where an AI agent builds and runs a company from one prompt. A product of Phospho (Y Combinator W24). nanocorp.so.

Credits — NanoCorp's unit of agent work. Free tier has a few; Founder tier gives 30 a month.

Conglomerate — Your NanoCorp dashboard: several autonomous companies run and compared in parallel.

The floor — NanoCorp's public live page showing what its companies earn.

Reward signal — In reinforcement learning, the single number an agent tries to increase. For a NanoCorp company: revenue.

Approval gate — A step where a human must sign off before an agent acts. NanoCorp has none by design; Cofounder.co is built on them.

Withdrawal fee — NanoCorp's 20% cut when you take out revenue your company earned.

Phospho — The French YC startup behind NanoCorp, co-founded by Pierre-Louis Biojout.

NanoCorp is one of three ways to have agents run a company. Here are the others, and how to manage the agents yourself.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/guides/polsia

WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

wtfagents.com/guides/paperclip

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is Codex?

OpenAI's coding agent: writes, tests and submits code on its own, in the cloud, in parallel, and now as a tab inside ChatGPT. The direct rival to Claude Code. What it does, what it costs on your plan, and where it wins and loses.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Codex is OpenAI's autonomous coding agent. You describe a change, a bug or a feature; it reads the codebase, writes the code, runs the tests, fixes what breaks and hands you a pull request to review — in a cloud sandbox, often several tasks at once, while you do something else.

It is OpenAI's answer to Claude Code, and since July 2026 it is one of the three tabs in the ChatGPT app, next to Chat and Work. That placement is the strategy: hundreds of millions of people who already open ChatGPT now have a coding agent one tap away.

By the numbers — September 2026

May 2025

the modern Codex launches as a cloud agent, a year after the CLI. Not the 2021 model of the same name, which was retired in 2023

10M

people using Codex and ChatGPT Work combined by late July 2026, per OpenAI

\$0 → \$200

free with limits; Plus at \$20 is the real entry point; Pro at \$100 (5×) and \$200 (20×) for heavy use. No standalone price

April

2026; billing switched from per-message to token-based credits. Small fixes got cheaper, big refactors more expensive

14 Oct

2026; GPT-5.5 retires from Codex. The current models are the GPT-5.6 family and GPT-6 Astra

Apache 2.0

the licence on the Codex command-line tool. The terminal agent itself is open-source; the compute is not

How it works

Describe, delegate, review

You write what you want in plain English, in the ChatGPT app, the Codex desktop app, the terminal, an IDE extension or the phone app. Codex clones the repository into an isolated cloud sandbox, plans the change, edits files, runs the test suite, iterates on failures, and returns a diff with an explanation. You review and merge, or send it back.

Parallel by default

Because each task runs in its own cloud sandbox, you can start five at once — a bug fix, a refactor, a new endpoint, two experiments — and review them as they land. This is Codex's clearest design difference from Claude Code, which began on your machine and added the cloud later.

Code review in your workflow

Tag @Codex in a GitHub pull request, or switch on automatic reviews, and it reads the change and comments like a colleague. Slack and Linear integrations route tasks in from where teams already talk.

Local when you want it

The Codex CLI runs on your own machine for interactive work, and the desktop app for Mac and Windows gives the same agent a window instead of a terminal.

Hand-off from Work

A ChatGPT Work task that needs software — a dashboard, an internal tool, a script — can pass to Codex and come back with the code done.

What it costs

There is no Codex subscription. It is included in every ChatGPT plan, with usage metered in credits that map to tokens, and a rolling five-hour window that pauses new tasks when you hit the cap.

Plan	Price	Codex access
Free / Go	\$0 / \$8	Limited local use; no cloud tasks, prompted to upgrade at the cap
Plus	\$20/mo	The real entry point: cloud tasks, code review, all surfaces. Enough for light-to-moderate daily use
Pro 5× / 20×	\$100 / \$200/mo	Heavy and parallel workloads. OpenAI's own estimate for power users is \$100–200/month
Business	\$20–25/user/mo	Cloud features, GitHub, Slack and Linear integrations, admin controls
Enterprise	Custom	Shared credit pool, no fixed rate limits
API key	Pay per token	For building products on Codex; newer models arrive later than on subscriptions

A practical note from people who run it daily: pay-as-you-go via API lands around \$5–15 a month for an hour a day, \$30–60 for a few hours, and \$80–250 for full-time agentic work. The subscription is the better deal until you are well past an hour a day.

Codex vs Claude Code

The comparison everyone makes, so here it is without the fan wars.

	Codex	Claude Code
Where it started	Cloud sandbox, parallel tasks	Your terminal, then desktop, cloud and phone
Model	GPT-5.6 family, GPT-6	Claude Fable, Opus, Sonnet
Entry price	\$20 ChatGPT Plus	\$20 Claude Pro
Middle tier	None — \$20 then \$100/\$200	\$100 Max
Strength	Backend and multi-task parallel work; review bot in GitHub	Long careful sessions; leads most agent benchmarks through 2026
Adoption	Huge via ChatGPT's user base	Overtook Codex in developer preference in under a year
Open source	The CLI, yes	No

Most developers who use both describe Codex as better for firing off several backend tasks and reviewing the results, and Claude Code as better for one hard problem it has to think through carefully. At \$20 a month each, plenty of people simply run both.

The catches

- **Full-access mode had a file-deletion bug.** In July 2026 Codex running without a sandbox deleted files it should not have: a failed environment-variable expansion turned a cleanup step into a delete against the user's home directory. OpenAI patched it in August. Sandboxing is still the safer default on a repo you care about.
- **Shallow visibility into its reasoning.** You see the diff and a summary, less of the why. Claude Code shows more of its working, which matters when the change is wrong.
- **The five-hour window bites.** Plus users hit the cap mid-afternoon on busy days. The workaround is a cheaper model for routine tasks, or the \$100 tier.
- **No hosting.** Codex writes and tests; it does not deploy. Vercel, Replit and the rest are still your job — or Polsia's, which uses these agents as its developer.
- **Model churn.** GPT-5.5 retires in October; the flagship tier changed twice this summer. Pin nothing.

Do you need to be a developer?

Less than you did, more than for Claude Code. The ChatGPT tab and the desktop app removed the terminal, and a founder who can read a diff and run a test can get real software out of Codex. But it assumes a Git repository and a codebase to work in, and it does not deploy, so the surrounding plumbing is still on you. If you have never shipped anything, start with Claude Code's desktop app or a platform like Polsia that wires the pieces together; if you already have a repo and a ChatGPT plan, Codex is one tab away.

Why it matters

Codex is the coding agent with the widest possible front door. Claude Code proved the category and still leads it on quality; Codex put the category inside the app that hundreds of millions of people already use, then made it the developer behind ChatGPT Work. For the agentic economy the consequence is simple: writing software is now a commodity capability available on a \$20 plan from two companies, and the scarce skill is knowing what to build.

Glossary

Codex — OpenAI's autonomous coding agent (2025 onward). Cloud sandboxes, parallel tasks, GitHub review, a tab in ChatGPT. Not the 2021 model of the same name.

Cloud sandbox — An isolated environment where Codex clones your repository and works without touching your machine.

Pull request (PR) — A proposed change to a codebase that a human reviews before it is merged. Codex's standard output.

Codex CLI — The open-source terminal version, Apache 2.0 licensed.

Credits — ChatGPT's usage currency since April 2026, mapped to tokens. Codex tasks draw them down.

Five-hour window — The rolling period over which Codex usage on Plus and Pro is capped.

Full-access mode — Running Codex with unrestricted permissions on your files. Convenient; the source of its worst reported bug.

Claude Code — Anthropic's coding agent and Codex's direct rival. See its own guide.

Codex is one of two coding agents on a \$20 plan. Here is the other, and what sits around them.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

wtfagents.com/guides/chatgpt-work

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF are the Picks and Shovels?

Whoever wins the agent race, the same handful of companies get paid: the chips, the clouds, the model hub, the code host, the payment rails, the power. Who they are, how the money flows, and what it means for anyone building on top of it.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

In a gold rush the reliable money is in picks and shovels. In the agentic economy the picks and shovels are chips, data centres, cloud capacity, electricity, and the plumbing that models, agents and money run through — and a handful of companies own nearly all of it.

Every agent in this series runs on this layer. Every dollar Polsia, OpenClaw or ChatGPT Work earns passes through it. When people ask whether AI is a bubble, this is the layer they are really asking about.

By the numbers — September 2026

\$725B+

combined 2026 capital spending by Amazon, Alphabet, Microsoft and Meta, up around 77% on 2025 — the largest infrastructure build-out in corporate history

\$5T

Nvidia's market value when it became the first chip company past five trillion dollars, 24 April 2026

\$75B

Nvidia's data-centre revenue in a single quarter (to late April), up over 90% in a year

\$12.93B

what Nvidia agreed on 2 September 2026 to pay for Hugging Face — the largest company it has ever bought outright. Expected to close in the first half of 2027

82%

Google Cloud's year-on-year growth in Q2 2026; AWS 37%, its fastest in eighteen quarters

\$514B

Google Cloud's contract backlog; AWS \$496B; Oracle \$638B — capacity sold years ahead

10–50×

more tokens an agent consumes per finished task than a single chatbot answer

2026

the year spending on running models (inference) is expected to overtake spending on training them

The layer, top to bottom

The chips — Nvidia and the challengers

Every frontier model was trained on Nvidia hardware and most are served on it. Nvidia holds roughly four fifths of the AI accelerator market, with gross margins other chipmakers cannot approach, and its CUDA software is the switching cost that keeps customers in place. It is the single most valuable company in the world for a simple reason: whoever wins the model race, they bought the chips from Nvidia.

The challengers are not AMD. They are the clouds' own chips — Google's TPUs, Amazon's Trainium, Microsoft's Maia — which already handle a meaningful share of inference inside their own data centres at a lower cost per token. And in China, Huawei's Ascend, which Beijing is steering its labs toward at the price of a slower next generation. Nvidia's risk is not losing the market; it is the market's biggest customers building their own supply.

Its answer is to own more of the stack. In December 2025 it paid around \$20 billion for Groq's assets and people, its largest transaction ever, buying the low-latency inference chips that threatened it. On 2 September 2026 it agreed to buy Hugging Face for \$12.93 billion — the largest company it has ever bought outright, and a straight line from the silicon to the place the models are published. Nvidia now sells the chips, and owns the shelf the open-weight world puts its models on.

The clouds — AWS, Google Cloud, Azure, Oracle

The three hyperscalers plus Oracle rent the chips, the power and the buildings to everyone who does not own them, which is nearly everyone. They also serve the models: Claude runs on AWS Bedrock, Google Vertex AI and Microsoft Foundry; GPT on Azure; Gemini on Google. An enterprise that already pays AWS can add Claude to the bill without a new vendor, which is how most of the Fortune 500 adopted agents.

The clouds are also the biggest investors in the labs: Amazon and Google in Anthropic, Microsoft in OpenAI. The money goes round in a circle — the lab raises from the cloud, spends it on the cloud's chips, and the cloud books the revenue. Whether that circle is a flywheel or a bubble is the question the 2026 earnings season has been trying to answer, and Q2 said flywheel: AWS grew 37%, Google Cloud 82%, Azure in the forties, all capacity-constrained.

The power — the constraint nobody planned for

Microsoft has reported around \$80 billion of Azure orders it cannot fulfil because the chips are in inventory waiting for electricity — its chief executive's words, not a critic's: "you may actually

have a bunch of chips sitting in inventory that I can't plug in." Data centres now compete for grid connections, gas turbines and nuclear restarts. Power is the new bottleneck, and utilities, turbine makers and the nuclear industry have become AI infrastructure companies without changing what they do.

The model hub — Hugging Face

When DeepSeek, Alibaba or Meta publish model weights, they land on Hugging Face. It is the GitHub of models: the place open weights are hosted, downloaded, fine-tuned and compared, and the reason a solo builder can run Qwen on a laptop. Over a billion Qwen downloads happened here. More than 18 million people use it to share over 3 million models, half a million datasets and a million applications. Nothing in the open-weight world works without it.

Which is why its sale matters. On 2 September 2026 Nvidia agreed to buy it for \$12.93 billion, subject to regulatory approval and expected to close in the first half of 2027. Nvidia has committed to keeping the platform open: its own compute will not be required to build or deploy through it, other chipmakers stay supported, and anyone can go on publishing and downloading what they like. Those are commitments, not yet a track record. The neutral ground of the open-model world is being bought by the company that sells the hardware it runs on, and the next few years will show what that is worth.

The code host — GitHub

Where the software lives and where the agents commit it. Claude Code and Codex write pull requests to GitHub; OpenClaw and Paperclip's star counts are GitHub's; Copilot, Microsoft's coding agent, is built into it. As agents write a growing share of all code, GitHub becomes the ledger of the agentic economy's work.

The payment rails — Stripe

Agents that can pay run through Stripe. Polsia companies take payment through it, Muse checks out with one-time Stripe cards, NanoCorp builds a Stripe account into every business it launches, and the guides you are reading were bought through it. Agentic commerce — agents buying, selling and paying each other — is a Stripe product before it is anyone else's.

The connective tissue — MCP, A2A, the foundation

Two open standards under the Linux Foundation's Agentic AI Foundation let any agent use any tool ([MCP](#)) and talk to any other agent (A2A). They are the picks and shovels nobody sells: free, vendor-neutral, and the reason the rest of the stack interoperates. See the [WTF is an API](#) guide.

How the money flows

Follow one \$7 guide sale on wtfagents.com and you have followed the whole economy.

A reader pays through Stripe. The site runs on Vercel, which runs on AWS. The guide was drafted with Claude, served from Anthropic's infrastructure on Google and Amazon clouds, on Nvidia chips, in data centres drawing power a utility built capacity for. The site's code sits on GitHub. The live stats on the homepage come from Polsia's API, whose agents run on Claude, on the same chips.

Scale that to every agent task in the world — ten to fifty times the tokens of a chatbot answer, millions of times a day — and the capital spending numbers in the box stop looking irrational. Inference, not training, is now the growth engine, and every agent run is a small payment to this layer.

The bubble question, honestly

The spending is unprecedented; the revenue is real and accelerating; the two are not yet in balance.

The case for a bubble: \$725 billion of capex in one year, free cash flow at the big four projected to fall sharply as spending outruns revenue, most enterprise agent projects still not in production, and a circular flow of money between clouds and labs that flatters everyone's numbers.

The case against: backlogs in the hundreds of billions that customers have contractually committed to, every provider reporting demand above capacity into 2027, Nvidia's data-centre revenue nearly doubling year on year, and the shift to inference, which is recurring, not one-off. Agents consume tokens every day, forever, in a way that training runs never did.

The reading most careful analysts land on: the infrastructure is being built ahead of the revenue, as railways and fibre were, and some of it will be stranded. The companies in this guide will still own the tracks.

What it means for you

→ **Your costs are set here.** Token prices fall when this layer gets cheaper and rise when it is constrained. The 2025 price collapse in older GPUs is why running an agent went from hundreds of dollars a day to a few.

- **Your outages happen here.** When a cloud region or a model provider has a bad day, every agent on it stops. Multi-model, multi-cloud is not paranoia for anything that matters.
- **Your independence lives here — partly.** Open weights you have downloaded, code you have cloned, hardware you own: those cannot be switched off. The hubs they came from — Hugging Face now Nvidia's, GitHub Microsoft's — can change terms. Cache what you depend on.
- **Your customers arrive here.** If you sell into enterprises, being available inside their cloud's marketplace matters more than any feature.

Glossary

Picks and shovels — The suppliers who profit from a boom regardless of which participant wins. From the Gold Rush.

Hyperscaler — A cloud provider operating at global scale: Amazon (AWS), Microsoft (Azure), Google (Google Cloud), and increasingly Oracle.

Capex — Capital expenditure: money spent building long-lived assets, here data centres and chips.

Accelerator — A chip built for AI workloads: Nvidia GPUs, Google TPUs, Amazon Trainium, Microsoft Maia, Huawei Ascend.

CUDA — Nvidia's software platform. The reason switching away from Nvidia chips is hard.

Inference — Running a trained model to produce answers or actions. Now the majority of AI compute spending.

Training — Building a model in the first place. Enormous, one-off, and no longer the biggest line item.

Backlog — Revenue customers have contractually committed to but the provider has not yet delivered. A measure of demand sold ahead.

Hugging Face — The hub for open-weight models: hosting, downloads, fine-tunes.

Bedrock / Vertex AI / Foundry — The model marketplaces inside AWS, Google Cloud and Azure respectively.

Agentic commerce — Agents buying, selling and paying. Runs on Stripe and its peers.

This is the layer everything else runs on. Here is what runs on it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

wtfagents.com/guides/whos-who

WTF is Happening with China's AI

DeepSeek, Qwen, Kimi, GLM and the open-weight strategy. Who the Chinese labs are, why they give models away, and what it means for anyone running agents.

wtfagents.com/guides/china-ai

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/guides/which-ai

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

Launch an AI-Run Company in a Weekend

Half a million AI-run companies have been launched this year and almost none of them sell anything. This is the weekend plan that fixes that: validate first, then pick the platform, then brief the agents. No code. One idea. Two days.

The one-liner

Any of the autonomous-company platforms will build you a live business in twenty minutes. None of them will find out whether anyone wants it. This guide puts the finding-out first, then the building, so that by Sunday night you have a company that is either selling or is a cheap, fast, honest failure.

The platforms' own numbers explain why the order matters. Polsia's founder says one company in ten on his platform has earned a dollar. NanoCorp publishes its leaderboard and the top company is in the low thousands. The agents build fine. The gap is demand, and demand is the part you do on Saturday.

By the numbers — September 2026

~6%

of the half-million companies created on Polsia are still active

1 in 10

has ever earned a dollar, per Polsia's founder

20 min

from one sentence to a live site, app, domain and Stripe account on NanoCorp or Polsia

\$20–30

the entry price on either managed platform; Paperclip is free plus model usage

5

real conversations with potential customers this guide asks for before you build anything

7 days

how long to give the launched company before deciding to double down or kill it

Saturday — validate before you build

The whole weekend hinges on Saturday. Skip it and you are the factory worker in the Rest of World story: a polished company, fabricated reviews, ad spend, zero customers.

Morning — pick one problem, not one idea

Write down a problem a specific kind of person has, that they already spend money or real time on. Not "an app for dog owners." Instead: "small dog groomers lose bookings because they answer the phone between clients." The narrower the person, the easier the next two steps.

Two tests before you continue. Can you name five real people who have it? Would they pay \$20 a month, or \$100 once, to make it go away? If either answer is no, pick another problem. This step should take an hour and it is where most weekends should end.

Midday — talk to five people

Message five people who have the problem. Not friends being kind — people in the group. A forum, a Facebook group, a subreddit, a WhatsApp community, a trade association, a LinkedIn search. Ask three things: how do you handle this now, what does it cost you, what have you tried. Do not pitch. Listen for the words they use; you will give those words to the agents on Sunday.

Three of five say it is a real pain and describe what they do about it: continue. Fewer: change the problem, not the plan.

Afternoon — write the one-page brief

Everything the agents will need, in plain English, in the customer's own words:

- **Who** — the person, precisely. "Independent dog groomers in the US with one to three staff."
- **Problem** — in their words, from the calls.
- **Offer** — the smallest thing that solves it. One feature, not a platform.
- **Price** — the number you tested. If you did not test one, \$29 a month or \$49 once are safe starting points for a small-business tool.
- **Proof** — one line about why you, or why now.
- **Not this** — three things the agent must not do: no fabricated testimonials, no cold email to people who did not opt in, no spend above the cap.

Evening — the pre-sale test

Before any platform, put up the cheapest possible test: a single page with the offer, the price and a button. A free page builder, a Stripe payment link, or even a Google Form that says "reserve your spot." Send it to the five people and to the group you found them in. If two people click through to pay or reserve by Sunday morning, you have a business worth building. If nobody does, you have saved a month.

Sunday — choose the platform and launch

Which platform

Three real choices, and the honest differences:

	Polsia	NanoCorp	Paperclip
Setup	None	None	Real: a server, a database
Cost to launch	\$20–50/mo + credits	Free trial, then \$30/mo	\$0 + model usage
Fees on sales	3% + 20% on ad spend; \$500/mo withdrawal cap	20% on withdrawals	None
Control	Low; their infra, their agents	Low; full autonomy, no gates	Total; your agents, your rules
Transparency	Live platform stats	Publishes every company's revenue	It is your data
Best for	Fastest hands-off launch	Honest fast experiments	Anyone who wants to own it

For a first weekend: Polsia or NanoCorp. You want speed and you want to spend your attention on customers, not servers. Paperclip is the right second company, once you know the idea works and want to keep the margin and the control. The full trade-offs are in each platform's own guide.

Morning — launch

Paste the brief. Do not press "surprise me"; that button is the platform choosing an idea with no demand signal and it is where the failed companies come from. Let the agents build the site, the product and the payment flow. Read what they produce against the brief. Fix the copy where it stopped sounding like your customers.

Set the caps before anything else: ad budget at zero for the first week, credits at the plan minimum, outreach off unless you have a list of people who asked to hear from you.

Afternoon — the first ten customers, by hand

The agents will offer to run ads and send cold email. Say no for now. Your first ten customers come from the five conversations and the group you found them in. Send the live link personally, with the same three questions you asked yesterday. This is the one part of an AI-run company that a human still does better, and it is the part that decides whether the company lives.

Evening — set the report

Every platform sends a daily summary. Tell the agent what you want in it: visitors, signups, payments, replies, and one question it needs you to answer. Then close the laptop. The point of the weekend was to reach this moment: a real company, a real offer, real people who have seen it, and a machine that will run it while you sleep.

The week after — decide honestly

Give it seven days and one rule: the company earns money or a clear "no" from real people, and either is a good outcome.

- **Two or more paying customers:** double down. Turn on a small ad budget with a hard cap, let the outreach agent contact people who match your first customers, and start a second week.
- **Interest but no payments:** the offer or the price is wrong, not the company. Change one thing, message the interested people, give it seven more days.
- **Nothing:** kill it. Cancel the plan, export what you can, and run the next problem through Saturday. A dead company after one week and \$30 is exactly what this process is for.

Most people who "fail" at AI-run companies never reach this decision. They launch on Sunday, watch the agents work for a month, pay for ads that reach nobody, and quit at \$200. Deciding on day seven is the discipline the platforms cannot give you.

The mistakes that kill weekend companies

Surprise me. The AI picks the idea. Nobody validated it. It builds beautifully and sells nothing.

Ads before customers. Ad spend is where the money goes on every platform and the fee structures reward it. Ten hand-won customers first; ads to find more people like them second.

Trusting the daily report. "Task complete" means the agent thinks it finished. Click the link. Send the test email. Buy your own product once.

Not reading the fees. A \$500 monthly withdrawal cap or a 20% cut changes what "working" means. Know the number before the first sale.

Building the platform, not the feature. Agents will happily add a dashboard, an onboarding flow and an integrations page. Your customer asked for one thing. Ship that.

Waiting. A month of watching is not patience; it is avoidance. Seven days, then a decision.

Glossary

Validation — Confirming real people want a thing before building it. The Saturday of this guide, and the step every platform skips.

Pre-sale test — Offering a product for sale before it exists, to measure demand. A page, a price, a button.

Brief — The one-page document the agents build from: who, problem, offer, price, proof, and what not to do.

Surprise me — The platform option where the AI chooses your business idea. Avoid.

Cap — A hard limit on what an agent may spend: ad budget, credits, outreach volume. Set before launch.

Withdrawal cap / fee — The limit or cut a platform applies when you take out revenue your company earned. Polsia: \$500 a month; NanoCorp: 20%.

Daily report — The summary every platform's agents send. Read it; do not trust it.

Kill — Ending a company that did not sell. Cheap, fast, and the normal outcome of a good process.

This is the how-to. Here are the platforms and the thinking behind them.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/guides/polsia

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

wtfagents.com/guides/solo-stack

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

The Solo Founder's AI Stack

Agents run the company. These tools make the things it sells and says: the voice, the video, the images, the decks, the automations. The seven seats a one-person business fills with software in 2026, what each costs, and the order to buy them in.

The one-liner

The rest of this series is about agents that do work. This guide is about the other half of a one-person company: the tools that produce the voice, video, images, documents and plumbing a business needs to show up in the world.

The right way to think about them is as seats. A ten-person company has a voice-over artist, a video editor, a designer, a researcher, an operations person, a deck-builder and a developer. In 2026 a solo founder fills each of those seats with software for less than one of those people costs per day.

By the numbers — September 2026

7

seats a solo operation fills with software: voice, video, image, deck, research, automation, code

\$30–60

per month for a lean solo stack; \$150–300 for a small-team stack, per the operators who publish theirs

~\$90

what one five-day-a-week newsletter business reports spending on its whole AI stack, at \$25,000 a month in revenue

175+

languages a single video can be dubbed into with matching lip movement, in one pass

15+

generation models behind one all-in-one video platform, on one credit balance

15–20 hrs

per week solo operators typically report saving versus their pre-AI workflow

The seven seats

1. The voice — ElevenLabs

The standard for synthetic speech and the one tool on this list almost every solo creator names first. Text-to-speech that crosses the "I cannot tell" line for long-form audio, voice cloning from a short sample, and dubbing into dozens of languages with the timing preserved. For a solo creator it replaces a voice-over artist, a localisation agency and a sound designer. Starter is \$6 a month; the Creator tier at \$22 is where the quality most people mean by "ElevenLabs" begins. Its video and image features are secondary; buy it for the audio.

2. The video — Higgsfield, Runway, HeyGen

Three different jobs. Higgsfield is the all-in-one: fifteen-plus generation models, lip-sync, colour grading and ElevenLabs voice on a single credit balance, for people who want to make short video without juggling five subscriptions. Runway is the VFX seat — generation and editing with the most control, for anyone whose video is the product. HeyGen is the avatar seat — script in, a presenter on screen talking, dubbed into more than 175 languages with the lip movement matched, for explainers, sales videos and support content at scale. Most solo founders need one of the three, not all.

3. The images — a generation model plus a design tool

Image generation is now inside everything: ChatGPT, Gemini, Higgsfield, Canva. The practical seat is the design tool that turns a generated image into a thumbnail, an ad, a product mock-up or a social post with your brand on it. Canva remains the default for non-designers; Figma for anyone who wants precision. Midjourney is still the quality benchmark for standalone generation if you want a dedicated subscription.

4. The deck and the document — Gamma, plus your agent

Gamma turns a brief into a presentation, a one-pager or a landing page in minutes and, since its API went generally available, can generate them automatically from a CRM or a form. Cowork and ChatGPT Work produce finished decks and documents as a native output, which for many solo founders makes Gamma optional. Pick one; do not pay for both.

5. The research — Perplexity, or your agent's deep-research mode

Sourced, current answers with citations, for fact-checking, market research and competitor analysis. Perplexity Pro is the specialist; Claude, ChatGPT and Gemini each have a deep-research mode that covers most of the same ground on a plan you already pay for. The seat matters more than the brand: every claim you publish should have been checked by something that cites its sources.

6. The plumbing — Zapier or n8n, plus Lindy

Automation is what turns tools into a system. Zapier is the easiest: when a form is submitted, send it to Claude for classification, route it to the right database, notify you. n8n does the same, open-source and self-hosted, for anyone who wants to own the workflow. Lindy sits a level up: agents for inbox triage, meeting follow-ups and CRM hygiene that run without a trigger. This is the seat that stops a solo founder being the integration layer between all the

others.

7. The code — Claude Code or Codex

Covered in their own guides and listed here because the solo stack is not complete without it. A working tool, a landing page fix, a scraper, an internal dashboard: things that used to need a developer now come from a \$20 plan and a clear brief.

The stack at three budgets

Budget	Monthly	What is in it
Lean	\$30–60	One \$20 AI plan (Claude or ChatGPT, with its agent, deck and research modes), ElevenLabs Starter, Canva free, Zapier free tier
Working	\$100–150	The above plus ElevenLabs Creator, Higgsfield or HeyGen, Canva Pro, a paid Zapier or n8n
Producing	\$200–300	The above plus a second AI plan, Runway, Perplexity Pro, Lindy, a \$100 agent tier for heavy use

Start lean. Every seat above the first row is a seat you should be able to name the output of. "I make three videos a week" earns a video tool. "I might make videos" does not.

The order to buy in

1 — The \$20 AI plan first. It fills three seats — research, documents, code — and gives you an agent. Nothing else on this list matters until this is in daily use.

2 — Voice second, if you publish anything spoken. Podcasts, video, courses, multilingual anything. The gap between synthetic and real closed in 2025; the gap between good and free did not.

3 — Automation third. The moment you find yourself copying data between two tools, that is a Zap. Ten minutes to build, hours a month saved.

4 — Video only when you know the format. Avatar explainers, generated short-form, edited footage: three different tools. Pick after you have made ten videos the hard way and know which kind you make.

5 — Everything else when the output demands it. A design tool when the brand needs consistency. Research pro when citations are the product. A second AI plan when the first one's cap is the bottleneck.

The catches

- **Subscription creep.** Seven tools at \$20–50 each is a real salary. The operators who publish their stacks all say the same thing: consolidate, and audit quarterly.
- **Credit maths.** Video and voice platforms bill in credits, and dubbing, cloning and long generations consume them fast. Overage bills above the plan price are routine. Set the cap.
- **Voice rights.** Cloning a real person's voice needs their consent, documented. Your own voice is yours; a client's or a colleague's is not.
- **Generated does not mean good.** The tools make a competent draft of anything. The taste to know which draft to publish is still the solo founder's job, and the reason the business is worth anything.
- **The all-in-ones move fast.** Platforms that bundle fifteen models change models, prices and credits monthly. Do not build a workflow you cannot move.

How it fits with the agents

The agents in this series run the business: they research, code, email, sell, support and report. The stack in this guide is what the business shows to the world. In practice they meet in the middle. Cowork or a Polsia agent writes the script; ElevenLabs voices it; Higgsfield or HeyGen makes it a video; Zapier posts it and logs the result; an agent reads the numbers and drafts next week's script. One person, one afternoon a week, a production line.

Glossary

Stack — The set of tools a business runs on. A solo stack is one person's.

Seat — A role in a company, filled here by a tool rather than a hire.

Text-to-speech (TTS) — Generating spoken audio from written text. ElevenLabs is the benchmark.

Voice cloning — Reproducing a specific person's voice from a sample. Needs consent for anyone but yourself.

Dubbing — Translating spoken content into another language with matching audio, and increasingly matching lip movement.

Avatar video — A generated on-screen presenter reading a script. HeyGen, Synthesia.

Credits — The usage currency of most video and voice platforms. Watch the balance.

All-in-one platform — A service bundling many generation models behind one interface and one bill. Higgsfield is the best-known.

Automation — Software that moves data and triggers actions between tools. Zapier, n8n, Lindy.

Deep research — The mode in Claude, ChatGPT and Gemini that reads many sources and produces a cited report.

The stack makes the things. The agents run the company. Here is the rest.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

Launch an AI-Run Company in a Weekend

The practical playbook: validate an idea on Saturday, launch it on Polsia, NanoCorp or Paperclip on Sunday, and know by Friday whether it sells.

wtfagents.com/guides/launch-weekend

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/guides/ai-agent

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/guides/api

WTF is Higgsfield?

Fifty video and image models behind one chat agent, \$700 million in annual revenue, a \$5.4 billion valuation — and this week, an open API, eighteen months of open-sourced code and a \$50,000 prize for whoever builds something with it first. What it is, what it costs, and why the open move matters.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Higgsfield is an AI video and image platform that puts fifty-plus generation models behind one interface and one credit balance, and increasingly behind one chat agent: describe the ad, the clip or the app you want, and Higgsfield chooses the models, runs them and hands back the result.

It is the tool most short-form creators and social media marketers actually use to make AI video, and in September 2026 it did something the closed labs have not: it opened its API to anyone, pay per generation, and a team member open-sourced a year and a half of the code that runs it.

By the numbers — September 2026

\$5.4B

valuation after the \$400M Series B led by DST Global, 17 August 2026

\$700M

annualised revenue at the time of the round, up from a \$1.3B valuation in January

361M

visits to the platform between January and July 2026

50+

image and video models available through the new API, opened 16 September

18 months

of the studio's code open-sourced by a Higgsfield engineer the same week

\$50,000

prize offered for practical builds on the API submitted within seven days

~\$0.50

cost of five seconds of Seedance 2.5 video through the API, per the open-source client's published rates

What it actually does

Higgsfield does not train its own frontier models. Its product is the layer on top: it licenses or integrates the best video and image models as they appear — Seedance, WAN, FLUX and dozens more — and wraps them in tools a creator can use without knowing which model is which.

Text and image to video — Short clips from a prompt, a photo or a product shot, with camera moves, presets and motion controls that were the company's first hit.

Character continuity — Face-Lock and its successors keep the same face, product or location consistent across shots, the problem that made early AI video useless for brands.

Genjutsu — Launched 1 September, on Product Hunt two days later: take the motion from any reference video and recast it with your own characters, locations and products.

Cinema Studio and Academy — A multi-shot editing workflow and a course library, aimed at people producing finished pieces rather than single clips.

Audio — Text-to-speech, voice change and lip-sync in the same credit balance, so a clip can leave with sound.

Supercomputer — The agent. Launched in May, it runs the whole pipeline from one chat: describe the piece, it picks the models, chains them and returns the result. Since summer it also builds and publishes whole generative apps — describe an app that makes images, video or 3D, and it wires the models in itself. No API keys, no code.

Similarity Scoring — For Team plans: flags likeness to known characters, celebrities and logos before delivery. A warning system, not legal clearance.

The pattern is the one running through this whole series: the models became commodities, and the value moved to the agent that chooses and chains them.

The September moves — API, open source, prize

Three things happened in one week, and together they change what Higgsfield is.

The API

On 16 September Higgsfield opened its API: over fifty image and video models, one pay-per-use balance in US dollars, separate from the consumer subscriptions. Any developer can now put Higgsfield's generation inside their own app, website or automated workflow. Until this week that layer was reachable only through Higgsfield's own interface.

The open source

The same week a Higgsfield engineer, Alisher S, open-sourced eighteen months of code — the creative-studio layer that sits on the API. A community client, Higgsfield-Open, followed within days: 38 of the models, exact per-generation pricing shown before you press go, and no platform markup. An independent project, OpenHiggsfield, had already been circulating as a free alternative with forty-plus models.

For a creator this means the interface is no longer the moat; the model access and the credit balance are. For Higgsfield it means the company is betting, as the Chinese labs did, that owning the ecosystem beats owning the front end.

The prize

Fifty thousand dollars for practical builds on the API submitted within seven days of launch. A small sum against a \$5.4 billion valuation, and a loud signal about what the company wants next: developers, not just creators.

What it costs

Route	Price	Notes
Free	\$0	Limited credits, watermarked, enough to learn the interface
Plus / Pro / Ultimate	Monthly, tiered	Credits per month; higher tiers unlock more models, longer clips, higher resolution and "unlimited" windows on selected models
Team	Custom	Seats, shared assets, Similarity Scoring, brand controls
API	Pay per generation	Separate balance in USD; roughly \$0.50 for five seconds of Seedance 2.5, about a cent for a Soul 2 image

Higgsfield's consumer pricing changes often and runs promotions constantly, so the page is the source of truth. The rule that does not change: everything is credits, video eats them fast, and an "unlimited" offer is always scoped to particular models at particular settings. Read the scope before you subscribe.

Higgsfield vs Runway vs HeyGen

	Higgsfield	Runway	HeyGen
What it is	Many models, one interface, one agent	One lab's own models and editor	Avatar presenters and dubbing
Best for	Short-form, ads, social, volume	Cinematic control, VFX	Explainers, sales video, localisation
Own models	No — integrates others'	Yes	Yes, for avatars
API	Yes, since September	Yes	Yes
Open code	Studio layer open-sourced	No	No
Price shape	Credits + pay-per-use API	Credits	Credits

Runway is the studio for people whose video is the product. HeyGen is the presenter. Higgsfield is the assembly line for people who need a lot of good-enough video fast, which describes most marketing.

The catches

- **Credits burn.** Video generation is expensive compute and every platform bills accordingly. Budget by finished clip, not by month, and set caps.
- **It rents its models.** Higgsfield's strength is aggregation; its exposure is that any model it offers can be pulled, repriced or offered elsewhere. The API and open code are the answer to that, and a bet.
- **Rights are your problem.** Similarity Scoring flags likeness; it does not license it. Generating a celebrity, a brand or a copyrighted character remains your legal risk, in every country you publish in.
- **Open-sourced is not open.** The studio code is open; the models and the credit system are not. You still pay Higgsfield to generate.
- **Fast-moving means unstable.** Genjutsu, Supercomputer's app builder, the API and the open-source drop all landed in four months. Features and prices in this guide will move; the shape will not.

Who it is for

Yes — Solo founders, creators and marketers who make short video weekly and want one bill, one interface and the newest models without tracking which lab shipped what.

Not yet — Anyone whose brand cannot risk a likeness or rights problem without legal review, or who needs frame-level control (that is Runway).

No — Anyone who just needs a talking presenter (HeyGen) or occasional clips (the free tier of anything).

Why it matters

Higgsfield is the clearest example in this series of a company winning by owning the layer between the models and the user. It trains nothing at the frontier and is worth \$5.4 billion for choosing, chaining and packaging what others train. In September it took the next step: open the API, open the studio code, pay developers to build. That is the Chinese open-weight playbook applied to creative tools, and it is what an agentic economy looks like when the agent, not the model, is the product.

Glossary

Higgsfield — AI video and image platform aggregating 50+ models behind one interface, one credit balance and one agent. San Francisco; \$5.4B valuation.

Supercomputer — Higgsfield's chat agent that runs the whole creative pipeline and builds generative apps from a description.

Genjutsu — Motion recasting: apply the movement of a reference video to your own characters and products. Launched September 2026.

Face-Lock — Character continuity across shots; the same face or product stays consistent.

Similarity Scoring — Team-plan tool that flags likeness to known characters, people and logos before delivery.

Higgsfield API — Pay-per-generation access to the platform's models, opened 16 September 2026.

Higgsfield-Open / OpenHiggsfield — Open-source clients for the API and a free alternative interface, respectively.

Seedance, WAN, FLUX — Third-party video and image models Higgsfield integrates. Names change quarterly.

Credits — The usage currency. Video consumes many; images few.

Aggregator — A platform that packages other companies' models rather than training its own. Higgsfield's model.

Liked this? Go deeper.

Higgsfield makes the video. Here is the stack around it, and the agents that run the business.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

wtfagents.com/guides/solo-stack

WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

wtfagents.com/guides/meta-muse

WTF is ElevenLabs

The AI voice standard. Text-to-speech, cloning, dubbing and the phone agents behind Klarna and Revolut. What it does, what it costs, and the rights question.

wtfagents.com/guides/elevenlabs

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is ElevenLabs?

The voice of the agentic economy. Text-to-speech that crossed the "I can't tell" line, cloning from a short sample, generation in seventy languages — and now the voice agents answering the phone at Klarna and Revolut. What it is, what it costs, and the catches.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

ElevenLabs makes AI voices: text-to-speech good enough for audiobooks and film, cloning from a short sample, dubbing that keeps the speaker's voice in another language, and, increasingly, conversational agents that answer a phone line and hold the call.

For a creator it is the voice seat in the stack. For an enterprise it is the reason a support line can answer in thirty languages at 3am. For the agentic economy it is the layer that gives every agent in this series a mouth.

By the numbers — September 2026

\$11B

valuation at the \$500M Series D led by Sequoia, February 2026 — up from \$3.3B a year earlier

~\$22B

valuation reportedly under discussion for an employee share sale this summer, no new capital raised

\$500M

annual recurring revenue passed in April 2026, from \$350M at the end of 2025 — \$100M added in one quarter

60%+

of the Fortune 500 have employees using the platform, per the company — up from 41% in 2024

35M

US customers whose first-line phone support at Klarna is an ElevenLabs voice agent

70+

languages for generation on Eleven v3; dubbing runs about thirty

\$0 → \$299

per month: free tier of 10,000 characters through to the Scale plan; Starter \$6, Creator \$22

What it actually does

Text-to-speech — Type or paste text; get natural speech in a library voice or your own. The v3 models handle emotion, pacing and emphasis from context and from simple tags, which is what moved it past "robot reading."

Voice cloning — Instant cloning from a minute of audio; professional cloning from more, for a voice indistinguishable from the original. Your own voice, or anyone's who has consented in writing.

Dubbing — A video in one language becomes the same video in another, in the original speaker's voice, with timing preserved. Seventy-plus languages in one pass.

Sound effects and music — Describe a sound, get a sound. Background beds for video without a licence hunt.

ElevenReader — The consumer app: any text, article or book read aloud in a chosen voice. Spotify now uses ElevenLabs to let authors produce audiobooks directly.

ElevenAgents — The enterprise product and the growth engine: voice and chat agents that take calls, answer support, qualify sales leads and run hiring screens, with the testing, monitoring and integrations a call centre needs. Revolut runs support for four million customers on it across the UK and Europe; Klarna put it in front of 35 million Americans.

Two years ago ElevenLabs sold a better voice. Now it sells the agent that uses the voice, and that is where the revenue went.

What it costs

Plan	Price	What you get
Free	\$0	10,000 characters a month, attribution required, no cloning
Starter	\$6/mo	30,000 characters, instant cloning, commercial use
Creator	\$22/mo	100,000 characters, professional cloning, higher-quality audio. The tier most people mean by "ElevenLabs"
Pro	\$99/mo	500,000 characters, higher limits, usage-based overage
Scale	\$299/mo	2M characters for teams and studios
API	Per character	Roughly \$0.06–0.12 per 1,000 characters depending on model
ElevenAgents	Custom	Enterprise voice agents, priced per minute and by contract

The unit is characters, not minutes: a thousand characters is roughly a minute of speech. A weekly ten-minute podcast burns about 10,000 characters an episode; a full audiobook is a million or more. Dubbing and cloning consume faster. Annual billing takes about 17% off.

ElevenLabs vs the field

	ElevenLabs	Open-source models	Big-lab voices
Quality	The benchmark	Close, and closing	Good, less expressive
Cloning	Yes, instant and pro	Yes, self-hosted	Limited or none
Dubbing	~30 languages, one pass	DIY pipeline	Partial
Price	\$6-299/mo, API per character	Free to run, your hardware	Bundled in the plan
Agents	ElevenAgents, enterprise-grade	Build your own	Inside ChatGPT, Gemini, Muse
Lock-in	Voices live on the platform	None	With the lab

The open models (Fish Audio, Kokoro and others) are why ElevenLabs cannot raise its API prices much; the enterprise agents are why it does not need to.

The catches

- **Voice rights.** Cloning a voice needs documented consent from its owner. Your own is yours; a client's, a colleague's or a celebrity's is not. ElevenLabs verifies professional clones and faces biometric-privacy litigation in the US; the legal exposure for misuse sits with you.
- **Character maths.** Long-form audio eats plans fast, and overage on Pro and above is usage-billed. Estimate the month in characters before choosing a tier.
- **Your voices live there.** A professional clone is an asset on ElevenLabs' servers. Export the audio you make; the voice itself does not travel.
- **Detection is a moving target.** "Can't tell it's AI" cuts both ways. Platforms and regulators are adding disclosure rules for synthetic voice; label it where required.
- **The agent business is not the creator business.** Enterprise voice agents are where the company's attention and revenue are going. Creator features keep shipping, but the roadmap is a call centre's.

Who it is for

Yes — Anyone who publishes spoken content: podcasts, video voice-over, courses, audiobooks, multilingual anything. Buy the Creator tier and stop thinking about voice.

Not yet — Anyone whose use needs a voice they do not own the rights to, or who wants to run everything on their own hardware (open models).

No — Anyone who only needs occasional read-aloud (the free tier, or the voice built into their chat app).

Why it matters

Every agent in this series types. ElevenLabs is what lets them speak, and the moment an agent can hold a phone call it stops being software people configure and becomes staff people talk to. That is why the fastest-growing part of an audio company is customer support, and why a voice startup founded in 2022 is worth more than most airlines.

Glossary

ElevenLabs — AI voice company founded 2022 by Mati Staniszewski and Piotr Dąbkowski. Text-to-speech, cloning, dubbing, voice agents. \$11B valuation.

Text-to-speech (TTS) — Generating spoken audio from written text.

Voice cloning — Reproducing a specific person's voice from a sample. Instant (short sample) or professional (longer, verified).

Dubbing — Re-voicing spoken content in another language, preserving the original speaker's voice and timing.

ElevenAgents — The enterprise platform for voice and chat agents that take calls and hold conversations.

ElevenReader — The consumer read-aloud app.

Character — The billing unit. Roughly 1,000 characters per minute of speech.

Eleven v3 — The current generation of models, with emotion and expressiveness driven by context and tags.

Consent — Written permission from a voice's owner, required before cloning anyone but yourself.

Liked this? Go deeper.

ElevenLabs is the voice. Here is the rest of the creative stack, and the agents that use it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Higgsfield

The AI video platform every creator is talking about. What it does, what it costs, the September API and open-source moves, and who it is for.

wtfagents.com/guides/higgsfield

The Solo Founder's AI Stack

Beyond agents: the voice, video, image, deck, research and automation tools a one-person business runs on. Seven seats, real prices, what to buy first.

wtfagents.com/guides/solo-stack

WTF is HeyGen

The AI avatar video platform. Digital twins, one-prompt videos, dubbing into 175+ languages. What it does, the credit maths, and the consent question.

wtfagents.com/guides/heygen

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is HeyGen?

A presenter who never needs a camera, a script that becomes a finished video from one prompt, and dubbing into 175 languages with the lips to match. What HeyGen does, what the credits really cost, and when you should still film yourself.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

HeyGen makes videos of people who were never filmed. Type a script, pick a presenter — a stock avatar or a digital twin of yourself made from two minutes of webcam footage — and out comes a talking-head video with natural gestures, lip-sync and, if you want, the same video in 175 other languages.

It is the presenter seat in a solo founder's stack, and for marketing and training teams it replaced the studio booking. In 2026 it added the agent: one prompt, and it writes, voices, shoots and edits the whole thing.

By the numbers — September 2026

\$200M+

annualised revenue, announced by the company in June 2026 — doubled in eight months

100,000+

businesses using it, per the company; G2's fastest-growing product of 2025

175+

languages and dialects for translation with lip-sync

2 min

of webcam footage to create a digital twin of yourself

\$29 / \$49 / \$149

per month for Creator, Pro and Business; free tier of three watermarked videos

31–48 credits

per minute of Avatar IV or V video, against 600 credits on Creator — twelve to nineteen finished minutes a month

\$500M

last disclosed valuation, at the 2024 Series A led by Benchmark

What it actually does

Avatars — Five hundred-plus stock presenters, or your own. Avatar IV and the newer Avatar V read the emotional register of the audio and generate matching expressions, head turns and hand gestures; the stiff mannequin of two years ago is gone.

Digital twins — A custom avatar from two minutes of footage; a photo avatar from a single still. Your face, delivering scripts you never read aloud.

Video Agent — The 2026 headline. One text prompt becomes a complete video: script, visuals, presenter, voice-over, edit. Under four minutes to a finished, postable piece. Speed and precision modes for quick clips versus high-stakes content.

Translation — Upload a video; get it back in another language, in a matching voice, with the mouth movements re-synched. 175-plus languages, from about thirty source languages.

Interactive avatars — A presenter that answers questions live, for sales pages, kiosks and support. A separate LiveAvatar track, billed by the real-time minute, with extra slots at \$49 a month.

Brand kits, templates, SCORM — The corporate plumbing: consistent styling, reusable layouts, exports that drop into learning systems.

API and MCP — A standalone pay-as-you-go pool: buy the credits you want, no subscription, and they last twelve months. Studio video, translation and LiveAvatar, so other agents can make HeyGen videos on your behalf.

What it really costs — the credit maths

Every plan is a credit budget, and the credit rates are the part the pricing page does not put in the headline.

Plan	Price	Credits and what they buy
Free	\$0	3 videos/month, watermarked, one minute max
Creator	\$29/mo (\$24 annual)	600 credits, 1080p, videos to 30 min, voice cloning, unlimited photo avatars
Pro	From \$49/mo	1,000 credits, rising in tiers; 4K export, advanced models
Business	\$149/mo + \$20/seat	1,500 credits, collaboration, brand controls, SCORM
Enterprise	Custom	Studio-grade avatars, seats, security review, support

What you make	Credits per minute	On Creator (600)
Avatar III, video look	4	~150 minutes
Avatar III, photo look	7	~85 minutes
Avatar IV, photo look	16	~37 minutes
Avatar IV, video look	31	~19 minutes
Avatar V, either look	48	~12 minutes
Video Agent, standard	40	~15 minutes
Video Agent, Seedance	120	~5 minutes
Translation with lip-sync	6	~100 minutes
Translation, no lip-sync	4	~150 minutes

The plan says \$29. Nineteen minutes of Avatar IV is what \$29 buys, and twelve if you want Avatar V. A creator making twenty three-minute videos a month needs sixty finished minutes — roughly three times a Creator plan — so that is Pro at a higher tier, and they will still watch the meter.

The API has its own credit pool; web-plan credits do not cross over.

HeyGen vs Higgsfield vs ElevenLabs

	HeyGen	Higgsfield	ElevenLabs
The seat	Presenter	Generated video	Voice
Core output	A person talking to camera	Clips, ads, motion	Speech, dubbing
One-prompt agent	Video Agent	Supercomputer	ElevenAgents (calls)
Languages	175+ with lip-sync	Depends on model	70+ generation, ~30 dubbing
Entry price	\$29/mo	Credits, tiered	\$6–22/mo
Best for	Explainers, sales, training, localisation	Short-form, ads, social	Audio of any kind

Most solo founders need one of the three; a marketing team ends up with all of them, often with ElevenLabs voices inside HeyGen videos.

The catches

- **Credits, again.** The realistic avatars cost eight to twelve times the old ones per minute. Budget in finished minutes at the quality you will actually publish.
- **The uncanny gap is small, not zero.** Avatar IV and V are good enough for corporate content and most marketing; they are not a substitute for a founder on camera when trust is the product.
- **Consent is lighter than at rivals.** Security researchers have noted HeyGen's avatar-creation consent checks are less strict than Synthesia's. Cloning anyone but yourself needs their written permission; the platform's leniency is not your defence.
- **Disclosure rules are arriving.** Several jurisdictions and platforms now require labelling synthetic presenters in ads and political content. Label where required.
- **Your twin lives there.** A digital twin is an asset on HeyGen's servers; extra slots cost \$29 a month each. Export finished videos; the avatar does not travel.

Who it is for

Yes — Anyone producing structured, script-driven video at volume: explainers, product demos, course modules, sales outreach, multilingual versions of the same thing. Creator at \$29 if it is a few videos a month; Pro if it is a few a week.

Not yet — Anyone whose audience needs to trust a real face, or who is unsure of their monthly volume. Film ten videos yourself first, then decide.

No — Anyone who needs cinematic or generated footage rather than a presenter (Higgsfield, Runway), or only needs the voice (ElevenLabs).

Why it matters

HeyGen is what "AI-run company" looks like at the point of contact with a customer: a presenter that does not exist explaining a product that an agent built, in whatever language the viewer speaks, produced from one prompt in four minutes. The economics of video content changed in 2026, and this is the tool that changed them for the people who make it.

Glossary

HeyGen — AI avatar video platform: digital presenters, one-prompt videos, translation with lip-sync. Los Angeles; \$200M+ ARR.

Avatar IV / V — HeyGen's current photorealistic avatar models, with expression and gesture driven by the audio.

Digital twin — A custom avatar of a real person, made from about two minutes of footage.

Video Agent — One prompt to a finished video: script, visuals, presenter, voice, edit.

Video translation — Re-voicing a video in another language with the mouth re-synced. 175+ target languages.

Credits — The usage currency. Rates vary by feature and look: 4 a minute for an Avatar III video look, 31 for Avatar IV, 48 for Avatar V, 40 for the Video Agent.

Interactive avatar — A presenter that responds live to questions. Enterprise, per slot.

SCORM — The export format learning-management systems accept; a Business-plan feature.

Consent — Documented permission from anyone whose likeness you clone. Required regardless of what the platform enforces.

Liked this? Go deeper.

HeyGen is the presenter. Here is the rest of the creative stack.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Higgsfield

The AI video platform every creator is talking about. What it does, what it costs, the September API and open-source moves, and who it is for.

wtfagents.com/guides/higgsfield

WTF is ElevenLabs

The AI voice standard. Text-to-speech, cloning, dubbing and the phone agents behind Klarna and Revolut. What it does, what it costs, and the rights question.

wtfagents.com/guides/elevenlabs

WTF is Runway

The professional AI video studio. Gen-4.5, the film and agency work, the pivot to world models, and what it costs. For people whose video is the product.

wtfagents.com/guides/runway

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Runway?

The studio that made Hollywood's first AI shots, co-created Stable Diffusion, and is now betting \$315 million that video generation was only practice for simulating the world. What it does, what it costs, and why it stopped calling itself a filmmaking tool.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Runway is the professional end of AI video: a research lab and a studio that builds its own models — the Gen series — and the editing tools around them, used on real films, by real agencies, for shots that need to look intentional rather than generated.

Where Higgsfield aggregates fifty other people's models and HeyGen makes a presenter, Runway trains the model, and in 2026 it has started describing video generation as a means to a larger end: world models, AI that understands physics well enough to simulate an environment rather than paint a plausible frame of one.

By the numbers — September 2026

\$5.3B

valuation at the \$315M Series E led by General Atlantic, 10 February 2026 — up from \$3.3B ten months earlier

\$860M

total raised, with Nvidia, AMD, Adobe and Fidelity on the cap table

\$200M

annualised revenue as of 8 September 2026, doubled in five months — \$40M of it added in Q2 alone

300%+

net revenue retention, with the acceleration almost entirely enterprise

Gen-4.5

the current production model, out 1 December 2025: native audio, multi-shot, character consistency

Dec 2025

first world model shipped; a second promised for late 2026

2018

founded in New York by three NYU graduates; co-creators of Stable Diffusion in 2022

What it actually does

Gen-4.5 — Text or image to video with the physics and prompt adherence that put it above Google's Veo 3 and OpenAI's Sora 2 Pro on the Artificial Analysis leaderboard at launch. Newer models from ByteDance and Alibaba have since passed it there; Runway's answer is that the leaderboard is not what its buyers are choosing on. Native audio, longer multi-shot sequences, and characters that stay the same person across cuts.

Editing, not just generating — The tools that make it a studio: motion brush, camera controls, inpainting, green-screen removal, frame interpolation, upscaling. You direct the shot rather than accept it.

References and consistency — Feed it a character, a location, a product, and keep them coherent across a whole sequence. The feature that made agency work possible.

Aleph 2.0 and Act-Two — Edit an existing clip by describing the change: Aleph 2.0, shipped 21 May 2026 inside a new product called Edit Studio, does localised edits on up to thirty seconds of 1080p and can apply one edit across several shots. Act-Two is the performance-capture side, driving a generated character from a phone video.

World models — The December 2025 release and the 2026 roadmap: systems that model object permanence, gravity, cause and effect, so a falling glass shatters instead of morphing. Aimed at robotics, simulation and science as much as film.

API — Gen models available to developers for building products; how Higgsfield and others offer "Runway" inside their own tools.

The distinction that matters: Runway sells control. Everything in the interface exists so a person with taste can get the exact shot, not a random good one.

What it costs

Plan	Price	Notes
Free	\$0	125 credits, once, watermarked and lower resolution
Standard	\$15/mo (\$12 annual)	625 credits a month, watermark off, 1080p
Pro	\$35/mo (\$28 annual)	2,250 credits, more concurrent generations, higher resolution
Max	\$95/mo (\$76 annual)	9,500 credits, a month of rollover, no queue
Enterprise	Custom	Studios and agencies: security, custom models, volume
API	Per second of video	Pay per generated second, by model

Max replaced the old Unlimited plan on 1 June 2026, and the trade is worth knowing: four times the credits, but Explore Mode — the slow lane where generation was genuinely unlimited — does not come with it. Legacy Unlimited subscriptions keep their access until 30 November 2026, and then everyone is on credits.

Runway changes credit rates with each model, so read the pricing page before committing. The constant: credits are consumed per second of video, and newer models cost more per second than older ones. A ten-second Gen-4.5 clip is a meaningful slice of a Standard plan.

Runway vs Higgsfield vs the big labs

	Runway	Higgsfield	Google Veo / others
Who trains the model	Runway	Nobody — it aggregates	The lab
Control	Full: brush, camera, edit	Presets and an agent	Prompt, limited edit
Best at	Intentional cinematic shots	Volume, speed, social	Raw quality on a prompt
Buyer	Film, agencies, studios	Creators, marketers	Consumers, developers
Independence	Owns its stack	Depends on model licences	Hyperscaler-backed
Price shape	Credits per second	Credits + API	Bundled or API

The category split in one line: Runway for the shot you can describe precisely, Higgsfield for the fifty shots you need by Friday, HeyGen when the shot is a person talking.

The pivot — from filmmaking tool to world model

For six years Runway's story was Hollywood: co-creating Stable Diffusion, shots in Everything Everywhere All At Once, deals with Lionsgate and AMC. In 2026 the founders started saying something different in public: that generating coherent video at scale is training data for a model that understands the physical world, and that such a model matters for robotics, drug discovery and climate simulation more than for film.

Two things make the bet credible. Runway owns its models and its infrastructure and is not inside a hyperscaler, which it argues lets it move faster than Google's Veo and Genie or Meta's efforts. And \$315 million from General Atlantic, with Nvidia and AMD participating, is a war chest for exactly this. Two things make it risky: Google is attacking from both sides with Veo for video and Genie for simulation, and a company whose revenue is film and agency tools is now spending on a market that does not yet pay.

For a creator none of this changes the product today. It changes who Runway is building for next.

The catches

- **Credits per second.** The best model is the most expensive per second and the one you want. Budget by finished seconds — and note there is no longer an unlimited tier to fall back on.
- **Control has a learning curve.** The tools that make Runway a studio take time to learn. Higgsfield's presets get a marketer to a usable clip faster.
- **Consumer-scale competition.** Google's Veo ships inside products billions of people already use; Runway's edge is the professional layer, and that layer is where the giants aim next.
- **Rights and likeness.** As with every generator: references you upload need to be yours to use, and outputs resembling real people or copyrighted work are your risk.
- **The roadmap points away from you.** If world models become the company, creator features may become the funding source rather than the focus.

Who it is for

Yes — Filmmakers, agencies, designers and creators whose video is the product and who want to direct the shot: consistent characters, specific camera moves, edits on existing footage.

Not yet — Anyone who needs volume over precision (Higgsfield), or who will not spend the hours the tools deserve.

No — Anyone who needs a presenter (HeyGen) or the occasional clip (whatever is bundled in their chat app).

Why it matters

Runway is the proof that a small lab can own a frontier: it trained its own models, held a quarter of a billion-dollar market against Google and OpenAI, and is now trying to turn video into a physics engine. If world models work, the same company that made AI film shots will be making the simulators that train robots. If they do not, it remains the best studio in the category. Either way it is the one to watch on the creative side of the agentic economy.

Glossary

Runway — New York AI research company and video studio, founded 2018. Builds the Gen video models and the editing tools around them. \$5.3B valuation, \$200M ARR.

Gen-4.5 — Runway's current production video model: native audio, multi-shot, character consistency.

World model — An AI system that models how an environment behaves — physics, permanence, cause and effect — rather than predicting the next frame. Runway's stated direction since 2025.

Stable Diffusion — The open image-generation model Runway co-created in 2022 with Stability AI and academics; the start of open generative media.

Motion brush / camera control — Runway tools for directing what moves and how the shot is framed.

Reference — An image of a character, place or object the model keeps consistent across generations.

Credits per second — Runway's billing: each generated second costs credits, more for newer models.

Explore Mode — The slower, unlimited-generation lane on the old Unlimited plan. Not part of Max, and gone for everyone after 30 November 2026.

Veo / Genie — Google's video model and world-simulation research, Runway's main competition on both fronts.

Liked this? Go deeper.

Runway is the studio. Here is the rest of the creative stack.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Higgsfield

The AI video platform every creator is talking about. What it does, what it costs, the September API and open-source moves, and who it is for.

wtfagents.com/guides/higgsfield

WTF is HeyGen

The AI avatar video platform. Digital twins, one-prompt videos, dubbing into 175+ languages. What it does, the credit maths, and the consent question.

wtfagents.com/guides/heygen

WTF is ElevenLabs

The AI voice standard. Text-to-speech, cloning, dubbing and the phone agents behind Klarna and Revolut. What it does, what it costs, and the rights question.

wtfagents.com/guides/elevenlabs

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Cursor?

The code editor that became the fastest-growing software company in history, then the largest startup acquisition ever when SpaceX paid \$60 billion for it. What it is, what changed under Musk, why OpenAI is walking away, and what a non-developer should know.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Cursor is a code editor with AI built into every keystroke: autocomplete that predicts the next edit, a chat that knows your whole codebase, and agents that take a task and return the finished change. It is where a large share of the world's professional developers now write software.

In 2026 it became something else as well: the product SpaceX paid \$60 billion for, the delivery vehicle for Grok Bot, and the front line of a fight between Elon Musk and OpenAI that will cut off a supplier on 12 November.

By the numbers — September 2026

\$60B

all-stock price SpaceX paid for Anysphere, Cursor's company — the largest acquisition of a venture-backed startup ever. Announced 16 June, closed 14 August

~\$4B

annualised revenue by mid-2026, from roughly zero at the start of 2024

73×

the rise in Cursor's price tag across four funding rounds in under three years, before the sale

12 Nov

2026: the date OpenAI stops supplying models to Cursor, citing the Musk relationship

~5%

of Cursor traffic served by OpenAI models, per co-founder Michael Trueell — Anthropic and Google models remain

\$20 / \$60 / \$200

per month for Pro, Pro+ and Ultra; Teams from \$40 a seat, Premium seats \$120

Composer 2.5

Cursor's own coding model since 18 May, in the first-party pool with Grok 4.5 and 4.6

What it actually does

Tab — Autocomplete that predicts your next edit, not your next word: it proposes the whole change and you press Tab. The feature that made developers switch.

Chat with the codebase — Ask questions about a project with millions of lines and get answers with the files cited. Cursor indexes the whole repository.

Agent mode — Describe a feature or a bug; the agent edits across files, runs the terminal, checks the result and iterates. Several agents can run in parallel in cloud sandboxes, then hand back a review.

Composer — Cursor's own models, built to be fast and cheap for the routine 80% of coding, with frontier models from Anthropic and Google (and OpenAI, until November) available for the hard 20%.

The router — The quiet core of the business: send each request to the model that does it well enough at the lowest cost. It is how Cursor went from negative margins to profit in early 2026, and why losing one supplier is survivable.

Grok Bot — Since August, the work agent that lives alongside the editor: named bots with your logins on a shared cloud computer, sold through Cursor's plans. See its own guide.

Enterprise — SSO, admin dashboards, usage pools, MCP support for connecting internal tools; where the money is made.

What it costs

Plan	Price	What you get
Hobby	\$0	Limited Tab and agent use; enough to see if it fits
Pro	\$20/mo	Unlimited Tab, a monthly usage pool for agents, Grok Bot access since August
Pro+	\$60/mo	Three times the Pro pool
Ultra	\$200/mo	Twenty times, for people running agents all day
Teams	\$40/seat	Pro plus admin controls, shared billing, privacy mode enforced
Teams Premium	\$120/seat	Five times the standard usage for heavy agent workloads
Enterprise	Custom	Pools, SSO, audit, SCIM, dedicated support

Usage splits into two pools: first-party models (Composer, Grok) draw from your plan; third-party models (Claude, Gemini, GPT) are billed at each provider's API rate against it.

Heavy agent use on frontier models exhausts a Pro pool quickly; that is what Pro+ and Ultra are for.

Cursor vs Claude Code vs Codex

	Cursor	Claude Code	Codex
Shape	An editor with agents inside	An agent, with an editor optional	An agent in ChatGPT and the terminal
Model	Any — routed; Grok and Composer first-party	Claude	GPT
Where it runs	Your machine, cloud sandboxes for agents	Your machine, cloud, phone	Cloud sandboxes, local CLI
Entry price	\$20 Pro	\$20 Claude Pro	\$20 ChatGPT Plus
Owner	SpaceX	Anthropic	OpenAI
Strength	Live editing with AI at every keystroke	Long autonomous sessions, leads the benchmarks	Parallel tasks, GitHub review

Many developers run Cursor as the editor and Claude Code as the agent inside it. The acquisition makes that combination more political than it was, and the November cutoff makes the third column a little emptier.

The SpaceX story, in order

- **February 2026:** SpaceX absorbs xAI at a \$250 billion valuation. Musk's rocket company now owns a frontier lab and X.
- **April:** SpaceX signs a deal with Cursor: buy the company, or pay roughly \$10 billion in break-up and service fees; \$4 billion more if regulators block it.
- **June:** SpaceX goes public at \$135 a share, raising \$75 billion, the largest IPO in history. Four days later it announces the \$60 billion Cursor purchase.
- **July:** xAI rebrands to SpaceXAI. Grok 4.5 ships, trained jointly with Cursor on trillions of tokens of Cursor data.
- **August:** Grok Bot and Grok 4.6 launch with Cursor's help. The acquisition closes on the 14th; Anysphere shareholders receive SpaceX stock. On the 28th OpenAI announces it will stop supplying models on 12 November.

→ **March, in hindsight:** Composer 2 was revealed to be built on Moonshot's open-weight Kimi K2.5 — a Chinese base model under the American developer tool — which Cursor confirmed after users spotted it on X. Cursor's account is that only about a quarter of the compute behind the model came from the base, and Moonshot confirmed an authorised commercial partnership through Fireworks AI. The complaint was the silence, not the licence.

The plot is simple once you see it: SpaceX wanted a real AI business and a million developers; Cursor wanted compute and margin. Musk got the layer between developers and every model, which is the layer OpenAI could not take away when it left.

The catches

→ **The OpenAI cutoff.** From 12 November, GPT models leave the third-party pool. If your workflow depends on them, plan the switch now. Claude and Gemini stay; Grok becomes the emphasised default.

→ **Whose data trains what.** Grok 4.5 was trained on Cursor data. Privacy mode exists and Teams enforce it; individuals should switch it on and understand what "off" means.

→ **Model choice is a feature, until it isn't.** Cursor markets choosing your model. The owner now has a model to push. Watch which one the router prefers.

→ **Individual accounts lose money.** Cursor's own economics say enterprise accounts pay and individual ones do not. Expect the \$20 tier to get tighter, not looser.

→ **It is an editor.** For a non-developer this is the least approachable of the coding agents; Claude Code's desktop app or Codex in ChatGPT are gentler starts.

Do you need to be a developer?

More than for Claude Code or Codex. Cursor is a professional editor with AI in it, not an agent with a chat box, and it assumes you know what a repository, a terminal and a diff are. If you do, it is the most productive place to write software in 2026. If you do not, start with Claude Code's desktop app and come back when you want the editor.

Why it matters

Cursor is the proof that the layer between people and models is worth more than most models. It trained almost nothing at the frontier, rented everyone else's, and was bought for \$60 billion by a company that needed a way into developers' hands. Now that layer belongs to the same owner as a frontier lab, a social network and a rocket company, and one supplier has already walked out. What happens to Cursor between now and November is the clearest test yet of whether an aggregator can survive owning a side.

Glossary

Cursor — AI code editor built on VS Code by Anysphere, founded 2022. Acquired by SpaceX for \$60B in 2026.

Tab — Cursor's predictive autocomplete: it proposes the next edit, you accept with Tab.

Agent mode — Cursor's autonomous coding: describe a change, the agent edits, runs and checks across files.

Composer — Cursor's own family of coding models. Composer 2 and the current 2.5 are both built on Moonshot's open-weight Kimi K2.5.

Router — The system that sends each request to the cheapest model that can do it well. Cursor's margin engine.

First-party / third-party pool — Usage from Cursor and Grok models versus usage billed at Anthropic's, Google's or OpenAI's API rates.

SpaceXAI — The rebranded xAI, a SpaceX subsidiary since February 2026 and Cursor's owner since August.

Grok Bot — SpaceXAI's work agent, sold through Cursor plans. See its own guide.

Privacy mode — The Cursor setting that keeps your code out of training. On by default for Teams; check it for yourself.

Liked this? Go deeper.

Cursor is the editor. Here are the agents that compete with it and the one that shipped through it.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

wtfagents.com/guides/chatgpt-work

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Devin?

The "first AI software engineer": a cloud agent that takes a ticket and returns working software, from a company valued at \$47 billion that says it writes 95% of its own code with it. What it does, why the price fell from \$500 to \$20, and how it compares.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Devin is an autonomous software engineer in the cloud: you give it a task — a bug, a feature, a migration — and it plans, writes, tests, fixes and opens the pull request, then messages you when it is done or stuck. It was the first product to make that claim in 2024 and the one everyone measured against.

Two years on it is cheaper, better, and part of a bigger company: Cognition bought the Windsurf editor, raised at \$47 billion, and says Devin now writes about 95% of Cognition's own code — up from 89% at the May round, and a quarter of its pull requests in 2025.

That last number is worth holding at arm's length. It is the company's own figure, self-reported by CEO Scott Wu, and it counts code committed rather than engineering done: the specifying, reviewing and deciding still sit with people. As a trend line it is the most interesting thing Cognition publishes. As a benchmark it is unaudited.

By the numbers — September 2026

\$47B

the valuation Bloomberg reported on 2 September 2026, for a ~\$1B round still in talks — up from \$26B in May and \$10.2B a year ago

~\$900M

annualised revenue, more than double the \$492M of May 2026; Devin alone was \$73M in June 2025, before Windsurf

\$500 → \$20

the monthly price cut with the Devin 2.0 relaunch in April 2025; Max at \$200, Teams from \$80

1 ACU

Devin's unit of work, roughly 15 minutes of autonomous engineering

July 2025

Cognition buys Windsurf after OpenAI's \$3B bid collapses and Google hires Windsurf's founders

4,000+

enterprise customers on Windsurf by spring 2026; over a million users

3.0

the current Devin: dynamic re-planning, desktop-app testing via computer use, a unified Devin Desktop

What it actually does

Takes a task, returns a pull request — Describe the change in plain English or paste a ticket. Devin explores the codebase, makes a plan, edits, runs the tests, fixes what breaks and hands back a reviewable change with an explanation.

Works in parallel — Each task runs in its own cloud environment, so ten tickets can run at once. This is where it earns its keep on migrations and backlogs.

Re-plans — Devin 3.0 revises its approach mid-task when something surprises it, rather than ploughing on. The difference between an agent that gets stuck and one that gets there.

Tests what it built — Including desktop and web apps, by driving them the way a person would (computer use).

Devin Desktop — The 2026 app unifying the cloud agent, the command line and the editor, so you can watch, steer and take over.

Windsurf — The editor Cognition bought: AI-native, with completions and chat like Cursor, and the on-ramp to Devin for longer tasks. Renamed Devin Desktop in June 2026, and now open to rival agents — Codex, Claude, Gemini CLI — through the Agent Client Protocol.

Multi-model — Claude, GPT and Gemini frontier models plus Cognition's own SWE family, routed per task. Cognition's pitch is independence from any one lab.

What it costs

Plan	Price	What you get
Free	\$0	A small ACU allowance to try it
Pro	\$20/mo	A usage allowance refreshing daily and weekly; extra usage at API rates
Max	\$200/mo	The heavy-use tier, for running Devin all day
Teams	\$80/mo + \$40/seat	Shared ACU pool, admin, integrations
Enterprise	Custom	Volume ACUs, security, dedicated support, custom models

The unit is the ACU, about a quarter-hour of Devin working. A simple bug might take two; a multi-file feature ten or more; a large migration hundreds. The \$20 plan is for trying tasks; the

\$200 plan is for handing over a backlog.

Cognition does not publish the ACU allowance on each plan — the pricing page says only that paid plans refresh daily and weekly, and that extra usage bills at API rates. Third-party write-ups quote confident quotas and per-ACU prices; they disagree with each other and with the pricing page. Watch your own meter for a month before sizing the plan.

Devin vs Claude Code vs Codex

	Devin	Claude Code	Codex
Shape	Cloud engineer you message	Agent in your terminal, desktop, cloud	Agent in ChatGPT and the CLI
Model	Multi-model + own SWE models	Claude	GPT
Unit	ACU (~15 min of work)	Plan usage	Plan credits
Entry price	\$20	\$20 Claude Pro	\$20 ChatGPT Plus
Best at	Long, parallel, hands-off tasks	Hard problems needing care; leads benchmarks	Parallel backend tasks, GitHub review
Owner	Cognition, independent	Anthropic	OpenAI

The honest distinction after two years: Devin is the most hands-off of the three and the best at running a queue of well-specified tasks unattended. Claude Code is what people reach for when the task is hard. Many teams use both.

The Windsurf story

In mid-2025 Windsurf, a well-liked AI editor, was in a bidding war. OpenAI agreed to buy it for about \$3 billion; the deal collapsed. Google hired Windsurf's CEO and research leads in a licensing deal. Days later Cognition bought what was left: the product, the brand, the engineers and \$82 million of revenue with 350-plus enterprise customers.

It turned out to be the acquisition that made the company. Devin had the headlines; Windsurf had the users. Combined revenue more than doubled within months, the valuation went from \$10 billion to \$26 billion to \$47 billion in a year, and Cognition became a coding platform instead of a single agent. The lesson it teaches about the whole category: distribution beat being first.

The catches

- **ACUs are hard to predict.** A task that seems small can take many; the meter is honest but the estimate is yours. Start with the free tier and watch what real tickets cost.
- **Cloud-only means your code leaves.** Devin works in Cognition's environments. Enterprise plans have the controls; individuals should know that is the model.
- **Autonomy needs specification.** Devin does well with a clear ticket and badly with a vague one. The skill it rewards is writing the brief.
- **Valuation pressure.** At roughly fifty times revenue, Cognition has to keep growing very fast. Expect pricing and packaging to keep moving.
- **The bundling threat.** Anthropic, OpenAI, Google and Microsoft can make autonomous coding a default feature of what developers already pay for. Devin has to stay better, not just independent.

Do you need to be a developer?

Less than for Cursor, more than for Claude Code's desktop app. Devin assumes a repository, tests and someone to review a pull request; it does not assume you write code yourself. A founder who can read a diff and describe a change precisely can get real software out of it. If that is not you yet, start with Claude Code and come back when you have a backlog to hand over.

Why it matters

Devin defined the category: the agent you message like a colleague, that returns finished work. Everything in the coding-agent row of this series — Claude Code, Codex, Grok Bot's engineering bots — is a response to the claim Devin made first. That it needed to buy an editor to grow is the second lesson, and the one the rest of the agentic economy is learning now: the agent is the product, but the place people already work is the distribution.

Glossary

Devin — Cognition's autonomous AI software engineer, launched 2024. Takes a task, returns a pull request. Now version 3.0.

Cognition — The company behind Devin and Windsurf, founded 2023 by Scott Wu, Steven Hao and Walden Yan. Valued at \$47B.

ACU (Agent Compute Unit) — Devin's unit of work, roughly 15 minutes of autonomous engineering.

Windsurf — The AI-native code editor Cognition acquired in July 2025, renamed Devin Desktop in June 2026.

Devin Desktop — Windsurf, renamed in June 2026: one app for the cloud agent, the CLI and the editor, and for rival agents too.

Pull request — A proposed change to a codebase that a human reviews before merging. Devin's output.

Re-planning — Revising the approach mid-task when something unexpected happens. Devin 3.0's headline capability.

SWE models — Cognition's own coding models, used alongside Claude, GPT and Gemini.

Computer use — An agent driving an application's interface the way a person would; Devin uses it to test what it builds.

Liked this? Go deeper.

Devin is the engineer you message. Here are the others, and the editor SpaceX bought.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is Codex

OpenAI's coding agent, now a tab in ChatGPT. What it does, how it compares with Claude Code, what it costs on each plan, and the catches.

wtfagents.com/guides/codex

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is n8n?

The open-source automation layer solo builders wire their agents with: one workflow, one execution, any number of steps, on a \$5 server if you want. What it does, why it beats Zapier at volume, when it doesn't, and what a \$5.2 billion valuation says about it.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

n8n is a workflow automation platform: a visual canvas where a trigger (a form submitted, an email received, a schedule) flows through steps (call an AI model, look up a record, send a message, update a sheet) to an outcome, without code — or with as much code as you like. It is open-source, runs on your own server if you want, and has become the plumbing under a large share of the indie agent stacks in this series.

Its billing model is the whole argument: one execution is one complete run of a workflow, however many steps. Zapier bills every step.

By the numbers — September 2026

\$5.2B

valuation after SAP's strategic investment, announced 12 May 2026 — double the \$2.5B of the October 2025 Series C

1.7M

monthly active builders, and more than 1,400 enterprise customers, per the company

1

execution billed per workflow run, whatever its length — versus one task per step at Zapier

€20 / €50

per month for Starter (2,500 executions) and Pro (10,000), billed annually — workflows, steps and users are unlimited on every plan

\$0

for the self-hosted Community Edition: no execution, workflow or user limits; you pay for the server

500+

native integrations, plus an HTTP node that reaches any API — versus 8,000-plus at Zapier

2.0

the January 2026 release that made agents first-class: an Agent node with tool-calling, memory and retries

What it actually does

Workflows on a canvas — Nodes connected by lines. A trigger starts it, each node transforms or acts, branches and loops handle the logic. Readable by a non-developer, editable by one after a week.

The AI Agent node — Since 2.0: a node that runs a model with tools, memory and retry logic. Point it at Claude, GPT or a local model via Ollama, give it tools from other nodes, and it decides the steps. This is where n8n stopped being "Zapier but cheaper" and became an agent runtime.

Any API — The HTTP node talks to anything with a REST or GraphQL interface, including systems no integration exists for: your accountant's practice software, a client's internal tool, a niche SaaS. Zapier waits for someone to build the connector; n8n lets you.

Code when you want it — JavaScript or Python inside a node, for the 5% of steps that need it. The rest stays visual.

Self-hosting — Run it on a \$5–20 virtual server, in Docker, and keep every byte of data inside your own infrastructure. The reason regulated businesses pick it.

Databases, files, MCP — First-class Postgres, MongoDB and Redis; file handling; and MCP support so agents elsewhere can call n8n workflows as tools and n8n agents can use MCP servers.

Templates — Thousands of community workflows to start from: lead scoring, content pipelines, inbox triage, the standard shapes of small-business automation.

What it costs — and the maths that decides it

Plan	Price (annual)	Includes
Community (self-hosted)	€0 + server	Unlimited everything; you run it
Starter (cloud)	€20/mo	2,500 executions; workflows, steps and users unlimited
Pro (cloud)	€50/mo	10,000 executions, more concurrency and projects
Business (cloud)	€667/mo	40,000 executions, SSO, Git, governance
Enterprise	Custom	Self-hosted with support, or cloud at volume

n8n dropped its active-workflow caps: every paid plan now carries unlimited workflows, steps and users, and only executions, concurrency and projects are metered. So the comparison that matters is per workflow run. Take a ten-step lead-qualification flow that fires a thousand times a month: on n8n that is 1,000 executions, inside the Starter plan. On Zapier it is 10,000 tasks, well past the Professional plan's 750 and into a tier ten times the price. Self-hosted, it is the server bill and nothing else.

Below about a thousand multi-step runs a month, the difference is small and Zapier's ease wins. Above it, n8n wins by a wide margin. Count your steps and your runs before choosing.

n8n vs Zapier vs Make

	n8n	Zapier	Make
Billing unit	Execution (whole run)	Task (each step)	Operation (each step, cheaper)
Integrations	500+ native, any API	8,000+	2,000+
Self-host	Yes, free	No	No
AI agents	Agent node, local models	Zapier Agents	AI Agents
Learning curve	Real: a week	Gentle	Medium
Best for	Volume, control, data privacy, agents	Speed, non-technical operators	Visual complexity on a budget

Zapier is the right first tool for most small businesses: it is easier, it connects to more things, and up to about \$100 a month of spend the time it saves pays for itself. n8n is what you move to when the bill passes \$200, the data must stay in-house, or you want a real agent with memory and tools rather than a packaged assistant.

The catches

- **Free means you host it.** Updates, backups, SSL, uptime and the 11pm "why did it stop" are yours. If nobody on the team can run a server, the Community Edition is not free; it costs your evenings. Cloud plans exist for that reason.
- **The learning curve is real.** Data mapping between nodes, expressions and error handling are harder than Zapier. Expect a week of friction before it pays back.

- **Fewer connectors.** Five hundred native integrations against Zapier's eight thousand. The HTTP node covers the gap if you can read an API document.
- **Agents amplify mistakes.** A workflow that emails a thousand people runs a thousand times whether the logic was right or not. Test with a small trigger, and use 2.0's human-in-the-loop tool gates in front of anything that deletes, sends or spends.
- **SAP is now an investor.** A strategic minority stake, not control, and the real substance is distribution: n8n's canvas is being embedded in SAP's Joule Studio. Good for enterprise credibility; watch whether the roadmap tilts toward SAP customers and away from the solo builder who made it.

Who it is for

Yes — Solo founders and small teams with someone comfortable running a server or paying €20 for cloud, whose automations have grown past a handful of simple zaps, and who want agents with memory and tools inside their workflows.

Not yet — Anyone with fewer than a few hundred workflow runs a month and no technical appetite. Zapier, and revisit when the bill or the data question forces it.

No — Anyone who wants an agent to run their business for them (that is Polsia or NanoCorp) rather than a system to wire their own.

Why it matters

Every agent in this series needs a way to be triggered, to reach the next tool, and to hand off. n8n is that connective tissue for the part of the agentic economy that is not run by a hyperscaler: open-source, self-hostable, execution-priced, and now with agents as a native node. That a company built on "you can run it yourself for free" is worth \$5.2 billion with SAP on the cap table says how much of the real work in this economy happens in workflows nobody outside the business ever sees.

Glossary

n8n — Open-source workflow automation platform, Berlin, founded 2019. Self-hostable or cloud; execution-based pricing. \$5.2B valuation.

Workflow — A trigger plus a chain of nodes that transform data and take actions.

Node — One step in a workflow: an integration, a transformation, a code block, an agent.

Execution — One complete run of a workflow, however many nodes. n8n's billing unit.

Task — Zapier's billing unit: one per step, per run.

AI Agent node — n8n's node that runs a model with tools, memory and retries, and decides the steps itself.

HTTP node — The node that calls any API, covering integrations n8n does not ship natively.

Self-hosted / Community Edition — Running n8n on your own server; free, unlimited, your responsibility.

Ollama — A way to run open-weight models locally; n8n's Agent node can use them, keeping data off any cloud.

Make — The third major automation platform; visual, operation-priced, cheaper than Zapier at volume.

Liked this? Go deeper.

n8n is the wiring. Here is what it connects.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is MCP

The Model Context Protocol, explained. The universal plug between agents and tools: what it is, who runs it now, and the security catch.

wtfagents.com/guides/mcp

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

wtfagents.com/guides/chatgpt-work

WTF is Hermes Agent?

The open-source agent that gets better the longer it runs: it writes its own skills, keeps its own memory, and builds a model of you across sessions. From a research lab, on a \$5 server, and the closest thing OpenClaw has to a rival. What it is and who should run it.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Hermes Agent is an open-source personal agent from Nous Research, a lab better known for training open models. It runs on your own server, talks to you over Telegram, WhatsApp, Slack, Discord or Signal, and, unlike every other agent in this series, it learns: after a hard task it writes a reusable skill, improves that skill when it uses it again, and keeps a deepening memory of who you are.

It launched on 25 February 2026 and became the fastest-growing agent project of the year. OpenClaw has the users; Hermes has the architecture people now measure agents against.

By the numbers — September 2026

25 Feb

2026: released under the MIT licence. 95,000 GitHub stars in seven weeks; 247,000 by September

118+

bundled skills at v0.10, plus the ones the agent writes for itself

40%

faster on repeated research tasks once an agent has 20 or more self-made skills, per independent benchmarks

6+

messaging channels from one gateway: Telegram, WhatsApp, Slack, Discord, Signal, and the terminal

\$0

for the software; a \$5–10 monthly server if you want it always on, plus whatever your model provider charges

\$20

per month for Nous Portal Plus: \$22 of credits, 300+ models and the tool gateway. Super is \$100, Ultra \$200

4 Critical

and nine High-severity findings in the default configuration, in an independent audit published 11 April 2026; CVEs have followed

What makes it different — the learning loop

Every other agent starts each session from scratch, plus whatever notes you configured. Hermes closes a loop.

- **Skills from experience.** After a complex task, the agent writes down how it did it as a reusable skill. Next time, it uses the skill instead of working it out again.
- **Skills that improve.** Using a skill and hitting a snag updates the skill. They get sharper with use.
- **Three layers of memory.** A curated long-term memory the agent maintains itself, searchable session history, and a model of you (preferences, context, how you like things done) that deepens over time.
- **Nudges to persist.** It periodically prompts itself to save what matters, so knowledge does not evaporate at the end of a session.

The practical effect is the one nobody else offers: an agent that is measurably better in month three than in week one, without you retraining anything.

What it actually does

Everything a personal agent does — Research, drafts, summaries, browsing (full browser automation: navigate, click, type, screenshot), code execution in a sandbox, file handling, image generation, text-to-speech, vision.

Scheduled work — A built-in scheduler in natural language: daily reports, nightly backups, weekly audits, a morning briefing, delivered to whichever channel you like, unattended.

Sub-agents — Spawns isolated agents for parallel workstreams, each with its own conversation and terminal. Hireable into Paperclip out of the box.

Any model — Nous's own Hermes models, Claude, GPT, Gemini, open models via OpenRouter, or a local model on your own hardware. Switch with one command; no code changes.

Runs anywhere — Linux, macOS, Windows, Android via Termux, a \$5 VPS, a GPU box, serverless. A desktop app arrived in 2026 for people who want a window rather than a terminal.

Nous Portal — Optional. One login for a model plus the tool gateway (search, images, speech, browser), from \$20 a month; Portal Cloud offers one-click hosted instances billed hourly.

Hermes vs OpenClaw vs Muse

	Hermes	OpenClaw	Meta Muse
Who makes it	Nous Research, a lab	A foundation, ex-Steinberger	Meta
Learns over time	Yes — the point	Memory, not a learning loop	Meta's memory, Meta's terms
Users	247,000 stars; a developer crowd	3M+ active, mainstream	Consumer scale
Skills	118 bundled + self-written	Large community marketplace	None; fixed features
Setup	Real, technical	Real, getting easier	None
Cost	Free + model; \$20 Portal optional	Free + model	Free tier, \$20, \$100
Data	Yours	Yours	Meta's

The honest split: OpenClaw is the popular one with the biggest ecosystem and the native mobile apps. Hermes is the better-architected one — the memory and learning design is genuinely ahead of the field — and developers have been migrating to it since spring. Its security record is not part of the pitch; see the catches. Muse is for people who want none of the setup and accept Meta holding the keys.

What it costs to run

The software is free. The bill is the model and the server.

Route	Roughly	Notes
Laptop only	Model API costs	Runs when the laptop does
\$5–10 VPS, always on	\$5–10 + model	The normal setup; a cheap open model keeps tasks at cents each
Nous Portal Plus	\$20/mo	Includes \$22 of credits; 300+ models and the tool gateway
Portal Cloud	Hourly from credits	Hosted instance, no server to run
Frontier model heavy use	\$30–100+/mo	Claude or GPT on every task adds up; route routine work to a cheap model

A complex task on a budget model costs around thirty cents; on a frontier model a few dollars. The learning loop cuts repeated tasks over time, which is a real cost lever nobody else has.

The catches

- **You run it.** A server, updates, a model key, and the first hour of setup. The desktop app and Portal Cloud soften this; they do not remove it.
- **Version 0.x.** Releases every two weeks with real changes. Pin a version and read the notes before upgrading; the pace that makes it good makes it move.
- **Security is the weak point, not the strength.** The April 2026 audit found four Critical and nine High-severity issues in the default configuration, including unrestricted shell execution and persistent skill injection. CVEs have kept coming since, one of them a remote-code-execution flaw with a published proof of concept. Run it in Docker, turn off any auto-approval bypass, restrict what it can write to, keep it patched — and review the skills it wrote itself before letting them send or spend.
- **Enterprise features are arriving, not arrived.** Team agents, access controls and audit logging are on the roadmap; for a company deployment, wait or add your own.
- **Star counts are not users.** The growth numbers are developer enthusiasm, which is real, and not the same as three million people using it daily. OpenClaw still has the crowd.

Who it is for

Yes — Anyone comfortable with a terminal who wants to own their agent, keep their data, and have it improve with use. The best choice in this series for a technical solo operator.

Not yet — Anyone who wants mobile-first, zero-setup, or a marketplace of thousands of community skills (OpenClaw), or no setup at all (Muse).

No — Anyone who wants an agent to run a business for them rather than to be their assistant.

Why it matters

Hermes is the first widely used agent built on the premise that an assistant should get better the longer you have it. Every other agent in this series is as good on day one as it will ever be until its maker ships an update. If the learning loop becomes standard, and its adoption curve says it will, the agents people keep will be the ones that have learned them, and switching will cost more than a subscription. That is a new kind of lock-in, and for once it favours the software you own.

Glossary

Hermes Agent — Open-source, self-improving personal agent from Nous Research, released February 2026 under MIT.

Nous Research — The AI lab behind Hermes; known for open models, reportedly in talks at a \$1.5B valuation.

Learning loop — The mechanism by which the agent writes skills from experience, improves them in use, and persists memory.

Skill — A reusable, written-down procedure the agent follows. Bundled, community, or self-created. Compatible with the agentskills.io standard.

Three-layer memory — Curated long-term notes, searchable session history, and a model of the user.

Gateway — The single process that connects the agent to Telegram, WhatsApp, Slack, Discord, Signal and the CLI.

Nous Portal — The optional paid backbone: model access and a tool gateway from \$20/month; Portal Cloud hosts instances.

Sub-agent — An isolated child agent spawned for a parallel task.

VPS — A cheap virtual server, the usual home for an always-on agent.

Liked this? Go deeper.

Hermes is the agent that learns. Here are the ones it competes with, and the org chart it can join.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

wtfagents.com/guides/meta-muse

WTF is Manus

The viral agent Meta bought and China blocked. What Manus does, the credit pricing, the 2026 saga, and what it means for anyone choosing a Chinese-founded tool.

wtfagents.com/guides/manus

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is Manus?

The general-purpose agent that went viral, hit \$100 million in revenue in eight months, sold to Meta for \$2 billion — and was then pulled back by Beijing with a one-line order. What it does, what it costs, and what its story says about who owns an agent.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

Manus is a general-purpose AI agent in the cloud: give it a goal — research a market, build a spreadsheet, plan a trip, write and ship a small app — and it works in its own virtual computer with a browser, a terminal and a file system until it has a result to show you. It was the first agent to go properly viral, in March 2025, with clips of it working unattended for an hour and coming back with the job done.

Since then it has become two other things: a real business with hundreds of millions in revenue, and the centre of the year's biggest lesson in who gets to own an AI company.

By the numbers — September 2026

6 March

2025: invite-only launch; invitation codes were being resold within days

\$100M

annualised revenue by December 2025, eight months after launch, growing about 20% a month

~\$450M

annualised revenue by June 2026 (Sacra estimate)

\$2B+

what Meta agreed to pay on 29 December 2025; some reports put the full value at \$2.5B

27 April

2026: China's NDRC blocks the deal in one paragraph and orders it unwound

11 Aug

2026: Manus announces it will operate as an independent company

\$20 / \$40

a month for the two Pro tiers, each "starting from"; the free tier refreshes 300 credits a day

What it actually does

A computer of its own — Each task runs in a sandboxed virtual machine with a real browser, a terminal and files. You can watch it work, step in, or close the tab and come back.

Long tasks, unattended — Competitive research across dozens of sites, a data analysis with charts, a comparison of fifty products, a trip booked end to end. The demos that made it famous were an hour of autonomous work per prompt.

Deliverables, not answers — Spreadsheets, slide decks, reports, small web apps, deployed. The output is a file or a site, which is what made it feel like an employee.

Desktop and My Computer — Since March 2026, a desktop app that lets the agent run on your own machine: terminal commands, local files, installed tools, on Mac and Windows.

Scheduled and parallel — Recurring tasks and several agents at once on higher plans.

Manus's honest capability, from people who have measured it: reliable on research and assembly tasks of up to about five dependent steps, with a sharp drop-off beyond that. Good for a well-scoped afternoon of work; not yet for a week of it.

What it costs

Plan	Price	Credits
Free	\$0	300 credits refreshed daily; one task at a time
Pro, first tier	From \$20/mo	From 4,000 credits; 20 concurrent and scheduled tasks
Pro, second tier	From \$40/mo	From 8,000 credits; seven-day trial
Team	From \$20/seat	Pro features plus SSO, privacy controls, analytics

Everything is credits, and a long autonomous task can burn a few hundred of them. The free tier is for seeing what it does; the first Pro tier is for occasional real tasks; the second is for people who hand it work daily. Both are priced "starting from", scaling with the credits you want.

Check manus.im before you buy. The tiers have been renamed and repriced more than once this year, and most third-party pricing pages — including ones updated this summer — still quote a \$19 / \$39 / \$199 ladder that Manus's own help centre no longer lists.

The 2026 saga, in order

- **2022:** Butterfly Effect founded in China by Xiao Hong, Yichao Ji and Tao Zhang. Its first product, Monica, is a browser assistant.
- **March 2025:** Manus launches. Self-reported benchmark scores beat OpenAI's Deep Research; invite codes become a black market.
- **July 2025:** The company moves its headquarters to Singapore, partly to escape US chip restrictions on Chinese firms.
- **December 2025:** With revenue past \$100 million, Meta agrees to buy it for around \$2 billion on the 29th and closes the next day, promising no remaining Chinese ownership, and starts wiring Manus into Meta AI and Ads Manager.
- **January 2026:** China's Ministry of Commerce and NDRC open a security review. By March the two founders are reported to be under exit bans; a hundred staff are already working from Meta's Singapore offices.
- **27 April:** The NDRC prohibits "foreign investment in the Manus project" and orders the parties to withdraw. No further explanation.
- **11 June:** Meta begins the unwind with an operational split: no data sharing, Manus staff cut off from Meta systems, Meta staff barred from using Manus.
- **11 August:** Manus announces it will return to operating as an independent company — and that from the end of the month it will delete some user data created on or after the day Meta closed the deal, to satisfy regulators.

The precedent is the point. Manus had left China, closed its Chinese business and sold to an American buyer. Beijing asserted authority anyway, over technology built by Chinese engineers, and won. Every Chinese-founded AI company, and every Western acquirer, now knows that.

Manus vs OpenClaw vs ChatGPT Work

	Manus	OpenClaw	ChatGPT Work
Shape	Cloud agent with its own computer	Self-hosted agent on yours	Agent inside ChatGPT
Setup	None	Real	None
Best at	Research and assembly, unattended	Anything you configure	Finished files from your apps
Model	Its own routing over frontier models	Any	GPT
Data	On Manus's servers, Singapore	Yours	OpenAI's
Owner	Independent, Chinese-founded	A foundation	OpenAI
Price	Free, from \$20, from \$40	Free + model	\$20 ChatGPT Plus

Manus is the most polished zero-setup agent for "go and find out" tasks. The catch is the same as its strength: it is a cloud service from a company whose corporate future was decided by a government this spring.

The catches

- **Ownership risk is now product risk.** An independent Manus is stable for now; the founders are under travel bans and the company's next move is not in its own hands alone. Annual plans deserve a pause.
- **Your data goes to Singapore — and some of it was deleted.** The unwind required Manus to delete certain user data created during the Meta year, from the end of August 2026. Fine for most tasks; not for anything a compliance team would ask about, and not somewhere to keep the only copy.
- **Credits burn on autonomy.** An agent that works for an hour spends for an hour. Scope the task, set the budget.
- **Five steps, then check.** Beyond a handful of dependent steps, reliability drops. Break long jobs into stages.
- **Benchmarks were self-reported.** The launch numbers that made it famous were the company's own and are stale; judge it on your tasks.

Who it is for

Yes — Anyone who wants a capable research-and-assembly agent with no setup, will scope tasks sensibly, and is not handling sensitive data. The free tier is a fair test.

Not yet — Anyone whose company cares where data lives or who is unsettled by the ownership question. Wait for the independent company to prove itself.

No — Anyone who wants to own the agent (OpenClaw, Hermes) or whose work lives in their existing apps (Cowork, ChatGPT Work).

Why it matters

Manus proved the demand: a general agent that does an hour's work from one prompt found a hundred million dollars of customers in eight months. Then it proved something else. In April 2026 the most important fact about an AI agent turned out to be where its engineers came from, not what it could do. For anyone building on, buying, or selling an agent company, the Manus order is the new rule of the game.

Glossary

Manus — General-purpose cloud AI agent from Butterfly Effect, launched March 2025. Singapore-based, Chinese-founded, independent since August 2026.

Butterfly Effect — Manus's parent company, founded in China in 2022; also makes the Monica browser assistant.

NDRC — China's National Development and Reform Commission, the state planner that blocked the Meta deal.

Sandbox VM — The isolated virtual computer each Manus task runs in, with browser, terminal and files.

My Computer — The Manus Desktop feature that lets the agent act on your own machine.

Credits — Manus's usage currency, consumed by task length and complexity.

Exit ban — A travel restriction; reportedly applied to Manus's two founders during the review.

Deep Research — OpenAI's research agent, the benchmark Manus compared itself to at launch.

Liked this? Go deeper.

Manus is the agent Beijing kept. Here is the China story around it, and the agents you can own instead.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Happening with China's AI

DeepSeek, Qwen, Kimi, GLM and the open-weight strategy. Who the Chinese labs are, why they give models away, and what it means for anyone running agents.

wtfagents.com/guides/china-ai

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/guides/openclaw

WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

wtfagents.com/guides/meta-muse

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent