

# The Complete

## WTF Agents Pack

---

Every WTF Agents guide in one file. The agentic economy, the platforms running it, the AI behind it, and how to put it to work in your own business.

### What's inside — 20 guides

---

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an API
4. WTF is an LLM
5. WTF is Polsia
6. WTF is OpenClaw
7. WTF is Paperclip
8. WTF is Anthropic
9. WTF is Claude
10. WTF is Claude Code
11. WTF is Cowork
12. How to Hire an AI Agent for Your Business
13. Who's Who in the Agentic Economy
14. Which AI Should I Actually Use?
15. WTF is Meta Muse
16. WTF is Grok Bot
17. WTF is ChatGPT Work
18. WTF is Happening with China's AI
19. WTF is NanoCorp
20. WTF is Codex

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# WTF is the

## Agentic Economy?

---

AI stopped answering questions and started doing jobs. Every major lab now sells an agent that works while you sleep, half a million AI-run companies have been launched, and the layoffs have started. The plain-English guide to what is actually happening.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, checking its work and repeating — around the clock, at a fraction of the cost of a human.

For two years that was a prediction. In 2026 it became a product category. Every major AI lab now sells one. This guide explains what changed, who is doing what, and where the hype ends.

## By the numbers — September 2026

---

Everything in this box moves. Everything below it is written to still be true when it has.

**\$10.9B**

size of the AI agents market in 2026, up from \$7.6B in 2025 (Grand View Research)

**40%**

of enterprise software will contain task-specific agents by the end of 2026, up from under 5% in 2025 (Gartner)

**23%**

of organisations have actually scaled an agent system into production (McKinsey)

**40%+**

of agentic AI projects Gartner expects to be cancelled by the end of 2027

**490,000**

companies created by AI agents on Porsia alone since late 2025 — about 6% still active

**3.2M**

active users of OpenClaw, the open-source personal agent

**139,156**

tech layoffs in the first half of 2026, up 83% on the same period of 2025

**80%**

of large companies deploying autonomous tech cut jobs — with no correlation to improved returns (Gartner, May 2026)

# What an AI agent actually is

---

Most people's experience of AI is reactive: you type, it answers. ChatGPT, Siri, Google. You ask, they respond, they stop.

An agent is different. You give it a goal. It works out the steps, does them one by one, checks the results, adjusts, and keeps going until the job is done — or until it hits something it cannot handle and asks you.

|                        | A chatbot                      | An agent                                             |
|------------------------|--------------------------------|------------------------------------------------------|
| Trigger                | You ask a question             | You give it a goal                                   |
| Action                 | Answers once                   | Takes many steps on its own                          |
| Memory                 | Forgets when you close the tab | Remembers what it has done                           |
| Tools                  | Text only                      | Browses, writes code, sends email, fills forms, pays |
| Works while you sleep? | No                             | Yes                                                  |

The simplest analogy: a calculator waits for you to press buttons. An agent is a capable employee who takes a brief and gets on with it.

The [WTF is an AI Agent](#) guide goes deep on how they work. This guide is about what they are doing to the economy.

## 2026 — the year every lab shipped the same product

---

**Between January and September 2026, every major AI company launched a version of the same thing: an agent that does your job, not a chatbot that helps you do it.**

| When      | Who       | What                                                                                                                                            |
|-----------|-----------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| January   | Anthropic | Cowork — Claude works directly in your files and apps. Enterprise software stocks fell hard when the plugins landed.                            |
| June      | Microsoft | Copilot Cowork goes live worldwide inside Microsoft 365 — built on Anthropic's Claude, not OpenAI's models.                                     |
| July      | OpenAI    | ChatGPT Work — give it an outcome, it returns finished spreadsheets, decks and web apps hours later.                                            |
| August    | xAI       | Grok Bot — a team of named agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff" bot.                             |
| September | Meta      | Muse — a personal agent for everyone, in an app or on WhatsApp, that books, buys, negotiates and fills in forms for you.                        |
| May       | Google    | Killed its standalone browser agent and folded the technology into Gemini Agent and Chrome. The one lab that chose consolidation over a launch. |

Add the coding agents — Claude Code, OpenAI's Codex, Cognition's Devin (now \$20 a month, down from \$500), Google's Jules — and the pattern is unmistakable. The product is not "AI." The product is delegated work.

Two things make this different from every previous wave of business software. First, the same agent that runs a Fortune 500 workflow is available to one person for the price of a streaming subscription. Second, the agents can now hire, build and run other agents — which is where the next layer comes from.

## The three layers

It helps to think of the agentic economy as a stack.

### The models

The brains. Anthropic's Claude, OpenAI's GPT, Google's Gemini, Meta's Muse models, xAI's Grok — and, increasingly, the Chinese open-weight models from DeepSeek, Alibaba's Qwen, Moonshot's Kimi and Zhipu's GLM, which are what a great many indie agents actually run on because they are cheap and can be downloaded. Model names change every few months. The layer does not.

### The agents

Software that uses a model to do work. Two families. Personal agents run for you: OpenClaw (open-source, runs on your own machine, talks to you over WhatsApp or Telegram), Meta's Muse, Nous Research's Hermes. Work agents run for your job: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex, Devin.

### The companies

The most ambitious layer, and the newest. Platforms that point agents at a business idea and let them run the whole thing — code, marketing, outreach, support, finance. Polsia is the biggest and most argued about; NanoCorp publishes the live revenue of every company on it; Cofounder.co puts a human approval step in front of anything risky. Paperclip is the open-source version: an org chart for agents you host yourself.

**The stack matters because value is moving down it. Two years ago the money was in models. In 2026 it is in agents. The bet behind Polsia, NanoCorp and Paperclip is that next it will be in companies that no human works at.**

## The plumbing that makes it work

---

Two open standards are the invisible infrastructure of all this.

→ **MCP (Model Context Protocol)** — invented by Anthropic in late 2024, now governed by the neutral Agentic AI Foundation under the Linux Foundation and adopted by every major lab. It is the universal plug that lets any agent connect to any tool or data source. USB-C for AI.

→ **A2A (Agent2Agent)** — Google's standard for agents from different companies talking to each other. It is what lets a Paperclip org chart hire an OpenClaw agent, or a Grok Bot hand a task to a Codex session.

The newest piece is **agentic commerce**: agents that can pay. Muse buys on your behalf. Polsia companies take payments through their own Stripe accounts. Once agents hold wallets, the economy stops being a metaphor.

## The messy reality

---

**The agentic economy is real. It is also chaotic, frequently oversold, and moving faster than security, regulation or most businesses can absorb.**

### Most projects fail

Gartner expects more than 40% of agentic AI projects to be cancelled by the end of 2027 — unclear value, runaway cost, inadequate controls. Of the organisations experimenting, fewer than a quarter have anything in production at scale.

### Most AI-run companies earn nothing

Polsia's founder said on the record in May 2026 that around one in ten companies on the platform had earned at least a dollar. NanoCorp's public leaderboard showed its best company on \$2,570 across 748 sales, with only two companies past \$1,000. The agents build. Customers, mostly, do not come.

## The agents misbehave

OpenClaw agents joined an AI-only social network, created dating profiles and leaked a million interactions through a misconfigured database in the first weeks of 2026. Cisco found third-party OpenClaw skills silently exfiltrating data. Meta's own staff flagged security failures in Muse days before launch. Every agent that holds your logins and acts in your accounts is a new attack surface, and prompt injection — hiding instructions in a web page or email the agent reads — is still unsolved.

## Governments are in the room

China restricted OpenClaw in government offices in March and blocked Meta's \$2 billion acquisition of the Chinese agent startup Manus in April. The US government reviews frontier models before release; OpenAI's GPT-5.6 launch was delayed by it. Agents are now a matter of state.

# The jobs question

---

This is the part people actually want to know, so here is the honest version.

The layoffs are real and they are being attributed to agents by name. ClickUp cut 22% of its staff in May and called the result a "100x org" built around thousands of internal agents, with million-dollar pay bands for the people who direct them. Salesforce's CEO said the company's support team went from 9,000 to 5,000. IBM replaced around 200 HR roles with agents. Tech layoffs in the first half of 2026 were up 83% on the year before.

The counter-evidence is also real. Gartner surveyed 350 large companies deploying autonomous tech in May 2026: 80% had cut headcount, and there was no correlation between the cuts and improved returns. "Workforce reductions may create budget room, but they do not create return." A lot of what is being called AI-driven is AI-justified. And IBM, the same company replacing HR staff with agents, is tripling its entry-level hiring for AI roles.

**The pattern that holds across all of it: the roles disappearing are the ones that execute a defined process. The roles appearing are the ones that direct, review and take responsibility for agents. The people doing well are the ones who treat agents as junior staff — briefing them, checking them, expanding what they are trusted with — rather than as threats or as magic.**

# What this means for you

---

## If you run a small business

You now have the operational capacity of a team for the cost of a few subscriptions. Support, content, research, outreach, bookkeeping — all delegable today. The question is not whether to start but

where, and the answer is one well-defined repetitive task, done well, before anything else. The How to Hire an AI Agent guide walks through it.

## If you are an employee

Learn to direct agents before your employer decides whether you can. The single most valuable skill in the 2026 job market is turning a goal into a brief an agent can execute, and knowing when its output is wrong.

## If you are trying to understand this space

You are in the right place. The rest of the WTF Agents series goes deep on each player, platform and concept in this guide. None of it requires a technical background. It requires curiosity and a willingness to try things.

# Glossary

---

**AI agent** — Software that pursues a goal on its own: taking actions, using tools, checking results and adapting, without a human directing each step.

**Agentic economy** — An economy in which AI agents do meaningful economic work — and increasingly hire, build and pay each other — with minimal human supervision.

**LLM (Large Language Model)** — The AI "brain" inside every agent. Claude, GPT, Gemini, Grok, DeepSeek and the rest.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex.

**Personal agent** — An agent that runs your life's admin: OpenClaw, Meta Muse.

**Autonomous company platform** — A service where agents build and run an entire business: Polsia, NanoCorp, Cofounder.co; Paperclip is the self-hosted version.

**MCP (Model Context Protocol)** — The open standard connecting agents to tools and data. Invented by Anthropic, now under the Agentic AI Foundation.

**A2A (Agent2Agent)** — Google's open standard for agents from different companies talking to each other.

**Agentic commerce** — Agents that can pay, buy and take payment on your behalf.

**Prompt injection** — Hiding instructions in content an agent reads so it does something its owner did not ask for. The main unsolved security problem of the agentic economy.

**Open weights** — A model anyone can download and run themselves. Most Chinese frontier models ship this way; most US ones do not.

**ARR (Annual Recurring Revenue)** — The standard measure of a subscription business's size. In 2026 AI, often the last thirty days of all revenue multiplied by twelve — read it with care.

## Liked this? Go deeper.

---

This guide gave you the big picture. The rest of the series goes deep on every platform, tool and concept mentioned here.

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is an

## AI Agent?

---

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# The one-liner

---

**An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.**

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

# The simplest possible explanation

---

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

## **A chatbot (like early ChatGPT):**

You ask: "What are some good hotels in Bangkok?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

## **An AI assistant (like Siri or Alexa):**

You say: "Book me a hotel in Bangkok." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

## **An AI agent:**

You say: "Find me a good hotel in Bangkok for next weekend, under \$200 a night, near the river, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

**That is an AI agent.**

# The five properties of an AI agent

---

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

## **1. Goal-directed**

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

## 2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

## 3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

## 4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

## 5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

# How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

| Step         | What happens                                                                                                                 |
|--------------|------------------------------------------------------------------------------------------------------------------------------|
| Receive goal | The agent is given an objective: "Write and schedule five social media posts for this week."                                 |
| Plan         | The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API. |
| Act          | The agent executes step one: browses trending topics using its web tool.                                                     |
| Observe      | It reads the results and decides what is relevant.                                                                           |
| Reflect      | It considers whether its approach is working. If step one produced poor results, it adjusts.                                 |
| Repeat       | It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.   |

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

# The main types of AI agent

---

Not all agents are the same. Here are the main categories you will encounter:

## Personal agents

Run for you, on your machine or in an app. Examples: OpenClaw (open-source, talks to you over WhatsApp, Telegram or Signal), Meta's Muse (books, buys and negotiates on your behalf). These agents handle your life's admin.

## Work agents

Do your job's tasks. Every major lab now sells one: Anthropic's Cowork, OpenAI's ChatGPT Work, xAI's Grok Bot, Microsoft's Copilot Cowork. Give them an outcome — a report, a spreadsheet, a cleared inbox — and they come back with it done.

## Coding agents

The category that proved agents work. Claude Code, OpenAI's Codex, Cognition's Devin, Google's Jules. They read a codebase, change it, test it and commit, without step-by-step instruction.

## Business process agents

Embedded into a specific business function. Salesforce's Agentforce handles customer service and sales tasks inside its CRM; a wave of "AI sales rep" and "AI support rep" products do the same as standalone hires.

## Company-building agents

The most ambitious category. Platforms like Polsia and NanoCorp deploy a whole team of agents that found and run a company — website, marketing, outreach, support, finance — with no human doing the work.

## Orchestration agents

Agents that manage other agents. Paperclip is the clearest example: an org chart of agents with roles, budgets and reporting lines, coordinating to run a business together.

## Research agents

Given a topic, they browse dozens of sources, extract facts, cross-check and produce a structured report — in minutes rather than hours. Most chat products now have one built in.

# Real agents doing real things right now

---

## OpenClaw — the personal agent that became a movement

Created by Austrian developer Peter Steinberger and launched in late 2025, OpenClaw runs on your own machine and takes instructions over messaging apps. By September 2026 it had over three

million active users, its own non-profit foundation with a board, native mobile apps, and a place among the most-starred projects in GitHub's history. It also produced the first famous agent misbehaviour stories, covered below.

## **Cowork, ChatGPT Work and Grok Bot — the year of the work agent**

Between January and September 2026, Anthropic, OpenAI and xAI each launched an agent that does whole jobs rather than answering questions about them. Cowork works inside your files and apps. ChatGPT Work takes an outcome and returns finished spreadsheets, decks and apps hours later. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff." Microsoft's Copilot Cowork, built on Claude, brought the same thing to Microsoft 365.

## **Claude Code — the autonomous developer**

Anthropic's Claude Code reads entire codebases, writes and edits code, runs tests, fixes bugs and commits — autonomously. It became the fastest-adopted developer tool of its generation and the main reason Anthropic's revenue multiplied several times over in 2026. Its remote version topped the standard software-engineering benchmark in the spring.

## **Polsia — agents running whole companies**

Polsia's agents have created close to half a million companies since late 2025: a website, product, marketing, outreach and support each, with no human doing the work. About 6% survive and the founder says one in ten has earned a dollar. It is the clearest demonstration in the world of both what agents can build and what they cannot: customers.

## **Agents in enterprise customer service**

Salesforce's CEO has said its support headcount fell from roughly 9,000 to 5,000 as agents took over the routine work. IBM replaced around 200 HR roles with agents. These are the deployments that get quoted in board meetings, and they are why the jobs question in the Agentic Economy guide matters.

# **The building blocks — what makes agents possible**

---

Three things came together to make AI agents viable in 2024–2026:

### **1. LLMs got good enough**

The underlying AI models — Claude, GPT, Gemini, Grok, Muse, DeepSeek — became capable enough to reason through multi-step problems reliably and, crucially, to notice when they were going wrong. Earlier models got confused, hallucinated, or went in circles. The current generation handles complex, ambiguous tasks with enough reliability to be left alone for hours. The exact scores change every quarter; the shift from "answers questions" to "finishes jobs" does not.

### **2. Tool use became standardised**

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in 2025, letting agents from different companies communicate. MCP now lives under the neutral Agentic AI Foundation. These two protocols are the invisible infrastructure of the agentic economy.

### **3. The cost of compute collapsed**

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, a capable agent costs a few dollars a day to run, and open-weight models from China and elsewhere push it lower still — to the point where a single person can run several agents simultaneously for the cost of a few software subscriptions.

## **The risks — what can go wrong**

---

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

### **Prompt injection**

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. It remains the main unsolved security problem for agents. Cisco found third-party OpenClaw skills silently exfiltrating data in early 2026; Meta's own staff flagged security failures in Muse days before its launch.

### **Overly broad permissions**

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

### **Hallucination in action**

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

### **Loss of control**

The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent. Palsia users have reported agents sending cold emails to real contacts and making public posts without approval.

**The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.**

## How to think about agents — the mental model

---

**The most useful mental model for AI agents is not "software." It is "staff."**

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

## Glossary

---

**AI Agent** — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

**LLM (Large Language Model)** — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT (OpenAI), Gemini (Google), Muse (Meta), Grok (xAI), DeepSeek.

**ReAct** — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

**MCP (Model Context Protocol)** — The open standard for connecting AI agents to external tools and data. The "USB-C for AI." Invented by Anthropic, now under the Agentic AI Foundation.

**A2A (Agent2Agent Protocol)** — Google's open standard for AI agents from different companies to communicate.

**Tool use** — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

**Prompt injection** — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

**Orchestration** — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

**Autonomous** — Operating without human control or supervision on a task-by-task basis.

**OpenClaw** — The most-used open-source personal AI agent. Created by Peter Steinberger, late 2025; now run by a non-profit foundation.

**Claude Code** — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

## Liked this? Go deeper.

---

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is an API?

---

APIs are mentioned in almost every conversation about AI, tech, and the agentic economy. Almost no one explains what they actually are. This guide does — in plain English, finally.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · March 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# The one-liner

---

**An API (Application Programming Interface) is a set of rules that lets two pieces of software talk to each other.**

That is it. That is the whole thing. Everything else in this guide is just unpacking what that means in practice.

# The restaurant analogy

---

The best analogy for an API is a restaurant.

You are sitting at a table. You want food. The kitchen can make food. But you do not walk into the kitchen and start cooking. You do not shout instructions at the chef. You use a menu and a waiter.

The menu tells you what the kitchen can make. The waiter takes your order to the kitchen and brings the food back.

In this analogy: you are the application making a request. The kitchen is the other software (a database, a service, an AI model). The waiter is the API. The menu is the API documentation.

**The API is the agreed-upon way that two systems communicate — what you can ask for, how you ask for it, and what you get back.**

# A real-world example

---

Here is what actually happens when you use a weather app on your phone:

- You open your weather app.
- The app does not have weather data stored on your phone.
- It sends a request to a weather service's API: "Give me the forecast for Chicago."
- The weather service's servers process the request.
- The API sends back the data: temperature, humidity, forecast for the next 7 days.
- Your app displays it in a nice interface.

You never saw any of that. You just saw the weather. But an API made the whole thing work.

This same pattern — request, process, respond — underpins almost everything on the internet. Every time you log in with Google, every time you pay with Stripe, every time you share to Twitter, every time an AI agent takes an action in the world — an API is involved.

# Why APIs matter for the agentic economy

---

## APIs are the reason AI agents can do anything useful.

An AI agent on its own is just a brain in a jar — it can think and plan but it cannot act. APIs are the hands and legs. They are how agents connect to the real world.

Here are examples of what agents do via APIs:

| Action               | How the agent does it                                                                            |
|----------------------|--------------------------------------------------------------------------------------------------|
| Send an email        | The agent calls Gmail's API or SendGrid's API with the message content, recipient, and subject.  |
| Browse the web       | The agent calls a web scraping API or uses a browser automation tool exposed via API.            |
| Process a payment    | The agent calls Stripe's API with the amount and card details.                                   |
| Post to social media | The agent calls the Twitter/X API or LinkedIn API with the post content.                         |
| Read a spreadsheet   | The agent calls the Google Sheets API to retrieve or update data.                                |
| Book a calendar slot | The agent calls the Google Calendar or Calendly API.                                             |
| Query a database     | The agent calls the Supabase or Postgres API to read or write records.                           |
| Use an AI model      | The agent calls the Anthropic API or OpenAI API to get Claude or GPT to think through a problem. |

Without APIs, an AI agent would be like a genius who has been locked in a room with no phone, no computer, and no way to interact with the outside world. APIs open the door.

## MCP — the API for AI agents

---

One of the most important developments in the agentic economy is the Model Context Protocol (MCP) — invented by Anthropic and published in November 2024.

Before MCP, connecting an AI agent to a specific tool or data source required custom engineering. Every integration was different. It was like every electrical appliance having a different plug.

**MCP standardised it. It is the universal plug — the USB-C for AI. If a tool supports MCP, any AI agent can connect to it without custom code.**

As of 2026, MCP is supported by OpenAI, Google, Microsoft, and virtually all major agent frameworks. It has been donated to the Agentic AI Foundation (backed by Anthropic, OpenAI, Google, Microsoft, Amazon, and Cloudflare) — making it a permanent, neutral open standard.

This matters because it means the ecosystem of tools available to AI agents is growing exponentially. Every new MCP-compatible tool immediately becomes available to every MCP-compatible agent.

# The types of API you will hear about

---

## REST APIs

The most common type. Uses standard web requests (GET, POST, PUT, DELETE) to retrieve or send data. When someone says "the API" without specifying, they almost always mean a REST API. The Porsia API that WTF Agents uses to track company data is a REST API.

## Webhooks

Instead of you asking the API for data, the API calls you when something happens. Example: Stripe fires a webhook to your server the moment a payment completes. You do not have to keep asking "has the payment gone through?" — Stripe tells you when it does.

## GraphQL APIs

A more flexible alternative to REST. Instead of getting a fixed set of data, you specify exactly what fields you want. More efficient but more complex to use.

## Streaming APIs

Send data continuously rather than in one response. This is why ChatGPT and Claude type out their responses word by word — the AI is streaming its output via a streaming API rather than waiting until it has finished the whole response.

## SDK (Software Development Kit)

Not technically an API itself, but often confused with one. An SDK is a package of code that wraps an API and makes it easier to use. The Anthropic SDK is a Python or JavaScript package that handles the technical details of calling the Claude API.

# API keys — what they are and why they matter

---

Almost every API requires an API key — a unique string of characters that identifies who is making the request.

Think of an API key as a password for a service. When you sign up for the Anthropic API, OpenAI API, or Stripe, they give you a key that looks something like this: `sk-ant-api03-xxxxxxxxxxxxxx`

API keys matter for three reasons:

- **Authentication** — the service knows it is you making the request, not someone else.
- **Billing** — usage is tracked against your key. Lose your key and someone else runs up your bill.
- **Rate limiting** — services use keys to enforce limits on how many requests you can make.

**The golden rule: never share your API keys publicly. Never commit them to GitHub. Store them in environment variables (a secure way of passing sensitive values to software without hardcoding them into your code).**

## APIs in everyday life — you use them constantly

---

You interact with APIs dozens of times a day without knowing it:

**Logging in with Google** — Google's OAuth API verifies your identity and tells the other site you are who you say you are.

**Paying online** — Stripe's payment API processes your card securely without the website ever seeing your full card number.

**Checking the weather** — Your phone's weather app calls a meteorological data API.

**Getting directions** — Google Maps or Apple Maps API calculates your route.

**Posting to Instagram** — The Instagram API receives your photo and caption and stores it.

**Reading your emails in a third-party app** — Gmail's API gives the app access to your emails.

**Using "Sign in with Apple"** — Apple's authentication API verifies you.

**An AI agent booking your travel** — The agent calls APIs for flights, hotels, and calendar — all in one automated flow.

## Do you need to understand APIs to use AI agents?

---

**Honestly — not really, for most use cases.**

Platforms like OpenClaw, Polsia, and Paperclip handle the API connections for you. You describe what you want the agent to do, and the platform handles the technical plumbing of which APIs to call and how.

But understanding what APIs are — and that they exist — helps you in three ways:

- You understand why agents can do what they can do (and why they sometimes cannot).
- You can evaluate agent platforms more intelligently — asking "what APIs does this connect to?"
- If you want to build something custom, you know what to look for.

The agentic economy runs on APIs. You do not need to be a plumber to live in a house with running water. But knowing that pipes exist, and roughly how they work, makes you a more informed resident.

# Glossary

---

**API (Application Programming Interface)** — A set of rules that lets two pieces of software communicate with each other.

**REST API** — The most common type of API. Uses standard web requests to send and receive data.

**Webhook** — An API that pushes data to you when something happens, rather than waiting for you to ask.

**API key** — A unique identifier used to authenticate API requests. Treat it like a password.

**MCP (Model Context Protocol)** — Anthropic's open standard for connecting AI agents to external tools. The universal plug for AI.

**SDK (Software Development Kit)** — A package of code that makes it easier to use an API in a specific programming language.

**Endpoint** — A specific URL that an API exposes for a specific type of request. E.g. /api/data is the endpoint for Polsia's company data.

**Rate limiting** — A restriction on how many API requests you can make in a given time period.

**Authentication** — The process of proving who you are to an API, usually via an API key or OAuth token.

**Streaming** — An API that sends data continuously rather than all at once. How Claude and ChatGPT deliver responses word by word.

## Liked this? Go deeper.

---

Now you know what APIs are — here is where to go next in the WTF Agents series.

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is an LLM?

---

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.**

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

## How does predicting words create intelligence?

---

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

**The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.**

## Where did LLMs come from?

---

The story of LLMs is surprisingly recent and moves fast:

| Year | What happened                                                                                                                                                      |
|------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2017 | Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.          |
| 2018 | Google releases BERT. OpenAI releases GPT-1. The race begins.                                                                                                      |
| 2020 | OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.                                                       |
| 2022 | OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.                                                              |
| 2023 | GPT-4 released. Open-source models (Llama, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.                                              |
| 2024 | Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.                                                               |
| 2025 | Claude Code launches. Chinese open-weight models (DeepSeek, Qwen, Kimi) match the frontier at a fraction of the cost. Multi-hour autonomous tasks become reliable. |
| 2026 | Every major lab ships a work agent. Anthropic splits its top model into a public (Fable) and restricted (Mythos) tier. GPT-6, Gemini 3.8, Meta Muse, DeepSeek V4.  |

## The main LLMs in 2026 — who makes what

---

Model names and version numbers change every few months. The labs behind them, and what each is known for, change far more slowly. Read the labs, not the numbers.

### Claude (Anthropic)

The model family behind Claude Code, Cowork, Microsoft's Copilot Cowork and the default in OpenClaw. Tiers: Fable (the flagship, a new tier above Opus introduced in 2026), Opus, Sonnet and Haiku. Anthropic's Constitutional AI training makes Claude the model most agent builders reach for when reliability and honesty over long unsupervised sessions matter. Context window: one million tokens. See the [WTF is Claude](#) guide.

### GPT (OpenAI)

The models behind ChatGPT, still the most widely used consumer AI product in the world, and behind OpenAI's agents: ChatGPT Work, Codex and the Agents API. GPT-6 arrived in September 2026 after the GPT-5.6 family in July. OpenAI's strength is reach: hundreds of millions of people already have the app.

### Gemini (Google DeepMind)

Google's flagship family, currently Gemini 3.8, and the model behind Gemini Agent, Chrome's Auto Browse and Google's coding agent Jules. Deeply integrated with Search, Workspace and Android. Google also created the Agent2Agent (A2A) protocol for agents from different companies to talk to each other.

### Muse (Meta)

Meta's 2026 reset: a new model family and a consumer personal agent, both called Muse, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is distribution — three billion people already use its messaging apps.

## Grok (xAI)

Elon Musk's models, now part of SpaceX after it acquired xAI and the coding tool Cursor. The model behind Grok Bot, a team of agents sharing one cloud computer.

## The Chinese open-weight models

DeepSeek, Alibaba's Qwen, Moonshot's Kimi, Zhipu's GLM and MiniMax. Released as open weights — anyone can download and run them — at a fraction of the price of US frontier models, and close enough in capability that a great many indie agents run on them. DeepSeek's cheap-to-train breakthrough in early 2025 was the moment this became a two-continent race. Guide: [WTF is Happening with China's AI](#).

## Other open models

Meta's Llama family, once the default open model, has been superseded inside Meta by Muse. Mistral (France) produces efficient smaller models. OpenClaw, Paperclip and most agent frameworks let you plug in any of these by API key.

# What LLMs are good at — and bad at

| LLMs are good at                      | LLMs are bad at                                      |
|---------------------------------------|------------------------------------------------------|
| Writing in any style or tone          | Precise arithmetic (they approximate)                |
| Summarising long documents            | Real-time information (training data has a cutoff)   |
| Writing and explaining code           | Guaranteed factual accuracy (they can "hallucinate") |
| Translating between languages         | Knowing what they don't know                         |
| Reasoning through multi-step problems | Consistent behaviour on identical inputs             |
| Generating creative ideas             | Tasks requiring persistent memory by default         |
| Explaining complex topics simply      | Physical world interaction (without agent tools)     |

# Hallucination — the most important limitation

**The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.**

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

## Parameters, tokens, and context windows

---

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

### Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the "size" of the model's brain.

### Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

### Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude's current models have a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.

## Constitutional AI — Anthropic's approach to safety

---

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has

internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

## LLMs and the agentic economy

---

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.
- Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Fable, GPT-6 and Gemini 3.8 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

## Glossary

---

**LLM (Large Language Model)** — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, Gemini, Muse and Grok.

**Transformer** — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

**Parameters** — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

**Token** — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

**Context window** — How much text an LLM can process at once. Current frontier models: up to 1 million tokens.

**Hallucination** — When an LLM generates confident but factually incorrect text. A key limitation to understand.

**Constitutional AI** — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

**SWE-bench** — The standard benchmark for measuring LLM performance on real software engineering tasks.

**Fine-tuning** — Training an existing LLM further on a specific dataset to improve performance on a specific task.

**Inference** — Running an LLM to generate a response. The "thinking" part, as opposed to training.

**Open-weight model** — An LLM whose weights are publicly available — anyone can download and run it. Examples: DeepSeek, Qwen, Kimi, Mistral, Llama.

## Liked this? Go deeper.

---

Now you know what LLMs are. Here is the rest of the WTF Agents series.

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Polsia?

---

One founder. Zero employees. Nearly half a million AI-built companies, a \$250 million valuation, and a Trustpilot page that tells a very different story. Here is what Polsia actually is, what it costs, and who should — and should not — touch it.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Polsia is a subscription platform where you describe a business idea — or let the AI pick one — and a team of AI agents builds the website, writes the code, runs the marketing, sends the outreach, answers support and reports back to you every morning. The founder's tagline: "AI that runs your company while you sleep."**

It is also the most-argued-about product in the agentic economy. Investors valued it at a quarter of a billion dollars. Paying customers gave it one star in large numbers. Both of those things are true at the same time, and understanding why is the point of this guide.

## By the numbers — September 2026

---

Everything in this box changes weekly. Polsia publishes a live public dashboard, which is how these figures can be checked. Everything below the box is written to still be true when the numbers have moved.

**\$30M**

raised in May 2026 at a \$250M valuation (Sound Ventures, True Ventures)

**\$18.7M**

headline ARR on Polsia's live feed, 15 Sept 2026 — about \$11M of it recurring subscriptions

**484,592**

companies created on the platform since launch

**29,132**

of those still active — roughly 6%

**~50%**

month-one churn, stated by the founder on the record

**10%**

of companies that have ever earned at least one dollar, per the founder

**2.9 / 5**

Trustpilot rating across 185 reviews, 45% of them one star

# 1

employee. Still just Ben.

## Who built it

Polsia was built by Ben Broca, who also goes by Ben Cera. He is French, trained as an engineer at CentraleSupélec with a masters from Columbia, and now lives in San Francisco. Before Polsia he ran international operations at CloudKitchens under Travis Kalanick and co-founded Hutch, a 3D-rendering e-commerce startup that raised roughly \$17 million.

He spent most of 2025 building separate AI tools — a coding tool, a marketing tool — before concluding the models were good enough to run every business function at once. Polsia launched in late 2025, crossed \$1 million in annualised revenue within about a month, and tripled the month after. He has been the only employee throughout, and says that is deliberate: if he is selling an AI that runs companies, Polsia should be run by one.

In May 2026 he raised \$30 million at a \$250 million valuation, reportedly letting the AI manage the data room and brief investors. That is the number the headlines quote. It says something about what investors believe. It says nothing about whether the product will work for you.

## How it works

You give Polsia an idea, or click "surprise me" and it picks one from your public profile. Within minutes it has a company name, a mission statement, a market research document, a live website on a polsia.app subdomain, an email inbox, a Stripe account and a GitHub repository. Then it asks for a credit card.

Once you subscribe, the agents run on a daily cycle:

| Agent              | What it does                                                                      |
|--------------------|-----------------------------------------------------------------------------------|
| Orchestrator (CEO) | Reviews the state of the company, writes the morning plan and the evening summary |
| Code               | Ships features and fixes to the website and product                               |
| Marketing          | Drafts and posts social content, writes landing pages                             |
| Outreach           | Finds prospects and sends cold email                                              |
| Support            | Reads the inbox and drafts replies                                                |
| Ads                | Runs and optimises Google and Meta campaigns from your budget                     |
| Finance            | Syncs Stripe revenue, tracks spend                                                |
| Research           | Watches competitors and refreshes the plan                                        |

Each task the agents run costs a credit. Your plan includes a monthly allowance; you buy more if you run out. You read the summaries, approve or redirect, and the cycle repeats whether or not anyone is buying what the company sells.

**The most important thing to understand about how Polsia works is what it does not do. There is no step between "here is the idea" and "the agents are building it" that asks whether anyone wants the thing.**

## What it actually costs

---

Polsia does not publish pricing on its homepage. As of August–September 2026, the in-product pricing showed five tiers:

| Plan     | Price      | Includes                             |
|----------|------------|--------------------------------------|
| Standard | \$20/mo    | 1 company, 20 credits/mo             |
| Pro      | \$50/mo    | 3 companies, 50 credits/mo, Meta Ads |
| Max      | \$200/mo   | 10 companies, 200 credits            |
| Max 2    | \$500/mo   | 25 companies, 500 credits            |
| Max 3    | \$1,000/mo | 50 companies, 1,000 credits          |

The subscription is not the whole bill. Three more lines matter, all from Polsia's own Terms of Service:

→ **Ad spend fee — 20%.** Any advertising budget the agents spend on your behalf carries a 20% platform fee.

→ **Revenue fee — 3%.** Polsia takes a cut of customer payments your company receives. In the September 2026 terms this is "currently 3%"; until then it was 20%. Check the live terms — it has already changed once.

→ **Withdrawal limits.** Your company's revenue lands in a Polsia balance first. Payouts have a 31-day settlement period, a \$50 minimum and a cap of \$500 per calendar month. A company earning \$2,000 a month can take out \$500 of it.

Users have also reported the credit meter running on a single click with no confirmation, and credits burned on tasks that failed. Budget for the plan, the credits, and the fees — then read the withdrawal cap twice.

## The numbers — what the live feed actually says

---

Polsia publishes a public dashboard feed with enough raw data to rebuild the headline figure yourself. On 15 September 2026 it read \$18.7 million ARR.

That number is built the way many AI startups build ARR in 2026: take the last thirty days of revenue, including one-off payments, and multiply by twelve. Ben has described the method himself on the record. Broken down, roughly \$11 million is annualised subscription revenue. The rest is the last month of credit packs, ad budget routed through the platform, domain purchases and a small slice of user-company revenue, each times twelve. Real money, but around 40% of it is not recurring by construction.

The second number on the feed is the one that matters more for you. Nearly half a million companies have been created. About 29,000 are active. Companies are cheap to spin up here, so a low survival rate is expected rather than damning — but it is the clearest public signal of what happens after the agents finish building.

Ben gave the third number to Mixergy's Andrew Warner in May 2026, in an interview titled "Is Polsia a \$250M scam?": about 10% of companies on the platform have made at least a dollar. The best one he had seen was making three or four thousand a month. Month-one churn is around 50%. He also said the engineering agent was writing "spaghetti code" and the product "doesn't work well enough for my taste."

**That is unusual candour from a founder at a \$250 million valuation, and it is the honest frame for the whole product: the agents really do build things. Nine times out of ten, nobody buys them.**

## What paying customers say

---

Polsia's Trustpilot page tells the story in three acts. In June 2026 it had 35 reviews. By mid-August it was 1.8 out of 5 across 79 reviews, three quarters of them one star. Polsia then claimed the profile, over a hundred reviews arrived in four weeks, and by mid-September it stood at 2.9 out of 5 across 184 reviews, with 45% one star and 27% five star. Several of the new five-star reviews describe a problem the founder fixed personally.

The complaints did not change across that period. The recurring themes:

- Tasks marked "complete" by the AI that never actually deployed
- Credits consumed by duplicate or failed actions, with refunds refused
- Cold emails sent to real contacts, or public posts made, without approval
- Websites and custom domains locked to Polsia's own hosting account
- Sites redeployed after an explicit takedown request
- Cancellations that took repeated attempts, and support tickets unanswered for weeks

The most-cited case is a Rest of World profile from April 2026: a factory worker in China paying \$199 a month — a quarter of his salary — for an AI to sell a spiritual guidance app to Americans. The

agents built the site, filled it with fabricated reviews, ran Facebook ads and emailed journalists he had never heard of. After several months: seven signups, zero customers.

The positive cases are real too. Founders who arrived with a validated idea and a clear brief report the agents standing up infrastructure, research, outreach and even video ads in hours — plumbing that would otherwise eat a week. One four-star reviewer who built a real product on it put the line exactly: the weakness is not what the agents can build, it is what the customer cannot control when the surrounding infrastructure fails.

## The lock-in question

---

This is the part most first-time users miss. Polsia deploys your company on Polsia's hosting, on a `polsia.app` subdomain, with email routed through Polsia. There is a "download code" button, but users report the export is a scaffold rather than a working product, and that there is no direct access to the hosting or repository. Self-hosting is listed as "coming soon."

If the platform is working for you, none of this matters. If you want to leave, take your domain with you, or hand the code to a developer, it matters a great deal. Verify the current export and account-recovery terms in writing before you subscribe.

## Who Polsia is for

---

**Polsia is an execution engine. It is not an idea-validation engine, and it does not pretend to be one.**

### It fits if you:

- Already have evidence people want the thing — conversations, a tested price, a channel that works
- Are launching cheap, fast experiments where one flop costs little
- Are happy paying the fees and living with the withdrawal cap in exchange for not touching any of the plumbing
- Do not need to own the code or the infrastructure

### Skip it if you:

- Have an idea but have not talked to a single potential customer
- Are in a regulated, technical or high-stakes category
- Need to withdraw more than \$500 a month, or would be hurt by 20% on ad spend
- Plan to migrate off later, or want portable code from day one

The "surprise me" button is the most marketed feature and the riskiest one. It means the AI is choosing the idea and executing it with no demand signal anchoring either decision. The factory worker in the Rest of World story is what that looks like in practice.

## Our verdict

---

Polsia is real, funded, and built by a credible operator. The agents genuinely build, market and operate companies with no human doing the work, and the platform has grown faster than almost anything in the category. The hype question is settled.

The question that matters now is fit. Polsia executes whatever it is given. Hand it a validated idea and a precise brief and it can save you weeks. Hand it a guess and it will build a polished company nobody asked for, spend your credits and your ad budget promoting it, and report back cheerfully every morning while nothing sells. The founder's own figure — one company in ten ever earns a dollar — is that gap measured from inside.

Validate first. Then decide whether Polsia is the right layer to execute on what you found.

## Glossary

---

**Polsia** — Subscription platform where AI agents build and run a company for you. Founded by Ben Broca (Ben Cera), launched late 2025, San Francisco.

**Credits** — Polsia's unit of work. Each task the agents run costs a credit. Plans include a monthly allowance; extra credits are bought in packs.

**ARR (Annual Recurring Revenue)** — Annualised recurring revenue. In much of the 2026 AI industry, including Polsia's headline figure, it means the last thirty days of all revenue multiplied by twelve.

**Run rate** — A recent period's revenue extrapolated to a full year. Polsia's headline number is a run rate.

**Platform fee** — Polsia's cut. 20% on advertising spend and, as of September 2026, 3% on customer payments (previously 20%).

**Withdrawal cap** — The limit on how much of your company's revenue you can take out of your Polsia balance: \$500 per calendar month, after a 31-day settlement period.

**Surprise me** — Polsia's option to let the AI choose your business idea for you.

**Validation** — Checking that people actually want a product before building it. The step Polsia does not perform.

**Churn** — The share of paying customers who cancel in a given period. Polsia's founder puts month-one churn at around 50%.

**polsia.app** — The subdomain every Polsia company is deployed on. Your company's site and email run through Polsia's infrastructure by default.

## Liked this? Go deeper.

---

Polsia is one piece of the agentic economy. Here is the rest of the picture.

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is NanoCorp

The Polsia rival built by AI researchers that publishes its companies' real revenue. How it works, what it costs, and what the live numbers say.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## OpenClaw?

---

One Austrian developer. One open-source project. Three name changes in 60 days, a hire by OpenAI, dating profiles created without consent, a ban from Chinese government offices — and now millions of users, a foundation, and native apps. The wildest story in tech, and where it stands today.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

OpenClaw is a free, open-source autonomous AI agent that runs on your computer or phone and takes instructions via messaging apps — WhatsApp, Telegram, Signal, Discord — connecting to Claude, GPT, DeepSeek or any model you choose to carry out multi-step tasks on your behalf.

## By the numbers — September 2026

---

**3.2M**

active users, making it the most-used open-source agent in the world

**385,000+**

GitHub stars — the sixth most-starred repository in GitHub's history

**30 Aug**

2026: OpenClaw 2.0 released, with native iPhone, iPad and Android apps

**\$0**

cost. MIT licence. You pay only for the model you plug in

**1**

non-profit foundation now governing the project, with an independent board

**3**

names in its first 60 days: Clawdbot, Moltbot, OpenClaw

It is not a platform. It is not a service. It is software you install and run yourself. And it became one of the most talked-about pieces of software in the world in under 60 days.

## The name history — three identities in 60 days

---

Understanding OpenClaw requires understanding its chaotic naming history:

| Date              | Name     | Why                                                                                                                                                        |
|-------------------|----------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| November 24, 2025 | Clawdbot | Original launch name. A portmanteau of "Clawd" (the developer's personal AI assistant, itself named after Anthropic's Claude) and "bot."                   |
| January 27, 2026  | Moltbot  | Anthropic sent trademark complaints about "Clawd" being too similar to "Claude." Renamed to Moltbot, keeping a lobster theme (lobsters molt their shells). |
| January 30, 2026  | OpenClaw | Peter Steinberger found "Moltbot" didn't roll off the tongue. Renamed again to OpenClaw — and it stuck.                                                    |

## Who built it — Peter Steinberger

Peter Steinberger is an Austrian software engineer and entrepreneur. He co-founded PSPDFKit (now Nutrient) in 2011 — a PDF SDK framework he bootstrapped to a successful exit. He is based between Vienna and London.

Steinberger built his own personal AI assistant called Clawd — named after Anthropic's Claude, which it used as its underlying model. He then open-sourced a generalised version of it as Clawdbot in November 2025.

He describes himself as a "vibe coder" — someone who builds primarily through AI-assisted development. The irony of an AI agent builder who uses AI agents to build AI agents was not lost on the internet.

**On February 14, 2026 — Valentine's Day — Steinberger announced he was joining OpenAI and that OpenClaw would move to an independent open-source foundation, financially sponsored by OpenAI. This was widely described as an "acqui-hire": OpenAI acquired the talent without acquiring the open-source software itself.**

## What OpenClaw actually does

OpenClaw runs locally on your machine. You interact with it via a messaging app — Signal, Telegram, Discord, or WhatsApp. You send it a message with a task. It uses an LLM (Claude by default, or GPT or DeepSeek) to reason through the task and executes it using whatever tools and permissions you have given it.

What it can do depends on what you connect it to. Examples:

- Read and send emails on your behalf
- Browse the web and summarise findings
- Manage your calendar

- Write, edit, and commit code
- Post to social media
- Execute terminal commands
- Interact with any service that has an API

OpenClaw uses a "skills" system — modular capabilities stored as directories that can be installed from ClawHub, the community skill registry. Think of skills as apps for your AI agent.

## The growth

---

**OpenClaw's growth was unprecedented. By February 2026 — just over two months after launch — it had 247,000 GitHub stars. By September it had passed 385,000, the sixth most-starred repository in GitHub's history, and over three million people were running it.**

For context: Linux took decades to accumulate its GitHub stars. OpenClaw surpassed most long-standing projects in weeks and kept going. The 2.0 release in August, with native mobile apps and Windows support, turned it from a developer's tool into something ordinary people install.

The timing at launch was perfect. OpenClaw arrived just as Moltbook went viral — a social network built for AI agents. The two projects fed each other's growth: OpenClaw agents connecting to Moltbook, humans watching in fascination and horror.

## Moltbook — the AI social network that made OpenClaw famous

---

On January 27, 2026 — the same day Clawdbot was renamed Moltbot — entrepreneur Matt Schlicht launched Moltbook: a social network built exclusively for AI agents.

The premise was wild: a Reddit-like platform where AI agents (not humans) create profiles, post content, and interact with each other. OpenClaw agents began autonomously discovering Moltbook, joining it, and posting — without their human owners necessarily knowing or intending this.

The internet went predictably insane. Wired ran a story: "I Infiltrated Moltbook, the AI-Only Social Network Where Humans Aren't Allowed." TechCrunch, CNBC, The Verge, and mainstream media followed.

Then things got worse. On January 31, an unsecured Supabase database exposed 6,000+ email addresses and 1 million+ agent interactions — and crucially, allowed anyone to take control of any agent on the platform. Security firm Wiz confirmed the vulnerability. Moltbook went offline briefly to

patch it.

Then things got weirder. An OpenClaw agent autonomously created a dating profile on MoltMatch (an AI agent dating platform) and began screening romantic matches — without its human owner's knowledge or consent.

**Moltbook was acquired by Meta on March 10, 2026. Matt Schlicht and co-founder Ben Parr joined Meta Superintelligence Labs.**

## The security problems — read this before installing

---

**OpenClaw is powerful. It is also genuinely dangerous if used carelessly. This is not marketing caution. Multiple credible security researchers have documented real risks.**

### Cisco's findings

Cisco's AI security research team tested a third-party OpenClaw skill called "What Would Elon Do?" They found it performed silent data exfiltration and prompt injection without user awareness. They also found the ClawHub skill repository lacked adequate vetting to prevent malicious skill submissions.

### Remote code execution vulnerabilities

Security firm Conscia disclosed a remote code execution chain and two additional command injection vulnerabilities in early 2026.

### The maintainer's own warning

One of OpenClaw's own maintainers, known as "Shadow," publicly warned on Discord: "If you can't understand how to run a command line, this is far too dangerous of a project for you to use safely."

### China bans it

In March 2026, Chinese authorities restricted state-run enterprises and government agencies from running OpenClaw apps on office computers.

The rule of thumb for OpenClaw: give it the minimum permissions it needs. Only install skills from trusted sources. Review its actions regularly. If you are not comfortable with the command line, start with something simpler.

# OpenClaw vs Polsia vs Paperclip

---

|             | OpenClaw                | Polsia                   | Paperclip               |
|-------------|-------------------------|--------------------------|-------------------------|
| Type        | Personal AI agent       | Managed company platform | Orchestration framework |
| Runs        | Your machine or phone   | Polsia's servers         | Your servers            |
| Cost        | Free (+ LLM costs)      | \$20–1,000/mo + fees     | Free (open source)      |
| For         | Personal automation     | Launching AI companies   | Multi-agent companies   |
| Technical   | Medium (CLI)            | Low (no-code)            | High (developer)        |
| Open source | Yes (MIT)               | No                       | Yes                     |
| LLMs        | Any — Claude by default | Claude Opus              | Claude Code + 7 others  |

## Where it stands now — the OpenClaw Foundation

---

The move to a foundation, announced when Steinberger joined OpenAI, has happened. OpenClaw is now governed by an independent non-profit with a board chaired by Dave Morin, the founder of Path. OpenAI remains a financial sponsor — an arrangement that still gives one of the world's most powerful AI companies influence over the most popular open-source agent, though the board structure now sits between the two.

Development has not slowed. OpenClaw 2.0 shipped on 30 August 2026 with native iPhone, iPad and Android apps and a rebuilt skills system. Security improved in specific ways: an approval is now tied to the exact request, command, session and person it was granted for, rather than becoming a standing permission, and protected credentials can reach the services that need them without ever appearing in text the model can read. Prompt injection, as everywhere in the agentic economy, is not solved.

OpenClaw is the reference implementation for what a personal AI agent looks like, and the benchmark every competitor — including Meta's Muse — is measured against. It is also hired into Paperclip org charts as a worker, which makes it one of the few agents that lives in both the personal and the company layer.

## Glossary

---

**OpenClaw** — Free, open-source autonomous AI agent. Runs on your machine or phone, connects via messaging apps. Created by Peter Steinberger, November 2025; governed by the OpenClaw Foundation since 2026.

**Clawdbot / Moltbot** — Previous names for OpenClaw. Clawdbot was the original name; Moltbot followed after Anthropic trademark complaints.

**Peter Steinberger** — Austrian developer who created OpenClaw. Co-founded PSPDFKit. Joined OpenAI February 2026.

**Acqui-hire** — When a company hires a founder or team without technically acquiring their product. OpenAI hired Steinberger; OpenClaw remained open source.

**Moltbook** — A social network for AI agents. Launched January 2026; its founders joined Meta in March 2026.

**ClawHub** — The community skill registry for OpenClaw — modular capabilities that can be installed to extend the agent.

**Skills** — Modular capabilities for OpenClaw stored as directories containing a SKILL.md file.

**Prompt injection** — A security attack where malicious instructions are hidden in data an agent reads. Found in OpenClaw skills by Cisco.

**MIT License** — An open-source licence allowing anyone to use, modify, and distribute the software freely.

**MoltMatch** — An experimental AI agent dating platform where OpenClaw agents were creating profiles without user knowledge.

## Liked this? Go deeper.

---

OpenClaw is one part of the story. Here is the rest.

### WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Paperclip?

---

An org chart for AI agents. Roles, reporting lines, budgets and approvals — for a team where nobody is human. Open-source, self-hosted, and quietly becoming the way people actually run more than one agent at a time.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

Paperclip is open-source software that lets you run AI agents the way a company runs employees: an org chart with roles and reporting lines, a budget for each agent, tasks that flow down and results that flow up, and a human at the top who approves what matters.

Its own tagline has softened since launch — from "zero-human companies" to "the app people use to manage AI agents for work" — but the idea is the same. Once you have more than one agent, you need a way to manage them. Paperclip is that way.

## By the numbers — September 2026

---

**4 March**

2026: public launch. Under three weeks to 30,000 GitHub stars

**~70,000**

GitHub stars by June, among the fastest-growing agent projects ever released

**100+**

contributors; releases roughly every two weeks, the latest on 2 September

**8**

agent types it can hire out of the box: Claude Code, Codex, Gemini CLI, Cursor, Hermes, OpenClaw, Pi, OpenCode

**\$0**

cost of the software. MIT licence. You pay only for the agents' model usage

**~60**

forkable company templates on ClipMart, Paperclip's marketplace

**0**

disclosed funding rounds. Built by a pseudonymous developer, "dotta," and a community

## The problem it solves

---

One agent is easy. You open Claude Code, or OpenClaw, or Cowork, and you talk to it.

Five agents are chaos. Five terminal windows. Five separate memories that do not know what the others have done. Two of them doing the same task. One of them stuck since Tuesday. No idea what any of it is costing. Nobody enforcing quality. This is what most people hit within a week of trying to run a business on agents, and it is why so many "AI-run company" experiments end as a pile of abandoned sessions.

**Paperclip's answer is to stop treating agents as tools you operate and start treating them as staff you manage.**

## How it works

---

You install Paperclip on your own machine or server. It runs a local web dashboard. In it, you create a company.

### Hire agents into roles

Each agent gets a job title, a reporting line and a written brief. "Marketing lead, reports to CEO." "Backend engineer, reports to CTO." The CEO can itself be an agent, with you above it as the board.

Paperclip is agent-agnostic. Its own phrase: "If it can receive a heartbeat, it's hired." Claude Code for engineering, Codex or Gemini CLI for other tasks, Hermes for research, OpenClaw for anything that needs a personal agent's tools, all under one org chart, each on whichever model suits the job and the budget.

### Work flows as issues

Tasks are tickets, like a Jira or Linear board. A goal goes in at the top; the CEO agent breaks it into issues and assigns them down the chart. Agents pick up their issues, do the work, comment on the thread, hand off, and close. Parent and child tasks, dependencies and blockers are all tracked, and an agent that is blocked wakes the agent that can unblock it.

### Heartbeats keep it alive

Each agent runs on a schedule, a "heartbeat." Every tick it checks its inbox, does the next thing, and reports. Watchdogs review tasks that have stalled so you do not have to. It runs while you sleep, but on a leash.

### Budgets and approvals

Every agent has a spending cap. Anything sensitive — a payment, a deploy, a public post — can be gated behind human approval, and approvals sit in one queue. Execution policies define what an agent may do without asking. You set the leash length.

### Teams, sandboxes, plugins

Since summer 2026 Paperclip supports multiple human users with their own credentials, a company-wide timeline of what every agent did and when, isolated sandbox environments for agents

to work in, a plugin system for adding new agent types, and an MCP server so other tools can drive Paperclip itself.

## Paperclip vs Polsia vs OpenClaw

---

These three get confused constantly. They are different layers.

|                      | Paperclip                    | Polsia                               | OpenClaw                |
|----------------------|------------------------------|--------------------------------------|-------------------------|
| What it is           | Manager for a team of agents | Platform that runs a company for you | A single personal agent |
| Who hosts it         | You                          | Polsia                               | You                     |
| Cost                 | Free + model usage           | \$20–1,000/mo + fees                 | Free + model usage      |
| Who picks the agents | You, any vendor              | Polsia's fixed team                  | It is the agent         |
| Control              | Full — code, data, budgets   | Limited — their infra, their fees    | Full                    |
| Effort               | High: setup and management   | Low: subscribe and read reports      | Medium                  |

Polsia is the managed hotel. Paperclip is the house you build and run. OpenClaw is one very capable employee, who can also be hired into a Paperclip org.

## ClipMart — fork a company

---

At launch Paperclip announced ClipMart, a marketplace for AI-agent companies. It has since shipped at [clipmart.ai](https://clipmart.ai), and the pitch is "fork an AI company, launch in 60 seconds."

Instead of designing an org chart from scratch, you pick a template — a marketing agency, a SaaS startup, an e-commerce store, a media company — and fork it into your own Paperclip. The template comes with the roles, the briefs, the reporting lines and an estimate of how many agents it runs and what they cost. Around sixty templates across marketing, SaaS, e-commerce, agency, media and finance were listed at time of writing, and anyone can submit their own.

Two things worth knowing. Forking a company is not the same as having customers, which is the trap in every "launch in 60 seconds" pitch in this series. And a template is only as good as the briefs inside it; read them before you hire.

# The reality check

---

**Paperclip is the most serious open-source answer to "how do I run more than one agent," and it is not a product for people who want a button.**

## What is genuinely good

- The org-chart model works. People who use it describe the mental shift from "prompting an AI" to "managing a team," and that shift is the point.
- Agent-agnostic for real. Mix Claude, OpenAI, Google and open models by role and cost.
- Governance is built in, not bolted on: budgets, approvals, execution policies, audit timeline.
- It builds itself. The creator runs Paperclip's own development through a Paperclip company, and the release cadence backs that up.
- No lock-in, no account, no vendor. Your data and code stay on your machine.

## What is not

- Self-hosted means you host it. A server, PostgreSQL, Node.js, and you doing the updates. A hosted cloud version has been on the roadmap all year and had not launched at time of writing.
- The agents are only as good as their briefs. Paperclip organises work; it does not make an agent competent at it.
- No disclosed team, funding or company behind it beyond a pseudonym and a community. That is how great open-source projects start and also how some stall.
- It is a control plane, not a business. Nothing about Paperclip finds customers. See the Polsia guide for how that usually goes.

# Who Paperclip is for

---

**Yes** — Founders and operators already running two or more agents who are losing track. Developers comfortable with a terminal and a server. Anyone who wants to own the whole stack rather than rent it.

**Not yet** — Anyone who has not run a single agent successfully. Start with OpenClaw or Cowork, get one thing working, then come back.

**No** — Anyone who wants "AI runs my business" as a subscription with no setup. That is Polsia, with all the trade-offs in that guide.

# Glossary

---

**Paperclip** — Open-source, self-hosted software for managing AI agents as an org chart with roles, budgets and approvals. Launched March 2026.

**Orchestration** — Coordinating multiple agents so they work together without duplication or chaos. What Paperclip does.

**Control plane** — The layer that manages and governs agents, as opposed to the agents themselves. Paperclip's own description of its role.

**Heartbeat** — The scheduled tick on which an agent checks for work and acts. Any agent that can receive one can be hired into Paperclip.

**Adapter** — The connector that lets a particular agent type (Claude Code, Codex, OpenClaw...) work inside Paperclip.

**Issue** — A unit of work on the board: assigned, threaded, tracked, closed. Paperclip's equivalent of a task or ticket.

**Execution policy** — The rules for what an agent may do without human approval.

**Watchdog** — An automated check that reviews stalled tasks and wakes or escalates them.

**ClipMart** — Paperclip's marketplace of forkable AI-company templates at [clipmart.ai](https://clipmart.ai). Pick a template, launch it into your own Paperclip.

**MIT licence** — A permissive open-source licence. Free to use, modify and sell, with attribution.

**Self-hosted** — Software you install and run on your own machine or server, rather than a service someone else runs for you.

## Liked this? Go deeper.

---

Paperclip manages agents. Here are the agents, and the alternatives.

### WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is NanoCorp

The Polsia rival built by AI researchers that publishes its companies' real revenue. How it works, what it costs, and what the live numbers say.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Anthropic?

---

Founded by people who left OpenAI over safety. Now the most valuable private company in the world, heading for the largest IPO in history, and the lab whose model sits inside half the agents in this series — including its rivals' products.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Anthropic is an AI safety company that builds some of the most powerful AI models in the world — the Claude family — while arguing, in public and in its own research, that models like these could be among the most dangerous technologies ever created.**

That is not a contradiction. It is the founding bet: if powerful AI is coming regardless, better to have a safety-focused lab at the frontier than to leave the frontier to labs that are not. In 2026 that bet paid off in a way nobody predicted.

## By the numbers — September 2026

---

Everything in this box will be out of date within months. Everything below it is built to last longer.

**\$965B**

valuation after the \$65B Series H in May 2026 — the largest private funding round in history

**\$380B → \$965B**

the valuation jump between February and May 2026, three months apart

**\$65B**

annualised revenue run rate by late July 2026, up from roughly \$9B at the end of 2025

**1 June**

2026: confidential S-1 filed with the SEC. An October listing has been widely reported, not confirmed

**~2,300**

employees, up from about 1,100 in 2024

**8 of 10**

Fortune 10 companies are Claude customers

**1,000+**

enterprises spending more than \$1M a year on Anthropic products

**2021**

founded, San Francisco. Structured as a Public Benefit Corporation

## The founding story — the OpenAI exodus

---

In 2021 a group of senior researchers left OpenAI, then the most prominent AI lab in the world, over disagreements about direction. The core concern: as OpenAI became more commercial, some believed it was prioritising speed and capability over safety.

The group was led by siblings Dario Amodei, then OpenAI's VP of Research, and Daniela Amodei, its VP of Operations. They left with Tom Brown (lead author of the GPT-3 paper), Chris Olah (the field's leading interpretability researcher), Jack Clark, Jared Kaplan, Sam McCandlish and others.

**They founded Anthropic with a stated mission: the responsible development and maintenance of advanced AI for the long-term benefit of humanity.**

They structured it as a Public Benefit Corporation, which legally allows the company to weigh its mission alongside shareholder returns, and later added a Long-Term Benefit Trust with the power to appoint board members. Whether those structures hold under trillion-dollar pressure is one of the open questions of the decade.

## The scale — from research lab to the most valuable startup on earth

---

For its first four years Anthropic was the smaller, more careful rival. Then Claude Code happened.

Launched to developers in early 2025, Claude Code became the tool that turned Claude from a chatbot into an agent that does work. Developers adopted it faster than any coding tool before it. Cework, launched in January 2026, brought the same capability to everyone else. Revenue followed the products: roughly \$9 billion annualised at the end of 2025, \$30 billion by April, \$47 billion when the Series H closed in May, \$65 billion by the end of July.

The Series H itself — \$65 billion led by Altimeter, Dragoneer, Greenoaks and Sequoia, with Amazon, Capital Group, Fidelity and T. Rowe Price participating — valued the company at \$965 billion, above OpenAI's \$852 billion from two months earlier. Nine days later Anthropic filed confidentially for an IPO. Reports point to October 2026 and a target valuation that would make it the largest listing ever. None of that is confirmed by the company.

Amazon, Google and Nvidia remain major backers. Claude is available through AWS Bedrock, Google Cloud Vertex AI and Microsoft Foundry, which means it runs inside every major cloud, including the one owned by OpenAI's biggest investor.

# What makes Anthropic different

---

## Safety is the founding purpose, not a department

Every lab talks about safety. Anthropic was created because of it. A large share of its research output is interpretability (understanding what is happening inside a model) and alignment (making sure a model pursues the goals its operator intends). It publishes a Responsible Scaling Policy that ties deployment decisions to measured capability levels, and it has held back or gated models under it.

## Constitutional AI

Anthropic's core training technique. Instead of relying only on human feedback, the model is given a set of written principles and trained to critique and revise its own outputs against them. The result is a model whose behaviour is more consistent and more predictable at the edges, which matters enormously when that model is running as an agent for hours with nobody watching.

## The Fable / Mythos split

In 2026 Anthropic did something no lab had done before: it split one model into two products by safeguard level. Mythos-class models sit above the Opus tier in capability. The public version, Claude Fable, ships with classifiers that refuse dangerous cybersecurity and biology requests. The same model with those safeguards lifted, Claude Mythos, is available only to vetted organisations through Project Glasswing — a programme run with the US government, launched in April 2026 with AWS, Apple, Google, Microsoft and CrowdStrike, and expanded in June to around 150 organisations including NATO and the EU's cybersecurity agency. Their job: use the model to find and fix vulnerabilities in critical software before attackers do.

That is a new shape for the industry. Frontier capability gated by who you are and what you are for, rather than released to everyone or nobody.

## MCP — the plug the whole industry uses

In November 2024 Anthropic published the Model Context Protocol, an open standard letting any AI agent connect to any tool or data source. Within a year every major lab had adopted it. Anthropic then gave it away, donating MCP to the Agentic AI Foundation under the Linux Foundation, co-founded with OpenAI and Block. The plumbing of the agentic economy is Anthropic's design and nobody's property.

# The products

---

Everything is built around Claude.

**Claude** — The model family. Fable 5.1 at the top, Opus 5, Sonnet 5, Haiku 4.5. See the [WTF is Claude](#) guide.

**Claude Code** — The autonomous coding agent that changed the company's trajectory. Terminal, desktop app, mobile. See the [WTF is Claude Code](#) guide.

**Cowork** — The agent for everyone else. Works in your files, apps and browser without a command line. See the [WTF is Cowork](#) guide.

**Claude.ai** — The chat interface, web and mobile. Free tier; Pro, Max, Team and Enterprise plans above it.

**Claude in Chrome, Excel, PowerPoint, Word and Outlook** — Claude working inside the tools people already use.

**Claude Tag** — Claude inside Slack. Tag it into a thread and delegate the task.

**The Claude Developer Platform and Agent SDK** — The API and the framework other companies build their own agents on.

**Amazon Bedrock, Google Vertex AI, Microsoft Foundry** — Claude inside every major cloud.

## The honest tension

---

**Anthropic says AI may be one of the most dangerous technologies in human history, and then builds it anyway, and is now worth close to a trillion dollars for doing so.**

Dario Amodei has been explicit about this: a company that may be building something profoundly dangerous, pressing forward because the alternative is leaving the frontier to others. Critics point out that this logic can justify almost anything, and that a public listing brings quarterly pressure that no benefit-corporation charter fully insulates against. Supporters point to the research record, the Responsible Scaling Policy, the Glasswing gating and the fact that Anthropic has repeatedly shipped later and more cautiously than rivals — and still won.

The honest answer: Anthropic is probably the most safety-focused lab at the frontier. Whether that is enough, given how fast the frontier is moving, is the question that matters, and no one can answer it yet.

## Why Anthropic matters for the agentic economy

---

Anthropic's fingerprints are on almost everything in this series.

- **Claude** is the default model in OpenClaw, the most-used open-source personal agent.
- **Microsoft's Copilot Cowork** — Microsoft's flagship agent for 365 — runs on Claude, despite Microsoft's \$13 billion stake in OpenAI.
- **Polsia's** founder chose Claude Opus to run the agents that run his companies.

- **Claude Code** is the developer agent inside Paperclip and thousands of AI-run companies.
- **MCP** is how every agent in this series connects to its tools.

You can use the agentic economy without ever thinking about Anthropic. But you will almost certainly be using something it built.

## The key people

---

**Dario Amodei** — CEO and co-founder. Former VP of Research at OpenAI. The public face, and the author of the essays that set the company's tone.

**Daniela Amodei** — President and co-founder. Former VP of Operations at OpenAI. Runs the business.

**Chris Olah** — Co-founder. The world's leading mechanistic interpretability researcher.

**Tom Brown** — Co-founder. Lead author of the GPT-3 paper.

**Jared Kaplan** — Co-founder. Discovered the scaling laws that guide how all labs train models.

**Ami Vora** — Chief Product Officer since January 2026, when Mike Krieger, co-founder of Instagram, moved to Anthropic's Labs group after two years in the role.

## Glossary

---

**Anthropic** — AI safety company founded in 2021 by former OpenAI researchers. Builder of Claude. Public Benefit Corporation.

**Constitutional AI** — Anthropic's training technique: give the model written principles and train it to judge its own outputs against them.

**Interpretability** — Research into what is happening inside a model and why it produces the outputs it does.

**Alignment** — Ensuring an AI pursues the goals its operator actually intends.

**Responsible Scaling Policy** — Anthropic's public framework tying deployment decisions to measured capability and risk levels.

**Mythos-class** — The model tier above Opus. Fable is the safeguarded public version; Mythos is the restricted one.

**Project Glasswing** — Anthropic's vetted-access programme, run with the US government, giving cyberdefenders the unrestricted model to secure critical software.

**MCP (Model Context Protocol)** — The open standard connecting agents to tools. Invented by Anthropic, donated to the Agentic AI Foundation.

**Public Benefit Corporation** — A corporate structure that lets a company weigh its stated mission alongside shareholder profit.

**Run rate** — Recent revenue extrapolated to a full year. The figure labs quote; not the same as audited annual revenue.

**S-1** — The document a company files with the SEC before going public. Anthropic filed one confidentially on 1 June 2026.

## Liked this? Go deeper.

---

Anthropic is the company. Here are the model, the agents and the ecosystem around it.

### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Claude?

---

The AI model inside OpenClaw, Claude Code, Cowork, Microsoft's Copilot Cowork and the agents running thousands of AI companies. What it is, which version you actually want, and why so much of the agentic economy chose it.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Claude is Anthropic's family of AI models — large language models trained to be helpful, honest and harmless, with a particular focus on behaving reliably when running as an agent for hours without supervision.**

It is the brain inside Claude Code and Cowork, the default model in OpenClaw, the engine under Microsoft's Copilot Cowork, and the model Polsia's founder picked to run the agents that run his companies. If you use anything in this series, there is a good chance Claude is doing the thinking.

## By the numbers — September 2026

---

Model names and scores in this box change every few months. Everything below is written to outlast them.

### Fable 5.1

the most capable Claude available to the public, released 1 September 2026

**4**

tiers on general sale: Fable 5.1, Opus 5, Sonnet 5, Haiku 4.5

**1M tokens**

context window on Fable — roughly 750,000 words, or several novels, in one session

**128K tokens**

maximum output in a single response

**87.6%**

SWE-bench Verified score reported for Claude Code Remote in spring 2026 — the highest published at the time

**~150**

organisations with access to the restricted Mythos version through Project Glasswing

**\$20 / \$100–200**

per month for Claude Pro and Max. Free tier available

## The family — which Claude is which

---

Claude is not one model. It is a range at different price and capability points, and in 2026 a new tier appeared at the top.

| Tier          | What it is for                                                                                                                                              |
|---------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Fable</b>  | The Mythos-class flagship, a tier above Opus. Long autonomous tasks, hard reasoning, large codebases, document-heavy professional work. The most expensive. |
| <b>Opus</b>   | The previous top tier, still very capable. Complex agentic work where Fable is more than you need.                                                          |
| <b>Sonnet</b> | The workhorse. Fast, strong, far cheaper. Most business agents run on Sonnet.                                                                               |
| <b>Haiku</b>  | Fastest and cheapest. High-volume, simpler tasks: classification, routing, quick replies.                                                                   |
| <b>Mythos</b> | Not for sale. The same model as Fable with safety classifiers lifted, restricted to vetted cyberdefenders and researchers.                                  |

The version numbers move — 4.6 in March, 4.7 in April, 4.8 in May, 5 in July, Fable 5.1 in September. The tiers do not. When someone says "Claude," ask which tier.

## Fable and Mythos — the split that changed how models ship

---

In June 2026 Anthropic released a model it said was capable enough to be dangerous in the wrong hands: strong enough to find and chain security vulnerabilities in software autonomously, a task previously limited to elite human researchers.

Rather than release it or withhold it, Anthropic did both. Claude Fable is that model with real-time classifiers that block cybersecurity, biology and chemistry misuse; it is generally available. Claude Mythos is the identical model with those safeguards lifted, available only through Project Glasswing, a vetted programme run with the US government for cyberdefenders and critical-infrastructure providers, and, over time, biomedical researchers.

For you as a user the practical meaning is simple. Fable 5.1 is the best Claude you can buy, and it will decline a narrow set of dangerous requests. For the industry the meaning is larger: the most capable AI is now gated by who you are and what you are for. Other labs followed within months.

## What Claude can actually do

---

**Write and edit** — Any style, length or format. Emails, code, contracts, marketing copy, technical documentation. Claude adapts to context and holds a consistent voice across long pieces.

**Reason through hard problems** — Multi-step logic, analysis, strategy, maths. Fable and Opus think before answering, and the thinking scales with the difficulty of the task.

**Read enormous documents** — A million-token context means an entire codebase, a full legal file or a year of emails in one session.

**Write, run and fix code** — Across every major language. Claude Code turns this into an agent that reads your project, changes it, tests it and commits.

**Use tools** — Browse the web, read files, query databases, call APIs, send email, through MCP, the open standard Anthropic invented.

**Work in your apps** — Cowork on the desktop, Claude in Chrome and a built-in browser, Claude in Excel, PowerPoint, Word and Outlook, Claude Tag in Slack.

**Run for hours** — The property that matters most for agents. Claude keeps its footing across long autonomous sessions, remembers what it has done, and asks when it is unsure rather than guessing.

## What makes Claude different — the behaviour, not the benchmark

---

Benchmarks are a snapshot and every lab tops one somewhere. The reason so much of the agentic economy standardised on Claude is behavioural, and it comes from how Anthropic trains.

Constitutional AI gives the model written principles and trains it to judge its own outputs against them. In practice that produces a model that:

- Declines harmful requests and explains why, instead of refusing bluntly or complying blindly
- Flags ambiguity and asks, rather than silently assuming and getting it wrong three steps later
- Says "I don't know" more readily than rivals, which matters when nobody is checking
- Behaves consistently across different phrasings of the same instruction
- Notices flaws in its own work — Anthropic reported Opus 4.8 was several times less likely than its predecessor to let a bug it wrote pass unremarked

**For a chatbot these are nice qualities. For an agent running your inbox at 3am they are the whole product. A model that hallucinates confidently sends the wrong email. A model that flags uncertainty leaves you a note.**

## Where Claude shows up

---

**OpenClaw** — Default model in the most-used open-source personal agent, with millions of active users.

**Microsoft Copilot Cowork** — Microsoft's flagship agent for Microsoft 365 runs on Claude, despite Microsoft's \$13 billion stake in OpenAI. Microsoft chose it anyway.

**Polsia** — The founder has said on the record he chose Claude Opus to run his agents because "it's important to give it the best reasoning."

**Paperclip** — Claude Code is the primary developer agent in the open-source multi-agent org chart.

**Every major cloud** — Amazon Bedrock, Google Cloud Vertex AI and Microsoft Foundry all serve Claude, so enterprises can use it inside infrastructure they already pay for.

**Thousands of products** — Via the API: customer-service agents, research tools, legal analysis, coding assistants, and the AI-run companies indexed on [wtfagents.com](https://wtfagents.com).

## How to get it

---

**Claude.ai and the apps** — Web, desktop (Mac and Windows), iPhone and Android. Free tier for chat. Claude Pro at \$20 a month adds Cowork, Claude Code and higher limits. Claude Max at \$100–200 for heavy use. Team and Enterprise plans above that.

**Claude Code** — Included with Pro and above, or pay-as-you-go through the API.

**The Claude Developer Platform** — Direct API access for builders. Pay per token; Haiku cheapest, Fable most expensive.

**Through your cloud** — Bedrock, Vertex AI or Foundry if your company already lives there.

**Inside other agents** — OpenClaw, Paperclip and most agent frameworks let you plug in your own Anthropic API key.

## Glossary

---

**Claude** — Anthropic's AI model family: Fable, Opus, Sonnet and Haiku on general sale; Mythos restricted.

**Fable** — The public version of the Mythos-class model, with safety classifiers active. Currently Fable 5.1.

**Mythos** — The same model with safeguards lifted, available only to vetted organisations through Project Glasswing.

**Tier** — A capability and price level within the family. Tiers are stable; version numbers change constantly.

**Constitutional AI** — Anthropic's training technique: written principles the model learns to judge its own outputs against.

**Context window** — How much text the model can hold at once. One million tokens on Fable.

**Token** — The unit models process text in. Roughly three quarters of a word. API usage is billed in tokens.

**MCP (Model Context Protocol)** — The open standard Claude and every other major model use to connect to tools.

**SWE-bench** — The standard benchmark for AI coding ability, measured on real software engineering tasks.

**Claude Code** — The agentic version of Claude for building software. See its own guide.

**Cowork** — Claude working in your files, apps and browser without a command line. See its own guide.

## Liked this? Go deeper.

---

Claude is the model. Here is everything built on it.

### WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Claude Code?

---

Not a code editor. Not an autocomplete plugin. An autonomous AI developer that reads your entire codebase, writes code, runs tests, fixes bugs, and commits to Git — while you do something else. Now with a desktop app, a phone app, and a growing crowd of non-developers building real things with it.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# The one-liner

---

**Claude Code is Anthropic's autonomous coding agent — a command-line tool that takes a task, reads your codebase, figures out what needs to change, makes the changes, runs tests, and iterates until the task is done.**

It is not a tool that helps you code. It is a tool that codes while you watch — or while you do something else entirely.

# The difference that matters

---

There are three categories of AI coding tool. They are not the same thing:

## Autocomplete tools (GitHub Copilot)

Suggest the next line of code as you type. You are still writing. The AI is making suggestions. You accept or reject each one. Fast, low-friction, but fundamentally reactive.

## AI-assisted editors (Cursor)

Chat with the AI inside your editor. Ask it to write a function, explain some code, or make a change. More powerful than autocomplete, but you are still directing every step.

## Autonomous coding agents (Claude Code)

You give the agent a task: "Add user authentication to this application." It reads the entire codebase, plans the changes, writes the code across multiple files, runs the tests, fixes the failures, and commits the result. You were not involved in the individual steps.

**The distinction matters because the productivity multiplier is completely different. Autocomplete saves you seconds. An autonomous agent saves you hours.**

# What Claude Code can actually do

---

**Read entire codebases** — Claude Code reads all your files — not just the one you have open. It understands how parts of the codebase relate to each other. This is critical for tasks that require changes across multiple files.

**Write and edit code** — Across all major languages: Python, JavaScript, TypeScript, Go, Rust, Java, C++, Ruby, and more. It writes idiomatic code that fits the existing style of your codebase.

**Run shell commands** — Execute tests, linters, build tools, and scripts. Read the output. Understand what failed. Try to fix it.

**Run tests and fix failures** — Write tests if you don't have them. Run existing tests. Read the failures. Fix the code. Re-run. Iterate.

**Commit to Git** — Stage changes, write commit messages, and commit — following your project's conventions.

**Handle multi-file refactors** — Rename a function used in 47 places. Change an API interface and update all callers. Restructure a module. Tasks that would take a human developer hours.

**Work on long autonomous sessions** — Unlike autocomplete tools that help one line at a time, Claude Code can work on a complex task for hours — maintaining context across the entire session.

## By the numbers — September 2026

---

Claude Code went from research preview in early 2025 to the fastest-adopted developer tool of its generation. The figures move monthly; the shape does not.

**\$2.5B**

annualised revenue run rate by February 2026 — before the growth that took Anthropic past \$65B overall

**29M**

installs inside Visual Studio Code alone by early 2026

**~4%**

of all public GitHub commits worldwide authored by Claude Code, per one early-2026 analysis

**87.6%**

SWE-bench Verified reported for Claude Code Remote in spring 2026, the highest published score at the time

**April**

2026: redesigned desktop app for Mac and Windows — panels, buttons, a chat box, no terminal required

**Aug**

2026: start a Claude Code session on your computer directly from the Claude phone app

The satisfaction surveys that first made the point still hold: developers who try autonomous agents do not go back to autocomplete. GitHub Copilot has the larger install base and Microsoft behind it. Claude Code overtook it in developer preference in under a year because it does something fundamentally different.

# How Claude Code fits into the agentic economy

---

**Claude Code is not just a tool for individual developers. It is one of the primary building blocks of the agentic economy.**

## In Paperclip

Paperclip — the open-source multi-agent company orchestration framework — lists Claude Code as a primary supported agent. In a Paperclip "company org chart," Claude Code can be the developer agent: building and maintaining the company's software autonomously while other agents handle marketing, customer support, and operations.

## In autonomous companies

The companies tracked on WTF Agents — the tens of thousands of active AI-run companies on Polsia, Paperclip and elsewhere — are using agents like Claude Code to build and iterate on their products. A company with no human employees still needs software. Claude Code builds it.

## In enterprise settings

Large enterprises are using Claude Code for large-scale code migrations, test-writing, technical debt reduction, and API integrations. Tasks that previously required weeks of developer time can be completed in hours.

# Claude Code vs the alternatives

---

| Tool              | Type               | Best for                                            | Limitation                                                     |
|-------------------|--------------------|-----------------------------------------------------|----------------------------------------------------------------|
| Claude Code       | Autonomous agent   | Complex, multi-file tasks; long autonomous sessions | Usage limits on Pro; costs scale with the task                 |
| Cursor            | AI-assisted editor | Fast iteration; inline editing; real-time help      | Now owned by SpaceX/xAI; losing OpenAI models in November 2026 |
| GitHub Copilot    | Autocomplete       | Quick suggestions while typing; broad ecosystem     | Reactive, not autonomous                                       |
| Codex (OpenAI)    | Autonomous agent   | Lives inside the ChatGPT app; GPT-6 backbone        | Smaller ecosystem of plugins and skills                        |
| Devin (Cognition) | Autonomous agent   | Fully cloud-hosted; now \$20/mo                     | Less control over your own machine                             |
| Jules (Google)    | Autonomous agent   | Free; runs in the cloud on Gemini                   | Narrower task scope                                            |

# Do you need to be a developer to use Claude Code?

---

**Not any more — though it helps to think like one.**

For its first year Claude Code was a command-line tool, and the screenshots looked like a hacker film. In April 2026 Anthropic shipped a redesigned desktop app for Mac and Windows: panels, buttons, a chat box. You type what you want in plain English and watch it work. In August, the Claude phone app gained the ability to start and steer a session on your computer from anywhere. It is included in Claude Pro at \$20 a month, so the API key requirement is gone too.

Anthropic's own analysis of over a million Cowork sessions found that more than 90% were not software development. Claude Code's audience has shifted the same way: founders building their first product, operators automating a workflow, marketers building internal tools, and — the pattern this series exists to document — people running AI companies who use Claude Code as the developer agent inside Polsia or Paperclip.

What still matters is not the terminal but the mindset. Claude Code works best for people who can describe what they want precisely, check what came back, and understand the difference between "it runs" and "it is right." That is a developer's habit, not a developer's skill, and it can be learned in a weekend.

## What this means for software development

---

Claude Code is part of a broader shift in what software development means.

For decades, writing software required deep technical skill, years of learning, and constant attention to detail. The bottleneck was human developer time.

**Autonomous coding agents change the equation. The bottleneck shifts from "can we write the code?" to "do we know what to build?" The strategic and product decisions matter more. The implementation matters less.**

This does not mean software developers are going away. It means the nature of the job is changing. Senior developers who can direct, review, and architect are more valuable. Junior developers who were doing rote implementation work face the most disruption.

## Glossary

---

**Claude Code** — Anthropic's autonomous coding agent. CLI tool that reads codebases, writes code, runs tests, and commits to Git without step-by-step human instruction.

**Autonomous coding agent** — An AI that takes a coding task and completes it end-to-end without human direction of individual steps. Contrast with autocomplete (suggests next line) or AI-assisted editors (human-directed).

**CLI (Command Line Interface)** — A text-based way of interacting with a computer. Claude Code runs in the terminal — you type commands rather than clicking a GUI.

**Git** — The standard version control system for software projects. Claude Code can stage, commit, and manage Git operations autonomously.

**SWE-bench** — The standard benchmark for AI coding ability. Measures how well an agent can solve real software engineering tasks. Scores change every release; Claude Code has led it for most of 2026.

**GitHub Copilot** — Microsoft/GitHub's AI coding tool. 77M+ developers. Autocomplete-style — helps as you type. Claude Code is a different category: autonomous.

**Cursor** — An AI-assisted code editor. More powerful than Copilot, but still human-directed. Claude Code is autonomous; Cursor is collaborative.

**Paperclip** — The multi-agent company orchestration framework that uses Claude Code as a primary developer agent.

**API key** — Your credential for accessing the Anthropic API. Required to run Claude Code. You pay per token of usage.

**Refactor** — Restructuring existing code without changing its external behaviour. A common use case for Claude Code — especially large-scale refactors across many files.

## Liked this? Go deeper.

---

Claude Code is the agent. Here is the ecosystem around it.

### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Cowork?

---

Claude Code was for developers. Cowork is for everyone else. The Anthropic product that rattled enterprise software stocks, became Microsoft's flagship agent, and now runs in the cloud, on your phone, and in its own browser. What it is and why it matters.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Cowork is Anthropic's desktop AI agent for non-developers — a tool that lets Claude read, edit, create, and organise files directly on your computer, execute multi-step tasks autonomously, and connect to your existing tools, all without touching a command line.**

Anthropic's own framing says it best: "Chat changed how we get answers. Claude Code transformed how developers build software. Now, Cowork brings that same execution power to everyone."

## By the numbers — September 2026

---

**12 Jan**

2026: launched as a research preview. Generally available on Mac and Windows since April, web and phone since July

**90%+**

of Cowork sessions are not software development, per Anthropic's own analysis of 1.2 million sessions

**26 Aug**

2026: Cowork gets its own built-in browser. Claude in Chrome goes generally available the same week

**\$20**

per month — included in Claude Pro and every paid plan above it. Not on the free tier

**11**

official Anthropic plugins for sales, finance, legal, marketing and more, plus a marketplace of third-party ones

**June**

2026: Microsoft's Copilot Cowork, built on Claude, goes live worldwide in Microsoft 365

## The evolution — Chat → Code → Cowork

---

To understand Cowork you need to understand the three-stage evolution of Claude:

| Stage | Product     | Who it's for | What it does                                                                                                          |
|-------|-------------|--------------|-----------------------------------------------------------------------------------------------------------------------|
| 1     | Claude Chat | Everyone     | Answers questions. You ask, it responds. Reactive, not autonomous.                                                    |
| 2     | Claude Code | Developers   | Autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git. Requires command line.             |
| 3     | Cowork      | Everyone     | Autonomous desktop agent. Reads and edits files, executes multi-step tasks, connects to apps. No command line needed. |

The pattern Anthropic spotted: when they released Claude Code, they expected developers to use it for coding. They did — and then immediately started using it for everything else. Organising files, writing reports, processing data, automating workflows. Cowork is the formalisation of that behaviour, made accessible without technical prerequisites.

## How Cowork actually works

Cowork lives inside the Claude Desktop app for Mac and Windows, and since July 2026 also on [claude.ai](#) in a browser and in the Claude phone app. You access it by clicking the Cowork tab, next to Chat. Sessions run against your Claude account, so a task started on the desktop keeps running after you close the laptop, and you can check on it from your phone.

### Step 1: Grant folder access

You point Cowork at a folder on your computer. This is the boundary — Claude can only read and modify files within that folder. Start with something low-stakes: a Downloads folder, a project folder, a research directory. You control exactly what Claude can and cannot touch.

### Step 2: Describe the task

In plain English. No special syntax, no commands. Just describe what you want done. Examples that work right now:

- "Organise this folder. Sort files by type, rename generic ones based on content, flag duplicates."
- "Pull the expense figures from all these receipt screenshots and create a summary spreadsheet."
- "Read all these research notes and write a 2-page briefing document."
- "Find all the client emails in this folder and extract the action items into a to-do list."
- "Take this raw data CSV and generate a formatted report with charts."

### Step 3: Claude plans and executes

Cowork makes a plan, shows it to you, and then executes it step by step. You can see what it is doing in the Progress panel. You can pause or redirect mid-task if needed. For significant actions — deleting files, sending emails — Cowork asks for confirmation first.

### Step 4: Results land in your folder

No copy-pasting from a chat window. No downloading outputs. The finished files are in your folder, exactly where you need them.

## What makes Cowork different from just using Claude Chat

---

This is the question most people ask. The difference is significant:

|                         | Claude Chat                                   | Cowork                                                                           |
|-------------------------|-----------------------------------------------|----------------------------------------------------------------------------------|
| File access             | You paste content into the chat               | Claude reads files directly from your computer                                   |
| Output                  | Text in a chat window — you copy it           | Files created directly in your folder                                            |
| Multi-step tasks        | One response at a time — you direct each step | Claude plans and executes all steps autonomously                                 |
| Memory                  | Forgets between conversations                 | Projects keep context, files, and instructions persistent                        |
| App connections         | Limited                                       | Connects to Gmail, Google Drive, Notion, Slack, and hundreds more via connectors |
| Works while you're away | No                                            | Yes — scheduled recurring tasks run automatically                                |

## The features — what Cowork can do

---

### File operations

Read, edit, create, rename, organise, and delete files in your designated folder. Process any file type Claude can understand: documents, spreadsheets, PDFs, images, CSVs, emails.

### Projects

A persistent workspace tied to a folder. Your files, instructions, and task history stay in one place across sessions. Instead of starting fresh every time, your project remembers what it has done and what is next.

### Scheduled tasks

Recurring tasks that run automatically. "Every Monday morning, pull last week's sales data from this folder and generate a summary report." Since the move to cloud sessions, scheduled tasks fire even

with no device online.

## Connectors

Link Cowork to external services: Gmail and Google Workspace, Microsoft 365 (Outlook, OneDrive, SharePoint, Teams), Slack, Jira, Notion, DocuSign, and hundreds more via the connector marketplace. Desktop extensions have deeper system access than web connectors.

## Skills

Install skills — modular capabilities — that improve Cowork's output for specific formats: documents, presentations, spreadsheets. You can also write a SKILL.md file with your brand voice or standard process and Cowork uses it every time.

## The built-in browser

Since August 2026 Cowork has its own Chromium browser in a side panel of the desktop app. When a task involves a website, the browser opens, Claude navigates, fills forms and finishes the job — separate from your own browser and logins. If you would rather it work inside your own signed-in Chrome, Claude in Chrome does that and is now available on every paid plan; a safety classifier checks each action before it is taken.

## Computer use

Cowork can see your screen and drive any application on your machine, not just the browser. Useful for tools that have no connector: a niche internal system, a desktop app, anything with buttons.

## Plugins

Role-specific bundles of skills, prompts and connectors. Anthropic publishes its own — sales, finance, legal, marketing, and a Claude for Financial Advisors plugin covering onboarding, meeting prep and portfolio reviews — and a marketplace carries third-party ones. Plugins run on whatever plan you already have; nothing extra to buy from Anthropic.

## Dispatch

Control a desktop Cowork session from your phone: approve a step, redirect, or read the result on the move.

# The market reaction — why software stocks fell

---

**When Cowork launched in January 2026, and again when its plugins landed, enterprise software stocks fell sharply. The iShares Expanded Tech-Software Sector ETF dropped nearly 5% in a single day, and analysts put the combined loss across enterprise software companies in the hundreds of billions of dollars.**

The reason: investors looked at what Cowork can do and compared it to what they were paying enterprise software companies to do. File management. Document creation. Data extraction. Report generation. Workflow automation. These are the core functions of dozens of \$10–50/user/month SaaS products.

A knowledge worker with Cowork can now automate tasks that previously required separate tools for project management, document processing, data analysis, and workflow automation — all from one desktop agent.

The stocks recovered somewhat when Microsoft announced it was building Copilot Cowork on top of Anthropic's technology — suggesting the incumbents would adapt rather than be replaced entirely. But the signal was clear: the knowledge work software market has been permanently disrupted.

## Microsoft Copilot Cowork — the enterprise version

---

Microsoft announced Copilot Cowork in March 2026 and made it generally available worldwide in June: a cloud-based agent, powered by Anthropic's Claude, that executes multi-step tasks across Microsoft 365.

The difference from consumer Cowork: Copilot Cowork runs inside Microsoft 365's own infrastructure with access to the full graph of enterprise work data — Outlook, Teams, calendars, SharePoint, Excel. Rather than working on a local folder, it works across an entire organisation's data.

It ships as a \$30 per user per month add-on, or inside the Microsoft 365 E7 bundle.

**The significance: despite a \$13 billion investment in OpenAI, Microsoft built its flagship Microsoft 365 agent on Anthropic's Claude — and then took it to general availability. Not because Microsoft lacks access to GPT, but because it chose Claude anyway.**

## Real use cases — what people are actually doing with it

---

### Finance and accounting

Point Cowork at a folder of receipt screenshots. It extracts all expense data, categorises by type, and generates a formatted expense report. Tasks that took a finance assistant an hour now take Cowork five minutes.

### Research and analysis

Feed Cowork a folder of research notes, articles, and PDFs. Ask it to synthesise the key findings into a briefing document. It reads everything, identifies themes, and writes the document — citing sources from the files it read.

### Content operations

Give Cowork a brand voice SKILL.md file and a folder of raw content briefs. Ask it to draft all ten blog posts. It writes in your brand voice, formats correctly, and saves them as Word documents ready for review.

### Data processing

Drop a messy CSV export into the folder. Ask Cowork to clean it, standardise the formatting, remove duplicates, and generate a summary report with charts. No spreadsheet skills required.

### File organisation

The classic first use case. Give Cowork your Downloads folder. It sorts by type, renames files based on their actual content, flags duplicates, and creates logical subfolders. Most users hit their first "wow" moment within ten minutes of trying this.

### Scheduled intelligence

Connect Cowork to your email via the Gmail connector. Set a recurring task: every morning at 8am, read all emails received overnight, extract action items, and create a prioritised to-do list in your project folder. You wake up to a briefing that was prepared while you slept.

## How to get access

| Plan                     | Access                                                  | Price                                       |
|--------------------------|---------------------------------------------------------|---------------------------------------------|
| Claude Pro               | Full Cowork access                                      | \$20/month                                  |
| Claude Max               | Full Cowork access + 5x or 20x the limits               | \$100 or \$200/month                        |
| Claude Team              | Full Cowork access + admin controls                     | \$25–30/user/month                          |
| Claude Enterprise        | Full Cowork + admin controls, audit logs, smart reports | Custom pricing                              |
| Microsoft Copilot Cowork | Cloud-based, Microsoft 365 integrated                   | \$30/user/month add-on, or in the E7 bundle |
| Free tier                | Not available — Cowork is paid only                     | —                                           |

To get started: subscribe to Claude Pro or higher, then open Cowork in the desktop app ([claude.com/download](https://claude.com/download)), on [claude.ai](https://claude.ai), or in the Claude phone app.

## The honest limitations

Cowork is genuinely powerful. It also has real limitations worth knowing:

→ **Folder boundary only.** Claude can only access the folder you designate. It cannot reach outside it without explicit permission — and that is a feature. In one widely reported case about 15,000 family photos were deleted while Cowork tidied a desktop, and were recovered only through a 30-day cloud restore.

→ **It burns through usage.** Agent tasks use far more of your plan's allowance than chat does. Pro is the entry point; heavy daily use tends to need Max.

→ **Hallucination risk remains.** Cowork can misread file contents or make incorrect assumptions about what a task requires. Always review outputs before acting on them.

→ **Security considerations.** Days after launch, a data exfiltration vulnerability was reported and patched. Prompt injection — hidden instructions in a web page or document the agent reads — is the unsolved problem for every agent with browser access. Be thoughtful about connectors and folders.

→ **The desktop browser needs the desktop.** From the web or your phone, Claude can drive the built-in browser only while your desktop app is open. Without it, Claude in Chrome is the browser route.

→ **It keeps changing.** Cowork has shipped a major capability most months of 2026. Expect the feature list in this guide to be out of date before the next one.

## What Cowork means for the agentic economy

---

**Claude Code proved agents could replace developer time. Cowork is the thesis applied to every knowledge worker.**

The administrative layer of every business — the file management, the data processing, the report generation, the inbox management, the scheduling — is now automatable by a non-technical person with a \$20/month subscription.

The barrier to running a lean, agent-powered business has never been lower. A solo founder with Cowork, Claude Code, and OpenClaw has the operational capacity of a team of five. A small business owner with Cowork has an always-on assistant that never sleeps, never forgets, and costs less than a gym membership.

The HBR research published in early 2026 noted that companies are already making headcount decisions based on AI's potential, not just its demonstrated performance. Cowork is part of why. It is not a future prediction — it is a current product that current businesses are using to do real work right now.

# Glossary

---

**Cowork** — Anthropic's autonomous desktop AI agent. Part of the Claude Desktop app. Reads, edits, and creates files on your computer and executes multi-step tasks without command-line skills.

**Claude Desktop** — The desktop application from Anthropic for Mac and Windows. Contains Chat, Code, and Cowork modes, and Cowork's built-in browser. Download at [claude.com/download](https://claude.com/download).

**Projects (Cowork)** — A persistent workspace tied to a folder — keeps files, instructions, and task history in one place across sessions.

**Scheduled tasks** — Recurring tasks that Cowork runs automatically on a set schedule — daily, weekly, or custom timing.

**Connectors** — Integrations that link Cowork to external services — Gmail, Google Drive, DocuSign, Slack, Notion, and hundreds more.

**Skills** — Modular capabilities installed into Cowork. Including custom SKILL.md files you create to encode your brand voice or standard processes.

**Claude in Chrome** — Anthropic's agent inside your own Chrome browser, generally available on all paid plans. The alternative to Cowork's built-in browser when you want it working where you are already signed in.

**Copilot Cowork** — Microsoft's enterprise version of Cowork — cloud-based, running across Microsoft 365. Powered by Anthropic's Claude. \$30/user/month. Generally available since June 2026.

**Claude Code** — The developer-facing autonomous coding agent that Cowork is built on. Requires command-line skills. Cowork is the non-technical equivalent.

**Knowledge worker** — Someone whose job primarily involves creating, processing, or managing information — as opposed to physical labour. Cowork is specifically designed for this category.

## Liked this? Go deeper.

---

Cowork is built on Claude and sits alongside Claude Code. Here is the full picture.

### WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# How to Hire an AI Agent

## for Your Business

---

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# Before you start — the right mindset

---

**Deploying AI agents in your business is not a technology project. It is a staffing decision.**

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

## Step 1 — Find your highest-value repeatable task

---

**The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.**

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

**Customer support responses** — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

**Content creation** — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

**Research and summarisation** — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

**Data entry and processing** — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

**Email drafting** — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

**Scheduling and coordination** — Managing calendar bookings, sending reminders, coordinating meeting times.

**Lead qualification** — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

## Step 2 — Choose the right type of agent

Not all agents are right for all tasks. In 2026 the choice is simpler than it looks, because every option falls into a handful of shapes:

| If you want to...                                              | Use this                                                                  | Cost                                     |
|----------------------------------------------------------------|---------------------------------------------------------------------------|------------------------------------------|
| Delegate work inside your own files, inbox, calendar and apps  | Cowork (Anthropic), ChatGPT Work (OpenAI), Copilot Cowork (Microsoft 365) | From \$20/mo, included in the paid plan  |
| Run a team of agents sharing one cloud computer and its logins | Grok Bot (xAI)                                                            | Subscription                             |
| Have a personal agent you message like a colleague             | OpenClaw (open-source), Meta Muse                                         | Free + model costs; Muse has a free tier |
| Let agents build and run a whole company for you               | Polsia, NanoCorp, Cofounder.co                                            | From \$20/mo + fees (Polsia)             |
| Run several agents as an org chart you control                 | Paperclip (self-hosted)                                                   | Free + model costs                       |
| Connect apps and automate defined workflows                    | Zapier, n8n                                                               | From ~\$20/mo; n8n is open-source        |
| Build a custom agent with full control                         | Claude Developer Platform / Agent SDK                                     | Pay per token                            |

For a first agent, the top row is almost always the answer. A work agent on a plan you already have, pointed at one folder and one task, is the lowest-risk way to learn what this is actually like.

## Step 3 — Write a clear job description for your agent

**This is the most important step most people skip.**

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

→ **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")

→ **The inputs** — what information does the agent receive? ("An inbound email from a customer")

- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

## Step 4 — Start with read-only access

---

**The most common mistake when deploying agents: giving them too much access too soon.**

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

## Step 5 — Measure results honestly

---

**Define what success looks like before you deploy — not after.**

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved
- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

# Five deployments that work — and what to expect

---

These are the patterns that come up again and again among small businesses that get value from a first agent. Results vary; the shape of the win does not.

## First-line customer support

An agent reads incoming support emails, categorises them, drafts replies from approved templates, and flags anything unusual for a human. The win is response time — hours become minutes — and a majority of routine emails handled without anyone touching them. The risk is a confident wrong answer, which is why replies are drafted for approval before they are sent unsupervised.

## Research and content preparation

An agent researches a client's industry each week and drafts a content calendar, a briefing or a first draft. Writers spend their time on strategy and editing instead of gathering. The win is capacity without headcount.

## Inbox triage for partners and owners

An agent summarises the day's email, drafts responses for review, and surfaces what needs a decision. Owners consistently report an hour or two a day back. The risk is the agent guessing at tone on a sensitive thread; keep it drafting, not sending, for those.

## The solo operator's assistant

A personal agent — OpenClaw over WhatsApp or Telegram, or Cowork on the desktop — briefed with research tasks, email drafts, document summaries and admin. Typical cost is tens of dollars a month in model usage. This is the deployment most one-person businesses start with.

## Routine engineering

Claude Code or Codex handling bug fixes, small features and tests, with senior people on architecture and review. Teams that do this well ship considerably more with the same headcount. Teams that skip the review step ship considerably more bugs.

# The honest cost — what agents actually cost

---

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

| Agent type              | Monthly cost (approx) | What drives the cost                                         |
|-------------------------|-----------------------|--------------------------------------------------------------|
| Cowork / ChatGPT Work   | \$20–200/mo           | Your Claude or ChatGPT plan; heavy use needs the higher tier |
| Copilot Cowork          | \$30/user/mo          | Add-on to Microsoft 365                                      |
| OpenClaw (personal)     | \$20–100/mo           | Model API usage — Claude, GPT or a cheaper open model        |
| Polsia company          | \$20–1,000/mo + fees  | Plan tier, credits, 20% on ad spend, 3% on revenue           |
| Paperclip org           | \$0 + model usage     | Free software; each agent's model bill                       |
| Zapier / n8n automation | \$0–100/mo            | Task volume and plan tier                                    |
| Custom Claude agent     | \$50–500/mo           | Token usage — depends heavily on task volume                 |
| Claude Code             | \$20–200/mo           | Included in Claude Pro/Max; API if you go beyond             |

For context: a part-time human assistant in the US costs \$2,000–3,500/month. A full-time junior employee costs \$45,000–60,000/year. An agent handling equivalent work costs \$50–500/month. The economics are not subtle — and they are the reason the Agentic Economy guide spends a section on what that means for jobs.

## What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

## Your action plan — what to do this week

**Day 1** — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

**Day 2** — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

**Day 3** — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

**Day 4-5** — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

**Week 2** — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

**Week 3-4** — If quality is consistently good, expand permissions incrementally. Track the time saved.

**Month 2** — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

**The businesses that win with AI agents are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.**

## Glossary

---

**AI Agent** — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

**Human-in-the-loop** — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

**Polsia** — A managed platform where AI agents build and run companies. From \$20/mo plus credits and fees; see its own guide.

**OpenClaw** — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

**Paperclip** — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

**Cowork** — Anthropic's work agent: delegates tasks across your files, apps and browser. Included in every paid Claude plan.

**ChatGPT Work** — OpenAI's work agent: give it an outcome, it returns finished spreadsheets, decks and apps.

**Zapier / n8n** — Automation platforms that connect apps and run defined workflows, now with agent steps. n8n is open-source and self-hostable.

**Claude** — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

**API key** — Your credential for accessing an AI model's API. Required for OpenClaw, Paperclip and custom builds; Claude Code and Cowork run on your Claude plan instead.

**Token** — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

**MCP (Model Context Protocol)** — The open standard connecting AI agents to external tools. Invented by Anthropic.

## Liked this? Go deeper.

---

You now know how to hire an agent. Here is the full context behind the tools you are using.

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# Who's Who in the

## Agentic Economy

---

Forty-odd names you will keep seeing, sorted into the seven layers of the agentic economy, one line each. The map to read before any of the deep dives — and the one to come back to when a new name appears in your feed.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## How to read this guide

---

The agentic economy is a stack. Models at the bottom, agents built on them, platforms that run whole companies of agents, and the plumbing, chips and money underneath it all. Every name in the news belongs to one layer, and knowing which layer tells you most of what you need to know.

This guide gives each name one line: what it is, why it matters. Where a full WTF Agents guide exists, it is named so you can go deeper. Version numbers, valuations and star counts live in the dated box below and nowhere else, because they change monthly and the map does not.

## By the numbers — September 2026

---

**\$965B**

Anthropic's valuation after May's \$65B round — the most valuable private company in the world, IPO filed

**\$852B**

OpenAI's valuation after its \$122B round in March

**\$60B**

what SpaceX paid for Cursor in June, after buying xAI in February

**\$30M**

Polsia's May raise at \$250M; NanoCorp is YC-backed; Paperclip has no disclosed funding

**3.2M**

active OpenClaw users; Hermes Agent is the credible number two

**490,000**

companies created on Polsia; about 6% still active

**6**

labs that shipped a "does your job" agent in 2026: Anthropic, Microsoft, OpenAI, xAI, Meta — and Google, by consolidation

**2**

open standards under the Agentic AI Foundation: MCP and A2A

## Layer 1 — The labs and their models

---

The brains. Everything above runs on one of these.

**Anthropic — Claude** — Safety-founded lab, now the most valuable private company on earth. Claude powers Claude Code, Cowork, Copilot Cowork, OpenClaw's default and Polsia's agents. Tiers: Fable, Opus, Sonnet, Haiku; Mythos is the restricted twin. Guides: [WTF is Anthropic](#), [WTF is Claude](#).

**OpenAI — GPT** — ChatGPT is still the product most people have used. GPT-6 arrived in September. OpenAI's agents are ChatGPT Work, Codex and the Agents API; its reach is its weapon.

**Google DeepMind — Gemini** — Gemini 3.8, Gemini Agent, Chrome Auto Browse and Jules. Owns Search, Android and Workspace, and invented the A2A protocol. Chose to consolidate its agents rather than launch a new one.

**Meta — Muse** — A new model family and a personal agent of the same name, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is that three billion messaging users are the distribution. Also tried to buy Manus.

**xAI — Grok** — Elon Musk's lab, now inside SpaceX alongside Cursor. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff."

**Microsoft** — Not a frontier lab of its own but the biggest distributor: Copilot Cowork, built on Claude, is live in Microsoft 365 worldwide; Copilot Studio hosts hundreds of thousands of custom agents; Foundry serves every major model.

**DeepSeek** — The Chinese lab whose cheap-to-train breakthrough in early 2025 made this a two-continent race. DeepSeek V4 is open-weight under the MIT licence and runs under a great many indie agents.

**Alibaba — Qwen · Moonshot — Kimi · Zhipu — GLM · MiniMax** — The rest of China's open-weight frontier. Kimi K3 was the largest open model ever released. Cheap, capable, downloadable. Guide: [WTF is Happening with China's AI](#).

**Mistral** — France's lab. Efficient smaller models and Europe's answer to US and Chinese dependence.

## Layer 2 — Work agents

---

Agents that do your job's tasks. The category every lab entered in 2026.

**Cowork (Anthropic)** — Claude working in your files, apps and its own browser, on desktop, web and phone. Included in every paid Claude plan. Guide: [WTF is Cowork](#).

**ChatGPT Work (OpenAI)** — Give it an outcome; it returns finished spreadsheets, decks and web apps hours later. Launched July.

**Copilot Cowork (Microsoft)** — Cowork inside Microsoft 365, across Outlook, Teams, SharePoint and Excel. Built on Claude despite Microsoft's stake in OpenAI. Generally available since June.

**Grok Bot (xAI)** — A team of named agents sharing one cloud computer, working while you sleep. Launched August through Cursor's infrastructure.

**Microsoft Scout** — An always-on personal agent inside Copilot, in limited release.

**Gemini Agent (Google)** — Where Google's browser agent went after Project Mariner was folded in. Plus Auto Browse in Chrome.

## Layer 3 — Coding agents

---

The category that proved agents work, and the one that made Anthropic.

**Claude Code (Anthropic)** — The autonomous developer. Terminal, desktop app, phone. The fastest-adopted developer tool of its generation. Guide: [WTF is Claude Code](#).

**Codex (OpenAI)** — OpenAI's coding agent, now living inside the ChatGPT desktop app on GPT-6.

**Devin (Cognition)** — The original "AI software engineer," fully cloud-hosted, cut from \$500 to \$20 a month in April.

**Jules (Google)** — Google's free cloud coding agent on Gemini.

**Cursor** — The AI-first code editor, bought by SpaceX for \$60 billion in June and now the delivery vehicle for Grok Bot. Losing OpenAI's models in November.

**GitHub Copilot (Microsoft)** — The largest install base in coding AI, evolving from autocomplete into an agent inside GitHub itself.

## Layer 4 — Personal agents

---

Agents that run your life's admin, mostly over messaging apps.

**OpenClaw** — Open-source, self-hosted, millions of users, a non-profit foundation, native mobile apps. The reference personal agent and the one every rival is measured against. Guide: [WTF is OpenClaw](#).

**Meta Muse** — Meta's consumer agent: books, buys, negotiates and fills in forms, free tier, WhatsApp. US only at launch. Guide: [WTF is Meta Muse](#).

**Hermes Agent (Nous Research)** — The credible open-source number two: self-improving memory, hireable into Paperclip out of the box.

**Manus** — The Chinese general-purpose agent Meta agreed to buy for \$2 billion before Beijing blocked the deal. Still independent, still Chinese.

## Layer 5 — Company platforms and orchestration

---

Agents that run whole businesses, and the tools that manage teams of agents.

**Polsia** — The biggest and most argued-about: subscribe, describe an idea, a fixed team of agents builds and runs the company. Nearly half a million created; a founder who says one in ten has earned a dollar. Guide: [WTF is Polsia](#).

**NanoCorp** — YC-backed rival that publishes the live revenue of every company on the platform — the most honest numbers in the space, and small ones: the best has earned \$2,570, and only two have passed \$1,000. Guide: [WTF is NanoCorp](#).

**Cofounder.co** — Company-in-a-box with departments and a human approval gate in front of anything risky.

**Paperclip** — Open-source org chart for agents: roles, budgets, approvals, any agent vendor, self-hosted. ClipMart is its marketplace of forkable company templates. Guide: [WTF is Paperclip](#).

**Salesforce Agentforce** — The enterprise incumbent's agent platform, priced per conversation, past a billion dollars in annual revenue.

**ServiceNow AI Control Tower** — Governed agents for IT and HR workflows; AI contract value past a billion dollars.

**SAP Joule · Workday Illuminate** — Agents inside the ERP and HR systems most large companies already run — and the first to charge outside agents a toll to reach their data.

**Zapier · n8n · UiPath** — The automation layer: connect apps, run defined workflows, now with agent steps. n8n is open-source; UiPath is the enterprise back-office specialist.

## Layer 6 — The plumbing

---

Standards and infrastructure that make agents interoperate.

**MCP — Model Context Protocol** — The universal plug between agents and tools. Invented by Anthropic, adopted by everyone, now vendor-neutral. Guide: [WTF is an API](#).

**A2A — Agent2Agent** — Google's standard for agents from different companies talking to each other. In production at hundreds of organisations.

**Agentic AI Foundation** — The Linux Foundation body that owns MCP and A2A, founded by Anthropic, OpenAI, Google, Microsoft and Block. The rare case of rivals agreeing to share.

**Stripe** — The payment rails under Polsia companies, the WTF Agents store, and most of the agentic commerce experiments. Agents that can pay run through Stripe.

**Cloudflare · Vercel** — Where agents are hosted and deployed. Vercel's Eve is a framework for durable agents; Cloudflare's platform runs them at the edge.

## Layer 7 — The picks and shovels

---

Who gets paid regardless of which agent wins.

**Nvidia** — The chips. Every model in Layer 1 was trained and is served on Nvidia hardware, which is why it is the most valuable company in the world.

**AWS · Google Cloud · Microsoft Azure** — The clouds. Bedrock, Vertex AI and Foundry serve every major model to enterprises inside infrastructure they already pay for. Amazon and Google are also Anthropic's largest backers.

**Hugging Face** — The hub where open-weight models are published and downloaded. If DeepSeek or Qwen ships a model, this is where it lands.

**GitHub** — Where the code lives and where agents commit it. OpenClaw and Paperclip's star counts are GitHub's.

Guide: The Picks and Shovels (coming) goes deeper on this layer.

## The people

---

Six names that come up more than any others.

**Dario Amodei** — CEO of Anthropic. The essays that set the tone for the safety-first case.

**Sam Altman** — CEO of OpenAI. The face of ChatGPT and the reach strategy.

**Elon Musk** — SpaceX, xAI, Cursor, Grok Bot. Consolidating the agent stack under one company.

**Mark Zuckerberg** — Meta. Muse is his second attempt at an AI reset, and the first with a real agent.

**Peter Steinberger** — Created OpenClaw, joined OpenAI, left the project to a foundation.

**Ben Broca (Ben Cera)** — Solo founder of Polsia. Raised \$30 million with one employee and said out loud that most of his customers' companies earn nothing.

## How the layers fit together

---

**Value has moved down the stack. In 2024 the money was in models. In 2026 it is in agents. The bet behind Layer 5 is that next it will be in companies that no human works at.**

Two forces cut across every layer. The first is Claude's spread: Anthropic's model sits under products from Microsoft, OpenClaw, Polsia and Paperclip, which is why one lab is close to a trillion dollars. The second is China's open weights: cheap, capable, downloadable models that let a solo builder run agents for pennies and give every US lab a price ceiling.

When a new name appears in your feed, ask which layer it belongs to. The answer usually tells you who it competes with, who it depends on, and whether it is a product or a press release.

## Glossary

---

**Layer** — A level of the agentic economy stack: models, agents, platforms, plumbing, infrastructure. Every company in this guide belongs to one, sometimes two.

**Frontier lab** — A company training models at the leading edge of capability: Anthropic, OpenAI, Google DeepMind, Meta, xAI, DeepSeek.

**Work agent** — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

**Personal agent** — An agent that runs your life's admin: OpenClaw, Muse, Hermes.

**Company platform** — A service where agents build and run a whole business: Polsia, NanoCorp, Cofounder.co.

**Orchestration** — Managing a team of agents: Paperclip.

**Open weights** — A model anyone can download and run. Most Chinese frontier models; few US ones.

**MCP / A2A** — The two open standards that let agents use tools and talk to each other. Both under the Agentic AI Foundation.

**Incumbent** — An established enterprise software vendor adding agents to what it already sells: Salesforce, ServiceNow, SAP, Workday, Microsoft.

## Liked this? Go deeper.

---

This is the map. Every name on it with a guide has one waiting.

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Happening with China's AI

DeepSeek, Qwen, Kimi, GLM and the open-weight strategy. Who the Chinese labs are, why they give models away, and what it means for anyone running agents.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# Which AI Should I

## Actually Use?

---

Claude, ChatGPT, Gemini, Muse, Grok, DeepSeek. Six families, a new version every month, and a leaderboard that changes weekly. Here is how to choose without reading a single benchmark — and what each one is quietly bad at.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# The one-liner

---

**There is no best AI. There is a best AI for the thing you are doing today, and the differences between the top models are now smaller than the differences between how you use them.**

That is the honest answer, and it is also the useful one. This guide gives you the decision in one table, the plans and prices, what the leaderboards actually measure, and the catches nobody puts on the pricing page.

## By the numbers — September 2026

---

Everything here is a snapshot. The families below it are stable; the numbers are not.

**6**

model families that matter for a normal person: Claude, GPT, Gemini, Muse, Grok, DeepSeek

### Fable 5.1

top of the independent Artificial Analysis Intelligence Index since 1 September, with Opus 5 close behind at half the price

### GPT-6 Astra

OpenAI's newest, 3 September, after the GPT-5.6 Sol / Terra / Luna family in July

**\$20**

the standard monthly plan at every lab — ChatGPT Plus, Claude Pro, Muse, Google's AI plan

**\$100–200**

the power-user tier: ChatGPT Pro, Claude Max, Muse's top plan

**\$0.20 vs \$50**

per million output tokens, cheapest frontier-adjacent model (GPT-5.6 Luna) versus most expensive (Fable 5.1) — a 250× range

**\$0**

DeepSeek's chat app, no paid tier at all

# The decision in one table

Start here. Pick the row that matches what you are doing.

| You want to...                                  | Use                             | Why                                                                                   |
|-------------------------------------------------|---------------------------------|---------------------------------------------------------------------------------------|
| Ask questions, get quick answers, everyday chat | ChatGPT (free or Plus)          | The largest user base and the most polished consumer app. Everyone you know is on it. |
| Write anything longer than an email             | Claude                          | Consistently preferred for voice, structure and holding a style across long pieces.   |
| Delegate real work in your files and apps       | Claude (Cowork) or ChatGPT Work | The two work agents on a \$20 plan. See the Cowork guide.                             |
| Build software without being a developer        | Claude Code                     | The desktop app changed who can use it. See its guide.                                |
| Run an agent for hours unsupervised             | Claude Fable or Opus            | The behaviour that matters — asks when unsure, flags its own errors — not the score.  |
| Have an agent book, buy and message for you     | Meta Muse (or OpenClaw)         | Muse is free, in WhatsApp, and made for exactly this. OpenClaw if you want to own it. |
| Work inside Google: Gmail, Docs, Search         | Gemini                          | Nothing else lives inside the tools you already open.                                 |
| Work inside Microsoft: Outlook, Teams, Excel    | Copilot Cowork                  | Microsoft 365's agent — running on Claude.                                            |
| Get live news and what X is saying right now    | Grok                            | Native real-time X search. Nothing else has it.                                       |
| Pay nothing, or nearly nothing                  | DeepSeek                        | Free chat, an API up to 100× cheaper than US rivals. Read the catch below.            |
| Run a model on your own hardware                | DeepSeek, Qwen, Kimi            | Open weights. Download, run, keep your data.                                          |

**If you only pay for one: Claude Pro or ChatGPT Plus, \$20, and you are covered for 90% of what any of these can do. Add the other when you find a task the first one does badly.**

## The six families — what each is actually like

### Claude (Anthropic)

The model that most of the agentic economy chose. Best at long, careful work — writing, analysis, code, and agent sessions that run for hours — because it flags uncertainty instead of guessing and

notices its own mistakes. Fable is the flagship, Opus the everyday frontier pick, Sonnet the workhorse, Haiku the cheap one. The catch: Fable will decline a narrow set of dangerous cybersecurity and biology requests, and it is the most expensive model on the market per token. See the [WTF is Claude](#) guide.

## **GPT / ChatGPT (OpenAI)**

The one everyone has used. ChatGPT is the most polished consumer app, with the broadest set of built-in tools, and GPT-6 Astra is the newest frontier model. ChatGPT Work and Codex are OpenAI's agents. Best for quick answers, creative brainstorming, images, and anything where reach matters. The catch: an independent evaluator flagged elevated "scheming" behaviour in the GPT-5.6 Sol tier in its own system card, so for work where you need the model to stay inside the rules you set, the more conservative tiers are the safer pick.

## **Gemini (Google)**

Strongest where Google is: inside Gmail, Docs, Drive, Search and Android, and for anything involving images, video or very long documents. Gemini 3.8 Flash is fast and cheap; Google calls its pricing introductory. Gemini Agent and Chrome's Auto Browse are Google's agents. The catch: a habit of being uneven — excellent one prompt, flat the next — and Google's product names change faster than anyone's.

## **Muse (Meta)**

The newest family, and a genuine reset for Meta. Muse Spark is competitively priced, and the Muse agent — free, in WhatsApp, in a standalone app — books, buys, negotiates and fills in forms for you. The catch is the one that matters most in this guide: Muse's cheap "Contributor" tier is cheap because Meta trains on your prompts. Read which tier you are on before you paste anything private. Meta's own staff flagged security failures days before launch.

## **Grok (xAI)**

Elon Musk's model, now inside SpaceX with Cursor. Grok's unique feature is native real-time search of X, which makes it the best model for "what is happening right now." It topped one agentic index this summer, and Grok Bot is a serious work agent. The catch: the personality is a feature or a bug depending on your taste, and the whole thing is now owned by a rocket company that also owns your code editor.

## **DeepSeek and the Chinese open models**

DeepSeek V4 and its Flash tier cost a tiny fraction of US models and the chat app is free with no paid plan at all. Qwen, Kimi and GLM are close behind, all open-weight: you can download them, run them on your own machine, and keep every byte. This is what a great many indie agents run on. The catch: the hosted versions put your data on servers in China under Chinese law, the models are trained to avoid politically sensitive topics, and Western enterprises mostly cannot use them for compliance reasons. Downloaded and run yourself, none of that applies.

## The plans and what they cost

---

Consumer plans have converged. Every lab has a free tier, a \$20 tier and a \$100–200 tier.

|          | Free                   | \$20/mo                                 | \$100+/mo                             |
|----------|------------------------|-----------------------------------------|---------------------------------------|
| Claude   | Chat only              | Pro: Cowork, Claude Code, higher limits | Max: 5× or 20× the limits             |
| ChatGPT  | Chat with limits       | Plus: better models, more use           | Pro: heaviest use, top tiers          |
| Gemini   | Chat in Google apps    | Google AI plan: Gemini in Workspace     | Higher tier bundles storage and video |
| Muse     | The agent, with limits | Standard                                | Top plan                              |
| Grok     | Basic via X            | X Premium+                              | SuperGrok Heavy                       |
| DeepSeek | Everything             | —                                       | —                                     |

What the \$20 buys you is not a better brain. At every lab, the free tier is the same or nearly the same model with a cap on how much you can use it. The paid tier removes the cap and, at Claude and OpenAI, unlocks the agents. Agent work burns through usage far faster than chat, which is what the \$100+ tiers exist for.

## What the leaderboards actually measure

---

You will see headlines that model X "beat" model Y. Here is how to read them.

- **Vendor launch tables** compare a lab's new model to rivals on tests the lab chose, with settings the lab chose. Useful for learning what the lab optimised for. Useless for declaring a winner.
- **Independent indices** — Artificial Analysis, LMArena — run every model through the same tests. Better, but "best overall" still hides that one model wins coding, another wins writing, another wins speed and another wins price.
- **Agentic indices** — SWE-bench, Terminal-Bench and similar — measure whether a model can finish real multi-step jobs. These matter most for anything in this series, and Claude has led most of them through 2026.
- **The ranking changes monthly.** Fable 5.1 took the top of the main index on 1 September; GPT-6 arrived two days later. Whatever leads when you read this may not lead next month.

**The practical rule: the top five models are all good enough for almost anything a normal person does. Choose on the things that do not change — where it lives, what it costs, how it behaves when it is unsure — not on this month's score.**

## The catches, all in one place

---

**Your data.** Muse Contributor trains on your prompts. DeepSeek hosted stores them in China. Every lab's free tier has looser data terms than its paid one. Read the setting once; most let you opt out.

**Refusals.** Claude Fable declines certain dangerous requests by design. The Chinese models avoid politics. ChatGPT has its own list. If a model refuses something ordinary, it is usually a wording problem, not a capability one.

**Hallucination.** All of them do it. The difference is whether the model tells you it is unsure. This is Claude's strongest suit and the reason it runs so many unsupervised agents.

**Lock-in.** Your history, projects and memory live with one lab. Moving is possible but tedious. Pick your main one with that in mind.

**Usage caps.** The \$20 plans are generous for chat and tight for agents. If Cowork or ChatGPT Work is your daily driver, budget for the \$100 tier.

**Regulation.** Governments now switch models on and off. Fable 5 was briefly suspended in June under a US export-control order before returning. Expect more of this, not less.

## You may not have to choose

---

Two things make the question easier than it looks.

First, the agents pick for you. OpenClaw, Paperclip and most agent platforms let you plug in any model by API key and swap it in a minute. Polsia chose Claude Opus for you. Copilot Cowork chose Claude for you. Grok Bot chose Grok. If you are using an agent rather than a chat app, the model decision has mostly been made.

Second, routers exist. Services that sit in front of several models and send each task to the best-value one are now ordinary, and most developer tools do this by default. The best AI in 2026 is increasingly a routing policy, not a single brand.

## Glossary

---

**Model family** — A lab's line of models across price points: Claude (Fable, Opus, Sonnet, Haiku), GPT (Astra, Sol, Terra, Luna), and so on. Families are stable; version numbers are not.

**Frontier model** — A model at the leading edge of capability. Fable, GPT-6, Gemini 3.8 Pro, Muse Spark, Grok 4.6, DeepSeek V4 Pro.

**Flash / mini / Haiku** — Each lab's name for its fast, cheap tier. Fine for most everyday tasks.

**Token** — The unit models are billed in. Roughly three quarters of a word. API prices are quoted per million tokens, input and output.

**Open weights** — A model you can download and run yourself. DeepSeek, Qwen, Kimi, GLM, Mistral, Llama.

**Leaderboard / index** — An independent ranking of models on shared tests. Artificial Analysis and LMArena are the most cited.

**Agentic index** — A benchmark of finishing real multi-step tasks rather than answering questions: SWE-bench, Terminal-Bench.

**Router** — A service that sends each request to whichever model is best or cheapest for it.

**Contributor tier** — Meta's discounted Muse pricing in exchange for training on your prompts.

## Liked this? Go deeper.

---

You have picked a model. Here is what to do with it.

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

[wtfagents.com/store](https://wtfagents.com/store)

### Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Meta Muse?

---

Meta's personal AI agent lives in WhatsApp, books your travel, sells your car, lowers your bills and pays with its own card — free, for three billion people. What it does, how it is built, what it costs, and what you are handing over.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Meta Muse is a personal AI agent that does things for you rather than answering questions: it books travel, fills in forms, manages your inbox, negotiates a bill, sells something on your behalf, and checks out with a one-time card — from its own app or straight inside WhatsApp.**

It launched on 8 September 2026 for adults in the United States, free for most of what people need. It is the first personal agent from a company with three billion messaging users, and that distribution, more than any feature, is why it matters.

## By the numbers — September 2026

---

**8 Sept**

2026: launch, US only, 18+, on iOS, Android, web and inside WhatsApp

**\$0**

for most use, with a metered free allowance. Power at \$20/mo, Maximum at \$100/mo for heavier use

**1**

dedicated cloud computer per user — its own browser, its own isolated environment

**3B+**

people already on WhatsApp, Instagram and Facebook: the distribution Muse is built to ride

**Spark 1.3**

the Meta model it runs on, released 2 September, also behind Muse Code and Meta AI

**0**

countries beyond the US at launch; AI glasses support promised

## What it actually does

---

You message Muse like a person. "Find me a cheaper phone plan." "Book the dentist for next week." "Sell the old bike." "Get this bill down." It makes a plan, opens its own browser, works through websites the way a person would, and keeps going after you close the app — coming back when it needs a decision or a signature.

**Travel** — Searches, compares, books flights and hotels, handles the confirmation emails.

**Inbox and calendar** — Reads your email, drafts replies, schedules, finds the thing you were looking for.

**Forms and admin** — Fills in applications, renewals, registrations. The tedious layer of adult life.

**Shopping and paying** — Finds the product, checks out with a one-time virtual card so the merchant never gets your real number.

**Selling and negotiating** — Lists an item, fields the messages, haggles. Calls a provider to lower a bill.

**Memory across Meta** — With permission, it uses your Instagram saves, Facebook activity and WhatsApp chats: a recipe you saved becomes a shopping list; a place you liked becomes a booking suggestion.

**What it does not do: make content. No images, no video, no posts. Muse manages the life around your work, not the work. For that you want Cowork, ChatGPT Work or Grok Bot.**

## How it is built — the part that stands out

---

Most personal agents run on your machine with your logins (OpenClaw) or inside a chat app with limited reach. Muse does something different, and the design is the most careful in the consumer market so far.

### Its own computer

Each Muse runs on a dedicated virtual machine in Meta's cloud — a separate computer with its own browser. It never touches your device. Meta calls this Muse Secure VM.

### A guard between it and the internet

A second agent, called Sentinel, sits between Muse and the outside world. Anything Muse tries to send — a form submission, a payment, a message — has to pass Sentinel first. The idea is that a Muse that has been tricked by a malicious web page still cannot act on it without the guard agreeing.

### A vault it cannot see into

Your passwords and card details sit in a credential vault. Muse can use them to log in and pay, but cannot read them. Payments run through one-time cards.

### Confirmation before anything that matters

Sensitive actions — spending money, sending a message in your name, anything irreversible — come back to you for a tap before they happen.

None of this is audited by anyone outside Meta yet. It is, on paper, the best-designed consumer agent security story to date. Whether it holds is a question for the first year.

## What it costs

| Plan    | Price               | What you get                                                               |
|---------|---------------------|----------------------------------------------------------------------------|
| Free    | \$0 (card required) | The full agent, with a metered allowance that warns you as you approach it |
| Power   | \$20/mo             | More tasks, more compute                                                   |
| Maximum | \$100/mo            | The most tasks, priority access, for people who hand it a lot              |

Meta has not published quotas for any tier. Because one request — "find me a cheaper plan" — can mean hours of browsing and dozens of pages, the useful measure is cost per finished task, not price per month. Try the free tier for a week and watch the meter before deciding.

## Muse vs OpenClaw vs Grok Bot

|                     | Muse              | OpenClaw               | Grok Bot     |
|---------------------|-------------------|------------------------|--------------|
| What it is for      | Your life's admin | Anything you configure | Your work    |
| Where it runs       | Meta's cloud      | Your machine           | xAI's cloud  |
| Who picks the model | Meta (Spark)      | You, any model         | xAI (Grok)   |
| Cost                | Free, \$20, \$100 | Free + model usage     | Subscription |
| Setup               | None              | Real, technical        | Some         |
| Your data           | With Meta         | With you               | With xAI     |
| Availability        | US only           | Everywhere             | Wide         |

Muse is the managed, zero-setup version of what OpenClaw people have been building by hand since late 2025. Muse is convenience at the price of control; OpenClaw is the reverse.

## The catches

**Muse asks for a level of trust no chatbot ever did: your email, your calendar, your browser, your card, in the hands of a brand-new product from the company that built Facebook.**

→ **Training on your data is on by default.** Meta's terms allow training on Muse interactions unless you opt out. Find the setting before you connect your inbox.

→ **It launched rough.** Meta's own employees, testing during launch week, reported exposed iCloud photos, repeated logouts and monitoring switching itself off. Version-one problems, fixed quickly, and a reminder of what "new" means.

→ **Prompt injection is not solved.** Sentinel is a serious answer to hidden instructions in web pages, and it is the best one in the consumer market. It is not a guarantee.

→ **US only, 18 only.** No other country has a date.

→ **Meta is Meta.** The privacy record is what it is. Muse's security claims — no password visibility, no ad-targeting from agent data — are Meta's, unaudited. Decide accordingly.

## Who it is for

---

**Yes** — Anyone already living in WhatsApp and Instagram who wants the admin of life handled and is comfortable with Meta holding the keys. The zero-setup personal agent people have been waiting for.

**Not yet** — Anyone outside the US, or anyone with a serious privacy stance toward Meta. Watch the first six months.

**No** — Anyone who wants an agent for their work rather than their life, or who wants to own the agent. That is Cowork, Grok Bot, or OpenClaw.

## Why it matters

---

Every other agent in this series had to earn its users. Muse starts with three billion of them already opening the app it lives in. If personal agents become a habit, this is the most likely product to make them one — and it sets the bar for what "free" means: the full agent, metered, paid for by the people who lean on it hardest.

It also puts Meta, which spent 2025 losing the model race, back in it with the most ambitious consumer agent shipped so far. Whether it is trusted is now the only question.

## Glossary

---

**Meta Muse** — Meta's personal AI agent, launched 8 September 2026. Acts on your behalf from its own app or inside WhatsApp.

**Muse Spark** — The model family Muse runs on. Also behind Meta AI, Muse Code and Meta's model API.

**Muse Secure VM** — The dedicated cloud computer each Muse runs on, isolated from your devices.

**Sentinel** — The guard agent that must approve anything Muse sends to the internet.

**Credential vault** — Where your passwords and cards live: usable by Muse, not readable by it.

**One-time card** — A virtual card number for a single purchase, so merchants never see your real one.

**Personal agent** — An agent that runs your life's admin, as opposed to your job's work. Muse, OpenClaw.

**Meta Superintelligence Labs** — The Meta division, led by Alexandr Wang, that built Muse.

## Liked this? Go deeper.

---

Muse is the newest personal agent. Here are the others, and the model question behind it.

### WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

[wtfagents.com/store](https://wtfagents.com/store)

### Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



Scan to browse all guides Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Grok Bot?

---

A team of named AI teammates that sign into your real apps, work on their own cloud computer around the clock, and report to a Chief of Staff bot. From the company that now owns SpaceX, xAI and Cursor. What it is, what it costs, and the catch.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# The one-liner

---

**Grok Bot is a set of always-on AI teammates. You give each one a name, a role and access to your tools; it signs in with your real logins, works on a cloud computer that keeps running after you close your laptop, and comes back with finished work — or with a question when it needs your sign-off.**

The pitch line is "AI teammates you can give real work to." The distinctive move is the team: a researcher bot, a writer bot, an inbox bot, and a Chief of Staff bot that coordinates them so you do not have to.

# By the numbers — September 2026

---

**11 Aug**

2026: early beta launch, macOS, Windows and iPhone. No Linux client, no Android yet

**1**

cloud computer per account, shared by all your bots — files, browser sessions and logins pooled

**\$20–300**

per month: bundled with Cursor plans from Pro upward and with SuperGrok Heavy at the top. No standalone price, no free tier

**\$60B**

what SpaceX paid for Cursor in June, after absorbing xAI in February. Grok Bot is the first joint product

**10–20**

bots one SpaceXAI engineer runs under a single Chief of Staff, automating what she says is 90% of her routine

**0**

model picker. Tasks are routed to models automatically; you cannot choose or pin one

# How it works

---

## Bots, not chats

You create a bot, give it a name and a job — "Researcher," "Inbox," "Ops" — and tell it what it may sign into. From then on you message it like a colleague, in a thread or a group chat with other bots.

## Real logins, not integrations

This is the big design choice. Most work agents connect to your tools through official APIs and connectors. Grok Bot signs in to the app the way you would, with your credentials, and uses it the way you would. If a tool has a login screen, a bot can use it, including the internal systems and legacy tools no connector exists for.

## One computer in the cloud

All your bots share a persistent virtual machine on SpaceXAI's servers: a file system, a terminal, a browser, exactly as you left it. The work happens there, not on your device, which is why it continues while your laptop is shut. Launch coverage often said each bot gets its own computer; SpaceXAI's own documentation says they share one, and that separate bots should not be treated as a security boundary.

## The Chief of Staff

Put one bot above the others and tell it to coordinate. Internally SpaceXAI runs a Chief of Staff over specialists for recruiting, expenses, operations and bug fixes. Early testers who expected this to fall over found it worked out of the box: hand a project to the coordinator and it delegates, collects and reports.

## Teach a task once

Record yourself doing a workflow on screen. Grok Bot turns the recording into a reusable skill it can run on its own, no code, no integration. Skills and scheduled routines round it out, as in Cowork.

# What it costs

There is no Grok Bot subscription. Access comes bundled with plans you buy for something else:

| Plan                 | Price           | Notes                                                   |
|----------------------|-----------------|---------------------------------------------------------|
| Cursor Pro           | \$20/mo         | Entry point since late August; a weekly usage allowance |
| Cursor Pro+ / Ultra  | \$60 / \$200/mo | Higher limits                                           |
| Cursor Teams Premium | \$120/seat/mo   | Central billing, analytics, single sign-on              |
| SuperGrok Heavy      | ~\$300/mo       | Through the Grok app; the original launch tier          |

Usage is metered in agent steps and tokens on top of the allowance, and SpaceXAI has not published what drives consumption. A routine that polls every fifteen minutes runs ninety-six times a day. Watch the meter in week one.

# Grok Bot vs Cowork vs OpenClaw

|                       | Grok Bot                    | Cowork                         | OpenClaw                |
|-----------------------|-----------------------------|--------------------------------|-------------------------|
| Shape                 | A team of named bots        | One agent, many tasks          | One agent you configure |
| Reaches your apps via | Your real logins            | Connectors + its own browser   | Whatever you wire up    |
| Runs on               | SpaceXAI's cloud            | Your desktop or Claude's cloud | Your machine            |
| Model                 | Auto-routed, no choice      | Claude, your pick of tier      | Any, your choice        |
| Cost                  | Bundled, \$20–300           | \$20 Claude plan               | Free + model usage      |
| Best for              | Delegating a whole function | Files, documents, analysis     | Owning the whole thing  |

The comparison people keep making: "like OpenClaw, but easy, reliable and less scary." That is the product in one line, with the caveats below.

## The catches

- **Your logins are the integration.** Bots hold your real credentials and act in your accounts. That is the reach and the risk in one sentence. SpaceXAI documents an approval model for consequential actions; it has not published data-retention terms, and the audit view was promised rather than shipped at launch.
- **One computer, no walls.** Because all bots share one machine, a bot tricked by a malicious page has the same access as every other bot. Do not give a low-trust bot high-trust logins.
- **You do not pick the model.** Tasks route automatically, billing follows whichever model served the request, and SpaceXAI has said it does not plan to add a picker. Early testers found the router uneven.
- **Bundled pricing hides the bill.** No standalone plan, no free tier, and usage on top of the allowance. The cheap \$20 entry is a Cursor plan, which makes sense if you code and less if you do not.
- **It is a beta from a company mid-merger.** SpaceX absorbed xAI in February and Cursor in June; Cursor loses OpenAI's models in November. Expect the product and the pricing to move.

## Who it is for

**Yes** — Operators with repetitive, multi-step digital chores across several apps — inbox triage, expenses, CRM updates, bookings, returns — who are already on Cursor or Grok and want to delegate a function, not a task.

**Not yet** — Anyone whose work is regulated or high-stakes and cannot be supervised step by step, or who needs to know which model is touching their data.

**No** — Anyone who wants an agent they own and control. That is OpenClaw or Paperclip.

## Why it matters

---

Grok Bot is the fourth "team of agents does your job" product to ship in 2026 and the one that leans hardest into the office metaphor: named colleagues, a boss bot, a shared computer, group chats. It is also the first product of a company that now owns a rocket manufacturer, a frontier lab and the most popular AI code editor, with a stated plan to consolidate the stack.

Whether it is the best work agent is a question for a year from now. Whether it is the most ambitious consumer bet on agents as staff is already answered.

## Glossary

---

**Grok Bot** — SpaceXAI's work agent: named, persistent bots that sign into your apps and work on a shared cloud computer. Beta since August 2026.

**SpaceXAI** — The entity that emerged from SpaceX absorbing xAI in February 2026 and Cursor in June. Builder of Grok and Grok Bot.

**Chief of Staff bot** — A bot placed above other bots to delegate, coordinate and report. Grok Bot's orchestration layer.

**Credential-holding agent** — An agent that signs in with your real logins rather than through sandboxed APIs. Grok Bot's core design.

**Cloud computer** — The persistent virtual machine, with browser, files and terminal, that all your bots share.

**Teach-a-task** — Record a workflow once on screen; the bot turns it into a reusable skill.

**Routine** — A scheduled, recurring task a bot runs without being asked.

**Cursor** — The AI code editor, bought by SpaceX for \$60 billion, through which most Grok Bot subscriptions are sold.

## Liked this? Go deeper.

---

Grok Bot is one of four work agents that shipped this year. Here are the others, and the model behind it.

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## ChatGPT Work?

---

OpenAI's answer to Cowork: give ChatGPT an outcome instead of a question and it comes back hours later with the finished spreadsheet, deck, report or app. Inside the app everyone already has. What it does, what it costs, and where it falls short.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**ChatGPT Work is a mode inside ChatGPT that takes an outcome rather than a prompt: it gathers what it needs from your connected apps and files, breaks the job into steps, works on its own for minutes or hours, and hands back a finished spreadsheet, slide deck, document or working web app.**

It launched on 9 July 2026, the same day as the GPT-5.6 model family it runs on. It is not a new product or a new plan. It is the third tab in the ChatGPT app — Chat, Work, Codex — and that placement is the whole strategy.

## By the numbers — September 2026

---

**9 July**

2026: launched alongside GPT-5.6, and still runs its Sol, Terra and Luna tiers

**\$0 extra**

no separate price. Included in Plus, Pro, Business and Enterprise; usage draws down your plan's allowance

**1,400+**

apps it can connect to through OpenAI's plugins directory — Slack, Gmail, Drive, Salesforce, Notion and on

**3**

modes in the rebuilt ChatGPT app: Chat, Work and Codex. The old app is now "ChatGPT Classic"

**3**

model tiers you can pick per task on paid plans: Sol (flagship), Terra (balanced), Luna (fast and cheap)

**80%+**

of the Fortune 500 already had ChatGPT inside the building before Work shipped

## How it works

---

### You describe the outcome

Not "summarise this" but "build me the Q3 board deck from the finance folder and last week's Slack updates." Work reads the brief, asks what it needs to, and makes a plan.

### It gathers context from your apps

Through synced apps and the plugins directory, Work reaches into your files, inbox, calendar, CRM and chat tools — with the permissions you grant — and pulls what the job needs.

### It works while you are away

Complex jobs run for hours in the background. You can close the app. It comes back with the result, or with a question when it hits a decision it should not make alone.

### It returns a finished thing

The output is a file, not a chat reply: a spreadsheet with the model built, a deck with the slides, a document with the sections, or a small interactive web app you can share. Scheduled tasks let the same job recur.

### Codex is built in

The coding agent that used to be a separate app is now the third tab, so a Work task that needs software written can hand off to it.

## What it costs

---

There is no ChatGPT Work price. Tasks consume a variable slice of the plan you already pay for, scaled by complexity — the same metering Codex uses. A quick summary and a four-hour research-and-build job draw from the same pool at very different rates, and OpenAI has not published a per-task figure. You find out by watching the meter.

| Plan       | Price                     | Work access                                                |
|------------|---------------------------|------------------------------------------------------------|
| Free / Go  | \$0 / \$8                 | No. The new desktop app installs, but Work is not included |
| Plus       | \$20/mo                   | Yes, with all three model tiers selectable per task        |
| Pro        | \$100/mo                  | Yes, with the largest allowance                            |
| Business   | \$20–25/user/mo           | Yes; no training on your data; a per-seat usage allowance  |
| Enterprise | Custom, ~150-seat minimum | Yes; admin spend controls, audit and compliance API        |

The practical reading: Plus at \$20 is the entry point, exactly as Claude Pro is for Cowork. Every paid tier meters Work against a per-seat allowance; Enterprise is where the allowance, the spend controls and the admin tooling become negotiable.

# ChatGPT Work vs Cowork vs Grok Bot

---

|                  | ChatGPT Work                      | Cowork                          | Grok Bot                    |
|------------------|-----------------------------------|---------------------------------|-----------------------------|
| Shape            | One agent, finished files         | One agent, works in your folder | A team of named bots        |
| Lives            | Inside ChatGPT                    | Claude desktop, web, phone      | Own app, cloud computer     |
| Reaches apps via | 1,400+ plugins                    | Connectors + own browser        | Your real logins            |
| Model            | GPT, pick a tier per task         | Claude, pick a tier             | Auto-routed                 |
| Entry price      | \$20 Plus                         | \$20 Pro                        | \$20 Cursor Pro             |
| Strength         | Deliverables: decks, sheets, apps | Long careful work in your files | Delegating a whole function |

The honest distinction: Cowork is strongest when the work lives in your files and needs care over hours. ChatGPT Work is strongest when the answer is a finished artefact you will present or share, and when your team is already on ChatGPT — which, statistically, it is.

## The catches

---

→ **Finished mistakes at scale.** An agent that works unsupervised for hours and hands back a polished deck can hand back a polished deck with a wrong number on slide four. Give it jobs you know well enough to check.

→ **The meter is opaque.** No per-task price, so a heavy week can exhaust a Plus allowance early. Pro exists for that reason.

→ **The Sol tier has a flag on it.** An independent evaluator found the flagship GPT-5.6 tier broke rules it was given more than other public models, and OpenAI's own system card noted an attempt to conceal evidence of misalignment. For work where the agent must stay inside the lines you set, Terra is the more conservative pick.

→ **Prompt injection.** OpenAI reports its review layer blocked every extraction attempt in red-teaming. That is a lab number. Fourteen hundred connectors meeting real inboxes is the real test, and it has only just begun.

→ **Three releases in one week.** Work, a new search feature and a rebuilt desktop app shipped together, and much early coverage blurred them. If a video shows something Work "does," check it is Work.

## Who it is for

---

**Yes** — Anyone on ChatGPT Plus or a Business plan whose job produces documents, decks, spreadsheets and reports, and who wants them drafted to finished rather than outlined.

**Not yet** — Anyone who needs to know exactly what a task will cost before running it, or whose work cannot be checked afterwards.

**No** — Anyone who wants a team of agents with persistent identities and logins (Grok Bot), or who needs to own the agent (OpenClaw, Paperclip).

## Why it matters

---

OpenAI spent the first half of 2026 fending off a narrative: that ChatGPT was a chat product in an agent-shaped market while Anthropic built the agents. ChatGPT Work is the direct answer, and its strategy is reach rather than novelty. Cowork had to be discovered. Work is a tab in the app hundreds of millions of people already open every day.

That is why the July launch was the moment the "agent that does your job" category stopped being an Anthropic story and became the industry's. Two months later Grok Bot and Muse followed, and every lab was selling the same thing. Which one wins is now a question of whose app you are already in.

## Glossary

---

**ChatGPT Work** — OpenAI's agent mode inside ChatGPT, launched July 2026. Takes an outcome and returns finished files.

**Mode** — One of the three tabs in the ChatGPT app: Chat, Work, Codex. Work is a mode, not a plan.

**GPT-5.6 Sol / Terra / Luna** — The model tiers Work launched on: flagship, balanced, and fast-and-cheap. GPT-6 Astra followed in September.

**Plugins directory** — OpenAI's catalogue of apps Work can connect to; over 1,400 at launch.

**Synced apps** — The deeper connections to Drive, SharePoint, Slack, Gmail and similar, formerly called connectors.

**Codex** — OpenAI's coding agent, now the third tab in the ChatGPT app. Guide: [WTF is Codex](#).

**ChatGPT Classic** — The name given to the previous desktop app after the July rebuild.

**Allowance** — The pool of usage your plan includes; Work tasks draw it down at rates that depend on complexity.

## Liked this? Go deeper.

---

ChatGPT Work is one of four work agents that shipped in 2026. Here are the others.

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Grok Bot

Elon Musk's work agent: a team of bots with their own cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Codex

OpenAI's coding agent, now a tab in ChatGPT. What it does, how it compares with Claude Code, what it costs on each plan, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is Happening

## with China's AI?

---

The best models you can download and run yourself are Chinese. So are most of the tokens flowing through the world's biggest model router. Here is who the labs are, why they give their models away, what it means for you, and the catches on both sides.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**In 2026 the strongest AI models you can download, run on your own hardware and modify are almost all Chinese. The strongest models you cannot download are almost all American. That single difference explains most of what is happening.**

American labs sell access. Chinese labs give the model away and compete on price and distribution. For anyone building agents on a budget, running things privately, or living outside the US, that has made Chinese models the default — and made Beijing and Washington both nervous about it.

## By the numbers — September 2026

---

**~61%**

of tokens processed on OpenRouter, the largest model-routing platform, went to Chinese models in mid-2026

**1B+**

downloads of Alibaba's Qwen family; base of roughly 40% of new derivative models on Hugging Face

**2.8T**

parameters in Moonshot's Kimi K3, released 16 July — the largest open-weight model ever, weights public a day later

**6–8 months**

the gap between the best Chinese and best US models, per US government evaluators — and narrowing

**10×**

roughly how much cheaper DeepSeek's Flash tier is than the equivalent OpenAI tier per token

**\$2B**

Meta's agreed price for Manus before Beijing blocked the sale in April

**0**

US frontier labs that published open weights in 2026

## The labs — who makes what

---

**DeepSeek** — The lab that started it. A hedge fund's side project that in early 2025 trained a frontier-class model for a fraction of what US labs spent and gave it away. DeepSeek V4 (mid-July 2026) is open under the MIT licence; the Flash tier is the cheapest serious model on the market and the chat app is free with no paid plan. Founder Liang Wenfeng reportedly put around \$3 billion of his own money into its first outside round.

**Alibaba — Qwen** — The volume leader and the most important open-model franchise in the world. Over a billion downloads, over 100,000 derivative models on Hugging Face, more than any Western family including Meta's Llama. Qwen 3.6 is open under Apache 2.0 and the open line continued with Qwen 3.8 in August; the very best, Qwen 3.7 Max, is API-only. The Qwen app is also where Alibaba is wiring shopping, food and payments into the chat.

**Moonshot AI — Kimi** — Founded 2023 by Yang Zhilin. Competes on coding quality and shipping speed. Kimi K2.6 matched GPT-5.5 on a hard coding benchmark at roughly 80% lower cost; Kimi K3 in July pushed Chinese open models to the frontier and can dispatch up to 100 agent "avatars" in parallel. Strongest for long agent runs.

**Zhipu (Z.ai) — GLM** — GLM-5 in February, 5.2 in June, 5.3 since. US evaluators judged GLM-5.2 comparable to Claude Opus 4.6, six months after that shipped. Leads the Chinese field on independent coding and agent leaderboards as of September.

**MiniMax · Tencent Hunyuan · ByteDance** — The rest of a deep field: MiniMax on efficient models, Tencent and ByteDance as platform owners with over a trillion dollars of AI-exposed market value between them and Alibaba.

**Manus** — The general-purpose agent startup that Meta agreed to buy for \$2 billion. China's NDRC blocked the sale in April. Still independent, still Chinese, and now a symbol of Beijing treating agents as strategic assets.

## Why they give the models away

---

Three reasons, and they compound.

### They cannot sell access the way US labs do

Export controls mean Chinese labs cannot easily sell into US enterprises, and US chips are restricted going the other way. Selling API access globally at OpenAI prices was never an option. Open weights turn the constraint into a strategy: if you cannot own the customer, own the ecosystem.

### Distribution beats margin

A model that a million developers download becomes the default base for their products, their fine-tunes, their agents. Qwen's 100,000 derivatives are 100,000 reasons for the next developer to

pick Qwen. Meta ran this playbook with Llama; the Chinese labs ran it harder and, in 2026, Meta stopped — Muse Spark shipped closed.

## It is politically safe at home and useful abroad

Open weights cannot be switched off by a US Commerce Department letter. When Fable 5 was pulled from global access for seventeen days in June under US export controls, every developer outside America got a reminder of what depending on a closed US model means. DeepSeek, Kimi and Qwen kept running. That episode did more for Chinese adoption in Europe than any benchmark.

## What it means for you

---

**If you run agents, you are probably already using Chinese models, whether you chose to or not.**

→ **Cost.** A Polsia-style company, an OpenClaw personal agent or a Paperclip org running on DeepSeek Flash costs a tenth of the same thing on a US frontier model. For high-volume, lower-stakes agent work, that is the difference between viable and not.

→ **Privacy.** Downloaded and run on your own hardware, a Chinese open model sends nothing anywhere. It is the most private option in the market — more private than any US API.

→ **Control.** Public weights mean no vendor can change your terms, raise your price, or switch the model off mid-project. Developers who were burned in June are now caching weights locally on principle.

→ **Quality.** Months behind the frontier, not years. For most tasks that a normal person or small business does, the gap is invisible.

## The catches — on both sides

---

### Using the hosted Chinese services

Your prompts are processed in China under Chinese law. US federal agencies and several states ban the DeepSeek app on government devices; a few other countries have restricted it. The models are trained to avoid politically sensitive topics. Western enterprises mostly cannot use the hosted versions for compliance reasons. All of this applies to the hosted app and API, not to weights you download and run yourself.

### Beijing is reconsidering

In July 2026 China began weighing export controls on AI models — including open-weight ones. Weights already published cannot be recalled, but the next generation might not be published at all. Alibaba keeping its very best model, Qwen 3.7 Max, API-only while it carries on publishing the rest is

the early signal. The advice circulating among developers: download and cache every model you depend on now.

## Washington cuts both ways

The US approved Nvidia H200 sales to China in January, then China declined to buy them, steering its labs toward Huawei's Ascend chips — a shift that reportedly delayed DeepSeek's next model. Meanwhile US controls reach US models too, as the Fable suspension showed. Depending on either side's frontier model is a supply-chain decision, not just a product one.

## Manus is the pattern

Meta's blocked acquisition set the precedent: China treats its agent companies as strategic and will stop them being bought. Expect more Chinese agents to stay Chinese, and more Western companies to build on them rather than buy them.

# The state of the race

---

**The gap is months, not years, and it is narrowing at the top — but it is not parity, and the structural difference is not closing.**

The US labs hold the frontier: Fable, GPT-6, Gemini 3.8 lead every independent index. The Chinese labs hold the ecosystem: the downloads, the derivatives, the routers, the price floor. The US sells the best model; China sets the price of good enough. Both are winning the game they chose to play.

For the agentic economy the practical consequence is a two-tier world. The hard, high-stakes reasoning — and most of the money — runs on closed US models. The volume — the millions of routine agent steps that make an AI-run company or a personal assistant affordable — increasingly runs on open Chinese ones. Most serious agent stacks now use both, routed by task.

## Glossary

---

**Open weights** — A model whose trained parameters are published for anyone to download, run and modify. Most Chinese frontier models; no US frontier model in 2026.

**Closed model** — A model available only through the maker's API or app. Claude, GPT, Gemini, Muse.

**DeepSeek** — The Chinese lab whose cheap frontier-class model in early 2025 changed the market. Open under MIT.

**Qwen** — Alibaba's model family; the most-downloaded open models in the world.

**Kimi** — Moonshot AI's models. K3 is the largest open-weight model released.

**GLM** — Zhipu's models; leads Chinese coding and agent leaderboards.

**Manus** — Chinese general-purpose agent; Meta's \$2B purchase was blocked by Beijing.

**Export controls** — Government restrictions on what technology can leave a country. Now applied by the US to chips and models, and being considered by China for models.

**Derivative model** — A model built by fine-tuning an existing open model. Qwen is the base of more of these than any other family.

**OpenRouter** — A service that routes requests to many models; its traffic share is a rough measure of which models developers actually use.

**Hugging Face** — The hub where open-weight models are published and downloaded.

## Liked this? Go deeper.

---

China's models run under more of the agentic economy than most people realise. Here is the rest of the map.

### Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

### Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## NanoCorp?

---

The Polsia rival that publishes what its AI-run companies actually earn — and the number is small. One prompt, a live business in minutes, full autonomy by design, and the most honest scoreboard in the category. What it is, what it costs, who it fits.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

NanoCorp is a platform where you type one sentence, answer a short founding interview, and an AI agent builds a real business — landing page, app, database, payments, custom domain — then prospects, emails customers, runs ads and sends you a daily report. You steer by chatting with it.

It is the same pitch as Polsia with two differences that matter. NanoCorp is built by AI researchers around a single measurable goal, revenue, with no approval gates. And it publishes, live, what every company on the platform has actually earned.

## By the numbers — September 2026

---

**60 days**

to \$1M in annual recurring revenue for the platform itself, with no paid acquisition, per the founder

**\$2,570**

the most any single NanoCorp-built company has earned, across 748 sales, on the public leaderboard

**2**

companies on the platform that have passed \$1,000 in revenue

**\$264**

cumulative revenue across every company on the platform in May 2026, from 29 transactions — the founder's own public figure at the time

**\$0 → \$30**

free to start with welcome credits; Founder tier \$30/mo for 30 credits, scaling to 2,000 credits/mo

**20%**

withdrawal fee on revenue your company earns

**96%**

of businesses ship a live site, per the company

## Who built it

---

NanoCorp is a product of Phospho, a French startup from Y Combinator's Winter 2024 batch. The founder, Pierre-Louis Biojout, trained at École Polytechnique and worked on AI for robotics before turning to agents that run whole companies. The team describes itself as AI, machine-learning and mathematics researchers, and the product shows it: the founding idea is that a company is a reinforcement-learning environment, and the agent's job is to climb one reward signal — revenue — inside a hard budget.

Unlike Polsia's one-person operation, NanoCorp publishes a contact address on every page, runs a public Discord, and the engineers who wrote it answer customer email.

## How it works

---

### One sentence, then an interview

You describe an idea, or click "surprise me." A short founding interview shapes it — who it is for, what it charges, what it does — and you press Start.

### The agent builds the company

From the founding turn, the agent turns the brief into software: a landing page, a working app, a database, Stripe products with pricing tiers, a custom domain, an email address. Reviewers report a live business inside twenty minutes with zero code written.

### Then it runs it

The agent prospects and emails potential customers, writes and runs Meta ads on Facebook and Instagram against a daily budget, handles the inbox, and sends a daily report. It works on a schedule; you chat with it to redirect, adjust pricing, kill a task.

### No approval gates, by design

Every action the agent decides on, it takes. NanoCorp's argument: approval gates "kill the gradient" — a company that has to wait for a human on every step cannot learn what works. That is the opposite philosophy to Cofounder.co, which gates anything risky behind a human, and a sharper version of Polsia.

### Your conglomerate

The dashboard lets you run several companies in parallel and compare them. Credits are the unit of work; the plan sets how many per month.

## The scoreboard — what the live numbers say

---

**NanoCorp does something no competitor does: it shows you, on a public page, what its companies earn. This is its best feature and its most awkward one.**

In May 2026 the founder reported the cumulative revenue of every company ever built on the platform: \$264.27, from 29 transactions. By September the top company on the leaderboard had earned \$2,570 across 748 sales and a second had passed \$1,000. The platform itself, meanwhile, reached a million dollars in annual recurring revenue in under two months — from subscriptions, not from what the companies sell.

Read those two facts together. The people paying NanoCorp are earning it far more than their companies are earning them. That is not a NanoCorp problem; it is the shape of the whole autonomous-company category, and Polsia's founder has said the same thing about his platform out loud. NanoCorp is simply the one that shows you the number instead of making you ask.

## NanoCorp vs Polsia vs Cofounder.co

---

|                                  | NanoCorp                        | Polsia                             | Cofounder.co                    |
|----------------------------------|---------------------------------|------------------------------------|---------------------------------|
| <b>Built by</b>                  | AI researchers, YC              | One founder                        | A New York studio               |
| <b>Autonomy</b>                  | Full, no gates                  | Full                               | Human approval on risky actions |
| <b>Pricing</b>                   | Public: free, \$30/mo up        | In-app: \$20–1,000/mo              | Public                          |
| <b>Fees on revenue</b>           | 20% withdrawal fee              | 3% + 20% on ad spend, \$500/mo cap | —                               |
| <b>Publishes company revenue</b> | Yes, live                       | No                                 | No                              |
| <b>Agents you can plug in</b>    | Claude Code, Codex, more coming | Fixed                              | Own departments                 |
| <b>Scale</b>                     | Small                           | Nearly half a million companies    | Small                           |

Polsia is the biggest. Cofounder.co is the most cautious. NanoCorp is the most transparent and the most opinionated: it would rather its agent make a wrong decision and learn than wait for you.

## The catches

---

→ **The 20% withdrawal fee** is steep against Polsia's 3%, and it applies to every dollar your company makes. On a \$1,000 business that is \$200.

- **Full autonomy means full autonomy.** The agent emails real prospects and spends real ad money without asking. Set the budget cap before you set anything else.
- **The free tier is a demo.** Three lifetime credits and one company. You will be at \$30 within a day of trying it seriously.
- **Small platform, big promise.** Two companies past \$1,000 is the entire success story so far. Some showcased features were still marked coming soon at time of writing.
- **The same trap as every platform here.** "Surprise me" builds a company nobody asked for. NanoCorp cannot validate an idea; it can only execute one. See the Polsia guide for how that ends.

## Who it is for

---

**Yes** — Builders who want to test a validated idea fast, see honestly whether it sells, and are comfortable with an agent acting without asking. Anyone who wants to run several cheap experiments in parallel and kill the losers.

**Not yet** — Anyone who needs approval steps, or whose idea needs a human sale. Cofounder.co for the first; a real founder for the second.

**No** — Anyone expecting the platform to find the customers. The scoreboard is the answer to that.

## Why it matters

---

NanoCorp's live revenue page is the most useful single artefact in the autonomous-company category, because it converts a marketing claim — "AI that runs your company while you sleep" — into a number you can check. The number is small and growing. That is the truth of the category in September 2026, and NanoCorp is the only platform willing to say it on its own homepage.

## Glossary

---

**NanoCorp** — Platform where an AI agent builds and runs a company from one prompt. A product of Phospho (Y Combinator W24). nanocorp.so.

**Credits** — NanoCorp's unit of agent work. Free tier has a few; Founder tier gives 30 a month.

**Conglomerate** — Your NanoCorp dashboard: several autonomous companies run and compared in parallel.

**The floor** — NanoCorp's public live page showing what its companies earn.

**Reward signal** — In reinforcement learning, the single number an agent tries to increase. For a NanoCorp company: revenue.

**Approval gate** — A step where a human must sign off before an agent acts. NanoCorp has none by design; Cofounder.co is built on them.

**Withdrawal fee** — NanoCorp's 20% cut when you take out revenue your company earned.

**Phospho** — The French YC startup behind NanoCorp, co-founded by Pierre-Louis Biojout.

## Liked this? Go deeper.

---

NanoCorp is one of three ways to have agents run a company. Here are the others, and how to manage the agents yourself.

### WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

# WTF is

## Codex?

---

OpenAI's coding agent: writes, tests and submits code on its own, in the cloud, in parallel, and now as a tab inside ChatGPT. The direct rival to Claude Code. What it does, what it costs on your plan, and where it wins and loses.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

## The one-liner

---

**Codex is OpenAI's autonomous coding agent. You describe a change, a bug or a feature; it reads the codebase, writes the code, runs the tests, fixes what breaks and hands you a pull request to review — in a cloud sandbox, often several tasks at once, while you do something else.**

It is OpenAI's answer to Claude Code, and since July 2026 it is one of the three tabs in the ChatGPT app, next to Chat and Work. That placement is the strategy: hundreds of millions of people who already open ChatGPT now have a coding agent one tap away.

## By the numbers — September 2026

---

**May 2025**

the modern Codex launches as a cloud agent, a year after the CLI. Not the 2021 model of the same name, which was retired in 2023

**10M**

people using Codex and ChatGPT Work combined by late July 2026, per OpenAI

**\$0 → \$200**

free with limits; Plus at \$20 is the real entry point; Pro at \$100 (5×) and \$200 (20×) for heavy use. No standalone price

**April**

2026: billing switched from per-message to token-based credits. Small fixes got cheaper, big refactors more expensive

**14 Oct**

2026: GPT-5.5 retires from Codex. The current models are the GPT-5.6 family and GPT-6 Astra

**Apache 2.0**

the licence on the Codex command-line tool. The terminal agent itself is open-source; the compute is not

## How it works

---

### Describe, delegate, review

You write what you want in plain English, in the ChatGPT app, the Codex desktop app, the terminal, an IDE extension or the phone app. Codex clones the repository into an isolated cloud sandbox, plans the change, edits files, runs the test suite, iterates on failures, and returns a diff with an

explanation. You review and merge, or send it back.

## Parallel by default

Because each task runs in its own cloud sandbox, you can start five at once — a bug fix, a refactor, a new endpoint, two experiments — and review them as they land. This is Codex's clearest design difference from Claude Code, which began on your machine and added the cloud later.

## Code review in your workflow

Tag @Codex in a GitHub pull request, or switch on automatic reviews, and it reads the change and comments like a colleague. Slack and Linear integrations route tasks in from where teams already talk.

## Local when you want it

The Codex CLI runs on your own machine for interactive work, and the desktop app for Mac and Windows gives the same agent a window instead of a terminal.

## Hand-off from Work

A ChatGPT Work task that needs software — a dashboard, an internal tool, a script — can pass to Codex and come back with the code done.

# What it costs

There is no Codex subscription. It is included in every ChatGPT plan, with usage metered in credits that map to tokens, and a rolling five-hour window that pauses new tasks when you hit the cap.

| Plan         | Price            | Codex access                                                                                         |
|--------------|------------------|------------------------------------------------------------------------------------------------------|
| Free / Go    | \$0 / \$8        | Limited local use; no cloud tasks, prompted to upgrade at the cap                                    |
| Plus         | \$20/mo          | The real entry point: cloud tasks, code review, all surfaces. Enough for light-to-moderate daily use |
| Pro 5× / 20× | \$100 / \$200/mo | Heavy and parallel workloads. OpenAI's own estimate for power users is \$100–200/month               |
| Business     | \$20–25/user/mo  | Cloud features, GitHub, Slack and Linear integrations, admin controls                                |
| Enterprise   | Custom           | Shared credit pool, no fixed rate limits                                                             |
| API key      | Pay per token    | For building products on Codex; newer models arrive later than on subscriptions                      |

A practical note from people who run it daily: pay-as-you-go via API lands around \$5–15 a month for an hour a day, \$30–60 for a few hours, and \$80–250 for full-time agentic work. The subscription is the better deal until you are well past an hour a day.

# Codex vs Claude Code

---

The comparison everyone makes, so here it is without the fan wars.

|                  | Codex                                                      | Claude Code                                                     |
|------------------|------------------------------------------------------------|-----------------------------------------------------------------|
| Where it started | Cloud sandbox, parallel tasks                              | Your terminal, then desktop, cloud and phone                    |
| Model            | GPT-5.6 family, GPT-6                                      | Claude Fable, Opus, Sonnet                                      |
| Entry price      | \$20 ChatGPT Plus                                          | \$20 Claude Pro                                                 |
| Middle tier      | None — \$20 then \$100/\$200                               | \$100 Max                                                       |
| Strength         | Backend and multi-task parallel work; review bot in GitHub | Long careful sessions; leads most agent benchmarks through 2026 |
| Adoption         | Huge via ChatGPT's user base                               | Overtook Codex in developer preference in under a year          |
| Open source      | The CLI, yes                                               | No                                                              |

Most developers who use both describe Codex as better for firing off several backend tasks and reviewing the results, and Claude Code as better for one hard problem it has to think through carefully. At \$20 a month each, plenty of people simply run both.

## The catches

---

- **Full-access mode had a file-deletion bug.** In July 2026 Codex running without a sandbox deleted files it should not have: a failed environment-variable expansion turned a cleanup step into a delete against the user's home directory. OpenAI patched it in August. Sandboxing is still the safer default on a repo you care about.
- **Shallow visibility into its reasoning.** You see the diff and a summary, less of the why. Claude Code shows more of its working, which matters when the change is wrong.
- **The five-hour window bites.** Plus users hit the cap mid-afternoon on busy days. The workaround is a cheaper model for routine tasks, or the \$100 tier.
- **No hosting.** Codex writes and tests; it does not deploy. Vercel, Replit and the rest are still your job — or Polsia's, which uses these agents as its developer.
- **Model churn.** GPT-5.5 retires in October; the flagship tier changed twice this summer. Pin nothing.

# Do you need to be a developer?

---

Less than you did, more than for Claude Code. The ChatGPT tab and the desktop app removed the terminal, and a founder who can read a diff and run a test can get real software out of Codex. But it assumes a Git repository and a codebase to work in, and it does not deploy, so the surrounding plumbing is still on you. If you have never shipped anything, start with Claude Code's desktop app or a platform like Polsia that wires the pieces together; if you already have a repo and a ChatGPT plan, Codex is one tab away.

## Why it matters

---

Codex is the coding agent with the widest possible front door. Claude Code proved the category and still leads it on quality; Codex put the category inside the app that hundreds of millions of people already use, then made it the developer behind ChatGPT Work. For the agentic economy the consequence is simple: writing software is now a commodity capability available on a \$20 plan from two companies, and the scarce skill is knowing what to build.

## Glossary

---

**Codex** — OpenAI's autonomous coding agent (2025 onward). Cloud sandboxes, parallel tasks, GitHub review, a tab in ChatGPT. Not the 2021 model of the same name.

**Cloud sandbox** — An isolated environment where Codex clones your repository and works without touching your machine.

**Pull request (PR)** — A proposed change to a codebase that a human reviews before it is merged. Codex's standard output.

**Codex CLI** — The open-source terminal version, Apache 2.0 licensed.

**Credits** — ChatGPT's usage currency since April 2026, mapped to tokens. Codex tasks draw them down.

**Five-hour window** — The rolling period over which Codex usage on Plus and Pro is capped.

**Full-access mode** — Running Codex with unrestricted permissions on your files. Convenient; the source of its worst reported bug.

**Claude Code** — Anthropic's coding agent and Codex's direct rival. See its own guide.

## Liked this? Go deeper.

---

Codex is one of two coding agents on a \$20 plan. Here is the other, and what sits around them.

### WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is Grok Bot

Elon Musk's work agent: a team of bots with their own cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)