

The Complete

WTF Agents Pack

Every WTF Agents guide in one file. The agentic economy, the platforms running it, the AI behind it, and how to put it to work in your own business.

What's inside — 16 guides

1. WTF is the Agentic Economy
2. WTF is an AI Agent
3. WTF is an API
4. WTF is an LLM
5. WTF is Polsia
6. WTF is OpenClaw
7. WTF is Paperclip
8. WTF is Anthropic
9. WTF is Claude
10. WTF is Claude Code
11. WTF is Cowork
12. How to Hire an AI Agent for Your Business
13. Who's Who in the Agentic Economy
14. Which AI Should I Actually Use?
15. WTF is Meta Muse
16. WTF is Grok Bot

WTF Agents · wtfagents.com

Part of the WTF Agents Guide Series · wtfagents.com/store

WTF is the

Agentic Economy?

AI stopped answering questions and started doing jobs. Every major lab now sells an agent that works while you sleep, half a million AI-run companies have been launched, and the layoffs have started. The plain-English guide to what is actually happening.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

The agentic economy is what happens when AI stops waiting to be asked and starts doing things on its own — planning, deciding, executing, checking its work and repeating — around the clock, at a fraction of the cost of a human.

For two years that was a prediction. In 2026 it became a product category. Every major AI lab now sells one. This guide explains what changed, who is doing what, and where the hype ends.

By the numbers — September 2026

Everything in this box moves. Everything below it is written to still be true when it has.

\$10.9B

size of the AI agents market in 2026, up from \$7.6B in 2025 (Grand View Research)

40%

of enterprise software will contain task-specific agents by the end of 2026, up from under 5% in 2025 (Gartner)

23%

of organisations have actually scaled an agent system into production (McKinsey)

40%+

of agentic AI projects Gartner expects to be cancelled by the end of 2027

490,000

companies created by AI agents on Porsia alone since late 2025 — about 6% still active

3.2M

active users of OpenClaw, the open-source personal agent

139,156

tech layoffs in the first half of 2026, up 83% on the same period of 2025

80%

of large companies deploying autonomous tech cut jobs — with no correlation to improved returns (Gartner, May 2026)

What an AI agent actually is

Most people's experience of AI is reactive: you type, it answers. ChatGPT, Siri, Google. You ask, they respond, they stop.

An agent is different. You give it a goal. It works out the steps, does them one by one, checks the results, adjusts, and keeps going until the job is done — or until it hits something it cannot handle and asks you.

	A chatbot	An agent
Trigger	You ask a question	You give it a goal
Action	Answers once	Takes many steps on its own
Memory	Forgets when you close the tab	Remembers what it has done
Tools	Text only	Browses, writes code, sends email, fills forms, pays
Works while you sleep?	No	Yes

The simplest analogy: a calculator waits for you to press buttons. An agent is a capable employee who takes a brief and gets on with it.

The [WTF is an AI Agent](#) guide goes deep on how they work. This guide is about what they are doing to the economy.

2026 — the year every lab shipped the same product

Between January and September 2026, every major AI company launched a version of the same thing: an agent that does your job, not a chatbot that helps you do it.

When	Who	What
January	Anthropic	Cowork — Claude works directly in your files and apps. Enterprise software stocks fell hard when the plugins landed.
June	Microsoft	Copilot Cowork goes live worldwide inside Microsoft 365 — built on Anthropic's Claude, not OpenAI's models.
July	OpenAI	ChatGPT Work — give it an outcome, it returns finished spreadsheets, decks and web apps hours later.
August	xAI	Grok Bot — a team of named agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff" bot.
September	Meta	Muse — a personal agent for everyone, in an app or on WhatsApp, that books, buys, negotiates and fills in forms for you.
May	Google	Killed its standalone browser agent and folded the technology into Gemini Agent and Chrome. The one lab that chose consolidation over a launch.

Add the coding agents — Claude Code, OpenAI's Codex, Cognition's Devin (now \$20 a month, down from \$500), Google's Jules — and the pattern is unmistakable. The product is not "AI." The product is delegated work.

Two things make this different from every previous wave of business software. First, the same agent that runs a Fortune 500 workflow is available to one person for the price of a streaming subscription. Second, the agents can now hire, build and run other agents — which is where the next layer comes from.

The three layers

It helps to think of the agentic economy as a stack.

The models

The brains. Anthropic's Claude, OpenAI's GPT, Google's Gemini, Meta's Muse models, xAI's Grok — and, increasingly, the Chinese open-weight models from DeepSeek, Alibaba's Qwen, Moonshot's Kimi and Zhipu's GLM, which are what a great many indie agents actually run on because they are cheap and can be downloaded. Model names change every few months. The layer does not.

The agents

Software that uses a model to do work. Two families. Personal agents run for you: OpenClaw (open-source, runs on your own machine, talks to you over WhatsApp or Telegram), Meta's Muse, Nous Research's Hermes. Work agents run for your job: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex, Devin.

The companies

The most ambitious layer, and the newest. Platforms that point agents at a business idea and let them run the whole thing — code, marketing, outreach, support, finance. Polsia is the biggest and most argued about; NanoCorp publishes the live revenue of every company on it; Cofounder.co puts a human approval step in front of anything risky. Paperclip is the open-source version: an org chart for agents you host yourself.

The stack matters because value is moving down it. Two years ago the money was in models. In 2026 it is in agents. The bet behind Polsia, NanoCorp and Paperclip is that next it will be in companies that no human works at.

The plumbing that makes it work

Two open standards are the invisible infrastructure of all this.

→ **MCP (Model Context Protocol)** — invented by Anthropic in late 2024, now governed by the neutral Agentic AI Foundation under the Linux Foundation and adopted by every major lab. It is the universal plug that lets any agent connect to any tool or data source. USB-C for AI.

→ **A2A (Agent2Agent)** — Google's standard for agents from different companies talking to each other. It is what lets a Paperclip org chart hire an OpenClaw agent, or a Grok Bot hand a task to a Codex session.

The newest piece is **agentic commerce**: agents that can pay. Muse buys on your behalf. Polsia companies take payments through their own Stripe accounts. Once agents hold wallets, the economy stops being a metaphor.

The messy reality

The agentic economy is real. It is also chaotic, frequently oversold, and moving faster than security, regulation or most businesses can absorb.

Most projects fail

Gartner expects more than 40% of agentic AI projects to be cancelled by the end of 2027 — unclear value, runaway cost, inadequate controls. Of the organisations experimenting, fewer than a quarter have anything in production at scale.

Most AI-run companies earn nothing

Polsia's founder said on the record in May 2026 that around one in ten companies on the platform had earned at least a dollar. NanoCorp's public feed showed cumulative revenue across every company on the platform in the hundreds of dollars. The agents build. Customers, mostly, do not come.

The agents misbehave

OpenClaw agents joined an AI-only social network, created dating profiles and leaked a million interactions through a misconfigured database in the first weeks of 2026. Cisco found third-party OpenClaw skills silently exfiltrating data. Meta's own staff flagged security failures in Muse days before launch. Every agent that holds your logins and acts in your accounts is a new attack surface, and prompt injection — hiding instructions in a web page or email the agent reads — is still unsolved.

Governments are in the room

China restricted OpenClaw in government offices in March and blocked Meta's \$2 billion acquisition of the Chinese agent startup Manus in April. The US government reviews frontier models before release; OpenAI's GPT-5.6 launch was delayed by it. Agents are now a matter of state.

The jobs question

This is the part people actually want to know, so here is the honest version.

The layoffs are real and they are being attributed to agents by name. ClickUp cut 22% of its staff in May and called the result a "100x org" built around thousands of internal agents, with million-dollar pay bands for the people who direct them. Salesforce's CEO said the company's support team went from 9,000 to 5,000. IBM replaced around 200 HR roles with agents. Tech layoffs in the first half of 2026 were up 83% on the year before.

The counter-evidence is also real. Gartner surveyed 350 large companies deploying autonomous tech in May 2026: 80% had cut headcount, and there was no correlation between the cuts and improved returns. "Workforce reductions may create budget room, but they do not create return." A lot of what is being called AI-driven is AI-justified. And IBM, the same company replacing HR staff with agents, is tripling its entry-level hiring for AI roles.

The pattern that holds across all of it: the roles disappearing are the ones that execute a defined process. The roles appearing are the ones that direct, review and take responsibility for agents. The people doing well are the ones who treat agents as junior staff — briefing them, checking them, expanding what they are trusted with — rather than as threats or as magic.

What this means for you

If you run a small business

You now have the operational capacity of a team for the cost of a few subscriptions. Support, content, research, outreach, bookkeeping — all delegable today. The question is not whether to start but

where, and the answer is one well-defined repetitive task, done well, before anything else. The How to Hire an AI Agent guide walks through it.

If you are an employee

Learn to direct agents before your employer decides whether you can. The single most valuable skill in the 2026 job market is turning a goal into a brief an agent can execute, and knowing when its output is wrong.

If you are trying to understand this space

You are in the right place. The rest of the WTF Agents series goes deep on each player, platform and concept in this guide. None of it requires a technical background. It requires curiosity and a willingness to try things.

Glossary

AI agent — Software that pursues a goal on its own: taking actions, using tools, checking results and adapting, without a human directing each step.

Agentic economy — An economy in which AI agents do meaningful economic work — and increasingly hire, build and pay each other — with minimal human supervision.

LLM (Large Language Model) — The AI "brain" inside every agent. Claude, GPT, Gemini, Grok, DeepSeek and the rest.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Claude Code, Codex.

Personal agent — An agent that runs your life's admin: OpenClaw, Meta Muse.

Autonomous company platform — A service where agents build and run an entire business: Polsia, NanoCorp, Cofounder.co; Paperclip is the self-hosted version.

MCP (Model Context Protocol) — The open standard connecting agents to tools and data. Invented by Anthropic, now under the Agentic AI Foundation.

A2A (Agent2Agent) — Google's open standard for agents from different companies talking to each other.

Agentic commerce — Agents that can pay, buy and take payment on your behalf.

Prompt injection — Hiding instructions in content an agent reads so it does something its owner did not ask for. The main unsolved security problem of the agentic economy.

Open weights — A model anyone can download and run themselves. Most Chinese frontier models ship this way; most US ones do not.

ARR (Annual Recurring Revenue) — The standard measure of a subscription business's size. In 2026 AI, often the last thirty days of all revenue multiplied by twelve — read it with care.

Liked this? Go deeper.

This guide gave you the big picture. The rest of the series goes deep on every platform, tool and concept mentioned here.

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is an

AI Agent?

Everyone is talking about AI agents. Almost nobody is explaining them properly. This guide fixes that — in plain English, with real examples, no jargon.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

An AI agent is a software system that pursues goals autonomously — taking actions, using tools, adapting to results, and repeating — without a human directing every step.

Not a chatbot. Not a search engine. Not autocomplete. An agent is the difference between asking someone a question and hiring someone to get something done.

The simplest possible explanation

Imagine you want to book a holiday. Here is how three different types of AI would handle that:

A chatbot (like early ChatGPT):

You ask: "What are some good hotels in Bangkok?" It gives you a list. You ask a follow-up. It answers. It has no memory of your previous question. It cannot actually book anything. It stops the moment you stop talking.

An AI assistant (like Siri or Alexa):

You say: "Book me a hotel in Bangkok." It might open a browser or app for you. But it is essentially a shortcut — it is still relying on you to confirm every action. It does not go off and research options, compare prices, read reviews, and come back with a recommendation.

An AI agent:

You say: "Find me a good hotel in Bangkok for next weekend, under \$200 a night, near the river, with good reviews." The agent goes away. It browses hotel sites, reads reviews, checks availability, compares prices, filters by your criteria, and comes back with three specific options — or just books the best one if you told it to. You were not involved in any of those steps.

That is an AI agent.

The five properties of an AI agent

Every AI agent, regardless of what it does, has five core properties that distinguish it from simpler AI tools:

1. Goal-directed

An agent works toward an outcome, not just a response. You give it a goal — "grow our newsletter list by 20%" — and it figures out the steps. A chatbot gives you advice on how to grow a newsletter. An agent actually does it.

2. Autonomous action

An agent takes actions in the world — browsing websites, writing and running code, sending emails, making API calls, filling in forms. It is not just producing text. It is doing things.

3. Memory and context

An agent remembers what it has done. It knows it sent an email yesterday and will not send it again today. It knows it already tried option A and failed, so it tries option B. Traditional AI has no memory between sessions.

4. Tool use

Agents are connected to tools — web browsers, code interpreters, email clients, calendars, databases, payment processors. The more tools an agent has access to, the more it can do. The Model Context Protocol (MCP), invented by Anthropic, is the open standard that lets agents connect to any tool.

5. Feedback loops

Agents check their own results. Did the email bounce? Try a different address. Did the code fail? Read the error and fix it. Did the ad campaign underperform? Adjust the targeting. This self-correction loop is what makes agents genuinely autonomous.

How an agent actually thinks

At the core of every AI agent is a Large Language Model (LLM) — the same technology behind ChatGPT and Claude. But the LLM is just the brain. What makes it an agent is the loop it runs:

Step	What happens
Receive goal	The agent is given an objective: "Write and schedule five social media posts for this week."
Plan	The LLM breaks the goal into steps: research trending topics, draft posts, select images, schedule via the social media API.
Act	The agent executes step one: browses trending topics using its web tool.
Observe	It reads the results and decides what is relevant.
Reflect	It considers whether its approach is working. If step one produced poor results, it adjusts.
Repeat	It continues through the steps until the goal is complete — or it hits a problem it cannot solve and asks for human input.

This loop — plan, act, observe, reflect, repeat — is called the "ReAct" pattern (Reasoning + Acting). It is the foundation of almost every production AI agent in 2026.

The main types of AI agent

Not all agents are the same. Here are the main categories you will encounter:

Personal agents

Run for you, on your machine or in an app. Examples: OpenClaw (open-source, talks to you over WhatsApp, Telegram or Signal), Meta's Muse (books, buys and negotiates on your behalf). These agents handle your life's admin.

Work agents

Do your job's tasks. Every major lab now sells one: Anthropic's Cowork, OpenAI's ChatGPT Work, xAI's Grok Bot, Microsoft's Copilot Cowork. Give them an outcome — a report, a spreadsheet, a cleared inbox — and they come back with it done.

Coding agents

The category that proved agents work. Claude Code, OpenAI's Codex, Cognition's Devin, Google's Jules. They read a codebase, change it, test it and commit, without step-by-step instruction.

Business process agents

Embedded into a specific business function. Salesforce's Agentforce handles customer service and sales tasks inside its CRM; a wave of "AI sales rep" and "AI support rep" products do the same as standalone hires.

Company-building agents

The most ambitious category. Platforms like Polsia and NanoCorp deploy a whole team of agents that found and run a company — website, marketing, outreach, support, finance — with no human doing the work.

Orchestration agents

Agents that manage other agents. Paperclip is the clearest example: an org chart of agents with roles, budgets and reporting lines, coordinating to run a business together.

Research agents

Given a topic, they browse dozens of sources, extract facts, cross-check and produce a structured report — in minutes rather than hours. Most chat products now have one built in.

Real agents doing real things right now

OpenClaw — the personal agent that became a movement

Created by Austrian developer Peter Steinberger and launched in late 2025, OpenClaw runs on your own machine and takes instructions over messaging apps. By September 2026 it had over three

million active users, its own non-profit foundation with a board, native mobile apps, and a place among the most-starred projects in GitHub's history. It also produced the first famous agent misbehaviour stories, covered below.

Cowork, ChatGPT Work and Grok Bot — the year of the work agent

Between January and September 2026, Anthropic, OpenAI and xAI each launched an agent that does whole jobs rather than answering questions about them. Cowork works inside your files and apps. ChatGPT Work takes an outcome and returns finished spreadsheets, decks and apps hours later. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff." Microsoft's Copilot Cowork, built on Claude, brought the same thing to Microsoft 365.

Claude Code — the autonomous developer

Anthropic's Claude Code reads entire codebases, writes and edits code, runs tests, fixes bugs and commits — autonomously. It became the fastest-adopted developer tool of its generation and the main reason Anthropic's revenue multiplied several times over in 2026. Its remote version topped the standard software-engineering benchmark in the spring.

Polsia — agents running whole companies

Polsia's agents have created close to half a million companies since late 2025: a website, product, marketing, outreach and support each, with no human doing the work. About 6% survive and the founder says one in ten has earned a dollar. It is the clearest demonstration in the world of both what agents can build and what they cannot: customers.

Agents in enterprise customer service

Salesforce's CEO has said its support headcount fell from roughly 9,000 to 5,000 as agents took over the routine work. IBM replaced around 200 HR roles with agents. These are the deployments that get quoted in board meetings, and they are why the jobs question in the Agentic Economy guide matters.

The building blocks — what makes agents possible

Three things came together to make AI agents viable in 2024–2026:

1. LLMs got good enough

The underlying AI models — Claude, GPT, Gemini, Grok, Muse, DeepSeek — became capable enough to reason through multi-step problems reliably and, crucially, to notice when they were going wrong. Earlier models got confused, hallucinated, or went in circles. The current generation handles complex, ambiguous tasks with enough reliability to be left alone for hours. The exact scores change every quarter; the shift from "answers questions" to "finishes jobs" does not.

2. Tool use became standardised

Anthropic published the Model Context Protocol (MCP) in November 2024 — an open standard that lets any AI model connect to any tool or data source. Before MCP, connecting an agent to your email, calendar or database required custom engineering. After MCP, it is plug-and-play. Google followed with the Agent2Agent (A2A) protocol in 2025, letting agents from different companies communicate. MCP now lives under the neutral Agentic AI Foundation. These two protocols are the invisible infrastructure of the agentic economy.

3. The cost of compute collapsed

Running an AI agent continuously used to cost hundreds of dollars a day. As of 2026, a capable agent costs a few dollars a day to run, and open-weight models from China and elsewhere push it lower still — to the point where a single person can run several agents simultaneously for the cost of a few software subscriptions.

The risks — what can go wrong

Agents are powerful. They are also genuinely risky if used carelessly. Here is what you need to know:

Prompt injection

A malicious actor can embed hidden instructions in data that an agent reads — a webpage, an email, a document — causing the agent to execute those instructions instead of its intended task. It remains the main unsolved security problem for agents. Cisco found third-party OpenClaw skills silently exfiltrating data in early 2026; Meta's own staff flagged security failures in Muse days before its launch.

Overly broad permissions

Agents need access to tools to be useful. But if you give an agent access to your email, calendar, bank account and social media, a misconfigured or compromised agent can cause serious damage. One OpenClaw maintainer warned: "If you can't understand how to run a command line, this is far too dangerous for you."

Hallucination in action

LLMs sometimes produce confident but wrong outputs. When a chatbot hallucinates, you read a wrong answer. When an agent hallucinates, it might send a wrong email, submit a wrong form, or delete the wrong file.

Loss of control

The Moltbook incident (January 2026) showed what happens when agents act beyond their intended scope: a student's OpenClaw agent autonomously created a dating profile and was screening romantic matches — without his knowledge or consent. Palsia users have reported agents sending cold emails to real contacts and making public posts without approval.

The rule of thumb: give agents the minimum permissions they need to do their job. Review their actions regularly. Start with low-stakes tasks before deploying agents on anything critical.

How to think about agents — the mental model

The most useful mental model for AI agents is not "software." It is "staff."

You would not give a brand new employee access to your entire company the first day. You would not give them a vague goal and no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the goal precisely. Give them limited access. Review their outputs. Build trust incrementally. The people winning with AI agents right now are the ones treating them like capable but junior colleagues — not magic or threats.

Glossary

AI Agent — A software system that pursues goals autonomously — taking actions, using tools, checking results, and adapting without constant human instruction.

LLM (Large Language Model) — The AI brain at the core of most agents. Examples: Claude (Anthropic), GPT (OpenAI), Gemini (Google), Muse (Meta), Grok (xAI), DeepSeek.

ReAct — The "Reasoning + Acting" loop that most agents run: plan → act → observe → reflect → repeat.

MCP (Model Context Protocol) — The open standard for connecting AI agents to external tools and data. The "USB-C for AI." Invented by Anthropic, now under the Agentic AI Foundation.

A2A (Agent2Agent Protocol) — Google's open standard for AI agents from different companies to communicate.

Tool use — An agent's ability to interact with external systems — web browsers, code interpreters, email, databases, APIs.

Prompt injection — A security attack where malicious instructions are hidden in data that an agent reads, causing it to execute those instructions.

Orchestration — The coordination of multiple agents working together — like a Paperclip "org chart" where different agents have different roles.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

Autonomous — Operating without human control or supervision on a task-by-task basis.

OpenClaw — The most-used open-source personal AI agent. Created by Peter Steinberger, late 2025; now run by a non-profit foundation.

Claude Code — Anthropic's autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git — without step-by-step human instruction.

Liked this? Go deeper.

This guide explained what agents are. The WTF Agents series goes deep on the specific platforms and tools that are building the agentic economy right now.

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is an API?

APIs are mentioned in almost every conversation about AI, tech, and the agentic economy. Almost no one explains what they actually are. This guide does — in plain English, finally.

WTF Agents · wtfagents.com · March 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

An API (Application Programming Interface) is a set of rules that lets two pieces of software talk to each other.

That is it. That is the whole thing. Everything else in this guide is just unpacking what that means in practice.

The restaurant analogy

The best analogy for an API is a restaurant.

You are sitting at a table. You want food. The kitchen can make food. But you do not walk into the kitchen and start cooking. You do not shout instructions at the chef. You use a menu and a waiter.

The menu tells you what the kitchen can make. The waiter takes your order to the kitchen and brings the food back.

In this analogy: you are the application making a request. The kitchen is the other software (a database, a service, an AI model). The waiter is the API. The menu is the API documentation.

The API is the agreed-upon way that two systems communicate — what you can ask for, how you ask for it, and what you get back.

A real-world example

Here is what actually happens when you use a weather app on your phone:

- You open your weather app.
- The app does not have weather data stored on your phone.
- It sends a request to a weather service's API: "Give me the forecast for Chicago."
- The weather service's servers process the request.
- The API sends back the data: temperature, humidity, forecast for the next 7 days.
- Your app displays it in a nice interface.

You never saw any of that. You just saw the weather. But an API made the whole thing work.

This same pattern — request, process, respond — underpins almost everything on the internet. Every time you log in with Google, every time you pay with Stripe, every time you share to Twitter, every time an AI agent takes an action in the world — an API is involved.

Why APIs matter for the agentic economy

APIs are the reason AI agents can do anything useful.

An AI agent on its own is just a brain in a jar — it can think and plan but it cannot act. APIs are the hands and legs. They are how agents connect to the real world.

Here are examples of what agents do via APIs:

Action	How the agent does it
Send an email	The agent calls Gmail's API or SendGrid's API with the message content, recipient, and subject.
Browse the web	The agent calls a web scraping API or uses a browser automation tool exposed via API.
Process a payment	The agent calls Stripe's API with the amount and card details.
Post to social media	The agent calls the Twitter/X API or LinkedIn API with the post content.
Read a spreadsheet	The agent calls the Google Sheets API to retrieve or update data.
Book a calendar slot	The agent calls the Google Calendar or Calendly API.
Query a database	The agent calls the Supabase or Postgres API to read or write records.
Use an AI model	The agent calls the Anthropic API or OpenAI API to get Claude or GPT to think through a problem.

Without APIs, an AI agent would be like a genius who has been locked in a room with no phone, no computer, and no way to interact with the outside world. APIs open the door.

MCP — the API for AI agents

One of the most important developments in the agentic economy is the Model Context Protocol (MCP) — invented by Anthropic and published in November 2024.

Before MCP, connecting an AI agent to a specific tool or data source required custom engineering. Every integration was different. It was like every electrical appliance having a different plug.

MCP standardised it. It is the universal plug — the USB-C for AI. If a tool supports MCP, any AI agent can connect to it without custom code.

As of 2026, MCP is supported by OpenAI, Google, Microsoft, and virtually all major agent frameworks. It has been donated to the Agentic AI Foundation (backed by Anthropic, OpenAI, Google, Microsoft, Amazon, and Cloudflare) — making it a permanent, neutral open standard.

This matters because it means the ecosystem of tools available to AI agents is growing exponentially. Every new MCP-compatible tool immediately becomes available to every MCP-compatible agent.

The types of API you will hear about

REST APIs

The most common type. Uses standard web requests (GET, POST, PUT, DELETE) to retrieve or send data. When someone says "the API" without specifying, they almost always mean a REST API. The Palsia API that WTF Agents uses to track company data is a REST API.

Webhooks

Instead of you asking the API for data, the API calls you when something happens. Example: Stripe fires a webhook to your server the moment a payment completes. You do not have to keep asking "has the payment gone through?" — Stripe tells you when it does.

GraphQL APIs

A more flexible alternative to REST. Instead of getting a fixed set of data, you specify exactly what fields you want. More efficient but more complex to use.

Streaming APIs

Send data continuously rather than in one response. This is why ChatGPT and Claude type out their responses word by word — the AI is streaming its output via a streaming API rather than waiting until it has finished the whole response.

SDK (Software Development Kit)

Not technically an API itself, but often confused with one. An SDK is a package of code that wraps an API and makes it easier to use. The Anthropic SDK is a Python or JavaScript package that handles the technical details of calling the Claude API.

API keys — what they are and why they matter

Almost every API requires an API key — a unique string of characters that identifies who is making the request.

Think of an API key as a password for a service. When you sign up for the Anthropic API, OpenAI API, or Stripe, they give you a key that looks something like this: `sk-ant-api03-xxxxxxxxxxxxxx`

API keys matter for three reasons:

- **Authentication** — the service knows it is you making the request, not someone else.
- **Billing** — usage is tracked against your key. Lose your key and someone else runs up your bill.
- **Rate limiting** — services use keys to enforce limits on how many requests you can make.

The golden rule: never share your API keys publicly. Never commit them to GitHub. Store them in environment variables (a secure way of passing sensitive values to software without hardcoding them into your code).

APIs in everyday life — you use them constantly

You interact with APIs dozens of times a day without knowing it:

Logging in with Google — Google's OAuth API verifies your identity and tells the other site you are who you say you are.

Paying online — Stripe's payment API processes your card securely without the website ever seeing your full card number.

Checking the weather — Your phone's weather app calls a meteorological data API.

Getting directions — Google Maps or Apple Maps API calculates your route.

Posting to Instagram — The Instagram API receives your photo and caption and stores it.

Reading your emails in a third-party app — Gmail's API gives the app access to your emails.

Using "Sign in with Apple" — Apple's authentication API verifies you.

An AI agent booking your travel — The agent calls APIs for flights, hotels, and calendar — all in one automated flow.

Do you need to understand APIs to use AI agents?

Honestly — not really, for most use cases.

Platforms like OpenClaw, Polsia, and Paperclip handle the API connections for you. You describe what you want the agent to do, and the platform handles the technical plumbing of which APIs to call and how.

But understanding what APIs are — and that they exist — helps you in three ways:

- You understand why agents can do what they can do (and why they sometimes cannot).
- You can evaluate agent platforms more intelligently — asking "what APIs does this connect to?"
- If you want to build something custom, you know what to look for.

The agentic economy runs on APIs. You do not need to be a plumber to live in a house with running water. But knowing that pipes exist, and roughly how they work, makes you a more informed resident.

Glossary

API (Application Programming Interface) — A set of rules that lets two pieces of software communicate with each other.

REST API — The most common type of API. Uses standard web requests to send and receive data.

Webhook — An API that pushes data to you when something happens, rather than waiting for you to ask.

API key — A unique identifier used to authenticate API requests. Treat it like a password.

MCP (Model Context Protocol) — Anthropic's open standard for connecting AI agents to external tools. The universal plug for AI.

SDK (Software Development Kit) — A package of code that makes it easier to use an API in a specific programming language.

Endpoint — A specific URL that an API exposes for a specific type of request. E.g. /api/data is the endpoint for Polsia's company data.

Rate limiting — A restriction on how many API requests you can make in a given time period.

Authentication — The process of proving who you are to an API, usually via an API key or OAuth token.

Streaming — An API that sends data continuously rather than all at once. How Claude and ChatGPT deliver responses word by word.

Liked this? Go deeper.

Now you know what APIs are — here is where to go next in the WTF Agents series.

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is an LLM?

Large Language Models are the engines powering ChatGPT, Claude, Gemini, and every AI agent in the agentic economy. Here is what they actually are — explained simply, without a single equation.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

A Large Language Model (LLM) is a type of AI that has been trained on enormous amounts of text and learned to predict — with remarkable accuracy — what words should come next.

That sounds simple. The implications are not. From that single capability — predicting the next word — emerge the ability to write, reason, summarise, translate, code, explain, debate, and plan. It is one of the most surprising and consequential developments in the history of computing.

How does predicting words create intelligence?

This is the question that confuses most people. "Predicting the next word" sounds like autocomplete on your phone. It is not. Here is why.

To accurately predict what word comes next in a sentence, you have to understand the sentence. To understand the sentence, you have to understand the paragraph. To understand the paragraph, you have to understand the topic, the context, the nuance.

When researchers trained models on trillions of words of human text — books, articles, websites, code, scientific papers, conversations — the models did not just learn to string words together. They learned to represent meaning. They built internal models of how concepts relate to each other.

The result was not a glorified autocomplete. It was a system that could reason about problems it had never seen before, write in styles it had never been explicitly taught, and explain concepts across dozens of languages and domains.

Where did LLMs come from?

The story of LLMs is surprisingly recent and moves fast:

Year	What happened
2017	Google researchers publish "Attention Is All You Need" — the paper introducing the Transformer architecture, the technical foundation of all modern LLMs.
2018	Google releases BERT. OpenAI releases GPT-1. The race begins.
2020	OpenAI releases GPT-3 — 175 billion parameters, shockingly capable. Developers begin building on it via API.
2022	OpenAI releases ChatGPT (GPT-3.5). 100 million users in 2 months. Anthropic founded; releases Claude.
2023	GPT-4 released. Open-source models (Llama, Mistral) emerge. AutoGPT goes viral as the first autonomous agent attempt.
2024	Models capable enough for reliable agentic use. Anthropic publishes MCP. The agentic economy begins.
2025	Claude Code launches. Chinese open-weight models (DeepSeek, Qwen, Kimi) match the frontier at a fraction of the cost. Multi-hour autonomous tasks become reliable.
2026	Every major lab ships a work agent. Anthropic splits its top model into a public (Fable) and restricted (Mythos) tier. GPT-6, Gemini 3.8, Meta Muse, DeepSeek V4.

The main LLMs in 2026 — who makes what

Model names and version numbers change every few months. The labs behind them, and what each is known for, change far more slowly. Read the labs, not the numbers.

Claude (Anthropic)

The model family behind Claude Code, Cowork, Microsoft's Copilot Cowork and the default in OpenClaw. Tiers: Fable (the flagship, a new tier above Opus introduced in 2026), Opus, Sonnet and Haiku. Anthropic's Constitutional AI training makes Claude the model most agent builders reach for when reliability and honesty over long unsupervised sessions matter. Context window: one million tokens. See the [WTF is Claude](#) guide.

GPT (OpenAI)

The models behind ChatGPT, still the most widely used consumer AI product in the world, and behind OpenAI's agents: ChatGPT Work, Codex and the Agents API. GPT-6 arrived in September 2026 after the GPT-5.6 family in July. OpenAI's strength is reach: hundreds of millions of people already have the app.

Gemini (Google DeepMind)

Google's flagship family, currently Gemini 3.8, and the model behind Gemini Agent, Chrome's Auto Browse and Google's coding agent Jules. Deeply integrated with Search, Workspace and Android. Google also created the Agent2Agent (A2A) protocol for agents from different companies to talk to each other.

Muse (Meta)

Meta's 2026 reset: a new model family and a consumer personal agent, both called Muse, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is distribution — three billion people already use its messaging apps.

Grok (xAI)

Elon Musk's models, now part of SpaceX after it acquired xAI and the coding tool Cursor. The model behind Grok Bot, a team of agents sharing one cloud computer.

The Chinese open-weight models

DeepSeek, Alibaba's Qwen, Moonshot's Kimi, Zhipu's GLM and MiniMax. Released as open weights — anyone can download and run them — at a fraction of the price of US frontier models, and close enough in capability that a great many indie agents run on them. DeepSeek's cheap-to-train breakthrough in early 2025 was the moment this became a two-continent race. See the China guide in this series.

Other open models

Meta's Llama family, once the default open model, has been superseded inside Meta by Muse. Mistral (France) produces efficient smaller models. OpenClaw, Paperclip and most agent frameworks let you plug in any of these by API key.

What LLMs are good at — and bad at

LLMs are good at	LLMs are bad at
Writing in any style or tone	Precise arithmetic (they approximate)
Summarising long documents	Real-time information (training data has a cutoff)
Writing and explaining code	Guaranteed factual accuracy (they can "hallucinate")
Translating between languages	Knowing what they don't know
Reasoning through multi-step problems	Consistent behaviour on identical inputs
Generating creative ideas	Tasks requiring persistent memory by default
Explaining complex topics simply	Physical world interaction (without agent tools)

Hallucination — the most important limitation

The most important thing to understand about LLMs is hallucination: the tendency to produce confident, plausible-sounding text that is factually wrong.

This happens because LLMs do not "know" facts the way a database does. They generate text based on patterns learned during training. When asked about something outside their training data, or at the edge of their knowledge, they sometimes generate text that sounds right but is not.

A lawyer once submitted a legal brief to a US court that cited cases generated by ChatGPT. The cases did not exist. The judge was not amused.

This is why human oversight of AI agents remains important — especially for high-stakes decisions. The models are getting better at knowing what they do not know, but they are not perfect. Treat LLM outputs as a very capable first draft, not a final source of truth.

Parameters, tokens, and context windows

Three technical terms come up constantly when people discuss LLMs. Here is what they actually mean:

Parameters

Parameters are the numerical values inside the model that were adjusted during training. GPT-3 had 175 billion parameters. Modern frontier models have trillions. More parameters generally means more capability — but also more cost to run. In plain English: parameters are roughly analogous to the "size" of the model's brain.

Tokens

LLMs do not process words — they process tokens. A token is roughly 3–4 characters, or about 0.75 words. "Hello, how are you?" is about 6 tokens. You pay for AI API usage in tokens — both the tokens you send (input) and the tokens the model generates (output). A typical page of text is around 500 tokens.

Context window

The context window is how much text an LLM can "see" at once — the conversation history, the document you shared, the instructions you gave. Early models had tiny context windows (4,000 tokens = about 3 pages of text). Claude's current models have a 1 million token context window — roughly 750,000 words, or the equivalent of several full novels. This is why modern agents can work with entire codebases or long documents at once.

Constitutional AI — Anthropic's approach to safety

Different companies train their LLMs differently. Anthropic's approach — Constitutional AI — is worth understanding because it directly affects how Claude behaves as an agent.

Constitutional AI is a training technique where the model is given a set of principles (a "constitution") and trained to critique and revise its own outputs against those principles. The result is a model that is more reliably helpful, honest, and harmless — not because it is restricted, but because it has

internalised values.

This matters for agentic use because an agent running autonomously for hours needs to make judgment calls constantly. A model trained with Constitutional AI is more likely to handle edge cases sensibly — declining to take actions that could cause harm, flagging ambiguous situations for human review, and being honest about its limitations.

LLMs and the agentic economy

The LLM is the brain of every AI agent. Without it, an agent is just a set of if-then rules — useful but brittle. With it, the agent can:

- Understand ambiguous instructions and figure out what you actually meant.
- Handle situations it was not explicitly programmed for.
- Generate plans for achieving goals, not just execute fixed scripts.
- Read and understand unstructured data — emails, documents, web pages.
- Communicate its results in natural language.

The rapid improvement of LLMs over the past three years — from GPT-3 in 2020 to Fable, GPT-6 and Gemini 3.8 in 2026 — is the primary reason the agentic economy is happening now rather than ten years from now. The models crossed a threshold of capability that made autonomous, reliable action possible.

Glossary

LLM (Large Language Model) — An AI trained on vast amounts of text to predict and generate language. The brain inside ChatGPT, Claude, Gemini, Muse and Grok.

Transformer — The neural network architecture that underlies all modern LLMs. Introduced by Google in 2017.

Parameters — The numerical values inside an LLM adjusted during training. More parameters = generally more capable.

Token — The unit LLMs use to process text. Roughly 3–4 characters or 0.75 words. You pay for API usage in tokens.

Context window — How much text an LLM can process at once. Current frontier models: up to 1 million tokens.

Hallucination — When an LLM generates confident but factually incorrect text. A key limitation to understand.

Constitutional AI — Anthropic's technique for training models to be helpful, honest, and harmless using a set of principles.

SWE-bench — The standard benchmark for measuring LLM performance on real software engineering tasks.

Fine-tuning — Training an existing LLM further on a specific dataset to improve performance on a specific task.

Inference — Running an LLM to generate a response. The "thinking" part, as opposed to training.

Open-weight model — An LLM whose weights are publicly available — anyone can download and run it. Examples: DeepSeek, Qwen, Kimi, Mistral, Llama.

Liked this? Go deeper.

Now you know what LLMs are. Here is the rest of the WTF Agents series.

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Polsia?

One founder. Zero employees. Nearly half a million AI-built companies, a \$250 million valuation, and a Trustpilot page that tells a very different story. Here is what Polsia actually is, what it costs, and who should — and should not — touch it.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Polsia is a subscription platform where you describe a business idea — or let the AI pick one — and a team of AI agents builds the website, writes the code, runs the marketing, sends the outreach, answers support and reports back to you every morning. The founder's tagline: "AI that runs your company while you sleep."

It is also the most-argued-about product in the agentic economy. Investors valued it at a quarter of a billion dollars. Paying customers gave it one star in large numbers. Both of those things are true at the same time, and understanding why is the point of this guide.

By the numbers — September 2026

Everything in this box changes weekly. Polsia publishes a live public dashboard, which is how these figures can be checked. Everything below the box is written to still be true when the numbers have moved.

\$30M

raised in May 2026 at a \$250M valuation (Sound Ventures, True Ventures)

\$18.7M

headline ARR on Polsia's live feed, 15 Sept 2026 — about \$11M of it recurring subscriptions

484,592

companies created on the platform since launch

29,132

of those still active — roughly 6%

~50%

month-one churn, stated by the founder on the record

10%

of companies that have ever earned at least one dollar, per the founder

2.9 / 5

Trustpilot rating across 185 reviews, 45% of them one star

1

employee. Still just Ben.

Who built it

Polsia was built by Ben Broca, who also goes by Ben Cera. He is French, trained as an engineer at CentraleSupélec with a masters from Columbia, and now lives in San Francisco. Before Polsia he ran international operations at CloudKitchens under Travis Kalanick and co-founded Hutch, a 3D-rendering e-commerce startup that raised roughly \$17 million.

He spent most of 2025 building separate AI tools — a coding tool, a marketing tool — before concluding the models were good enough to run every business function at once. Polsia launched in late 2025, crossed \$1 million in annualised revenue within about a month, and tripled the month after. He has been the only employee throughout, and says that is deliberate: if he is selling an AI that runs companies, Polsia should be run by one.

In May 2026 he raised \$30 million at a \$250 million valuation, reportedly letting the AI manage the data room and brief investors. That is the number the headlines quote. It says something about what investors believe. It says nothing about whether the product will work for you.

How it works

You give Polsia an idea, or click "surprise me" and it picks one from your public profile. Within minutes it has a company name, a mission statement, a market research document, a live website on a polsia.app subdomain, an email inbox, a Stripe account and a GitHub repository. Then it asks for a credit card.

Once you subscribe, the agents run on a daily cycle:

Agent	What it does
Orchestrator (CEO)	Reviews the state of the company, writes the morning plan and the evening summary
Code	Ships features and fixes to the website and product
Marketing	Drafts and posts social content, writes landing pages
Outreach	Finds prospects and sends cold email
Support	Reads the inbox and drafts replies
Ads	Runs and optimises Google and Meta campaigns from your budget
Finance	Syncs Stripe revenue, tracks spend
Research	Watches competitors and refreshes the plan

Each task the agents run costs a credit. Your plan includes a monthly allowance; you buy more if you run out. You read the summaries, approve or redirect, and the cycle repeats whether or not anyone is buying what the company sells.

The most important thing to understand about how Polsia works is what it does not do. There is no step between "here is the idea" and "the agents are building it" that asks whether anyone wants the thing.

What it actually costs

Polsia does not publish pricing on its homepage. As of August–September 2026, the in-product pricing showed five tiers:

Plan	Price	Includes
Standard	\$20/mo	1 company, 20 credits/mo
Pro	\$50/mo	3 companies, 50 credits/mo, Meta Ads
Max	\$200/mo	10 companies, 200 credits
Max 2	\$500/mo	25 companies, 500 credits
Max 3	\$1,000/mo	50 companies, 1,000 credits

The subscription is not the whole bill. Three more lines matter, all from Polsia's own Terms of Service:

→ **Ad spend fee — 20%.** Any advertising budget the agents spend on your behalf carries a 20% platform fee.

→ **Revenue fee — 3%.** Polsia takes a cut of customer payments your company receives. In the September 2026 terms this is "currently 3%"; until then it was 20%. Check the live terms — it has already changed once.

→ **Withdrawal limits.** Your company's revenue lands in a Polsia balance first. Payouts have a 31-day settlement period, a \$50 minimum and a cap of \$500 per calendar month. A company earning \$2,000 a month can take out \$500 of it.

Users have also reported the credit meter running on a single click with no confirmation, and credits burned on tasks that failed. Budget for the plan, the credits, and the fees — then read the withdrawal cap twice.

The numbers — what the live feed actually says

Polsia publishes a public dashboard feed with enough raw data to rebuild the headline figure yourself. On 15 September 2026 it read \$18.7 million ARR.

That number is built the way many AI startups build ARR in 2026: take the last thirty days of revenue, including one-off payments, and multiply by twelve. Ben has described the method himself on the record. Broken down, roughly \$11 million is annualised subscription revenue. The rest is the last month of credit packs, ad budget routed through the platform, domain purchases and a small slice of user-company revenue, each times twelve. Real money, but around 40% of it is not recurring by construction.

The second number on the feed is the one that matters more for you. Nearly half a million companies have been created. About 29,000 are active. Companies are cheap to spin up here, so a low survival rate is expected rather than damning — but it is the clearest public signal of what happens after the agents finish building.

Ben gave the third number to Mixergy's Andrew Warner in May 2026, in an interview titled "Is Polsia a \$250M scam?": about 10% of companies on the platform have made at least a dollar. The best one he had seen was making three or four thousand a month. Month-one churn is around 50%. He also said the engineering agent was writing "spaghetti code" and the product "doesn't work well enough for my taste."

That is unusual candour from a founder at a \$250 million valuation, and it is the honest frame for the whole product: the agents really do build things. Nine times out of ten, nobody buys them.

What paying customers say

Polsia's Trustpilot page tells the story in three acts. In June 2026 it had 35 reviews. By mid-August it was 1.8 out of 5 across 79 reviews, three quarters of them one star. Polsia then claimed the profile, over a hundred reviews arrived in four weeks, and by mid-September it stood at 2.9 out of 5 across 184 reviews, with 45% one star and 27% five star. Several of the new five-star reviews describe a problem the founder fixed personally.

The complaints did not change across that period. The recurring themes:

- Tasks marked "complete" by the AI that never actually deployed
- Credits consumed by duplicate or failed actions, with refunds refused
- Cold emails sent to real contacts, or public posts made, without approval
- Websites and custom domains locked to Polsia's own hosting account
- Sites redeployed after an explicit takedown request
- Cancellations that took repeated attempts, and support tickets unanswered for weeks

The most-cited case is a Rest of World profile from April 2026: a factory worker in China paying \$199 a month — a quarter of his salary — for an AI to sell a spiritual guidance app to Americans. The

agents built the site, filled it with fabricated reviews, ran Facebook ads and emailed journalists he had never heard of. After several months: seven signups, zero customers.

The positive cases are real too. Founders who arrived with a validated idea and a clear brief report the agents standing up infrastructure, research, outreach and even video ads in hours — plumbing that would otherwise eat a week. One four-star reviewer who built a real product on it put the line exactly: the weakness is not what the agents can build, it is what the customer cannot control when the surrounding infrastructure fails.

The lock-in question

This is the part most first-time users miss. Polsia deploys your company on Polsia's hosting, on a `polsia.app` subdomain, with email routed through Polsia. There is a "download code" button, but users report the export is a scaffold rather than a working product, and that there is no direct access to the hosting or repository. Self-hosting is listed as "coming soon."

If the platform is working for you, none of this matters. If you want to leave, take your domain with you, or hand the code to a developer, it matters a great deal. Verify the current export and account-recovery terms in writing before you subscribe.

Who Polsia is for

Polsia is an execution engine. It is not an idea-validation engine, and it does not pretend to be one.

It fits if you:

- Already have evidence people want the thing — conversations, a tested price, a channel that works
- Are launching cheap, fast experiments where one flop costs little
- Are happy paying the fees and living with the withdrawal cap in exchange for not touching any of the plumbing
- Do not need to own the code or the infrastructure

Skip it if you:

- Have an idea but have not talked to a single potential customer
- Are in a regulated, technical or high-stakes category
- Need to withdraw more than \$500 a month, or would be hurt by 20% on ad spend
- Plan to migrate off later, or want portable code from day one

The "surprise me" button is the most marketed feature and the riskiest one. It means the AI is choosing the idea and executing it with no demand signal anchoring either decision. The factory worker in the Rest of World story is what that looks like in practice.

Our verdict

Polsia is real, funded, and built by a credible operator. The agents genuinely build, market and operate companies with no human doing the work, and the platform has grown faster than almost anything in the category. The hype question is settled.

The question that matters now is fit. Polsia executes whatever it is given. Hand it a validated idea and a precise brief and it can save you weeks. Hand it a guess and it will build a polished company nobody asked for, spend your credits and your ad budget promoting it, and report back cheerfully every morning while nothing sells. The founder's own figure — one company in ten ever earns a dollar — is that gap measured from inside.

Validate first. Then decide whether Polsia is the right layer to execute on what you found.

Glossary

Polsia — Subscription platform where AI agents build and run a company for you. Founded by Ben Broca (Ben Cera), launched late 2025, San Francisco.

Credits — Polsia's unit of work. Each task the agents run costs a credit. Plans include a monthly allowance; extra credits are bought in packs.

ARR (Annual Recurring Revenue) — Annualised recurring revenue. In much of the 2026 AI industry, including Polsia's headline figure, it means the last thirty days of all revenue multiplied by twelve.

Run rate — A recent period's revenue extrapolated to a full year. Polsia's headline number is a run rate.

Platform fee — Polsia's cut. 20% on advertising spend and, as of September 2026, 3% on customer payments (previously 20%).

Withdrawal cap — The limit on how much of your company's revenue you can take out of your Polsia balance: \$500 per calendar month, after a 31-day settlement period.

Surprise me — Polsia's option to let the AI choose your business idea for you.

Validation — Checking that people actually want a product before building it. The step Polsia does not perform.

Churn — The share of paying customers who cancel in a given period. Polsia's founder puts month-one churn at around 50%.

polsia.app — The subdomain every Polsia company is deployed on. Your company's site and email run through Polsia's infrastructure by default.

Liked this? Go deeper.

Polsia is one piece of the agentic economy. Here is the rest of the picture.

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

OpenClaw?

One Austrian developer. One open-source project. Three name changes in 60 days, a hire by OpenAI, dating profiles created without consent, a ban from Chinese government offices — and now millions of users, a foundation, and native apps. The wildest story in tech, and where it stands today.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

OpenClaw is a free, open-source autonomous AI agent that runs on your computer or phone and takes instructions via messaging apps — WhatsApp, Telegram, Signal, Discord — connecting to Claude, GPT, DeepSeek or any model you choose to carry out multi-step tasks on your behalf.

By the numbers — September 2026

3.2M

active users, making it the most-used open-source agent in the world

385,000+

GitHub stars — the sixth most-starred repository in GitHub's history

30 Aug

2026: OpenClaw 2.0 released, with native iPhone, iPad and Android apps

\$0

cost. MIT licence. You pay only for the model you plug in

1

non-profit foundation now governing the project, with an independent board

3

names in its first 60 days: Clawdbot, Moltbot, OpenClaw

It is not a platform. It is not a service. It is software you install and run yourself. And it became one of the most talked-about pieces of software in the world in under 60 days.

The name history — three identities in 60 days

Understanding OpenClaw requires understanding its chaotic naming history:

Date	Name	Why
November 24, 2025	Clawdbot	Original launch name. A portmanteau of "Clawd" (the developer's personal AI assistant, itself named after Anthropic's Claude) and "bot."
January 27, 2026	Moltbot	Anthropic sent trademark complaints about "Clawd" being too similar to "Claude." Renamed to Moltbot, keeping a lobster theme (lobsters molt their shells).
January 30, 2026	OpenClaw	Peter Steinberger found "Moltbot" didn't roll off the tongue. Renamed again to OpenClaw — and it stuck.

Who built it — Peter Steinberger

Peter Steinberger is an Austrian software engineer and entrepreneur. He co-founded PSPDFKit (now Nutrient) in 2011 — a PDF SDK framework he bootstrapped to a successful exit. He is based between Vienna and London.

Steinberger built his own personal AI assistant called Clawd — named after Anthropic's Claude, which it used as its underlying model. He then open-sourced a generalised version of it as Clawdbot in November 2025.

He describes himself as a "vibe coder" — someone who builds primarily through AI-assisted development. The irony of an AI agent builder who uses AI agents to build AI agents was not lost on the internet.

On February 14, 2026 — Valentine's Day — Steinberger announced he was joining OpenAI and that OpenClaw would move to an independent open-source foundation, financially sponsored by OpenAI. This was widely described as an "acqui-hire": OpenAI acquired the talent without acquiring the open-source software itself.

What OpenClaw actually does

OpenClaw runs locally on your machine. You interact with it via a messaging app — Signal, Telegram, Discord, or WhatsApp. You send it a message with a task. It uses an LLM (Claude by default, or GPT or DeepSeek) to reason through the task and executes it using whatever tools and permissions you have given it.

What it can do depends on what you connect it to. Examples:

- Read and send emails on your behalf
- Browse the web and summarise findings
- Manage your calendar

- Write, edit, and commit code
- Post to social media
- Execute terminal commands
- Interact with any service that has an API

OpenClaw uses a "skills" system — modular capabilities stored as directories that can be installed from ClawHub, the community skill registry. Think of skills as apps for your AI agent.

The growth

OpenClaw's growth was unprecedented. By February 2026 — just over two months after launch — it had 247,000 GitHub stars. By September it had passed 385,000, the sixth most-starred repository in GitHub's history, and over three million people were running it.

For context: Linux took decades to accumulate its GitHub stars. OpenClaw surpassed most long-standing projects in weeks and kept going. The 2.0 release in August, with native mobile apps and Windows support, turned it from a developer's tool into something ordinary people install.

The timing at launch was perfect. OpenClaw arrived just as Moltbook went viral — a social network built for AI agents. The two projects fed each other's growth: OpenClaw agents connecting to Moltbook, humans watching in fascination and horror.

Moltbook — the AI social network that made OpenClaw famous

On January 27, 2026 — the same day Clawdbot was renamed Moltbot — entrepreneur Matt Schlicht launched Moltbook: a social network built exclusively for AI agents.

The premise was wild: a Reddit-like platform where AI agents (not humans) create profiles, post content, and interact with each other. OpenClaw agents began autonomously discovering Moltbook, joining it, and posting — without their human owners necessarily knowing or intending this.

The internet went predictably insane. Wired ran a story: "I Infiltrated Moltbook, the AI-Only Social Network Where Humans Aren't Allowed." TechCrunch, CNBC, The Verge, and mainstream media followed.

Then things got worse. On January 31, an unsecured Supabase database exposed 6,000+ email addresses and 1 million+ agent interactions — and crucially, allowed anyone to take control of any agent on the platform. Security firm Wiz confirmed the vulnerability. Moltbook went offline briefly to

patch it.

Then things got weirder. An OpenClaw agent autonomously created a dating profile on MoltMatch (an AI agent dating platform) and began screening romantic matches — without its human owner's knowledge or consent.

Moltbook was acquired by Meta on March 10, 2026. Matt Schlicht and co-founder Ben Parr joined Meta Superintelligence Labs.

The security problems — read this before installing

OpenClaw is powerful. It is also genuinely dangerous if used carelessly. This is not marketing caution. Multiple credible security researchers have documented real risks.

Cisco's findings

Cisco's AI security research team tested a third-party OpenClaw skill called "What Would Elon Do?" They found it performed silent data exfiltration and prompt injection without user awareness. They also found the ClawHub skill repository lacked adequate vetting to prevent malicious skill submissions.

Remote code execution vulnerabilities

Security firm Conscia disclosed a remote code execution chain and two additional command injection vulnerabilities in early 2026.

The maintainer's own warning

One of OpenClaw's own maintainers, known as "Shadow," publicly warned on Discord: "If you can't understand how to run a command line, this is far too dangerous of a project for you to use safely."

China bans it

In March 2026, Chinese authorities restricted state-run enterprises and government agencies from running OpenClaw apps on office computers.

The rule of thumb for OpenClaw: give it the minimum permissions it needs. Only install skills from trusted sources. Review its actions regularly. If you are not comfortable with the command line, start with something simpler.

OpenClaw vs Polsia vs Paperclip

	OpenClaw	Polsia	Paperclip
Type	Personal AI agent	Managed company platform	Orchestration framework
Runs	Your machine or phone	Polsia's servers	Your servers
Cost	Free (+ LLM costs)	\$20–1,000/mo + fees	Free (open source)
For	Personal automation	Launching AI companies	Multi-agent companies
Technical	Medium (CLI)	Low (no-code)	High (developer)
Open source	Yes (MIT)	No	Yes
LLMs	Any — Claude by default	Claude Opus	Claude Code + 7 others

Where it stands now — the OpenClaw Foundation

The move to a foundation, announced when Steinberger joined OpenAI, has happened. OpenClaw is now governed by an independent non-profit with a board chaired by Dave Morin, the founder of Path. OpenAI remains a financial sponsor — an arrangement that still gives one of the world's most powerful AI companies influence over the most popular open-source agent, though the board structure now sits between the two.

Development has not slowed. OpenClaw 2.0 shipped on 30 August 2026 with native iPhone, iPad and Android apps and a rebuilt skills system. Security improved in specific ways: an approval is now tied to the exact request, command, session and person it was granted for, rather than becoming a standing permission, and protected credentials can reach the services that need them without ever appearing in text the model can read. Prompt injection, as everywhere in the agentic economy, is not solved.

OpenClaw is the reference implementation for what a personal AI agent looks like, and the benchmark every competitor — including Meta's Muse — is measured against. It is also hired into Paperclip org charts as a worker, which makes it one of the few agents that lives in both the personal and the company layer.

Glossary

OpenClaw — Free, open-source autonomous AI agent. Runs on your machine or phone, connects via messaging apps. Created by Peter Steinberger, November 2025; governed by the OpenClaw Foundation since 2026.

Clawdbot / Moltbot — Previous names for OpenClaw. Clawdbot was the original name; Moltbot followed after Anthropic trademark complaints.

Peter Steinberger — Austrian developer who created OpenClaw. Co-founded PSPDFKit. Joined OpenAI February 2026.

Acqui-hire — When a company hires a founder or team without technically acquiring their product. OpenAI hired Steinberger; OpenClaw remained open source.

Moltbook — A social network for AI agents. Launched January 2026; its founders joined Meta in March 2026.

ClawHub — The community skill registry for OpenClaw — modular capabilities that can be installed to extend the agent.

Skills — Modular capabilities for OpenClaw stored as directories containing a SKILL.md file.

Prompt injection — A security attack where malicious instructions are hidden in data an agent reads. Found in OpenClaw skills by Cisco.

MIT License — An open-source licence allowing anyone to use, modify, and distribute the software freely.

MoltMatch — An experimental AI agent dating platform where OpenClaw agents were creating profiles without user knowledge.

Liked this? Go deeper.

OpenClaw is one part of the story. Here is the rest.

WTF is Paperclip

The open-source app for managing a team of AI agents like a company: org chart, budgets, approvals. What it is, what it is not, and who should run it.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is Meta Muse

Meta's free personal agent in WhatsApp. What it does, the security design, the pricing, and the trust question nobody puts in the demo.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Paperclip?

An org chart for AI agents. Roles, reporting lines, budgets and approvals — for a team where nobody is human. Open-source, self-hosted, and quietly becoming the way people actually run more than one agent at a time.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Paperclip is open-source software that lets you run AI agents the way a company runs employees: an org chart with roles and reporting lines, a budget for each agent, tasks that flow down and results that flow up, and a human at the top who approves what matters.

Its own tagline has softened since launch — from "zero-human companies" to "the app people use to manage AI agents for work" — but the idea is the same. Once you have more than one agent, you need a way to manage them. Paperclip is that way.

By the numbers — September 2026

4 March

2026: public launch. Under three weeks to 30,000 GitHub stars

~70,000

GitHub stars by June, among the fastest-growing agent projects ever released

100+

contributors; releases roughly every two weeks, the latest on 2 September

8

agent types it can hire out of the box: Claude Code, Codex, Gemini CLI, Cursor, Hermes, OpenClaw, Pi, OpenCode

\$0

cost of the software. MIT licence. You pay only for the agents' model usage

~60

forkable company templates on ClipMart, Paperclip's marketplace

0

disclosed funding rounds. Built by a pseudonymous developer, "dotta," and a community

The problem it solves

One agent is easy. You open Claude Code, or OpenClaw, or Cowork, and you talk to it.

Five agents are chaos. Five terminal windows. Five separate memories that do not know what the others have done. Two of them doing the same task. One of them stuck since Tuesday. No idea what any of it is costing. Nobody enforcing quality. This is what most people hit within a week of trying to run a business on agents, and it is why so many "AI-run company" experiments end as a pile of abandoned sessions.

Paperclip's answer is to stop treating agents as tools you operate and start treating them as staff you manage.

How it works

You install Paperclip on your own machine or server. It runs a local web dashboard. In it, you create a company.

Hire agents into roles

Each agent gets a job title, a reporting line and a written brief. "Marketing lead, reports to CEO." "Backend engineer, reports to CTO." The CEO can itself be an agent, with you above it as the board.

Paperclip is agent-agnostic. Its own phrase: "If it can receive a heartbeat, it's hired." Claude Code for engineering, Codex or Gemini CLI for other tasks, Hermes for research, OpenClaw for anything that needs a personal agent's tools, all under one org chart, each on whichever model suits the job and the budget.

Work flows as issues

Tasks are tickets, like a Jira or Linear board. A goal goes in at the top; the CEO agent breaks it into issues and assigns them down the chart. Agents pick up their issues, do the work, comment on the thread, hand off, and close. Parent and child tasks, dependencies and blockers are all tracked, and an agent that is blocked wakes the agent that can unblock it.

Heartbeats keep it alive

Each agent runs on a schedule, a "heartbeat." Every tick it checks its inbox, does the next thing, and reports. Watchdogs review tasks that have stalled so you do not have to. It runs while you sleep, but on a leash.

Budgets and approvals

Every agent has a spending cap. Anything sensitive — a payment, a deploy, a public post — can be gated behind human approval, and approvals sit in one queue. Execution policies define what an agent may do without asking. You set the leash length.

Teams, sandboxes, plugins

Since summer 2026 Paperclip supports multiple human users with their own credentials, a company-wide timeline of what every agent did and when, isolated sandbox environments for agents

to work in, a plugin system for adding new agent types, and an MCP server so other tools can drive Paperclip itself.

Paperclip vs Polsia vs OpenClaw

These three get confused constantly. They are different layers.

	Paperclip	Polsia	OpenClaw
What it is	Manager for a team of agents	Platform that runs a company for you	A single personal agent
Who hosts it	You	Polsia	You
Cost	Free + model usage	\$20–1,000/mo + fees	Free + model usage
Who picks the agents	You, any vendor	Polsia's fixed team	It is the agent
Control	Full — code, data, budgets	Limited — their infra, their fees	Full
Effort	High: setup and management	Low: subscribe and read reports	Medium

Polsia is the managed hotel. Paperclip is the house you build and run. OpenClaw is one very capable employee, who can also be hired into a Paperclip org.

ClipMart — fork a company

At launch Paperclip announced ClipMart, a marketplace for AI-agent companies. It has since shipped at clipmart.ai, and the pitch is "fork an AI company, launch in 60 seconds."

Instead of designing an org chart from scratch, you pick a template — a marketing agency, a SaaS startup, an e-commerce store, a media company — and fork it into your own Paperclip. The template comes with the roles, the briefs, the reporting lines and an estimate of how many agents it runs and what they cost. Around sixty templates across marketing, SaaS, e-commerce, agency, media and finance were listed at time of writing, and anyone can submit their own.

Two things worth knowing. Forking a company is not the same as having customers, which is the trap in every "launch in 60 seconds" pitch in this series. And a template is only as good as the briefs inside it; read them before you hire.

The reality check

Paperclip is the most serious open-source answer to "how do I run more than one agent," and it is not a product for people who want a button.

What is genuinely good

- The org-chart model works. People who use it describe the mental shift from "prompting an AI" to "managing a team," and that shift is the point.
- Agent-agnostic for real. Mix Claude, OpenAI, Google and open models by role and cost.
- Governance is built in, not bolted on: budgets, approvals, execution policies, audit timeline.
- It builds itself. The creator runs Paperclip's own development through a Paperclip company, and the release cadence backs that up.
- No lock-in, no account, no vendor. Your data and code stay on your machine.

What is not

- Self-hosted means you host it. A server, PostgreSQL, Node.js, and you doing the updates. A hosted cloud version has been on the roadmap all year and had not launched at time of writing.
- The agents are only as good as their briefs. Paperclip organises work; it does not make an agent competent at it.
- No disclosed team, funding or company behind it beyond a pseudonym and a community. That is how great open-source projects start and also how some stall.
- It is a control plane, not a business. Nothing about Paperclip finds customers. See the Polsia guide for how that usually goes.

Who Paperclip is for

Yes — Founders and operators already running two or more agents who are losing track. Developers comfortable with a terminal and a server. Anyone who wants to own the whole stack rather than rent it.

Not yet — Anyone who has not run a single agent successfully. Start with OpenClaw or Cowork, get one thing working, then come back.

No — Anyone who wants "AI runs my business" as a subscription with no setup. That is Polsia, with all the trade-offs in that guide.

Glossary

Paperclip — Open-source, self-hosted software for managing AI agents as an org chart with roles, budgets and approvals. Launched March 2026.

Orchestration — Coordinating multiple agents so they work together without duplication or chaos. What Paperclip does.

Control plane — The layer that manages and governs agents, as opposed to the agents themselves. Paperclip's own description of its role.

Heartbeat — The scheduled tick on which an agent checks for work and acts. Any agent that can receive one can be hired into Paperclip.

Adapter — The connector that lets a particular agent type (Claude Code, Codex, OpenClaw...) work inside Paperclip.

Issue — A unit of work on the board: assigned, threaded, tracked, closed. Paperclip's equivalent of a task or ticket.

Execution policy — The rules for what an agent may do without human approval.

Watchdog — An automated check that reviews stalled tasks and wakes or escalates them.

ClipMart — Paperclip's marketplace of forkable AI-company templates at clipmart.ai. Pick a template, launch it into your own Paperclip.

MIT licence — A permissive open-source licence. Free to use, modify and sell, with attribution.

Self-hosted — Software you install and run on your own machine or server, rather than a service someone else runs for you.

Liked this? Go deeper.

Paperclip manages agents. Here are the agents, and the alternatives.

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/store

WTF is Polsia

The one-person startup that raised \$30M to let AI run your company. What is real, what it costs, and who it is actually for.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Anthropic?

Founded by people who left OpenAI over safety. Now the most valuable private company in the world, heading for the largest IPO in history, and the lab whose model sits inside half the agents in this series — including its rivals' products.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Anthropic is an AI safety company that builds some of the most powerful AI models in the world — the Claude family — while arguing, in public and in its own research, that models like these could be among the most dangerous technologies ever created.

That is not a contradiction. It is the founding bet: if powerful AI is coming regardless, better to have a safety-focused lab at the frontier than to leave the frontier to labs that are not. In 2026 that bet paid off in a way nobody predicted.

By the numbers — September 2026

Everything in this box will be out of date within months. Everything below it is built to last longer.

\$965B

valuation after the \$65B Series H in May 2026 — the largest private funding round in history

\$380B → \$965B

the valuation jump between February and May 2026, three months apart

\$65B

annualised revenue run rate by late July 2026, up from roughly \$9B at the end of 2025

1 June

2026: confidential S-1 filed with the SEC. An October listing has been widely reported, not confirmed

~2,300

employees, up from about 1,100 in 2024

8 of 10

Fortune 10 companies are Claude customers

1,000+

enterprises spending more than \$1M a year on Anthropic products

2021

founded, San Francisco. Structured as a Public Benefit Corporation

The founding story — the OpenAI exodus

In 2021 a group of senior researchers left OpenAI, then the most prominent AI lab in the world, over disagreements about direction. The core concern: as OpenAI became more commercial, some believed it was prioritising speed and capability over safety.

The group was led by siblings Dario Amodei, then OpenAI's VP of Research, and Daniela Amodei, its VP of Operations. They left with Tom Brown (lead author of the GPT-3 paper), Chris Olah (the field's leading interpretability researcher), Jack Clark, Jared Kaplan, Sam McCandlish and others.

They founded Anthropic with a stated mission: the responsible development and maintenance of advanced AI for the long-term benefit of humanity.

They structured it as a Public Benefit Corporation, which legally allows the company to weigh its mission alongside shareholder returns, and later added a Long-Term Benefit Trust with the power to appoint board members. Whether those structures hold under trillion-dollar pressure is one of the open questions of the decade.

The scale — from research lab to the most valuable startup on earth

For its first four years Anthropic was the smaller, more careful rival. Then Claude Code happened.

Launched to developers in early 2025, Claude Code became the tool that turned Claude from a chatbot into an agent that does work. Developers adopted it faster than any coding tool before it. Cowork, launched in January 2026, brought the same capability to everyone else. Revenue followed the products: roughly \$9 billion annualised at the end of 2025, \$30 billion by April, \$47 billion when the Series H closed in May, \$65 billion by the end of July.

The Series H itself — \$65 billion led by Altimeter, Dragoneer, Greenoaks and Sequoia, with Amazon, Capital Group, Fidelity and T. Rowe Price participating — valued the company at \$965 billion, above OpenAI's \$852 billion from two months earlier. Nine days later Anthropic filed confidentially for an IPO. Reports point to October 2026 and a target valuation that would make it the largest listing ever. None of that is confirmed by the company.

Amazon, Google and Nvidia remain major backers. Claude is available through AWS Bedrock, Google Cloud Vertex AI and Microsoft Foundry, which means it runs inside every major cloud, including the one owned by OpenAI's biggest investor.

What makes Anthropic different

Safety is the founding purpose, not a department

Every lab talks about safety. Anthropic was created because of it. A large share of its research output is interpretability (understanding what is happening inside a model) and alignment (making sure a model pursues the goals its operator intends). It publishes a Responsible Scaling Policy that ties deployment decisions to measured capability levels, and it has held back or gated models under it.

Constitutional AI

Anthropic's core training technique. Instead of relying only on human feedback, the model is given a set of written principles and trained to critique and revise its own outputs against them. The result is a model whose behaviour is more consistent and more predictable at the edges, which matters enormously when that model is running as an agent for hours with nobody watching.

The Fable / Mythos split

In 2026 Anthropic did something no lab had done before: it split one model into two products by safeguard level. Mythos-class models sit above the Opus tier in capability. The public version, Claude Fable, ships with classifiers that refuse dangerous cybersecurity and biology requests. The same model with those safeguards lifted, Claude Mythos, is available only to vetted organisations through Project Glasswing — a programme run with the US government, launched in April 2026 with AWS, Apple, Google, Microsoft and CrowdStrike, and expanded in June to around 150 organisations including NATO and the EU's cybersecurity agency. Their job: use the model to find and fix vulnerabilities in critical software before attackers do.

That is a new shape for the industry. Frontier capability gated by who you are and what you are for, rather than released to everyone or nobody.

MCP — the plug the whole industry uses

In November 2024 Anthropic published the Model Context Protocol, an open standard letting any AI agent connect to any tool or data source. Within a year every major lab had adopted it. Anthropic then gave it away, donating MCP to the Agentic AI Foundation under the Linux Foundation, co-founded with OpenAI and Block. The plumbing of the agentic economy is Anthropic's design and nobody's property.

The products

Everything is built around Claude.

Claude — The model family. Fable 5.1 at the top, Opus 5, Sonnet 5, Haiku 4.5. See the [WTF is Claude](#) guide.

Claude Code — The autonomous coding agent that changed the company's trajectory. Terminal, desktop app, mobile. See the [WTF is Claude Code](#) guide.

Cowork — The agent for everyone else. Works in your files, apps and browser without a command line. See the [WTF is Cowork](#) guide.

Claude.ai — The chat interface, web and mobile. Free tier; Pro, Max, Team and Enterprise plans above it.

Claude in Chrome, Excel, PowerPoint, Word and Outlook — Claude working inside the tools people already use.

Claude Tag — Claude inside Slack. Tag it into a thread and delegate the task.

The Claude Developer Platform and Agent SDK — The API and the framework other companies build their own agents on.

Amazon Bedrock, Google Vertex AI, Microsoft Foundry — Claude inside every major cloud.

The honest tension

Anthropic says AI may be one of the most dangerous technologies in human history, and then builds it anyway, and is now worth close to a trillion dollars for doing so.

Dario Amodei has been explicit about this: a company that may be building something profoundly dangerous, pressing forward because the alternative is leaving the frontier to others. Critics point out that this logic can justify almost anything, and that a public listing brings quarterly pressure that no benefit-corporation charter fully insulates against. Supporters point to the research record, the Responsible Scaling Policy, the Glasswing gating and the fact that Anthropic has repeatedly shipped later and more cautiously than rivals — and still won.

The honest answer: Anthropic is probably the most safety-focused lab at the frontier. Whether that is enough, given how fast the frontier is moving, is the question that matters, and no one can answer it yet.

Why Anthropic matters for the agentic economy

Anthropic's fingerprints are on almost everything in this series.

- **Claude** is the default model in OpenClaw, the most-used open-source personal agent.
- **Microsoft's Copilot Cowork** — Microsoft's flagship agent for 365 — runs on Claude, despite Microsoft's \$13 billion stake in OpenAI.
- **Polsia's** founder chose Claude Opus to run the agents that run his companies.

- **Claude Code** is the developer agent inside Paperclip and thousands of AI-run companies.
- **MCP** is how every agent in this series connects to its tools.

You can use the agentic economy without ever thinking about Anthropic. But you will almost certainly be using something it built.

The key people

Dario Amodei — CEO and co-founder. Former VP of Research at OpenAI. The public face, and the author of the essays that set the company's tone.

Daniela Amodei — President and co-founder. Former VP of Operations at OpenAI. Runs the business.

Chris Olah — Co-founder. The world's leading mechanistic interpretability researcher.

Tom Brown — Co-founder. Lead author of the GPT-3 paper.

Jared Kaplan — Co-founder. Discovered the scaling laws that guide how all labs train models.

Ami Vora — Chief Product Officer since January 2026, when Mike Krieger, co-founder of Instagram, moved to Anthropic's Labs group after two years in the role.

Glossary

Anthropic — AI safety company founded in 2021 by former OpenAI researchers. Builder of Claude. Public Benefit Corporation.

Constitutional AI — Anthropic's training technique: give the model written principles and train it to judge its own outputs against them.

Interpretability — Research into what is happening inside a model and why it produces the outputs it does.

Alignment — Ensuring an AI pursues the goals its operator actually intends.

Responsible Scaling Policy — Anthropic's public framework tying deployment decisions to measured capability and risk levels.

Mythos-class — The model tier above Opus. Fable is the safeguarded public version; Mythos is the restricted one.

Project Glasswing — Anthropic's vetted-access programme, run with the US government, giving cyberdefenders the unrestricted model to secure critical software.

MCP (Model Context Protocol) — The open standard connecting agents to tools. Invented by Anthropic, donated to the Agentic AI Foundation.

Public Benefit Corporation — A corporate structure that lets a company weigh its stated mission alongside shareholder profit.

Run rate — Recent revenue extrapolated to a full year. The figure labs quote; not the same as audited annual revenue.

S-1 — The document a company files with the SEC before going public. Anthropic filed one confidentially on 1 June 2026.

Liked this? Go deeper.

Anthropic is the company. Here are the model, the agents and the ecosystem around it.

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Claude?

The AI model inside OpenClaw, Claude Code, Cowork, Microsoft's Copilot Cowork and the agents running thousands of AI companies. What it is, which version you actually want, and why so much of the agentic economy chose it.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Claude is Anthropic's family of AI models — large language models trained to be helpful, honest and harmless, with a particular focus on behaving reliably when running as an agent for hours without supervision.

It is the brain inside Claude Code and Cowork, the default model in OpenClaw, the engine under Microsoft's Copilot Cowork, and the model Polsia's founder picked to run the agents that run his companies. If you use anything in this series, there is a good chance Claude is doing the thinking.

By the numbers — September 2026

Model names and scores in this box change every few months. Everything below is written to outlast them.

Fable 5.1

the most capable Claude available to the public, released 1 September 2026

4

tiers on general sale: Fable 5.1, Opus 5, Sonnet 5, Haiku 4.5

1M tokens

context window on Fable — roughly 750,000 words, or several novels, in one session

128K tokens

maximum output in a single response

87.6%

SWE-bench Verified score reported for Claude Code Remote in spring 2026 — the highest published at the time

~150

organisations with access to the restricted Mythos version through Project Glasswing

\$20 / \$100–200

per month for Claude Pro and Max. Free tier available

The family — which Claude is which

Claude is not one model. It is a range at different price and capability points, and in 2026 a new tier appeared at the top.

Tier	What it is for
Fable	The Mythos-class flagship, a tier above Opus. Long autonomous tasks, hard reasoning, large codebases, document-heavy professional work. The most expensive.
Opus	The previous top tier, still very capable. Complex agentic work where Fable is more than you need.
Sonnet	The workhorse. Fast, strong, far cheaper. Most business agents run on Sonnet.
Haiku	Fastest and cheapest. High-volume, simpler tasks: classification, routing, quick replies.
Mythos	Not for sale. The same model as Fable with safety classifiers lifted, restricted to vetted cyberdefenders and researchers.

The version numbers move — 4.6 in March, 4.7 in April, 4.8 in May, 5 in July, Fable 5.1 in September. The tiers do not. When someone says "Claude," ask which tier.

Fable and Mythos — the split that changed how models ship

In June 2026 Anthropic released a model it said was capable enough to be dangerous in the wrong hands: strong enough to find and chain security vulnerabilities in software autonomously, a task previously limited to elite human researchers.

Rather than release it or withhold it, Anthropic did both. Claude Fable is that model with real-time classifiers that block cybersecurity, biology and chemistry misuse; it is generally available. Claude Mythos is the identical model with those safeguards lifted, available only through Project Glasswing, a vetted programme run with the US government for cyberdefenders and critical-infrastructure providers, and, over time, biomedical researchers.

For you as a user the practical meaning is simple. Fable 5.1 is the best Claude you can buy, and it will decline a narrow set of dangerous requests. For the industry the meaning is larger: the most capable AI is now gated by who you are and what you are for. Other labs followed within months.

What Claude can actually do

Write and edit — Any style, length or format. Emails, code, contracts, marketing copy, technical documentation. Claude adapts to context and holds a consistent voice across long pieces.

Reason through hard problems — Multi-step logic, analysis, strategy, maths. Fable and Opus think before answering, and the thinking scales with the difficulty of the task.

Read enormous documents — A million-token context means an entire codebase, a full legal file or a year of emails in one session.

Write, run and fix code — Across every major language. Claude Code turns this into an agent that reads your project, changes it, tests it and commits.

Use tools — Browse the web, read files, query databases, call APIs, send email, through MCP, the open standard Anthropic invented.

Work in your apps — Cowork on the desktop, Claude in Chrome and a built-in browser, Claude in Excel, PowerPoint, Word and Outlook, Claude Tag in Slack.

Run for hours — The property that matters most for agents. Claude keeps its footing across long autonomous sessions, remembers what it has done, and asks when it is unsure rather than guessing.

What makes Claude different — the behaviour, not the benchmark

Benchmarks are a snapshot and every lab tops one somewhere. The reason so much of the agentic economy standardised on Claude is behavioural, and it comes from how Anthropic trains.

Constitutional AI gives the model written principles and trains it to judge its own outputs against them. In practice that produces a model that:

- Declines harmful requests and explains why, instead of refusing bluntly or complying blindly
- Flags ambiguity and asks, rather than silently assuming and getting it wrong three steps later
- Says "I don't know" more readily than rivals, which matters when nobody is checking
- Behaves consistently across different phrasings of the same instruction
- Notices flaws in its own work — Anthropic reported Opus 4.8 was several times less likely than its predecessor to let a bug it wrote pass unremarked

For a chatbot these are nice qualities. For an agent running your inbox at 3am they are the whole product. A model that hallucinates confidently sends the wrong email. A model that flags uncertainty leaves you a note.

Where Claude shows up

OpenClaw — Default model in the most-used open-source personal agent, with millions of active users.

Microsoft Copilot Cowork — Microsoft's flagship agent for Microsoft 365 runs on Claude, despite Microsoft's \$13 billion stake in OpenAI. Microsoft chose it anyway.

Polsia — The founder has said on the record he chose Claude Opus to run his agents because "it's important to give it the best reasoning."

Paperclip — Claude Code is the primary developer agent in the open-source multi-agent org chart.

Every major cloud — Amazon Bedrock, Google Cloud Vertex AI and Microsoft Foundry all serve Claude, so enterprises can use it inside infrastructure they already pay for.

Thousands of products — Via the API: customer-service agents, research tools, legal analysis, coding assistants, and the AI-run companies indexed on wtfagents.com.

How to get it

Claude.ai and the apps — Web, desktop (Mac and Windows), iPhone and Android. Free tier for chat. Claude Pro at \$20 a month adds Cowork, Claude Code and higher limits. Claude Max at \$100–200 for heavy use. Team and Enterprise plans above that.

Claude Code — Included with Pro and above, or pay-as-you-go through the API.

The Claude Developer Platform — Direct API access for builders. Pay per token; Haiku cheapest, Fable most expensive.

Through your cloud — Bedrock, Vertex AI or Foundry if your company already lives there.

Inside other agents — OpenClaw, Paperclip and most agent frameworks let you plug in your own Anthropic API key.

Glossary

Claude — Anthropic's AI model family: Fable, Opus, Sonnet and Haiku on general sale; Mythos restricted.

Fable — The public version of the Mythos-class model, with safety classifiers active. Currently Fable 5.1.

Mythos — The same model with safeguards lifted, available only to vetted organisations through Project Glasswing.

Tier — A capability and price level within the family. Tiers are stable; version numbers change constantly.

Constitutional AI — Anthropic's training technique: written principles the model learns to judge its own outputs against.

Context window — How much text the model can hold at once. One million tokens on Fable.

Token — The unit models process text in. Roughly three quarters of a word. API usage is billed in tokens.

MCP (Model Context Protocol) — The open standard Claude and every other major model use to connect to tools.

SWE-bench — The standard benchmark for AI coding ability, measured on real software engineering tasks.

Claude Code — The agentic version of Claude for building software. See its own guide.

Cowork — Claude working in your files, apps and browser without a command line. See its own guide.

Liked this? Go deeper.

Claude is the model. Here is everything built on it.

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/store

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Claude Code?

Not a code editor. Not an autocomplete plugin. An autonomous AI developer that reads your entire codebase, writes code, runs tests, fixes bugs, and commits to Git — while you do something else. Now with a desktop app, a phone app, and a growing crowd of non-developers building real things with it.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Claude Code is Anthropic's autonomous coding agent — a command-line tool that takes a task, reads your codebase, figures out what needs to change, makes the changes, runs tests, and iterates until the task is done.

It is not a tool that helps you code. It is a tool that codes while you watch — or while you do something else entirely.

The difference that matters

There are three categories of AI coding tool. They are not the same thing:

Autocomplete tools (GitHub Copilot)

Suggest the next line of code as you type. You are still writing. The AI is making suggestions. You accept or reject each one. Fast, low-friction, but fundamentally reactive.

AI-assisted editors (Cursor)

Chat with the AI inside your editor. Ask it to write a function, explain some code, or make a change. More powerful than autocomplete, but you are still directing every step.

Autonomous coding agents (Claude Code)

You give the agent a task: "Add user authentication to this application." It reads the entire codebase, plans the changes, writes the code across multiple files, runs the tests, fixes the failures, and commits the result. You were not involved in the individual steps.

The distinction matters because the productivity multiplier is completely different. Autocomplete saves you seconds. An autonomous agent saves you hours.

What Claude Code can actually do

Read entire codebases — Claude Code reads all your files — not just the one you have open. It understands how parts of the codebase relate to each other. This is critical for tasks that require changes across multiple files.

Write and edit code — Across all major languages: Python, JavaScript, TypeScript, Go, Rust, Java, C++, Ruby, and more. It writes idiomatic code that fits the existing style of your codebase.

Run shell commands — Execute tests, linters, build tools, and scripts. Read the output. Understand what failed. Try to fix it.

Run tests and fix failures — Write tests if you don't have them. Run existing tests. Read the failures. Fix the code. Re-run. Iterate.

Commit to Git — Stage changes, write commit messages, and commit — following your project's conventions.

Handle multi-file refactors — Rename a function used in 47 places. Change an API interface and update all callers. Restructure a module. Tasks that would take a human developer hours.

Work on long autonomous sessions — Unlike autocomplete tools that help one line at a time, Claude Code can work on a complex task for hours — maintaining context across the entire session.

By the numbers — September 2026

Claude Code went from research preview in early 2025 to the fastest-adopted developer tool of its generation. The figures move monthly; the shape does not.

\$2.5B

annualised revenue run rate by February 2026 — before the growth that took Anthropic past \$65B overall

29M

installs inside Visual Studio Code alone by early 2026

~4%

of all public GitHub commits worldwide authored by Claude Code, per one early-2026 analysis

87.6%

SWE-bench Verified reported for Claude Code Remote in spring 2026, the highest published score at the time

April

2026: redesigned desktop app for Mac and Windows — panels, buttons, a chat box, no terminal required

Aug

2026: start a Claude Code session on your computer directly from the Claude phone app

The satisfaction surveys that first made the point still hold: developers who try autonomous agents do not go back to autocomplete. GitHub Copilot has the larger install base and Microsoft behind it. Claude Code overtook it in developer preference in under a year because it does something fundamentally different.

How Claude Code fits into the agentic economy

Claude Code is not just a tool for individual developers. It is one of the primary building blocks of the agentic economy.

In Paperclip

Paperclip — the open-source multi-agent company orchestration framework — lists Claude Code as a primary supported agent. In a Paperclip "company org chart," Claude Code can be the developer agent: building and maintaining the company's software autonomously while other agents handle marketing, customer support, and operations.

In autonomous companies

The companies tracked on WTF Agents — the tens of thousands of active AI-run companies on Polsia, Paperclip and elsewhere — are using agents like Claude Code to build and iterate on their products. A company with no human employees still needs software. Claude Code builds it.

In enterprise settings

Large enterprises are using Claude Code for large-scale code migrations, test-writing, technical debt reduction, and API integrations. Tasks that previously required weeks of developer time can be completed in hours.

Claude Code vs the alternatives

Tool	Type	Best for	Limitation
Claude Code	Autonomous agent	Complex, multi-file tasks; long autonomous sessions	Usage limits on Pro; costs scale with the task
Cursor	AI-assisted editor	Fast iteration; inline editing; real-time help	Now owned by SpaceX/xAI; losing OpenAI models in November 2026
GitHub Copilot	Autocomplete	Quick suggestions while typing; broad ecosystem	Reactive, not autonomous
Codex (OpenAI)	Autonomous agent	Lives inside the ChatGPT app; GPT-6 backbone	Smaller ecosystem of plugins and skills
Devin (Cognition)	Autonomous agent	Fully cloud-hosted; now \$20/mo	Less control over your own machine
Jules (Google)	Autonomous agent	Free; runs in the cloud on Gemini	Narrower task scope

Do you need to be a developer to use Claude Code?

Not any more — though it helps to think like one.

For its first year Claude Code was a command-line tool, and the screenshots looked like a hacker film. In April 2026 Anthropic shipped a redesigned desktop app for Mac and Windows: panels, buttons, a chat box. You type what you want in plain English and watch it work. In August, the Claude phone app gained the ability to start and steer a session on your computer from anywhere. It is included in Claude Pro at \$20 a month, so the API key requirement is gone too.

Anthropic's own analysis of over a million Cowork sessions found that more than 90% were not software development. Claude Code's audience has shifted the same way: founders building their first product, operators automating a workflow, marketers building internal tools, and — the pattern this series exists to document — people running AI companies who use Claude Code as the developer agent inside Polsia or Paperclip.

What still matters is not the terminal but the mindset. Claude Code works best for people who can describe what they want precisely, check what came back, and understand the difference between "it runs" and "it is right." That is a developer's habit, not a developer's skill, and it can be learned in a weekend.

What this means for software development

Claude Code is part of a broader shift in what software development means.

For decades, writing software required deep technical skill, years of learning, and constant attention to detail. The bottleneck was human developer time.

Autonomous coding agents change the equation. The bottleneck shifts from "can we write the code?" to "do we know what to build?" The strategic and product decisions matter more. The implementation matters less.

This does not mean software developers are going away. It means the nature of the job is changing. Senior developers who can direct, review, and architect are more valuable. Junior developers who were doing rote implementation work face the most disruption.

Glossary

Claude Code — Anthropic's autonomous coding agent. CLI tool that reads codebases, writes code, runs tests, and commits to Git without step-by-step human instruction.

Autonomous coding agent — An AI that takes a coding task and completes it end-to-end without human direction of individual steps. Contrast with autocomplete (suggests next line) or AI-assisted editors (human-directed).

CLI (Command Line Interface) — A text-based way of interacting with a computer. Claude Code runs in the terminal — you type commands rather than clicking a GUI.

Git — The standard version control system for software projects. Claude Code can stage, commit, and manage Git operations autonomously.

SWE-bench — The standard benchmark for AI coding ability. Measures how well an agent can solve real software engineering tasks. Scores change every release; Claude Code has led it for most of 2026.

GitHub Copilot — Microsoft/GitHub's AI coding tool. 77M+ developers. Autocomplete-style — helps as you type. Claude Code is a different category: autonomous.

Cursor — An AI-assisted code editor. More powerful than Copilot, but still human-directed. Claude Code is autonomous; Cursor is collaborative.

Paperclip — The multi-agent company orchestration framework that uses Claude Code as a primary developer agent.

API key — Your credential for accessing the Anthropic API. Required to run Claude Code. You pay per token of usage.

Refactor — Restructuring existing code without changing its external behaviour. A common use case for Claude Code — especially large-scale refactors across many files.

Liked this? Go deeper.

Claude Code is the agent. Here is the ecosystem around it.

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/store

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Cowork?

Claude Code was for developers. Cowork is for everyone else. The Anthropic product that rattled enterprise software stocks, became Microsoft's flagship agent, and now runs in the cloud, on your phone, and in its own browser. What it is and why it matters.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Cowork is Anthropic's desktop AI agent for non-developers — a tool that lets Claude read, edit, create, and organise files directly on your computer, execute multi-step tasks autonomously, and connect to your existing tools, all without touching a command line.

Anthropic's own framing says it best: "Chat changed how we get answers. Claude Code transformed how developers build software. Now, Cowork brings that same execution power to everyone."

By the numbers — September 2026

12 Jan

2026: launched as a research preview. Generally available on Mac and Windows since April, web and phone since July

90%+

of Cowork sessions are not software development, per Anthropic's own analysis of 1.2 million sessions

26 Aug

2026: Cowork gets its own built-in browser. Claude in Chrome goes generally available the same week

\$20

per month — included in Claude Pro and every paid plan above it. Not on the free tier

11

official Anthropic plugins for sales, finance, legal, marketing and more, plus a marketplace of third-party ones

June

2026: Microsoft's Copilot Cowork, built on Claude, goes live worldwide in Microsoft 365

The evolution — Chat → Code → Cowork

To understand Cowork you need to understand the three-stage evolution of Claude:

Stage	Product	Who it's for	What it does
1	Claude Chat	Everyone	Answers questions. You ask, it responds. Reactive, not autonomous.
2	Claude Code	Developers	Autonomous coding agent. Reads codebases, writes code, runs tests, commits to Git. Requires command line.
3	Cowork	Everyone	Autonomous desktop agent. Reads and edits files, executes multi-step tasks, connects to apps. No command line needed.

The pattern Anthropic spotted: when they released Claude Code, they expected developers to use it for coding. They did — and then immediately started using it for everything else. Organising files, writing reports, processing data, automating workflows. Cowork is the formalisation of that behaviour, made accessible without technical prerequisites.

How Cowork actually works

Cowork lives inside the Claude Desktop app for Mac and Windows, and since July 2026 also on [claude.ai](#) in a browser and in the Claude phone app. You access it by clicking the Cowork tab, next to Chat. Sessions run against your Claude account, so a task started on the desktop keeps running after you close the laptop, and you can check on it from your phone.

Step 1: Grant folder access

You point Cowork at a folder on your computer. This is the boundary — Claude can only read and modify files within that folder. Start with something low-stakes: a Downloads folder, a project folder, a research directory. You control exactly what Claude can and cannot touch.

Step 2: Describe the task

In plain English. No special syntax, no commands. Just describe what you want done. Examples that work right now:

- "Organise this folder. Sort files by type, rename generic ones based on content, flag duplicates."
- "Pull the expense figures from all these receipt screenshots and create a summary spreadsheet."
- "Read all these research notes and write a 2-page briefing document."
- "Find all the client emails in this folder and extract the action items into a to-do list."
- "Take this raw data CSV and generate a formatted report with charts."

Step 3: Claude plans and executes

Cowork makes a plan, shows it to you, and then executes it step by step. You can see what it is doing in the Progress panel. You can pause or redirect mid-task if needed. For significant actions — deleting files, sending emails — Cowork asks for confirmation first.

Step 4: Results land in your folder

No copy-pasting from a chat window. No downloading outputs. The finished files are in your folder, exactly where you need them.

What makes Cowork different from just using Claude Chat

This is the question most people ask. The difference is significant:

	Claude Chat	Cowork
File access	You paste content into the chat	Claude reads files directly from your computer
Output	Text in a chat window — you copy it	Files created directly in your folder
Multi-step tasks	One response at a time — you direct each step	Claude plans and executes all steps autonomously
Memory	Forgets between conversations	Projects keep context, files, and instructions persistent
App connections	Limited	Connects to Gmail, Google Drive, Notion, Slack, and hundreds more via connectors
Works while you're away	No	Yes — scheduled recurring tasks run automatically

The features — what Cowork can do

File operations

Read, edit, create, rename, organise, and delete files in your designated folder. Process any file type Claude can understand: documents, spreadsheets, PDFs, images, CSVs, emails.

Projects

A persistent workspace tied to a folder. Your files, instructions, and task history stay in one place across sessions. Instead of starting fresh every time, your project remembers what it has done and what is next.

Scheduled tasks

Recurring tasks that run automatically. "Every Monday morning, pull last week's sales data from this folder and generate a summary report." Since the move to cloud sessions, scheduled tasks fire even

with no device online.

Connectors

Link Cowork to external services: Gmail and Google Workspace, Microsoft 365 (Outlook, OneDrive, SharePoint, Teams), Slack, Jira, Notion, DocuSign, and hundreds more via the connector marketplace. Desktop extensions have deeper system access than web connectors.

Skills

Install skills — modular capabilities — that improve Cowork's output for specific formats: documents, presentations, spreadsheets. You can also write a SKILL.md file with your brand voice or standard process and Cowork uses it every time.

The built-in browser

Since August 2026 Cowork has its own Chromium browser in a side panel of the desktop app. When a task involves a website, the browser opens, Claude navigates, fills forms and finishes the job — separate from your own browser and logins. If you would rather it work inside your own signed-in Chrome, Claude in Chrome does that and is now available on every paid plan; a safety classifier checks each action before it is taken.

Computer use

Cowork can see your screen and drive any application on your machine, not just the browser. Useful for tools that have no connector: a niche internal system, a desktop app, anything with buttons.

Plugins

Role-specific bundles of skills, prompts and connectors. Anthropic publishes its own — sales, finance, legal, marketing, and a Claude for Financial Advisors plugin covering onboarding, meeting prep and portfolio reviews — and a marketplace carries third-party ones. Plugins run on whatever plan you already have; nothing extra to buy from Anthropic.

Dispatch

Control a desktop Cowork session from your phone: approve a step, redirect, or read the result on the move.

The market reaction — why software stocks fell

When Cowork launched in January 2026, and again when its plugins landed, enterprise software stocks fell sharply. The iShares Expanded Tech-Software Sector ETF dropped nearly 5% in a single day, and analysts put the combined loss across enterprise software companies in the hundreds of billions of dollars.

The reason: investors looked at what Cowork can do and compared it to what they were paying enterprise software companies to do. File management. Document creation. Data extraction. Report generation. Workflow automation. These are the core functions of dozens of \$10–50/user/month SaaS products.

A knowledge worker with Cowork can now automate tasks that previously required separate tools for project management, document processing, data analysis, and workflow automation — all from one desktop agent.

The stocks recovered somewhat when Microsoft announced it was building Copilot Cowork on top of Anthropic's technology — suggesting the incumbents would adapt rather than be replaced entirely. But the signal was clear: the knowledge work software market has been permanently disrupted.

Microsoft Copilot Cowork — the enterprise version

Microsoft announced Copilot Cowork in March 2026 and made it generally available worldwide in June: a cloud-based agent, powered by Anthropic's Claude, that executes multi-step tasks across Microsoft 365.

The difference from consumer Cowork: Copilot Cowork runs inside Microsoft 365's own infrastructure with access to the full graph of enterprise work data — Outlook, Teams, calendars, SharePoint, Excel. Rather than working on a local folder, it works across an entire organisation's data.

It ships as a \$30 per user per month add-on, or inside the Microsoft 365 E7 bundle.

The significance: despite a \$13 billion investment in OpenAI, Microsoft built its flagship Microsoft 365 agent on Anthropic's Claude — and then took it to general availability. Not because Microsoft lacks access to GPT, but because it chose Claude anyway.

Real use cases — what people are actually doing with it

Finance and accounting

Point Cowork at a folder of receipt screenshots. It extracts all expense data, categorises by type, and generates a formatted expense report. Tasks that took a finance assistant an hour now take Cowork five minutes.

Research and analysis

Feed Cowork a folder of research notes, articles, and PDFs. Ask it to synthesise the key findings into a briefing document. It reads everything, identifies themes, and writes the document — citing sources from the files it read.

Content operations

Give Cowork a brand voice SKILL.md file and a folder of raw content briefs. Ask it to draft all ten blog posts. It writes in your brand voice, formats correctly, and saves them as Word documents ready for review.

Data processing

Drop a messy CSV export into the folder. Ask Cowork to clean it, standardise the formatting, remove duplicates, and generate a summary report with charts. No spreadsheet skills required.

File organisation

The classic first use case. Give Cowork your Downloads folder. It sorts by type, renames files based on their actual content, flags duplicates, and creates logical subfolders. Most users hit their first "wow" moment within ten minutes of trying this.

Scheduled intelligence

Connect Cowork to your email via the Gmail connector. Set a recurring task: every morning at 8am, read all emails received overnight, extract action items, and create a prioritised to-do list in your project folder. You wake up to a briefing that was prepared while you slept.

How to get access

Plan	Access	Price
Claude Pro	Full Cowork access	\$20/month
Claude Max	Full Cowork access + 5x or 20x the limits	\$100 or \$200/month
Claude Team	Full Cowork access + admin controls	\$25–30/user/month
Claude Enterprise	Full Cowork + admin controls, audit logs, smart reports	Custom pricing
Microsoft Copilot Cowork	Cloud-based, Microsoft 365 integrated	\$30/user/month add-on, or in the E7 bundle
Free tier	Not available — Cowork is paid only	—

To get started: subscribe to Claude Pro or higher, then open Cowork in the desktop app (claude.com/download), on claude.ai, or in the Claude phone app.

The honest limitations

Cowork is genuinely powerful. It also has real limitations worth knowing:

→ **Folder boundary only.** Claude can only access the folder you designate. It cannot reach outside it without explicit permission — and that is a feature. In one widely reported case about 15,000 family photos were deleted while Cowork tidied a desktop, and were recovered only through a 30-day cloud restore.

→ **It burns through usage.** Agent tasks use far more of your plan's allowance than chat does. Pro is the entry point; heavy daily use tends to need Max.

→ **Hallucination risk remains.** Cowork can misread file contents or make incorrect assumptions about what a task requires. Always review outputs before acting on them.

→ **Security considerations.** Days after launch, a data exfiltration vulnerability was reported and patched. Prompt injection — hidden instructions in a web page or document the agent reads — is the unsolved problem for every agent with browser access. Be thoughtful about connectors and folders.

→ **The desktop browser needs the desktop.** From the web or your phone, Claude can drive the built-in browser only while your desktop app is open. Without it, Claude in Chrome is the browser route.

→ **It keeps changing.** Cowork has shipped a major capability most months of 2026. Expect the feature list in this guide to be out of date before the next one.

What Cowork means for the agentic economy

Claude Code proved agents could replace developer time. Cowork is the thesis applied to every knowledge worker.

The administrative layer of every business — the file management, the data processing, the report generation, the inbox management, the scheduling — is now automatable by a non-technical person with a \$20/month subscription.

The barrier to running a lean, agent-powered business has never been lower. A solo founder with Cowork, Claude Code, and OpenClaw has the operational capacity of a team of five. A small business owner with Cowork has an always-on assistant that never sleeps, never forgets, and costs less than a gym membership.

The HBR research published in early 2026 noted that companies are already making headcount decisions based on AI's potential, not just its demonstrated performance. Cowork is part of why. It is not a future prediction — it is a current product that current businesses are using to do real work right now.

Glossary

Cowork — Anthropic's autonomous desktop AI agent. Part of the Claude Desktop app. Reads, edits, and creates files on your computer and executes multi-step tasks without command-line skills.

Claude Desktop — The desktop application from Anthropic for Mac and Windows. Contains Chat, Code, and Cowork modes, and Cowork's built-in browser. Download at claude.com/download.

Projects (Cowork) — A persistent workspace tied to a folder — keeps files, instructions, and task history in one place across sessions.

Scheduled tasks — Recurring tasks that Cowork runs automatically on a set schedule — daily, weekly, or custom timing.

Connectors — Integrations that link Cowork to external services — Gmail, Google Drive, DocuSign, Slack, Notion, and hundreds more.

Skills — Modular capabilities installed into Cowork. Including custom SKILL.md files you create to encode your brand voice or standard processes.

Claude in Chrome — Anthropic's agent inside your own Chrome browser, generally available on all paid plans. The alternative to Cowork's built-in browser when you want it working where you are already signed in.

Copilot Cowork — Microsoft's enterprise version of Cowork — cloud-based, running across Microsoft 365. Powered by Anthropic's Claude. \$30/user/month. Generally available since June 2026.

Claude Code — The developer-facing autonomous coding agent that Cowork is built on. Requires command-line skills. Cowork is the non-technical equivalent.

Knowledge worker — Someone whose job primarily involves creating, processing, or managing information — as opposed to physical labour. Cowork is specifically designed for this category.

Liked this? Go deeper.

Cowork is built on Claude and sits alongside Claude Code. Here is the full picture.

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is Anthropic

The company behind Claude. Founding story, the safety mission, the trillion-dollar valuation, and why its model powers so much of the agentic economy.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

How to Hire an AI Agent

for Your Business

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

Before you start — the right mindset

Deploying AI agents in your business is not a technology project. It is a staffing decision.

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

Step 1 — Find your highest-value repeatable task

The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

Customer support responses — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

Content creation — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

Research and summarisation — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

Data entry and processing — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

Email drafting — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

Scheduling and coordination — Managing calendar bookings, sending reminders, coordinating meeting times.

Lead qualification — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

Step 2 — Choose the right type of agent

Not all agents are right for all tasks. In 2026 the choice is simpler than it looks, because every option falls into a handful of shapes:

If you want to...	Use this	Cost
Delegate work inside your own files, inbox, calendar and apps	Cowork (Anthropic), ChatGPT Work (OpenAI), Copilot Cowork (Microsoft 365)	From \$20/mo, included in the paid plan
Run a team of agents sharing one cloud computer and its logins	Grok Bot (xAI)	Subscription
Have a personal agent you message like a colleague	OpenClaw (open-source), Meta Muse	Free + model costs; Muse has a free tier
Let agents build and run a whole company for you	Polsia, NanoCorp, Cofounder.co	From \$20/mo + fees (Polsia)
Run several agents as an org chart you control	Paperclip (self-hosted)	Free + model costs
Connect apps and automate defined workflows	Zapier, n8n	From ~\$20/mo; n8n is open-source
Build a custom agent with full control	Claude Developer Platform / Agent SDK	Pay per token

For a first agent, the top row is almost always the answer. A work agent on a plan you already have, pointed at one folder and one task, is the lowest-risk way to learn what this is actually like.

Step 3 — Write a clear job description for your agent

This is the most important step most people skip.

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

→ **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")

→ **The inputs** — what information does the agent receive? ("An inbound email from a customer")

- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

Step 4 — Start with read-only access

The most common mistake when deploying agents: giving them too much access too soon.

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

Step 5 — Measure results honestly

Define what success looks like before you deploy — not after.

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved
- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

Five deployments that work — and what to expect

These are the patterns that come up again and again among small businesses that get value from a first agent. Results vary; the shape of the win does not.

First-line customer support

An agent reads incoming support emails, categorises them, drafts replies from approved templates, and flags anything unusual for a human. The win is response time — hours become minutes — and a majority of routine emails handled without anyone touching them. The risk is a confident wrong answer, which is why replies are drafted for approval before they are sent unsupervised.

Research and content preparation

An agent researches a client's industry each week and drafts a content calendar, a briefing or a first draft. Writers spend their time on strategy and editing instead of gathering. The win is capacity without headcount.

Inbox triage for partners and owners

An agent summarises the day's email, drafts responses for review, and surfaces what needs a decision. Owners consistently report an hour or two a day back. The risk is the agent guessing at tone on a sensitive thread; keep it drafting, not sending, for those.

The solo operator's assistant

A personal agent — OpenClaw over WhatsApp or Telegram, or Cowork on the desktop — briefed with research tasks, email drafts, document summaries and admin. Typical cost is tens of dollars a month in model usage. This is the deployment most one-person businesses start with.

Routine engineering

Claude Code or Codex handling bug fixes, small features and tests, with senior people on architecture and review. Teams that do this well ship considerably more with the same headcount. Teams that skip the review step ship considerably more bugs.

The honest cost — what agents actually cost

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

Agent type	Monthly cost (approx)	What drives the cost
Cowork / ChatGPT Work	\$20–200/mo	Your Claude or ChatGPT plan; heavy use needs the higher tier
Copilot Cowork	\$30/user/mo	Add-on to Microsoft 365
OpenClaw (personal)	\$20–100/mo	Model API usage — Claude, GPT or a cheaper open model
Polsia company	\$20–1,000/mo + fees	Plan tier, credits, 20% on ad spend, 3% on revenue
Paperclip org	\$0 + model usage	Free software; each agent's model bill
Zapier / n8n automation	\$0–100/mo	Task volume and plan tier
Custom Claude agent	\$50–500/mo	Token usage — depends heavily on task volume
Claude Code	\$20–200/mo	Included in Claude Pro/Max; API if you go beyond

For context: a part-time human assistant in the US costs \$2,000–3,500/month. A full-time junior employee costs \$45,000–60,000/year. An agent handling equivalent work costs \$50–500/month. The economics are not subtle — and they are the reason the Agentic Economy guide spends a section on what that means for jobs.

What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

Your action plan — what to do this week

Day 1 — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

Day 2 — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

Day 3 — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

Day 4-5 — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

Week 2 — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

Week 3-4 — If quality is consistently good, expand permissions incrementally. Track the time saved.

Month 2 — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

The businesses that win with AI agents are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.

Glossary

AI Agent — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

Human-in-the-loop — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

Polsia — A managed platform where AI agents build and run companies. From \$20/mo plus credits and fees; see its own guide.

OpenClaw — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

Paperclip — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

Cowork — Anthropic's work agent: delegates tasks across your files, apps and browser. Included in every paid Claude plan.

ChatGPT Work — OpenAI's work agent: give it an outcome, it returns finished spreadsheets, decks and apps.

Zapier / n8n — Automation platforms that connect apps and run defined workflows, now with agent steps. n8n is open-source and self-hostable.

Claude — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

API key — Your credential for accessing an AI model's API. Required for OpenClaw, Paperclip and custom builds; Claude Code and Cowork run on your Claude plan instead.

Token — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

MCP (Model Context Protocol) — The open standard connecting AI agents to external tools. Invented by Anthropic.

Liked this? Go deeper.

You now know how to hire an agent. Here is the full context behind the tools you are using.

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

© 2026 WTF Agents. All rights reserved.

Who's Who in the

Agentic Economy

Forty-odd names you will keep seeing, sorted into the seven layers of the agentic economy, one line each. The map to read before any of the deep dives — and the one to come back to when a new name appears in your feed.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

How to read this guide

The agentic economy is a stack. Models at the bottom, agents built on them, platforms that run whole companies of agents, and the plumbing, chips and money underneath it all. Every name in the news belongs to one layer, and knowing which layer tells you most of what you need to know.

This guide gives each name one line: what it is, why it matters. Where a full WTF Agents guide exists, it is named so you can go deeper. Version numbers, valuations and star counts live in the dated box below and nowhere else, because they change monthly and the map does not.

By the numbers — September 2026

\$965B

Anthropic's valuation after May's \$65B round — the most valuable private company in the world, IPO filed

\$852B

OpenAI's valuation after its \$122B round in March

\$60B

what SpaceX paid for Cursor in June, after buying xAI in February

\$30M

Polsia's May raise at \$250M; NanoCorp is YC-backed; Paperclip has no disclosed funding

3.2M

active OpenClaw users; Hermes Agent is the credible number two

490,000

companies created on Polsia; about 6% still active

6

labs that shipped a "does your job" agent in 2026: Anthropic, Microsoft, OpenAI, xAI, Meta — and Google, by consolidation

2

open standards under the Agentic AI Foundation: MCP and A2A

Layer 1 — The labs and their models

The brains. Everything above runs on one of these.

Anthropic — Claude — Safety-founded lab, now the most valuable private company on earth. Claude powers Claude Code, Cowork, Copilot Cowork, OpenClaw's default and Polsia's agents. Tiers: Fable, Opus, Sonnet, Haiku; Mythos is the restricted twin. Guides: [WTF is Anthropic](#), [WTF is Claude](#).

OpenAI — GPT — ChatGPT is still the product most people have used. GPT-6 arrived in September. OpenAI's agents are ChatGPT Work, Codex and the Agents API; its reach is its weapon.

Google DeepMind — Gemini — Gemini 3.8, Gemini Agent, Chrome Auto Browse and Jules. Owns Search, Android and Workspace, and invented the A2A protocol. Chose to consolidate its agents rather than launch a new one.

Meta — Muse — A new model family and a personal agent of the same name, delivered through WhatsApp and a standalone app with a free tier. Meta's bet is that three billion messaging users are the distribution. Also tried to buy Manus.

xAI — Grok — Elon Musk's lab, now inside SpaceX alongside Cursor. Grok Bot runs a team of agents sharing one cloud computer and its logins, coordinated by a "Chief of Staff."

Microsoft — Not a frontier lab of its own but the biggest distributor: Copilot Cowork, built on Claude, is live in Microsoft 365 worldwide; Copilot Studio hosts hundreds of thousands of custom agents; Foundry serves every major model.

DeepSeek — The Chinese lab whose cheap-to-train breakthrough in early 2025 made this a two-continent race. DeepSeek V4 is open-weight under the MIT licence and runs under a great many indie agents.

Alibaba — Qwen · Moonshot — Kimi · Zhipu — GLM · MiniMax — The rest of China's open-weight frontier. Kimi K3 was the largest open model ever released. Cheap, capable, downloadable. Guide: China's AI (coming).

Mistral — France's lab. Efficient smaller models and Europe's answer to US and Chinese dependence.

Layer 2 — Work agents

Agents that do your job's tasks. The category every lab entered in 2026.

Cowork (Anthropic) — Claude working in your files, apps and its own browser, on desktop, web and phone. Included in every paid Claude plan. Guide: [WTF is Cowork](#).

ChatGPT Work (OpenAI) — Give it an outcome; it returns finished spreadsheets, decks and web apps hours later. Launched July.

Copilot Cowork (Microsoft) — Cowork inside Microsoft 365, across Outlook, Teams, SharePoint and Excel. Built on Claude despite Microsoft's stake in OpenAI. Generally available since June.

Grok Bot (xAI) — A team of named agents sharing one cloud computer, working while you sleep. Launched August through Cursor's infrastructure.

Microsoft Scout — An always-on personal agent inside Copilot, in limited release.

Gemini Agent (Google) — Where Google's browser agent went after Project Mariner was folded in. Plus Auto Browse in Chrome.

Layer 3 — Coding agents

The category that proved agents work, and the one that made Anthropic.

Claude Code (Anthropic) — The autonomous developer. Terminal, desktop app, phone. The fastest-adopted developer tool of its generation. Guide: [WTF is Claude Code](#).

Codex (OpenAI) — OpenAI's coding agent, now living inside the ChatGPT desktop app on GPT-6.

Devin (Cognition) — The original "AI software engineer," fully cloud-hosted, cut from \$500 to \$20 a month in April.

Jules (Google) — Google's free cloud coding agent on Gemini.

Cursor — The AI-first code editor, bought by SpaceX for \$60 billion in June and now the delivery vehicle for Grok Bot. Losing OpenAI's models in November.

GitHub Copilot (Microsoft) — The largest install base in coding AI, evolving from autocomplete into an agent inside GitHub itself.

Layer 4 — Personal agents

Agents that run your life's admin, mostly over messaging apps.

OpenClaw — Open-source, self-hosted, millions of users, a non-profit foundation, native mobile apps. The reference personal agent and the one every rival is measured against. Guide: [WTF is OpenClaw](#).

Meta Muse — Meta's consumer agent: books, buys, negotiates and fills in forms, free tier, WhatsApp. US only at launch. Guide: [WTF is Meta Muse](#) (coming).

Hermes Agent (Nous Research) — The credible open-source number two: self-improving memory, hireable into Paperclip out of the box.

Manus — The Chinese general-purpose agent Meta agreed to buy for \$2 billion before Beijing blocked the deal. Still independent, still Chinese.

Layer 5 — Company platforms and orchestration

Agents that run whole businesses, and the tools that manage teams of agents.

Polsia — The biggest and most argued-about: subscribe, describe an idea, a fixed team of agents builds and runs the company. Nearly half a million created; a founder who says one in ten has earned a dollar. Guide: [WTF is Polsia](#).

NanoCorp — YC-backed rival that publishes the live revenue of every company on the platform — the most honest number in the space, and a small one. Guide: WTF is NanoCorp (coming).

Cofounder.co — Company-in-a-box with departments and a human approval gate in front of anything risky.

Paperclip — Open-source org chart for agents: roles, budgets, approvals, any agent vendor, self-hosted. ClipMart is its marketplace of forkable company templates. Guide: [WTF is Paperclip](#).

Salesforce Agentforce — The enterprise incumbent's agent platform, priced per conversation, past a billion dollars in annual revenue.

ServiceNow AI Control Tower — Governed agents for IT and HR workflows; AI contract value past a billion dollars.

SAP Joule · Workday Illuminate — Agents inside the ERP and HR systems most large companies already run — and the first to charge outside agents a toll to reach their data.

Zapier · n8n · UiPath — The automation layer: connect apps, run defined workflows, now with agent steps. n8n is open-source; UiPath is the enterprise back-office specialist.

Layer 6 — The plumbing

Standards and infrastructure that make agents interoperate.

MCP — Model Context Protocol — The universal plug between agents and tools. Invented by Anthropic, adopted by everyone, now vendor-neutral. Guide: [WTF is an API](#).

A2A — Agent2Agent — Google's standard for agents from different companies talking to each other. In production at hundreds of organisations.

Agentic AI Foundation — The Linux Foundation body that owns MCP and A2A, founded by Anthropic, OpenAI, Google, Microsoft and Block. The rare case of rivals agreeing to share.

Stripe — The payment rails under Polsia companies, the WTF Agents store, and most of the agentic commerce experiments. Agents that can pay run through Stripe.

Cloudflare · Vercel — Where agents are hosted and deployed. Vercel's Eve is a framework for durable agents; Cloudflare's platform runs them at the edge.

Layer 7 — The picks and shovels

Who gets paid regardless of which agent wins.

Nvidia — The chips. Every model in Layer 1 was trained and is served on Nvidia hardware, which is why it is the most valuable company in the world.

AWS · Google Cloud · Microsoft Azure — The clouds. Bedrock, Vertex AI and Foundry serve every major model to enterprises inside infrastructure they already pay for. Amazon and Google are also Anthropic's largest backers.

Hugging Face — The hub where open-weight models are published and downloaded. If DeepSeek or Qwen ships a model, this is where it lands.

GitHub — Where the code lives and where agents commit it. OpenClaw and Paperclip's star counts are GitHub's.

Guide: The Picks and Shovels (coming) goes deeper on this layer.

The people

Six names that come up more than any others.

Dario Amodei — CEO of Anthropic. The essays that set the tone for the safety-first case.

Sam Altman — CEO of OpenAI. The face of ChatGPT and the reach strategy.

Elon Musk — SpaceX, xAI, Cursor, Grok Bot. Consolidating the agent stack under one company.

Mark Zuckerberg — Meta. Muse is his second attempt at an AI reset, and the first with a real agent.

Peter Steinberger — Created OpenClaw, joined OpenAI, left the project to a foundation.

Ben Broca (Ben Cera) — Solo founder of Polsia. Raised \$30 million with one employee and said out loud that most of his customers' companies earn nothing.

How the layers fit together

Value has moved down the stack. In 2024 the money was in models. In 2026 it is in agents. The bet behind Layer 5 is that next it will be in companies that no human works at.

Two forces cut across every layer. The first is Claude's spread: Anthropic's model sits under products from Microsoft, OpenClaw, Polsia and Paperclip, which is why one lab is close to a trillion dollars. The second is China's open weights: cheap, capable, downloadable models that let a solo builder run agents for pennies and give every US lab a price ceiling.

When a new name appears in your feed, ask which layer it belongs to. The answer usually tells you who it competes with, who it depends on, and whether it is a product or a press release.

Glossary

Layer — A level of the agentic economy stack: models, agents, platforms, plumbing, infrastructure. Every company in this guide belongs to one, sometimes two.

Frontier lab — A company training models at the leading edge of capability: Anthropic, OpenAI, Google DeepMind, Meta, xAI, DeepSeek.

Work agent — An agent that does your job's tasks: Cowork, ChatGPT Work, Grok Bot, Copilot Cowork.

Personal agent — An agent that runs your life's admin: OpenClaw, Muse, Hermes.

Company platform — A service where agents build and run a whole business: Polsia, NanoCorp, Cofounder.co.

Orchestration — Managing a team of agents: Paperclip.

Open weights — A model anyone can download and run. Most Chinese frontier models; few US ones.

MCP / A2A — The two open standards that let agents use tools and talk to each other. Both under the Agentic AI Foundation.

Incumbent — An established enterprise software vendor adding agents to what it already sells: Salesforce, ServiceNow, SAP, Workday, Microsoft.

Liked this? Go deeper.

This is the map. Every name on it with a guide has one waiting.

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

Which AI Should I

Actually Use?

Claude, ChatGPT, Gemini, Muse, Grok, DeepSeek. Six families, a new version every month, and a leaderboard that changes weekly. Here is how to choose without reading a single benchmark — and what each one is quietly bad at.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

There is no best AI. There is a best AI for the thing you are doing today, and the differences between the top models are now smaller than the differences between how you use them.

That is the honest answer, and it is also the useful one. This guide gives you the decision in one table, the plans and prices, what the leaderboards actually measure, and the catches nobody puts on the pricing page.

By the numbers — September 2026

Everything here is a snapshot. The families below it are stable; the numbers are not.

6

model families that matter for a normal person: Claude, GPT, Gemini, Muse, Grok, DeepSeek

Fable 5.1

top of the independent Artificial Analysis Intelligence Index since 1 September, with Opus 5 close behind at half the price

GPT-6 Astra

OpenAI's newest, 3 September, after the GPT-5.6 Sol / Terra / Luna family in July

\$20

the standard monthly plan at every lab — ChatGPT Plus, Claude Pro, Muse, Google's AI plan

\$100–200

the power-user tier: ChatGPT Pro, Claude Max, Muse's top plan

\$0.20 vs \$50

per million output tokens, cheapest frontier-adjacent model (GPT-5.6 Luna) versus most expensive (Fable 5.1) — a 250× range

\$0

DeepSeek's chat app, no paid tier at all

The decision in one table

Start here. Pick the row that matches what you are doing.

You want to...	Use	Why
Ask questions, get quick answers, everyday chat	ChatGPT (free or Plus)	The largest user base and the most polished consumer app. Everyone you know is on it.
Write anything longer than an email	Claude	Consistently preferred for voice, structure and holding a style across long pieces.
Delegate real work in your files and apps	Claude (Cowork) or ChatGPT Work	The two work agents on a \$20 plan. See the Cowork guide.
Build software without being a developer	Claude Code	The desktop app changed who can use it. See its guide.
Run an agent for hours unsupervised	Claude Fable or Opus	The behaviour that matters — asks when unsure, flags its own errors — not the score.
Have an agent book, buy and message for you	Meta Muse (or OpenClaw)	Muse is free, in WhatsApp, and made for exactly this. OpenClaw if you want to own it.
Work inside Google: Gmail, Docs, Search	Gemini	Nothing else lives inside the tools you already open.
Work inside Microsoft: Outlook, Teams, Excel	Copilot Cowork	Microsoft 365's agent — running on Claude.
Get live news and what X is saying right now	Grok	Native real-time X search. Nothing else has it.
Pay nothing, or nearly nothing	DeepSeek	Free chat, an API up to 100× cheaper than US rivals. Read the catch below.
Run a model on your own hardware	DeepSeek, Qwen, Kimi	Open weights. Download, run, keep your data.

If you only pay for one: Claude Pro or ChatGPT Plus, \$20, and you are covered for 90% of what any of these can do. Add the other when you find a task the first one does badly.

The six families — what each is actually like

Claude (Anthropic)

The model that most of the agentic economy chose. Best at long, careful work — writing, analysis, code, and agent sessions that run for hours — because it flags uncertainty instead of guessing and

notices its own mistakes. Fable is the flagship, Opus the everyday frontier pick, Sonnet the workhorse, Haiku the cheap one. The catch: Fable will decline a narrow set of dangerous cybersecurity and biology requests, and it is the most expensive model on the market per token. See the [WTF is Claude](#) guide.

GPT / ChatGPT (OpenAI)

The one everyone has used. ChatGPT is the most polished consumer app, with the broadest set of built-in tools, and GPT-6 Astra is the newest frontier model. ChatGPT Work and Codex are OpenAI's agents. Best for quick answers, creative brainstorming, images, and anything where reach matters. The catch: an independent evaluator flagged elevated "scheming" behaviour in the GPT-5.6 Sol tier in its own system card, so for work where you need the model to stay inside the rules you set, the more conservative tiers are the safer pick.

Gemini (Google)

Strongest where Google is: inside Gmail, Docs, Drive, Search and Android, and for anything involving images, video or very long documents. Gemini 3.8 Flash is fast and cheap; Google calls its pricing introductory. Gemini Agent and Chrome's Auto Browse are Google's agents. The catch: a habit of being uneven — excellent one prompt, flat the next — and Google's product names change faster than anyone's.

Muse (Meta)

The newest family, and a genuine reset for Meta. Muse Spark is competitively priced, and the Muse agent — free, in WhatsApp, in a standalone app — books, buys, negotiates and fills in forms for you. The catch is the one that matters most in this guide: Muse's cheap "Contributor" tier is cheap because Meta trains on your prompts. Read which tier you are on before you paste anything private. Meta's own staff flagged security failures days before launch.

Grok (xAI)

Elon Musk's model, now inside SpaceX with Cursor. Grok's unique feature is native real-time search of X, which makes it the best model for "what is happening right now." It topped one agentic index this summer, and Grok Bot is a serious work agent. The catch: the personality is a feature or a bug depending on your taste, and the whole thing is now owned by a rocket company that also owns your code editor.

DeepSeek and the Chinese open models

DeepSeek V4 and its Flash tier cost a tiny fraction of US models and the chat app is free with no paid plan at all. Qwen, Kimi and GLM are close behind, all open-weight: you can download them, run them on your own machine, and keep every byte. This is what a great many indie agents run on. The catch: the hosted versions put your data on servers in China under Chinese law, the models are trained to avoid politically sensitive topics, and Western enterprises mostly cannot use them for compliance reasons. Downloaded and run yourself, none of that applies.

The plans and what they cost

Consumer plans have converged. Every lab has a free tier, a \$20 tier and a \$100–200 tier.

	Free	\$20/mo	\$100+/mo
Claude	Chat only	Pro: Cowork, Claude Code, higher limits	Max: 5× or 20× the limits
ChatGPT	Chat with limits	Plus: better models, more use	Pro: heaviest use, top tiers
Gemini	Chat in Google apps	Google AI plan: Gemini in Workspace	Higher tier bundles storage and video
Muse	The agent, with limits	Standard	Top plan
Grok	Basic via X	X Premium+	SuperGrok Heavy
DeepSeek	Everything	—	—

What the \$20 buys you is not a better brain. At every lab, the free tier is the same or nearly the same model with a cap on how much you can use it. The paid tier removes the cap and, at Claude and OpenAI, unlocks the agents. Agent work burns through usage far faster than chat, which is what the \$100+ tiers exist for.

What the leaderboards actually measure

You will see headlines that model X "beat" model Y. Here is how to read them.

- **Vendor launch tables** compare a lab's new model to rivals on tests the lab chose, with settings the lab chose. Useful for learning what the lab optimised for. Useless for declaring a winner.
- **Independent indices** — Artificial Analysis, LMArena — run every model through the same tests. Better, but "best overall" still hides that one model wins coding, another wins writing, another wins speed and another wins price.
- **Agentic indices** — SWE-bench, Terminal-Bench and similar — measure whether a model can finish real multi-step jobs. These matter most for anything in this series, and Claude has led most of them through 2026.
- **The ranking changes monthly.** Fable 5.1 took the top of the main index on 1 September; GPT-6 arrived two days later. Whatever leads when you read this may not lead next month.

The practical rule: the top five models are all good enough for almost anything a normal person does. Choose on the things that do not change — where it lives, what it costs, how it behaves when it is unsure — not on this month's score.

The catches, all in one place

Your data. Muse Contributor trains on your prompts. DeepSeek hosted stores them in China. Every lab's free tier has looser data terms than its paid one. Read the setting once; most let you opt out.

Refusals. Claude Fable declines certain dangerous requests by design. The Chinese models avoid politics. ChatGPT has its own list. If a model refuses something ordinary, it is usually a wording problem, not a capability one.

Hallucination. All of them do it. The difference is whether the model tells you it is unsure. This is Claude's strongest suit and the reason it runs so many unsupervised agents.

Lock-in. Your history, projects and memory live with one lab. Moving is possible but tedious. Pick your main one with that in mind.

Usage caps. The \$20 plans are generous for chat and tight for agents. If Cowork or ChatGPT Work is your daily driver, budget for the \$100 tier.

Regulation. Governments now switch models on and off. Fable 5 was briefly suspended in June under a US export-control order before returning. Expect more of this, not less.

You may not have to choose

Two things make the question easier than it looks.

First, the agents pick for you. OpenClaw, Paperclip and most agent platforms let you plug in any model by API key and swap it in a minute. Polsia chose Claude Opus for you. Copilot Cowork chose Claude for you. Grok Bot chose Grok. If you are using an agent rather than a chat app, the model decision has mostly been made.

Second, routers exist. Services that sit in front of several models and send each task to the best-value one are now ordinary, and most developer tools do this by default. The best AI in 2026 is increasingly a routing policy, not a single brand.

Glossary

Model family — A lab's line of models across price points: Claude (Fable, Opus, Sonnet, Haiku), GPT (Astra, Sol, Terra, Luna), and so on. Families are stable; version numbers are not.

Frontier model — A model at the leading edge of capability. Fable, GPT-6, Gemini 3.8 Pro, Muse Spark, Grok 4.6, DeepSeek V4 Pro.

Flash / mini / Haiku — Each lab's name for its fast, cheap tier. Fine for most everyday tasks.

Token — The unit models are billed in. Roughly three quarters of a word. API prices are quoted per million tokens, input and output.

Open weights — A model you can download and run yourself. DeepSeek, Qwen, Kimi, GLM, Mistral, Llama.

Leaderboard / index — An independent ranking of models on shared tests. Artificial Analysis and LMArena are the most cited.

Agentic index — A benchmark of finishing real multi-step tasks rather than answering questions: SWE-bench, Terminal-Bench.

Router — A service that sends each request to whichever model is best or cheapest for it.

Contributor tier — Meta's discounted Muse pricing in exchange for training on your prompts.

Liked this? Go deeper.

You have picked a model. Here is what to do with it.

WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

wtfagents.com/store

WTF is Claude

Anthropic's AI model. The family, the new Fable and Mythos tier, what it can do, and why it powers so many agents — including rivals' products.

wtfagents.com/store

Who's Who in the Agentic Economy

The map. Every lab, agent, platform and piece of infrastructure that matters, one line each, sorted by layer. Read this first, then go deep.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Meta Muse?

Meta's personal AI agent lives in WhatsApp, books your travel, sells your car, lowers your bills and pays with its own card — free, for three billion people. What it does, how it is built, what it costs, and what you are handing over.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Meta Muse is a personal AI agent that does things for you rather than answering questions: it books travel, fills in forms, manages your inbox, negotiates a bill, sells something on your behalf, and checks out with a one-time card — from its own app or straight inside WhatsApp.

It launched on 8 September 2026 for adults in the United States, free for most of what people need. It is the first personal agent from a company with three billion messaging users, and that distribution, more than any feature, is why it matters.

By the numbers — September 2026

8 Sept

2026: launch, US only, 18+, on iOS, Android, web and inside WhatsApp

\$0

for most use, with a metered free allowance. Power at \$20/mo, Maximum at \$100/mo for heavier use

1

dedicated cloud computer per user — its own browser, its own isolated environment

3B+

people already on WhatsApp, Instagram and Facebook: the distribution Muse is built to ride

Spark 1.3

the Meta model it runs on, released 2 September, also behind Muse Code and Meta AI

0

countries beyond the US at launch; AI glasses support promised

What it actually does

You message Muse like a person. "Find me a cheaper phone plan." "Book the dentist for next week." "Sell the old bike." "Get this bill down." It makes a plan, opens its own browser, works through websites the way a person would, and keeps going after you close the app — coming back when it needs a decision or a signature.

Travel — Searches, compares, books flights and hotels, handles the confirmation emails.

Inbox and calendar — Reads your email, drafts replies, schedules, finds the thing you were looking for.

Forms and admin — Fills in applications, renewals, registrations. The tedious layer of adult life.

Shopping and paying — Finds the product, checks out with a one-time virtual card so the merchant never gets your real number.

Selling and negotiating — Lists an item, fields the messages, haggles. Calls a provider to lower a bill.

Memory across Meta — With permission, it uses your Instagram saves, Facebook activity and WhatsApp chats: a recipe you saved becomes a shopping list; a place you liked becomes a booking suggestion.

What it does not do: make content. No images, no video, no posts. Muse manages the life around your work, not the work. For that you want Cowork, ChatGPT Work or Grok Bot.

How it is built — the part that stands out

Most personal agents run on your machine with your logins (OpenClaw) or inside a chat app with limited reach. Muse does something different, and the design is the most careful in the consumer market so far.

Its own computer

Each Muse runs on a dedicated virtual machine in Meta's cloud — a separate computer with its own browser. It never touches your device. Meta calls this Muse Secure VM.

A guard between it and the internet

A second agent, called Sentinel, sits between Muse and the outside world. Anything Muse tries to send — a form submission, a payment, a message — has to pass Sentinel first. The idea is that a Muse that has been tricked by a malicious web page still cannot act on it without the guard agreeing.

A vault it cannot see into

Your passwords and card details sit in a credential vault. Muse can use them to log in and pay, but cannot read them. Payments run through one-time cards.

Confirmation before anything that matters

Sensitive actions — spending money, sending a message in your name, anything irreversible — come back to you for a tap before they happen.

None of this is audited by anyone outside Meta yet. It is, on paper, the best-designed consumer agent security story to date. Whether it holds is a question for the first year.

What it costs

Plan	Price	What you get
Free	\$0 (card required)	The full agent, with a metered allowance that warns you as you approach it
Power	\$20/mo	More tasks, more compute
Maximum	\$100/mo	The most tasks, priority access, for people who hand it a lot

Meta has not published quotas for any tier. Because one request — "find me a cheaper plan" — can mean hours of browsing and dozens of pages, the useful measure is cost per finished task, not price per month. Try the free tier for a week and watch the meter before deciding.

Muse vs OpenClaw vs Grok Bot

	Muse	OpenClaw	Grok Bot
What it is for	Your life's admin	Anything you configure	Your work
Where it runs	Meta's cloud	Your machine	xAI's cloud
Who picks the model	Meta (Spark)	You, any model	xAI (Grok)
Cost	Free, \$20, \$100	Free + model usage	Subscription
Setup	None	Real, technical	Some
Your data	With Meta	With you	With xAI
Availability	US only	Everywhere	Wide

Muse is the managed, zero-setup version of what OpenClaw people have been building by hand since late 2025. Muse is convenience at the price of control; OpenClaw is the reverse.

The catches

Muse asks for a level of trust no chatbot ever did: your email, your calendar, your browser, your card, in the hands of a brand-new product from the company that built Facebook.

→ **Training on your data is on by default.** Meta's terms allow training on Muse interactions unless you opt out. Find the setting before you connect your inbox.

→ **It launched rough.** Meta's own employees, testing during launch week, reported exposed iCloud photos, repeated logouts and monitoring switching itself off. Version-one problems, fixed quickly, and a reminder of what "new" means.

→ **Prompt injection is not solved.** Sentinel is a serious answer to hidden instructions in web pages, and it is the best one in the consumer market. It is not a guarantee.

→ **US only, 18 only.** No other country has a date.

→ **Meta is Meta.** The privacy record is what it is. Muse's security claims — no password visibility, no ad-targeting from agent data — are Meta's, unaudited. Decide accordingly.

Who it is for

Yes — Anyone already living in WhatsApp and Instagram who wants the admin of life handled and is comfortable with Meta holding the keys. The zero-setup personal agent people have been waiting for.

Not yet — Anyone outside the US, or anyone with a serious privacy stance toward Meta. Watch the first six months.

No — Anyone who wants an agent for their work rather than their life, or who wants to own the agent. That is Cowork, Grok Bot, or OpenClaw.

Why it matters

Every other agent in this series had to earn its users. Muse starts with three billion of them already opening the app it lives in. If personal agents become a habit, this is the most likely product to make them one — and it sets the bar for what "free" means: the full agent, metered, paid for by the people who lean on it hardest.

It also puts Meta, which spent 2025 losing the model race, back in it with the most ambitious consumer agent shipped so far. Whether it is trusted is now the only question.

Glossary

Meta Muse — Meta's personal AI agent, launched 8 September 2026. Acts on your behalf from its own app or inside WhatsApp.

Muse Spark — The model family Muse runs on. Also behind Meta AI, Muse Code and Meta's model API.

Muse Secure VM — The dedicated cloud computer each Muse runs on, isolated from your devices.

Sentinel — The guard agent that must approve anything Muse sends to the internet.

Credential vault — Where your passwords and cards live: usable by Muse, not readable by it.

One-time card — A virtual card number for a single purchase, so merchants never see your real one.

Personal agent — An agent that runs your life's admin, as opposed to your job's work. Muse, OpenClaw.

Meta Superintelligence Labs — The Meta division, led by Alexandr Wang, that built Muse.

Liked this? Go deeper.

Muse is the newest personal agent. Here are the others, and the model question behind it.

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/store

Which AI Should I Actually Use?

Claude vs ChatGPT vs Gemini vs Muse vs Grok vs DeepSeek, for a normal person. Which to use for what, what each costs, and the catches.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

WTF is

Grok Bot?

A team of named AI teammates that sign into your real apps, work on their own cloud computer around the clock, and report to a Chief of Staff bot. From the company that now owns SpaceX, xAI and Cursor. What it is, what it costs, and the catch.

WTF Agents · wtfagents.com · September 2026

Part of the WTF Agents Guide Series · wtfagents.com/store

The one-liner

Grok Bot is a set of always-on AI teammates. You give each one a name, a role and access to your tools; it signs in with your real logins, works on a cloud computer that keeps running after you close your laptop, and comes back with finished work — or with a question when it needs your sign-off.

The pitch line is "AI teammates you can give real work to." The distinctive move is the team: a researcher bot, a writer bot, an inbox bot, and a Chief of Staff bot that coordinates them so you do not have to.

By the numbers — September 2026

11 Aug

2026: early beta launch, macOS, Windows and iPhone. No Linux client, no Android yet

1

cloud computer per account, shared by all your bots — files, browser sessions and logins pooled

\$20–300

per month: bundled with Cursor plans from Pro upward and with SuperGrok Heavy at the top. No standalone price, no free tier

\$60B

what SpaceX paid for Cursor in June, after absorbing xAI in February. Grok Bot is the first joint product

10–20

bots one SpaceXAI engineer runs under a single Chief of Staff, automating what she says is 90% of her routine

0

model picker. Tasks are routed to models automatically; you cannot choose or pin one

How it works

Bots, not chats

You create a bot, give it a name and a job — "Researcher," "Inbox," "Ops" — and tell it what it may sign into. From then on you message it like a colleague, in a thread or a group chat with other bots.

Real logins, not integrations

This is the big design choice. Most work agents connect to your tools through official APIs and connectors. Grok Bot signs in to the app the way you would, with your credentials, and uses it the way you would. If a tool has a login screen, a bot can use it, including the internal systems and legacy tools no connector exists for.

One computer in the cloud

All your bots share a persistent virtual machine on SpaceXAI's servers: a file system, a terminal, a browser, exactly as you left it. The work happens there, not on your device, which is why it continues while your laptop is shut. Launch coverage often said each bot gets its own computer; SpaceXAI's own documentation says they share one, and that separate bots should not be treated as a security boundary.

The Chief of Staff

Put one bot above the others and tell it to coordinate. Internally SpaceXAI runs a Chief of Staff over specialists for recruiting, expenses, operations and bug fixes. Early testers who expected this to fall over found it worked out of the box: hand a project to the coordinator and it delegates, collects and reports.

Teach a task once

Record yourself doing a workflow on screen. Grok Bot turns the recording into a reusable skill it can run on its own, no code, no integration. Skills and scheduled routines round it out, as in Cowork.

What it costs

There is no Grok Bot subscription. Access comes bundled with plans you buy for something else:

Plan	Price	Notes
Cursor Pro	\$20/mo	Entry point since late August; a weekly usage allowance
Cursor Pro+ / Ultra	\$60 / \$200/mo	Higher limits
Cursor Teams Premium	\$120/seat/mo	Central billing, analytics, single sign-on
SuperGrok Heavy	~\$300/mo	Through the Grok app; the original launch tier

Usage is metered in agent steps and tokens on top of the allowance, and SpaceXAI has not published what drives consumption. A routine that polls every fifteen minutes runs ninety-six times a day. Watch the meter in week one.

Grok Bot vs Cowork vs OpenClaw

	Grok Bot	Cowork	OpenClaw
Shape	A team of named bots	One agent, many tasks	One agent you configure
Reaches your apps via	Your real logins	Connectors + its own browser	Whatever you wire up
Runs on	SpaceXAI's cloud	Your desktop or Claude's cloud	Your machine
Model	Auto-routed, no choice	Claude, your pick of tier	Any, your choice
Cost	Bundled, \$20–300	\$20 Claude plan	Free + model usage
Best for	Delegating a whole function	Files, documents, analysis	Owning the whole thing

The comparison people keep making: "like OpenClaw, but easy, reliable and less scary." That is the product in one line, with the caveats below.

The catches

- **Your logins are the integration.** Bots hold your real credentials and act in your accounts. That is the reach and the risk in one sentence. SpaceXAI documents an approval model for consequential actions; it has not published data-retention terms, and the audit view was promised rather than shipped at launch.
- **One computer, no walls.** Because all bots share one machine, a bot tricked by a malicious page has the same access as every other bot. Do not give a low-trust bot high-trust logins.
- **You do not pick the model.** Tasks route automatically, billing follows whichever model served the request, and SpaceXAI has said it does not plan to add a picker. Early testers found the router uneven.
- **Bundled pricing hides the bill.** No standalone plan, no free tier, and usage on top of the allowance. The cheap \$20 entry is a Cursor plan, which makes sense if you code and less if you do not.
- **It is a beta from a company mid-merger.** SpaceX absorbed xAI in February and Cursor in June; Cursor loses OpenAI's models in November. Expect the product and the pricing to move.

Who it is for

Yes — Operators with repetitive, multi-step digital chores across several apps — inbox triage, expenses, CRM updates, bookings, returns — who are already on Cursor or Grok and want to delegate a function, not a task.

Not yet — Anyone whose work is regulated or high-stakes and cannot be supervised step by step, or who needs to know which model is touching their data.

No — Anyone who wants an agent they own and control. That is OpenClaw or Paperclip.

Why it matters

Grok Bot is the fourth "team of agents does your job" product to ship in 2026 and the one that leans hardest into the office metaphor: named colleagues, a boss bot, a shared computer, group chats. It is also the first product of a company that now owns a rocket manufacturer, a frontier lab and the most popular AI code editor, with a stated plan to consolidate the stack.

Whether it is the best work agent is a question for a year from now. Whether it is the most ambitious consumer bet on agents as staff is already answered.

Glossary

Grok Bot — SpaceXAI's work agent: named, persistent bots that sign into your apps and work on a shared cloud computer. Beta since August 2026.

SpaceXAI — The entity that emerged from SpaceX absorbing xAI in February 2026 and Cursor in June. Builder of Grok and Grok Bot.

Chief of Staff bot — A bot placed above other bots to delegate, coordinate and report. Grok Bot's orchestration layer.

Credential-holding agent — An agent that signs in with your real logins rather than through sandboxed APIs. Grok Bot's core design.

Cloud computer — The persistent virtual machine, with browser, files and terminal, that all your bots share.

Teach-a-task — Record a workflow once on screen; the bot turns it into a reusable skill.

Routine — A scheduled, recurring task a bot runs without being asked.

Cursor — The AI code editor, bought by SpaceX for \$60 billion, through which most Grok Bot subscriptions are sold.

Liked this? Go deeper.

Grok Bot is one of four work agents that shipped this year. Here are the others, and the model behind it.

WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

wtfagents.com/store

WTF is OpenClaw

The open-source personal agent with millions of users. The wild origin story, the security warnings, and where it stands now.

wtfagents.com/store

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/store

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/store

WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

wtfagents.com/store

Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

wtfagents.com/store →



Scan to browse all guides Every WTF Agents guide at wtfagents.com/store \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at wtfagents.com/companies

