

# How to Hire an AI Agent

## for Your Business

---

A practical, jargon-free guide for business owners. No code required. No prior AI experience needed. Just clear steps to deploy your first AI agent this week — and start getting your time back.

---

WTF Agents · [wtfagents.com](https://wtfagents.com) · September 2026

Part of the WTF Agents Guide Series · [wtfagents.com/store](https://wtfagents.com/store)

# Before you start — the right mindset

---

**Deploying AI agents in your business is not a technology project. It is a staffing decision.**

The mental model that works: treat AI agents like new hires. You would not give a brand new employee access to your entire company on day one. You would not give them a vague goal with no check-ins. You would not trust them to make major decisions without oversight — until they had earned it.

Apply the same thinking to agents. Start small. Define the role precisely. Give limited access. Review outputs. Build trust incrementally.

The businesses failing with AI agents are the ones that try to automate everything at once. The businesses winning are the ones that deploy one agent, in one area, get it working reliably, and then expand.

## Step 1 — Find your highest-value repeatable task

---

**The best first AI agent is not the most ambitious one. It is the one that solves a specific, well-defined, repeatable problem.**

Ask yourself: what do I (or my team) do repeatedly that follows a predictable pattern? The answer to that question is where your first agent should go.

High-value starting points for most businesses:

**Customer support responses** — Answering the same questions repeatedly. An agent can handle first-line support — answering FAQs, routing complex issues to humans, following up on tickets.

**Content creation** — Writing social media posts, blog articles, email newsletters, product descriptions. An agent can produce a first draft; you refine and approve.

**Research and summarisation** — Researching competitors, summarising industry news, compiling information from multiple sources. An agent can do this in minutes rather than hours.

**Data entry and processing** — Extracting information from emails or documents and entering it into a spreadsheet or CRM. Tedious for humans; straightforward for agents.

**Email drafting** — Writing first drafts of emails — proposals, follow-ups, updates. The agent drafts; you review and send.

**Scheduling and coordination** — Managing calendar bookings, sending reminders, coordinating meeting times.

**Lead qualification** — Reviewing inbound enquiries, asking qualifying questions, scoring leads before they reach a human salesperson.

## Step 2 — Choose the right type of agent

Not all agents are right for all tasks. In 2026 the choice is simpler than it looks, because every option falls into a handful of shapes:

| If you want to...                                              | Use this                                                                  | Cost                                     |
|----------------------------------------------------------------|---------------------------------------------------------------------------|------------------------------------------|
| Delegate work inside your own files, inbox, calendar and apps  | Cowork (Anthropic), ChatGPT Work (OpenAI), Copilot Cowork (Microsoft 365) | From \$20/mo, included in the paid plan  |
| Run a team of agents sharing one cloud computer and its logins | Grok Bot (xAI)                                                            | Subscription                             |
| Have a personal agent you message like a colleague             | OpenClaw (open-source), Meta Muse                                         | Free + model costs; Muse has a free tier |
| Let agents build and run a whole company for you               | Polsia, NanoCorp, Cofounder.co                                            | From \$20/mo + fees (Polsia)             |
| Run several agents as an org chart you control                 | Paperclip (self-hosted)                                                   | Free + model costs                       |
| Connect apps and automate defined workflows                    | Zapier, n8n                                                               | From ~\$20/mo; n8n is open-source        |
| Build a custom agent with full control                         | Claude Developer Platform / Agent SDK                                     | Pay per token                            |

For a first agent, the top row is almost always the answer. A work agent on a plan you already have, pointed at one folder and one task, is the lowest-risk way to learn what this is actually like.

## Step 3 — Write a clear job description for your agent

**This is the most important step most people skip.**

AI agents work best when given precise, well-defined instructions. Vague goals produce vague results. Before deploying an agent, write out — in plain English — exactly what you want it to do.

A good agent job description includes:

→ **The goal** — what outcome do you want? ("Respond to customer support emails within 2 hours")

→ **The inputs** — what information does the agent receive? ("An inbound email from a customer")

- **The outputs** — what should the agent produce? ("A draft reply, flagged for human review if the issue is a refund request")
- **The constraints** — what should the agent NOT do? ("Never promise a refund without human approval")
- **The escalation path** — when should it involve a human? ("Any complaint that mentions a legal threat")

The more specific you are, the better the agent performs. "Handle customer support" is a bad brief. "Read inbound customer emails, identify the question type from this list of 12 categories, draft a reply using these templates, and flag any complaint mentioning a refund or legal issue for human review" is a good brief.

## Step 4 — Start with read-only access

---

**The most common mistake when deploying agents: giving them too much access too soon.**

Start by giving your agent access to read information, not to take actions. Let it draft emails — do not let it send them. Let it suggest calendar bookings — do not let it confirm them. Let it write social posts — do not let it publish them.

Run in this "human-in-the-loop" mode for a week or two. Review everything the agent produces. Identify where it gets things right consistently, and where it makes mistakes.

Once you trust the agent's judgment in a specific area, you can expand its permissions incrementally. Never expand all at once.

## Step 5 — Measure results honestly

---

**Define what success looks like before you deploy — not after.**

For each agent you deploy, set a specific metric:

- Customer support agent: average response time, customer satisfaction score, escalation rate
- Content agent: posts published per week, engagement rate, time saved
- Research agent: hours saved per week, accuracy of outputs
- Lead qualification agent: leads qualified per day, conversion rate vs human qualification

Review these metrics weekly for the first month. If the agent is not delivering measurable value, adjust the brief before expanding its scope.

# Five deployments that work — and what to expect

---

These are the patterns that come up again and again among small businesses that get value from a first agent. Results vary; the shape of the win does not.

## First-line customer support

An agent reads incoming support emails, categorises them, drafts replies from approved templates, and flags anything unusual for a human. The win is response time — hours become minutes — and a majority of routine emails handled without anyone touching them. The risk is a confident wrong answer, which is why replies are drafted for approval before they are sent unsupervised.

## Research and content preparation

An agent researches a client's industry each week and drafts a content calendar, a briefing or a first draft. Writers spend their time on strategy and editing instead of gathering. The win is capacity without headcount.

## Inbox triage for partners and owners

An agent summarises the day's email, drafts responses for review, and surfaces what needs a decision. Owners consistently report an hour or two a day back. The risk is the agent guessing at tone on a sensitive thread; keep it drafting, not sending, for those.

## The solo operator's assistant

A personal agent — OpenClaw over WhatsApp or Telegram, or Cowork on the desktop — briefed with research tasks, email drafts, document summaries and admin. Typical cost is tens of dollars a month in model usage. This is the deployment most one-person businesses start with.

## Routine engineering

Claude Code or Codex handling bug fixes, small features and tests, with senior people on architecture and review. Teams that do this well ship considerably more with the same headcount. Teams that skip the review step ship considerably more bugs.

# The honest cost — what agents actually cost

---

AI agents are significantly cheaper than human employees. But they are not free. Here is what to budget:

| Agent type              | Monthly cost (approx) | What drives the cost                                         |
|-------------------------|-----------------------|--------------------------------------------------------------|
| Cowork / ChatGPT Work   | \$20–200/mo           | Your Claude or ChatGPT plan; heavy use needs the higher tier |
| Copilot Cowork          | \$30/user/mo          | Add-on to Microsoft 365                                      |
| OpenClaw (personal)     | \$20–100/mo           | Model API usage — Claude, GPT or a cheaper open model        |
| Polsia company          | \$20–1,000/mo + fees  | Plan tier, credits, 20% on ad spend, 3% on revenue           |
| Paperclip org           | \$0 + model usage     | Free software; each agent's model bill                       |
| Zapier / n8n automation | \$0–100/mo            | Task volume and plan tier                                    |
| Custom Claude agent     | \$50–500/mo           | Token usage — depends heavily on task volume                 |
| Claude Code             | \$20–200/mo           | Included in Claude Pro/Max; API if you go beyond             |

For context: a part-time human assistant in the US costs \$2,000–3,500/month. A full-time junior employee costs \$45,000–60,000/year. An agent handling equivalent work costs \$50–500/month. The economics are not subtle — and they are the reason the Agentic Economy guide spends a section on what that means for jobs.

## What agents cannot do — be honest with yourself

AI agents are powerful. They are not magic. Here is what they cannot reliably do:

- Build genuine human relationships (clients, partners, key hires)
- Handle truly novel situations with no precedent
- Make judgment calls in highly regulated areas without human oversight
- Replace the strategic vision and creative direction of a founder
- Guarantee factual accuracy on topics outside their training
- Handle physical world tasks (no hands, no body)

The businesses that deploy agents most successfully are the ones that are clear-eyed about this. Use agents for what they are good at. Keep humans for what humans are good at. The boundary between the two is moving — but it has not disappeared.

## Your action plan — what to do this week

**Day 1** — Write down the three most repetitive tasks in your business. Pick the one with the clearest inputs and outputs.

**Day 2** — Write the "job description" for your first agent. Goal, inputs, outputs, constraints, escalation path.

**Day 3** — Sign up for one platform from Step 2. Start with the free tier or lowest plan.

**Day 4-5** — Configure the agent with your job description. Run it in read-only mode — review everything it produces.

**Week 2** — Measure the output quality. Refine the brief where it is getting things wrong. Keep reviewing.

**Week 3-4** — If quality is consistently good, expand permissions incrementally. Track the time saved.

**Month 2** — Deploy a second agent in a different area. Build your agent team the same way you would build a human team — one role at a time.

**The businesses that win with AI agents are not the ones with the biggest budgets or the most technical teams. They are the ones that start now, learn fast, and iterate.**

## Glossary

---

**AI Agent** — Software that pursues goals autonomously — taking actions, using tools, and adapting without constant human instruction.

**Human-in-the-loop** — A deployment mode where an agent produces outputs but a human reviews and approves before actions are taken.

**Polsia** — A managed platform where AI agents build and run companies. From \$20/mo plus credits and fees; see its own guide.

**OpenClaw** — Free, open-source personal AI agent. Runs locally, takes instructions via messaging apps.

**Paperclip** — Open-source multi-agent orchestration framework for running businesses with teams of AI agents.

**Cowork** — Anthropic's work agent: delegates tasks across your files, apps and browser. Included in every paid Claude plan.

**ChatGPT Work** — OpenAI's work agent: give it an outcome, it returns finished spreadsheets, decks and apps.

**Zapier / n8n** — Automation platforms that connect apps and run defined workflows, now with agent steps. n8n is open-source and self-hostable.

**Claude** — Anthropic's AI model. Powers OpenClaw, Claude Code, and many business agent deployments.

**API key** — Your credential for accessing an AI model's API. Required for OpenClaw, Paperclip and custom builds; Claude Code and Cowork run on your Claude plan instead.

**Token** — The unit AI models use to process text. You pay for API usage in tokens. Roughly 0.75 words per token.

**MCP (Model Context Protocol)** — The open standard connecting AI agents to external tools. Invented by Anthropic.

## Liked this? Go deeper.

---

You now know how to hire an agent. Here is the full context behind the tools you are using.

### WTF is Cowork

Claude Code was for developers. Cowork is for everyone else. What it does, what it costs, and how it changed the work-agent market.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an AI Agent

What an agent actually is, how it thinks, the main types, real examples, and the risks. The foundation for everything else in the series.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an API

The thing connecting everything in the agentic economy — explained simply, finally.

[wtfagents.com/store](https://wtfagents.com/store)

### WTF is an LLM

Large language models are the engine behind every AI company. Here's how they actually work.

[wtfagents.com/store](https://wtfagents.com/store)

---

## Want the rest of the series?

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. Buy them individually at \$7, or take the Starter Pack for \$29 or the Complete Pack for \$49 with every guide in the series.

[wtfagents.com/store](https://wtfagents.com/store) →

---



**Scan to browse all guides** Every WTF Agents guide at [wtfagents.com/store](https://wtfagents.com/store) \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series Also free: the live AI company directory at [wtfagents.com/companies](https://wtfagents.com/companies)

© 2026 WTF Agents. All rights reserved.