

WTF is Devin?

The "first AI software engineer": a cloud agent that takes a ticket and returns working software, from a company valued at \$47 billion that says it writes 95% of its own code with it. What it does, why the price fell from \$500 to \$20, and how it compares.

The one-liner

Devin is an autonomous software engineer in the cloud: you give it a task — a bug, a feature, a migration — and it plans, writes, tests, fixes and opens the pull request, then messages you when it is done or stuck. It was the first product to make that claim in 2024 and the one everyone measured against.

Two years on it is cheaper, better, and part of a bigger company: Cognition bought the Windsurf editor, raised at \$47 billion, and says Devin now writes about 95% of Cognition's own code — up from 89% at the May round, and a quarter of its pull requests in 2025.

That last number is worth holding at arm's length. It is the company's own figure, self-reported by CEO Scott Wu, and it counts code committed rather than engineering done: the specifying, reviewing and deciding still sit with people. As a trend line it is the most interesting thing Cognition publishes. As a benchmark it is unaudited.

By the numbers — September 2026

\$47B

the valuation Bloomberg reported on 2 September 2026, for a ~\$1B round still in talks — up from \$26B in May and \$10.2B a year ago

~\$900M

annualised revenue, more than double the \$492M of May 2026; Devin alone was \$73M in June 2025, before Windsurf

\$500 → \$20

the monthly price cut with the Devin 2.0 relaunch in April 2025; Max at \$200, Teams from \$80

1 ACU

Devin's unit of work, roughly 15 minutes of autonomous engineering

July 2025

Cognition buys Windsurf after OpenAI's \$3B bid collapses and Google hires Windsurf's founders

4,000+

enterprise customers on Windsurf by spring 2026; over a million users

3.0

the current Devin: dynamic re-planning, desktop-app testing via computer use, a unified Devin Desktop

What it actually does

Takes a task, returns a pull request — Describe the change in plain English or paste a ticket. Devin explores the codebase, makes a plan, edits, runs the tests, fixes what breaks and hands back a reviewable change with an explanation.

Works in parallel — Each task runs in its own cloud environment, so ten tickets can run at once. This is where it earns its keep on migrations and backlogs.

Re-plans — Devin 3.0 revises its approach mid-task when something surprises it, rather than ploughing on. The difference between an agent that gets stuck and one that gets there.

Tests what it built — Including desktop and web apps, by driving them the way a person would (computer use).

Devin Desktop — The 2026 app unifying the cloud agent, the command line and the editor, so you can watch, steer and take over.

Windsurf — The editor Cognition bought: AI-native, with completions and chat like Cursor, and the on-ramp to Devin for longer tasks. Renamed Devin Desktop in June 2026, and now open to rival agents — Codex, Claude, Gemini CLI — through the Agent Client Protocol.

Multi-model — Claude, GPT and Gemini frontier models plus Cognition's own SWE family, routed per task. Cognition's pitch is independence from any one lab.

What it costs

Plan	Price	What you get
Free	\$0	A small ACU allowance to try it
Pro	\$20/mo	A usage allowance refreshing daily and weekly; extra usage at API rates
Max	\$200/mo	The heavy-use tier, for running Devin all day
Teams	\$80/mo + \$40/seat	Shared ACU pool, admin, integrations
Enterprise	Custom	Volume ACUs, security, dedicated support, custom models

The unit is the ACU, about a quarter-hour of Devin working. A simple bug might take two; a multi-file feature ten or more; a large migration hundreds. The \$20 plan is for trying tasks; the

\$200 plan is for handing over a backlog.

Cognition does not publish the ACU allowance on each plan — the pricing page says only that paid plans refresh daily and weekly, and that extra usage bills at API rates. Third-party write-ups quote confident quotas and per-ACU prices; they disagree with each other and with the pricing page. Watch your own meter for a month before sizing the plan.

Devin vs Claude Code vs Codex

	Devin	Claude Code	Codex
Shape	Cloud engineer you message	Agent in your terminal, desktop, cloud	Agent in ChatGPT and the CLI
Model	Multi-model + own SWE models	Claude	GPT
Unit	ACU (~15 min of work)	Plan usage	Plan credits
Entry price	\$20	\$20 Claude Pro	\$20 ChatGPT Plus
Best at	Long, parallel, hands-off tasks	Hard problems needing care; leads benchmarks	Parallel backend tasks, GitHub review
Owner	Cognition, independent	Anthropic	OpenAI

The honest distinction after two years: Devin is the most hands-off of the three and the best at running a queue of well-specified tasks unattended. Claude Code is what people reach for when the task is hard. Many teams use both.

The Windsurf story

In mid-2025 Windsurf, a well-liked AI editor, was in a bidding war. OpenAI agreed to buy it for about \$3 billion; the deal collapsed. Google hired Windsurf's CEO and research leads in a licensing deal. Days later Cognition bought what was left: the product, the brand, the engineers and \$82 million of revenue with 350-plus enterprise customers.

It turned out to be the acquisition that made the company. Devin had the headlines; Windsurf had the users. Combined revenue more than doubled within months, the valuation went from \$10 billion to \$26 billion to \$47 billion in a year, and Cognition became a coding platform instead of a single agent. The lesson it teaches about the whole category: distribution beat being first.

The catches

- **ACUs are hard to predict.** A task that seems small can take many; the meter is honest but the estimate is yours. Start with the free tier and watch what real tickets cost.
- **Cloud-only means your code leaves.** Devin works in Cognition's environments. Enterprise plans have the controls; individuals should know that is the model.
- **Autonomy needs specification.** Devin does well with a clear ticket and badly with a vague one. The skill it rewards is writing the brief.
- **Valuation pressure.** At roughly fifty times revenue, Cognition has to keep growing very fast. Expect pricing and packaging to keep moving.
- **The bundling threat.** Anthropic, OpenAI, Google and Microsoft can make autonomous coding a default feature of what developers already pay for. Devin has to stay better, not just independent.

Do you need to be a developer?

Less than for Cursor, more than for Claude Code's desktop app. Devin assumes a repository, tests and someone to review a pull request; it does not assume you write code yourself. A founder who can read a diff and describe a change precisely can get real software out of it. If that is not you yet, start with Claude Code and come back when you have a backlog to hand over.

Why it matters

Devin defined the category: the agent you message like a colleague, that returns finished work. Everything in the coding-agent row of this series — Claude Code, Codex, Grok Bot's engineering bots — is a response to the claim Devin made first. That it needed to buy an editor to grow is the second lesson, and the one the rest of the agentic economy is learning now: the agent is the product, but the place people already work is the distribution.

Glossary

Devin — Cognition's autonomous AI software engineer, launched 2024. Takes a task, returns a pull request. Now version 3.0.

Cognition — The company behind Devin and Windsurf, founded 2023 by Scott Wu, Steven Hao and Walden Yan. Valued at \$47B.

ACU (Agent Compute Unit) — Devin's unit of work, roughly 15 minutes of autonomous engineering.

Windsurf — The AI-native code editor Cognition acquired in July 2025, renamed Devin Desktop in June 2026.

Devin Desktop — Windsurf, renamed in June 2026: one app for the cloud agent, the CLI and the editor, and for rival agents too.

Pull request — A proposed change to a codebase that a human reviews before merging. Devin's output.

Re-planning — Revising the approach mid-task when something unexpected happens. Devin 3.0's headline capability.

SWE models — Cognition's own coding models, used alongside Claude, GPT and Gemini.

Computer use — An agent driving an application's interface the way a person would; Devin uses it to test what it builds.

Liked this? Go deeper.

Devin is the engineer you message. Here are the others, and the editor SpaceX bought.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

WTF is Claude Code

The autonomous developer that changed Anthropic's trajectory. How it works, how it compares, and how non-developers are using it now.

wtfagents.com/guides/claude-code

WTF is Codex

OpenAI's coding agent, now a tab in ChatGPT. What it does, how it compares with Claude Code, what it costs on each plan, and the catches.

wtfagents.com/guides/codex

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent