

WTF is n8n?

The open-source automation layer solo builders wire their agents with: one workflow, one execution, any number of steps, on a \$5 server if you want. What it does, why it beats Zapier at volume, when it doesn't, and what a \$5.2 billion valuation says about it.

WTF Agents · wtfagents.com · September 2026
Part of the WTF Agents Guide Series · wtfagents.com

The one-liner

n8n is a workflow automation platform: a visual canvas where a trigger (a form submitted, an email received, a schedule) flows through steps (call an AI model, look up a record, send a message, update a sheet) to an outcome, without code — or with as much code as you like. It is open-source, runs on your own server if you want, and has become the plumbing under a large share of the indie agent stacks in this series.

Its billing model is the whole argument: one execution is one complete run of a workflow, however many steps. Zapier bills every step.

By the numbers — September 2026

\$5.2B

valuation after SAP's strategic investment, announced 12 May 2026 — double the \$2.5B of the October 2025 Series C

1.7M

monthly active builders, and more than 1,400 enterprise customers, per the company

1

execution billed per workflow run, whatever its length — versus one task per step at Zapier

€20 / €50

per month for Starter (2,500 executions) and Pro (10,000), billed annually — workflows, steps and users are unlimited on every plan

\$0

for the self-hosted Community Edition: no execution, workflow or user limits; you pay for the server

500+

native integrations, plus an HTTP node that reaches any API — versus 8,000-plus at Zapier

2.0

the January 2026 release that made agents first-class: an Agent node with tool-calling, memory and retries

What it actually does

Workflows on a canvas — Nodes connected by lines. A trigger starts it, each node transforms or acts, branches and loops handle the logic. Readable by a non-developer, editable by one after a week.

The AI Agent node — Since 2.0: a node that runs a model with tools, memory and retry logic. Point it at Claude, GPT or a local model via Ollama, give it tools from other nodes, and it decides the steps. This is where n8n stopped being "Zapier but cheaper" and became an agent runtime.

Any API — The HTTP node talks to anything with a REST or GraphQL interface, including systems no integration exists for: your accountant's practice software, a client's internal tool, a niche SaaS. Zapier waits for someone to build the connector; n8n lets you.

Code when you want it — JavaScript or Python inside a node, for the 5% of steps that need it. The rest stays visual.

Self-hosting — Run it on a \$5–20 virtual server, in Docker, and keep every byte of data inside your own infrastructure. The reason regulated businesses pick it.

Databases, files, MCP — First-class Postgres, MongoDB and Redis; file handling; and MCP support so agents elsewhere can call n8n workflows as tools and n8n agents can use MCP servers.

Templates — Thousands of community workflows to start from: lead scoring, content pipelines, inbox triage, the standard shapes of small-business automation.

What it costs — and the maths that decides it

Plan	Price (annual)	Includes
Community (self-hosted)	€0 + server	Unlimited everything; you run it
Starter (cloud)	€20/mo	2,500 executions; workflows, steps and users unlimited
Pro (cloud)	€50/mo	10,000 executions, more concurrency and projects
Business (cloud)	€667/mo	40,000 executions, SSO, Git, governance
Enterprise	Custom	Self-hosted with support, or cloud at volume

n8n dropped its active-workflow caps: every paid plan now carries unlimited workflows, steps and users, and only executions, concurrency and projects are metered. So the comparison that matters is per workflow run. Take a ten-step lead-qualification flow that fires a thousand times a month: on n8n that is 1,000 executions, inside the Starter plan. On Zapier it is 10,000 tasks, well past the Professional plan's 750 and into a tier ten times the price. Self-hosted, it is the server bill and nothing else.

Below about a thousand multi-step runs a month, the difference is small and Zapier's ease wins. Above it, n8n wins by a wide margin. Count your steps and your runs before choosing.

n8n vs Zapier vs Make

	n8n	Zapier	Make
Billing unit	Execution (whole run)	Task (each step)	Operation (each step, cheaper)
Integrations	500+ native, any API	8,000+	2,000+
Self-host	Yes, free	No	No
AI agents	Agent node, local models	Zapier Agents	AI Agents
Learning curve	Real: a week	Gentle	Medium
Best for	Volume, control, data privacy, agents	Speed, non-technical operators	Visual complexity on a budget

Zapier is the right first tool for most small businesses: it is easier, it connects to more things, and up to about \$100 a month of spend the time it saves pays for itself. n8n is what you move to when the bill passes \$200, the data must stay in-house, or you want a real agent with memory and tools rather than a packaged assistant.

The catches

- **Free means you host it.** Updates, backups, SSL, uptime and the 11pm "why did it stop" are yours. If nobody on the team can run a server, the Community Edition is not free; it costs your evenings. Cloud plans exist for that reason.
- **The learning curve is real.** Data mapping between nodes, expressions and error handling are harder than Zapier. Expect a week of friction before it pays back.

- **Fewer connectors.** Five hundred native integrations against Zapier's eight thousand. The HTTP node covers the gap if you can read an API document.
- **Agents amplify mistakes.** A workflow that emails a thousand people runs a thousand times whether the logic was right or not. Test with a small trigger, and use 2.0's human-in-the-loop tool gates in front of anything that deletes, sends or spends.
- **SAP is now an investor.** A strategic minority stake, not control, and the real substance is distribution: n8n's canvas is being embedded in SAP's Joule Studio. Good for enterprise credibility; watch whether the roadmap tilts toward SAP customers and away from the solo builder who made it.

Who it is for

Yes — Solo founders and small teams with someone comfortable running a server or paying €20 for cloud, whose automations have grown past a handful of simple zaps, and who want agents with memory and tools inside their workflows.

Not yet — Anyone with fewer than a few hundred workflow runs a month and no technical appetite. Zapier, and revisit when the bill or the data question forces it.

No — Anyone who wants an agent to run their business for them (that is Polsia or NanoCorp) rather than a system to wire their own.

Why it matters

Every agent in this series needs a way to be triggered, to reach the next tool, and to hand off. n8n is that connective tissue for the part of the agentic economy that is not run by a hyperscaler: open-source, self-hostable, execution-priced, and now with agents as a native node. That a company built on "you can run it yourself for free" is worth \$5.2 billion with SAP on the cap table says how much of the real work in this economy happens in workflows nobody outside the business ever sees.

Glossary

n8n — Open-source workflow automation platform, Berlin, founded 2019. Self-hostable or cloud; execution-based pricing. \$5.2B valuation.

Workflow — A trigger plus a chain of nodes that transform data and take actions.

Node — One step in a workflow: an integration, a transformation, a code block, an agent.

Execution — One complete run of a workflow, however many nodes. n8n's billing unit.

Task — Zapier's billing unit: one per step, per run.

AI Agent node — n8n's node that runs a model with tools, memory and retries, and decides the steps itself.

HTTP node — The node that calls any API, covering integrations n8n does not ship natively.

Self-hosted / Community Edition — Running n8n on your own server; free, unlimited, your responsibility.

Ollama — A way to run open-weight models locally; n8n's Agent node can use them, keeping data off any cloud.

Make — The third major automation platform; visual, operation-priced, cheaper than Zapier at volume.

Liked this? Go deeper.

n8n is the wiring. Here is what it connects.

GO DEEPER

Get the rest of the series

Every WTF Agents guide is written the same way — plain English, no hype, no jargon. \$7 each, or take a bundle.

Starter Pack \$29

Complete Pack \$49

wtfagents.com →



Scan to browse all guides · wtfagents.com \$7 each · Starter Pack \$29 · Complete Pack \$49 — every guide in the series

How to Hire an AI Agent for Your Business

The practical guide. Pick the task, pick the agent, write the brief, start read-only, measure honestly. No code, no prior AI experience.

wtfagents.com/guides/hire-agent

WTF is MCP

The Model Context Protocol, explained. The universal plug between agents and tools: what it is, who runs it now, and the security catch.

wtfagents.com/guides/mcp

WTF is Grok Bot

Elon Musk's work agent: a team of bots sharing one cloud computer and your real logins, coordinated by a Chief of Staff. What it does, costs and risks.

wtfagents.com/guides/grok-bot

WTF is the Agentic Economy

The big picture. What actually happened in 2026, who the players are, what is real and what is hype. Start here.

wtfagents.com/guides/agentic-economy

WTF is ChatGPT Work

OpenAI's work agent, inside ChatGPT. Give it an outcome, get finished files back. What it does, what it costs on your plan, and the catches.

wtfagents.com/guides/chatgpt-work